

Compilation of all my lecture notes

toc:

- Computer System Overview
 - What Operating Systems Do
 - Computer-System Organization
 - Competing for Resources
 - Computer-System Operation
 - Instruction Execution
 - Interrupts
 - Instruction Cycle with Interrupts
 - Program Status Word
 - Flow Control With and Without Interrupts
 - Multiple Interrupts
 - Storage Structure
 - Von Neumann Architecture
 - Main Memory
 - ROM, EPROM, and EEPROM
 - ROM
 - EPROM
 - EEPROM
 - Secondary Storage
 - Tertiary Storage
 - Storage Hierarchy
 - Caching
 - I/O Structure
 - Direct Memory Access Structure
- Operating System Structures
 - Computer System Architecture
 - Operating System Operations
 - Multiprogramming and Multitasking
 - Dual Mode
 - Transition from User to Kernel Mode
 - Resource Management
 - Process Management
 - Memory Management
 - File-System Management
 - Mass-Storage Management
 - Caching Management
 - Cache Coherence
 - I/O Subsystem Management
 - Operating System Services
 - For the User
 - For the System

- User Operating System Interfaces
- System call
- Process Description and Control
 - What is a Process
 - Process States
 - Process Trace
 - Simple Two-State Process Model
 - Operations on Processes
 - Creation
 - fork
 - CreateProcess
 - Tree of Processes in Linux
 - Termination
 - Five-State Process Model
 - Suspended Processes
 - Process Description and Control
 - Memory Tables
 - I/O Tables
 - File Tables
 - Process Tables
 - Process Control
 - Process Location
 - Process Attributes
 - Process Control
 - Creation
 - Switching
 - Execution of the Operating System
 - Separate Kernel
 - User Processes
 - Process-Based OS
 - Summary
- Threads
 - Processes and Threads
 - Threaded Approaches
 - What is a Thread
 - Why Multithread?
 - Multicore Programming
 - Concurrency vs Parallelism
 - Amdahl's Law
 - Thread States and Operations
 - Thread Synchronization
 - Multithreading Models
 - User Threads
 - Kernel Threads
 - Combined User and Kernel level threading
- Concurrency: Mutual Exclusion and Synchronization

- Background for Concurrency
- Critical-Section Problem
 - Single Core
 - OS
- Mutual Exclusion
 - Software Approaches
 - Dekker's Algo
 - Peterson's Solution
 - Hardware Support Approaches
 - Memory Barrier
 - Hardware Instructions
 - Atomic Variables
 - OS & Programming Languages Approaches
 - Mutex Locks
 - Semaphores
 - Monitors
- Deadlock and Starvation
 - Principles of Deadlocks
 - Deadlock
 - Deadlock in Multithreaded Applications
 - Deadlock Characterization
 - Resource Allocation Graph
 - Methods for Handling Deadlocks
 - Prevention
 - Deny Hold and Wait
 - Deny No Preemption
 - Deny Circular Wait
 - Deadlock Avoidance
 - Safe State
 - Avoidance Algos
 - Resource-Allocation Graph Scheme
 - Banker's Algorithm
 - Deadlock Detection
 - Recovery from Deadlock
- CPU Scheduling
 - Recap for CPU Scheduling
 - Basics for CPU Scheduling
 - Scheduling Criteria
 - Scheduling Algos
 - Multi-Processor Scheduling
 - Multicore
 - Multithreaded Cores
 - Load Balancing
 - Processor Affinity
 - NUMA Systems
 - Heterogeneous Multiprocessing

- Real-Time CPU Scheduling
 - Priority Based Scheduling
- Algo Eval
 - Deterministic
 - Queueing Models
 - Simulations
 - Real Implementation
- Main Memory
 - Background for Memory Management
 - Memory Protection
 - Logical vs Physical Address Space
 - Memory Management Unit
 - Dynamic Loading
 - Dynamic Linking
 - Contiguous Memory Allocation
 - Memory Protection in Contiguous Memory Allocation
 - Memory Allocation
 - Fragmentation
 - Paging
 - Paging Hardware
 - Example
 - Calculating Internal Fragmentation
 - Free Frames
 - Implementation of Page Table
 - TLB
 - Effective Access Time
 - Memory Protection in Page Table
 - Shared Pages
 - Swapping
- Virtual Memory
 - Background for Virtual Memory
 - Demand Paging
 - Copy on Write
 - Page Replacement
 - Page Replacement Algorithms
 - Page Replacement Algo Evaluation
 - Different Types
 - FIFO
 - Optimal Algo
 - Least Recently Used (LRU) Algorithm
 - Second-Chance Algo (Clock) - LRU Approximation
 - Enhanced Second-Chance Algo
 - Counting Algos
 - Allocation of Frames
 - Thrashing
 - Demand Paging

- Working-Set Model
- Allocating Kernel Memory
- Storage Management
 - Mass-Storage Structure
 - Overview of Mass Storage Structure
 - Hard Disks
 - Nonvolatile Memory Devices
 - Disk Attachment
 - Address Mapping
 - HDD Scheduling
 - First Come First Serve
 - SCAN
 - C-SCAN
 - NVM Scheduling
 - RAID Structure
 - I/O Systems
 - Overview
 - I/O Hardware
- Security and Protection
 - Security
 - How Systems Fail
 - What is Security
 - Info Security Strats
 - Protection
 - Goal of protection
 - Principles of Protection
- Exam Review
 - Exam Housekeeping
 - QnA
 - Final Exam Practice

Computer System Overview

What Operating Systems Do

Components to a computer structure:

- user
- application programs
- operating system
- computer hardware

OS sits in between the user and the hardware, handling how the application programs interact with the hardware.

OS Goals:

- execute user programs
- convenient and easy to use
 - not hard to use
- efficiently use hardware

user pov:

- want
 - easy to use
 - good performance
 - secure
- don't care about efficiency

Smartphones and tablets are good from the user pov since they are optimized for usability and batter life.

The touch screens and voice interactions allow for better usability.

Even though in reality the phone could be pretty wasteful when it comes to the resources.

Profesor's Soapbox:

- we should be more concerned with the lack of privacy that comes with the convenience
- she doesn't like the argument that if you don't have anything to hide then you shouldn't care
- "privacy is our right"
- she wants us to take into account privacy and security regarding user data

Embedded Computers:

- little/no user interface
- ex
 - home devices
 - motion sensor lights
 - computerized water systems
 - computerized heating system
 - car computers
 - numeric keypads
 - presto scanner
- run primarily without user intervention

System pov:

- OS must keep all users happy in the case of shared computers

- OS is the resource allocator and control program
 - efficiently manage the execution of user programs
- I/O devices are important in this respect
 - hard to manage and work with efficiently

Operating System definition

- no universally accepted definition
 - “we define the operating system as we go”
- is a resource allocator
 - manage all resources
 - decides between conflicting requests for efficient and fair resource use
- is a control program
- controls execution of programs to prevent errors and improper use of the computer
- everything a vendor ships when you order an operating system is a good approximation but varies wildly
- the one program running at all times on the computer is the kernel which is part of the operating system
- everything else is either
 - a system program that ships with the operating system
 - or
 - an application program

Mobile operating systems often include not only a core kernel but also middleware

- a set of software frameworks that provide additional services to application developers

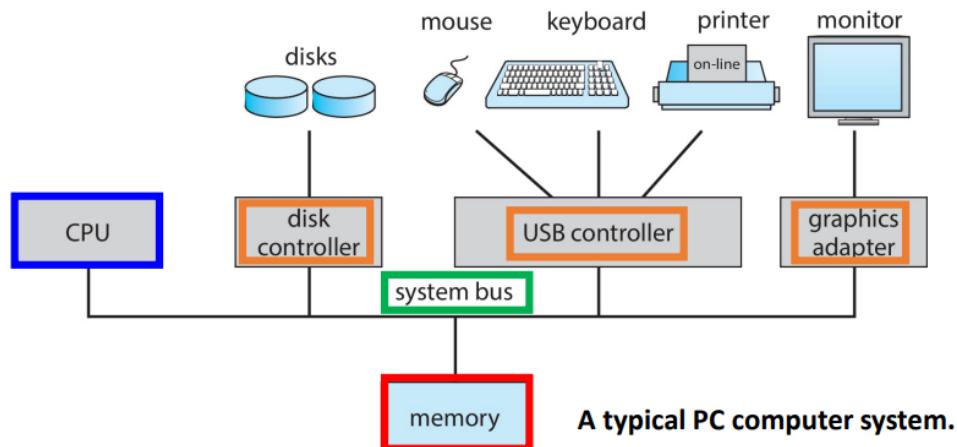
kernel is the core of the system

not all computer systems have some sort of user interaction

Computer-System Organization

Computer System Organization

- Consists of one or more **CPUs** and a number of device **controllers** connected through a common **system bus** that provides access between components and shared **memory**.



cpu talks to the memory in order to get instructions

memory is typically volatile so it loses all information once it shuts down

system bus:

- provides communication among processors, main memory, and I/O modules

Device controllers are things like the disk controller and usb controller.

There's local memory on the controller that acts as a buffer storage.

the controller moves data between the peripheral devices and the buffer storage.

There are software versions of these drivers too.

Competing for Resources

CPU and controllers can execute in parallel, competing for memory cycles.

They both need to go through the memory

- cpu fetches instructions from memory
- I/O device sends information through memory
- modern computers use DMA (direct memory address) that we speak on later

Computer-System Operation

How does the processor execute an instruction?

There are registers inside of the cpu itself

- this is based off of the architecture (design) of the processor itself
- very fast storage to read and write to
- used to save, store, and use instructions
- very expensive, fast, and limited

Memory Address Register (MAR) specifies the address in memory for the next read or write

Memory Buffer Register (MBR) contains the data to be written/read into/from memory

I/O Address Register (I/O AR) specifies a particular I/O device we're working with

I/O Buffer Register (I/O BR) used for the exchange of data between an I/O module and the processor.

CPU is waiting for the input

Instructions always have to come from memory. It can only be loaded from the memory.

I/O doesn't input instructions, it inputs/outputs data.

CPU can pull directly from IO device

Memory and I/O:

- memory is a set of locations defined in sequentially numbered addresses
- each location contains a bit pattern that can be interpreted as either an instruction or data
- an I/O module transfers data from external devices to processor and memory, and vice versa

Instruction Execution

A program to be executed by a processor consists of a set of instructions stored in memory.

The processor reads (fetches) instructions from memory one at a time and executes each instruction.

The processing required for a single instruction is called an instruction cycle.

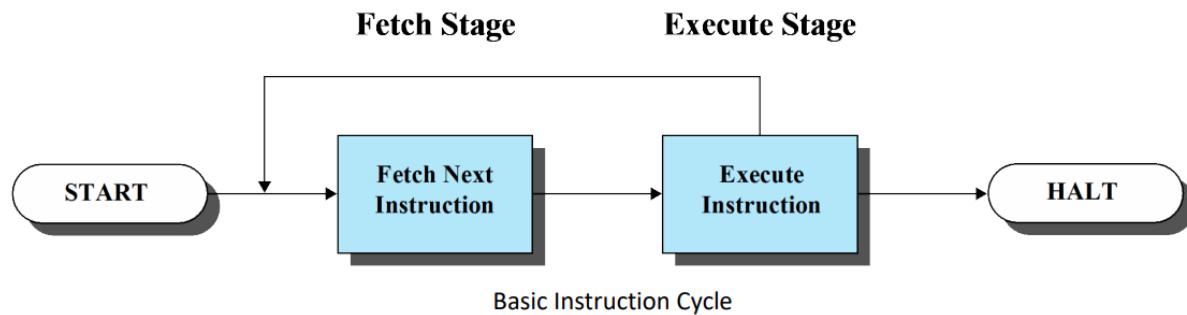
Fetch and Execute is an atomic operation:

- If you fetch an operation then you can't do anything else until you execute that operation
- they come together

Program halts if:

- processor turns off
- unrecoverable error occurs

- program instruction that halts the processor is encountered
 - tell the processor to wait for an input or output



Instruction Fetch and Execute from CPU pov:

- processor fetches an instruction from memory
- program counter (PC) holds the address of the next instruction to be fetched
- PC is incremented after each fetch
 - on to the next instruction
 - unless the instruction loaded tells us to change the PC differently
- fetched instruction is loaded into the instruction register (IR)
- the processor interprets the instruction and performs the required action:
 - processor-memory: data transfer
 - Processor-I/O: data transfer
 - Data processing: arithmetic or logic operations
 - Control: an instruction may specify that the sequence of execution be altered.

Hypoethical 16 bit machine

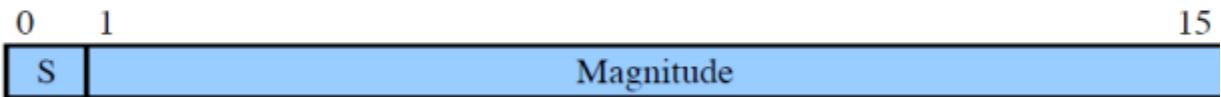
opcode - operation code that tells the processor what to do

address - data needed

S - sign of the data (0 = +ve, 1 = -ve)



(a) Instruction format



(b) Integer format

Program counter (PC) = Address of instruction
 Instruction register (IR) = Instruction being executed
 Accumulator (AC) = Temporary storage

(c) Internal CPU registers

0001 = Load AC from memory
 0010 = Store AC to memory
 0101 = Add to AC from memory

(d) Partial list of opcodes

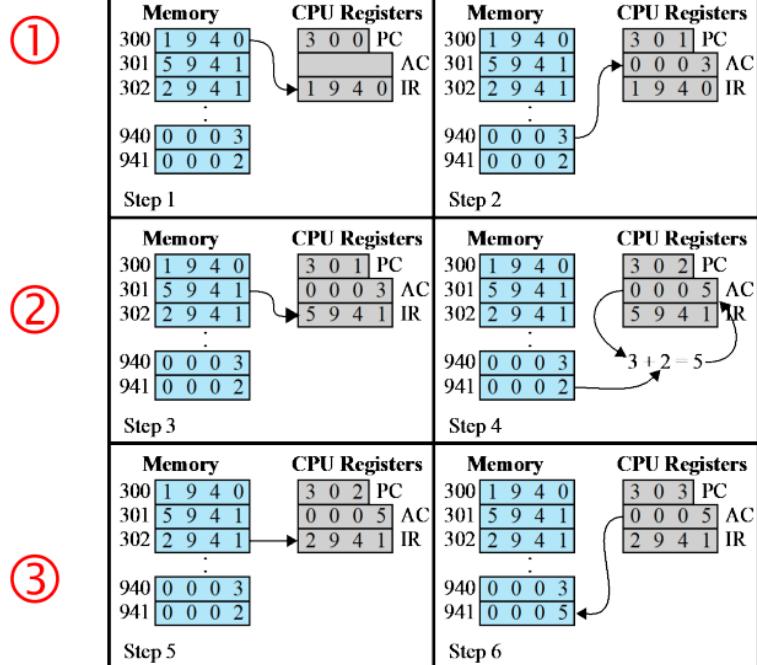
These are all displayed in hex

Example of Program Execution

Contents of memory and registers in hexadecimal.

0001 = Load AC from memory
 0010 = Store AC to memory
 0101 = Add to AC from memory

Question: What do you think the code is?



1:

- fetch instruction from **300**
- **1940** is loaded to **IR**
 - **1 = 001** = load AC from memory
 - **940** is where we're grabbing from
- increment **PC**
- load **0003** from **940** into **AC**

Interrupts

An interrupt is a signal for the cpu to pay attention to the device.

You can process user inputs very quickly so that it doesn't feel laggy.

Most IO devices are much slower than the cpu so it's less efficient for the cpu to just wait on the IO device so it's better for the IO device to just cut in whenever it needs to.

Most operating systems are interrupt driven, I/O devices will cut in whenever they need to so that the cpu doesn't have to pause and wait on them.

lect 1 end

There are different kinds of interrupts:

- hardware and software
 - whenever hardware or software wants to seek the attention of the processor
 - hardware sends *signal* through a peripheral device to interrupt the processor
 - software executes a specific instruction to interrupt the processor
- vectored and non-vectored
 - interrupt transfers control to a specific address which contains the interrupt vector
 - the interrupt vector has the address of the ISR
 - it's like a low level database that tells you what to do for a given interrupt
 - interrupt vector table contains addresses that inform the interrupt handler as to where to find the interrupt service routines (ISR)
 - vectored: manufacturer of the processor predefines this vector address
 - it's hardcoded/built-in
 - non-vectored: vector address is not predefined
 - address of the required ISR for the interrupt is provided by the interrupting device
- maskable and non-maskable
 - non-maskable: cannot be ignored
 - reserved for events such as
 - unrecoverable memory errors
 - hardware failure
 - system crash
 - force shutdown or cutting power

- non-maskable, hardware interrupt
- needs to have the option to shutdown gracefully and minimize damage to hardware
- overheating
 - heat sensors tell the processor to stop so that the heat can dissipate
- don't ignore
- attend to immediately
- maskable: can be ignored
 - turned off by CPU before the execution of critical instruction sequences that cannot be interrupted
 - non-essential for core system function
 - used by device controllers to request service
 - keyboard
 - printer
 - usually only ignored for a while
 - could potentially be ignored forever
 - almost never happens

These different types are can overlap.

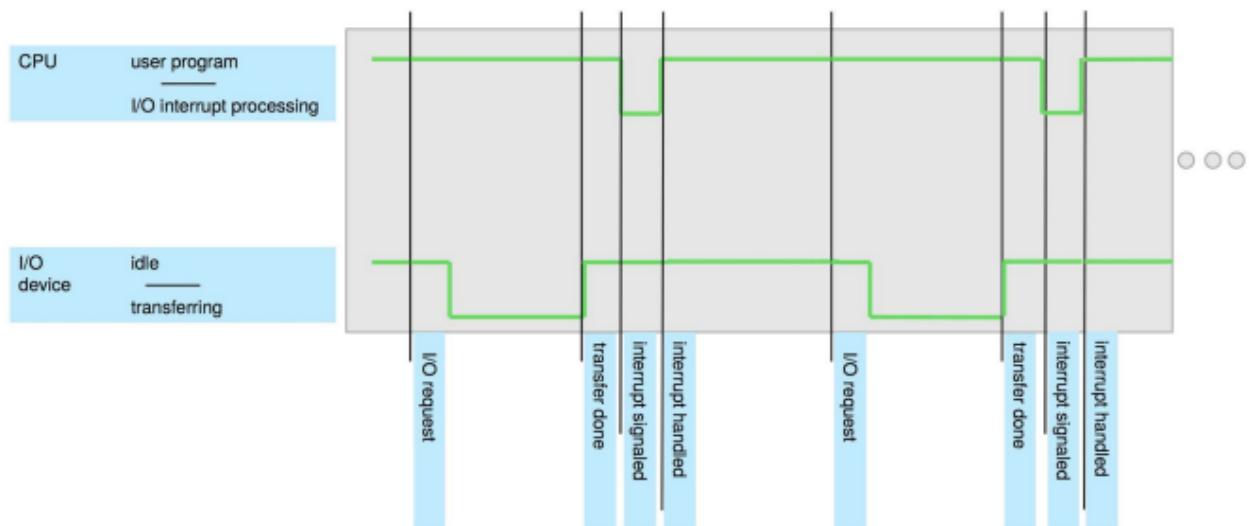
We might have a vectored, non-maskable, software interrupt or a non-vectored, maskable, hardware interrupt.

When the cpu is interrupted, it stops what it's doing and transfers execution to a fixed location.

Fixed location usually has a starting address where the service routine for the interrupt is located.

The interrupt service routine executes.

After the routine completes, the cpu goes back to what it was doing before.

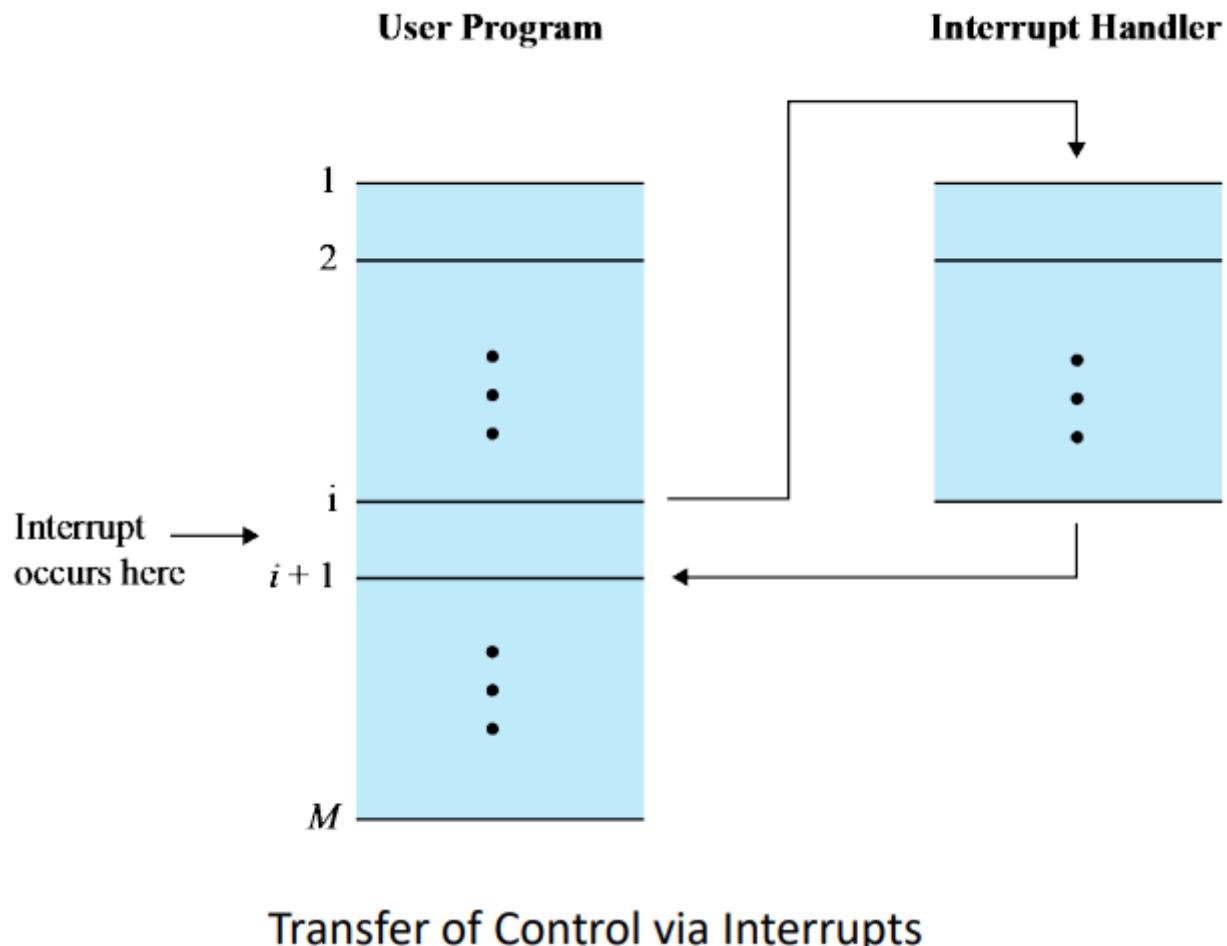


Interrupt timeline for a single program doing output.

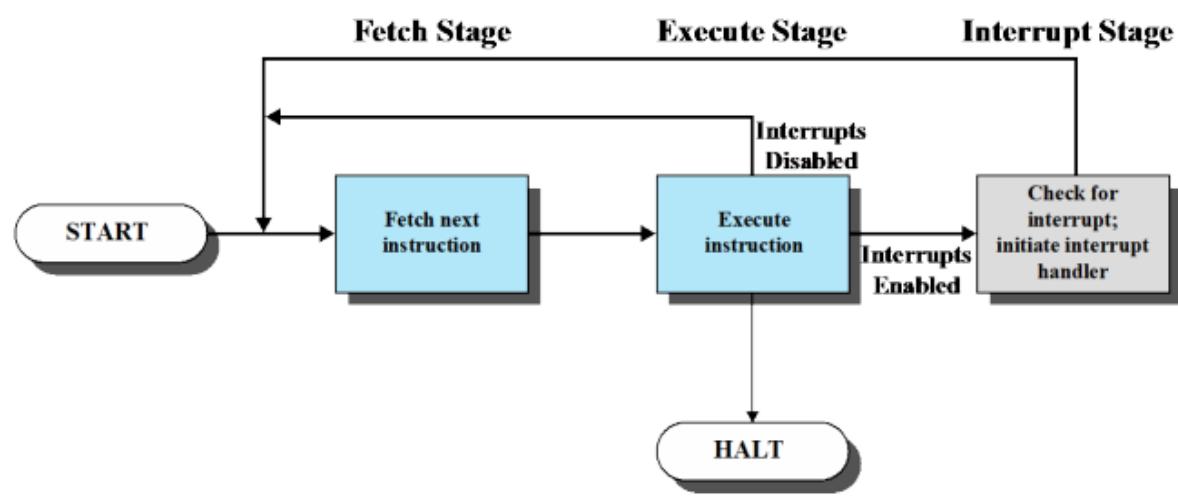
For the user program, an interrupt suspends the normal sequence of execution.

When the interrupt processing is completed, execution resumes.

In the above image we see the I/O device making the request, sending it to the processor, then the processor processing the interrupt shortly thereafter then signalling that it's done when it's done.



Instruction Cycle with Interrupts



Instruction Cycle with Interrupts

Interrupt stage is added to the instruction cycle.

At the execute stage, the processor checks to see if any interrupts have occurred.

If none, proceed to fetch stage as usual.

Otherwise, the processor executes an interrupt handler routine instead of going to the fetch stage continue the program execution. After that it continues.

An interrupt handler routine is generally part of the OS.

The routine:

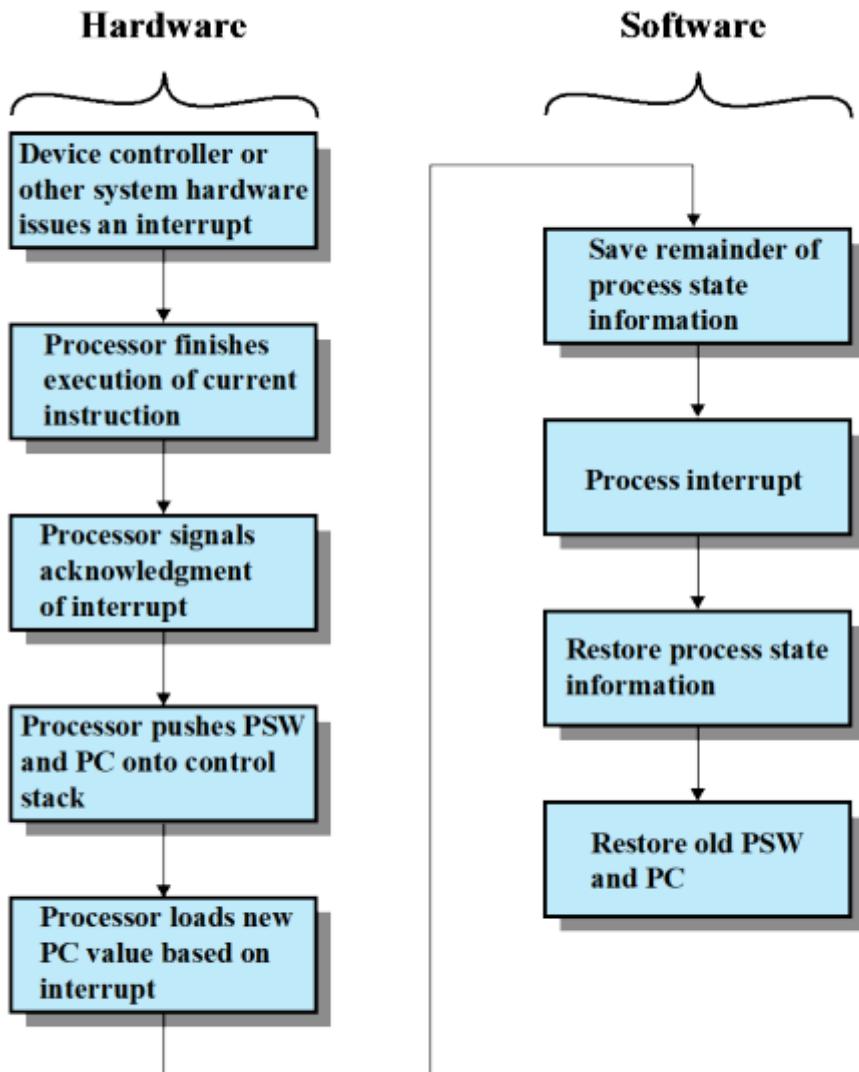
- determines the nature of the interrupt
- performs whatever actions are needed
 - e.g. find out what I/O module generated the interrupt and branch to a program that will write more data out to that I/O module
- allows the processor to continue executing the user program after it finishes

Program Status Word

An interrupt can trigger lots of events in the processor hardware and software.

Program Status Word (PSW) contains status information about the currently running process:

- memory usage
- codes
- status information
 - interrupt enable/disable bit
 - kernel/user-mode bit

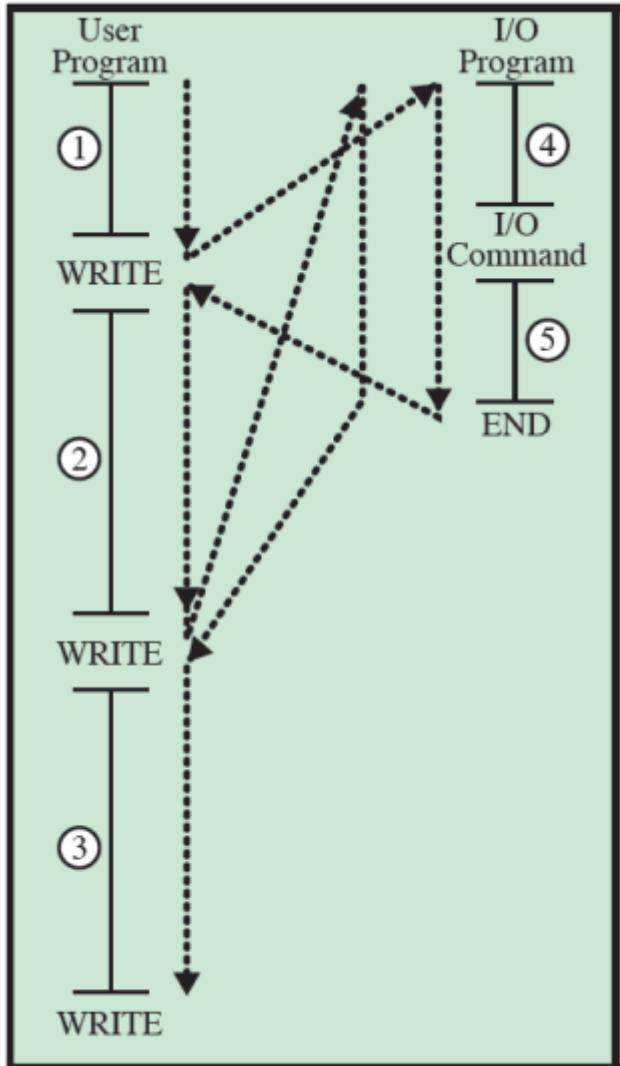


Simple Interrupt Processing – IO Device
PSW: Program Status Word.

We cover this in more detail later on.

Flow Control With and Without Interrupts

Let's take an example of flow control without interrupts.



(a) No interrupts

the black dotted arrow shows how the execution flows

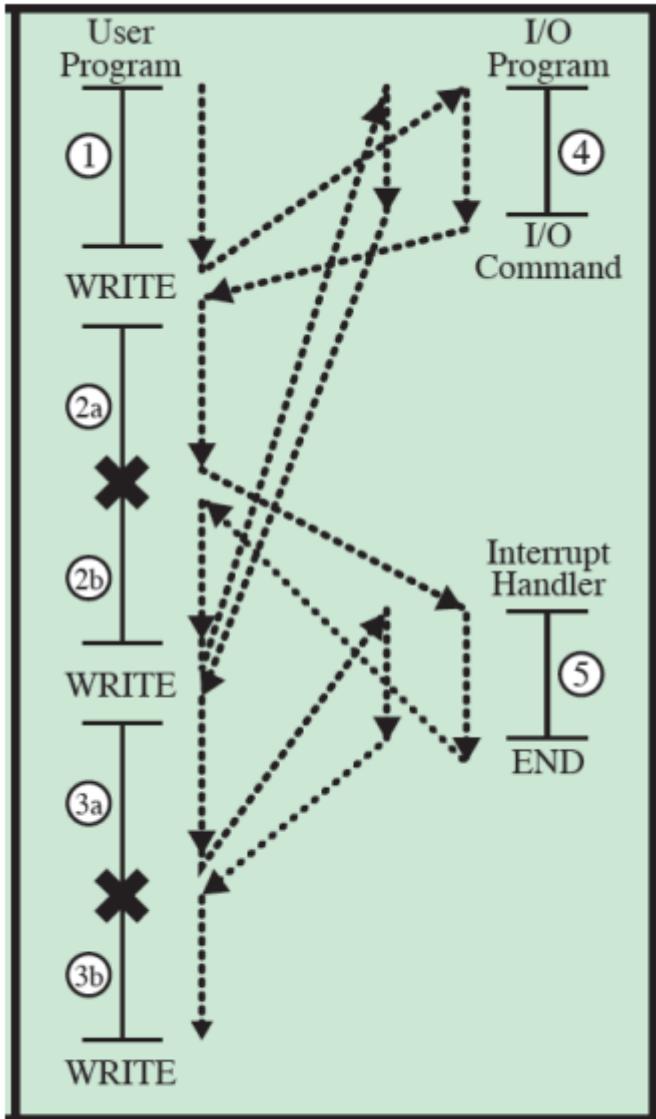
1, 2, 3 refer are segments of code that don't involve I/O

The write calls involve I/O and call an I/O routine/program pictured on the right:

- **4** prepares for the actual I/O operation
- then there's the I/O command
- **5** complete the I/O operation
- then it ends and we return

The write calls call to an I/O routine

Suppose now there are interrupts with short waits.

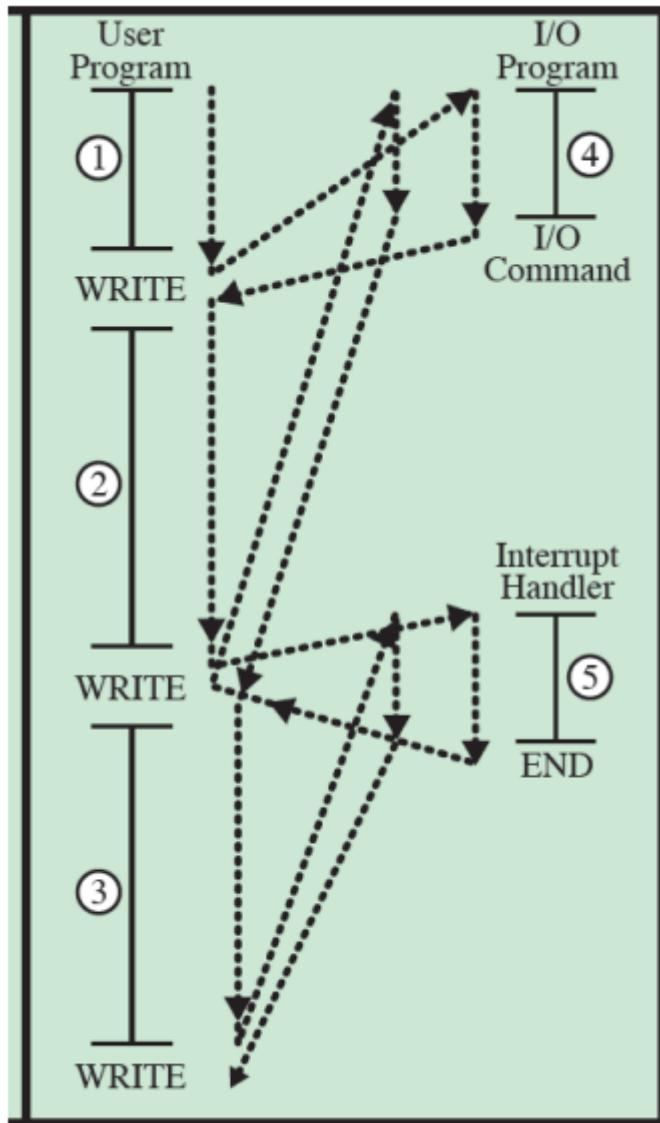


(b) Interrupts; short I/O wait

Now we see an interrupt in the segments 2 and 3 dividing them into 2a & 2b and 3a & 3b respectively.

At the dividing points we have the execution move to the interrupt handler before going back to the user program.

Suppose now there are still interrupts but now with long waits



(c) Interrupts; long I/O wait

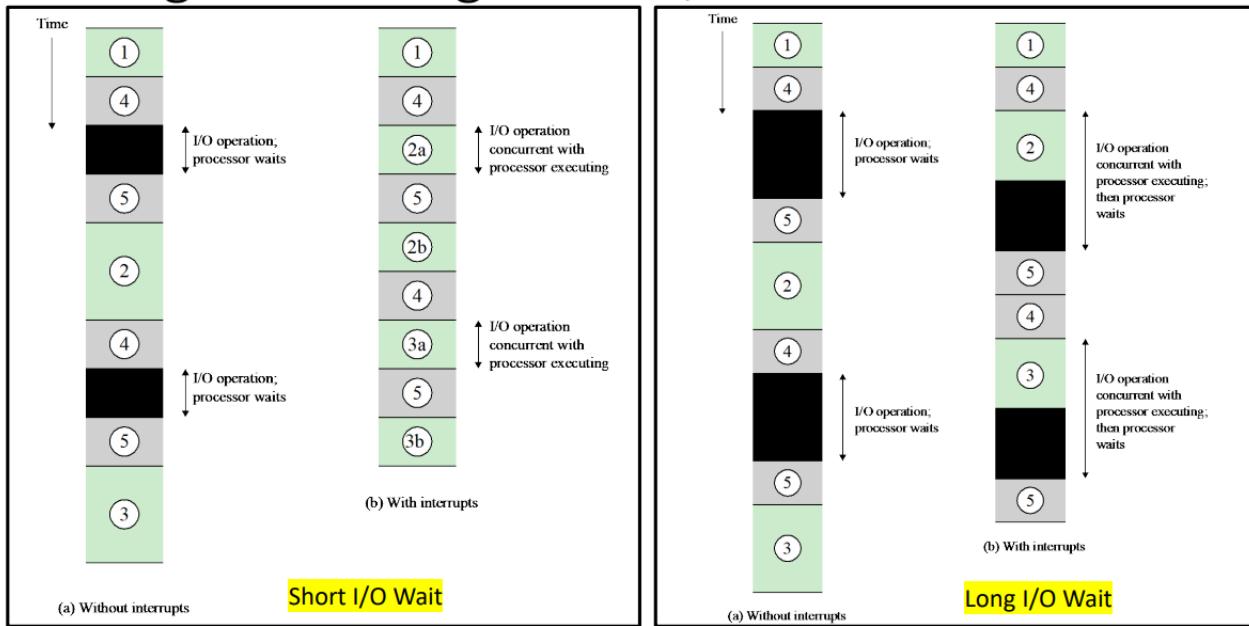
The interrupts no longer occur during the segments that don't involve calling to another routine.

Instead they happen when the user program calls to another routine, the write calls.

So above we went over different techniques for handling I/O operations

- (a) : programmed I/O
- (b) and (c) : interrupt-driven I/O

There is another third technique that we can go over later today called direct memory access (DMA)



Using interrupts it takes less time and is more efficient as the processor doesn't need to idle.

We have to be aware of the overhead as we get faster however.

A cpu core can execute 1 instruction at a time. Multi-core CPUs can execute multiple cores at a time. Multi-threading adds to this.

In order to utilize this we need to program in a way to take advantage of parallel execution.

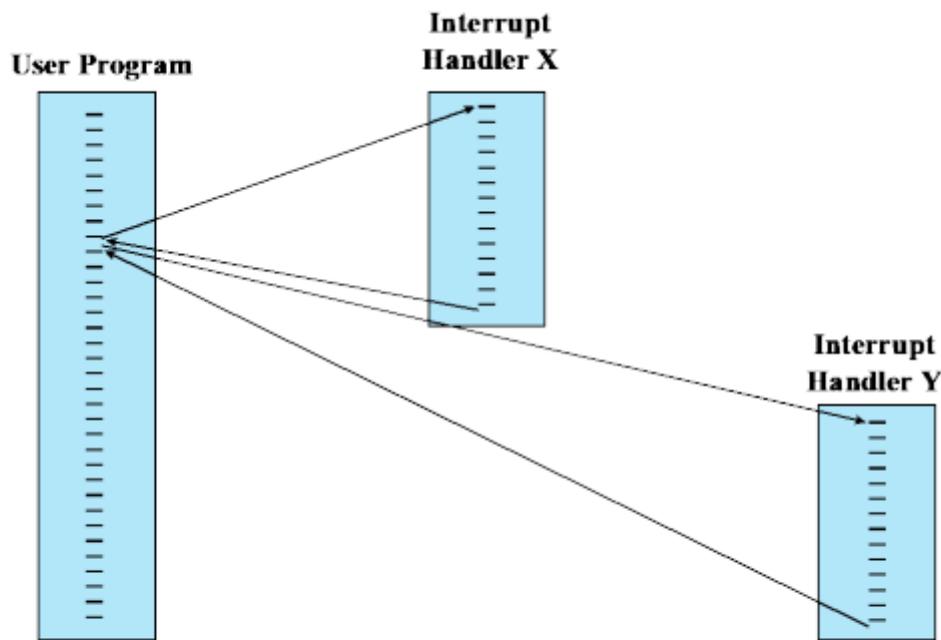
Multiple Interrupts

Suppose there is an interrupt during the processing of another interrupt, we can either:

- disable interrupts while an interrupt is being processed
 - hard to work with if there's un-maskable interrupts
- use a priority scheme

With disabled interrupts we might have an interrupt while we're handling another interrupt.

We finish handling the first interrupt before attending to the other interrupt.

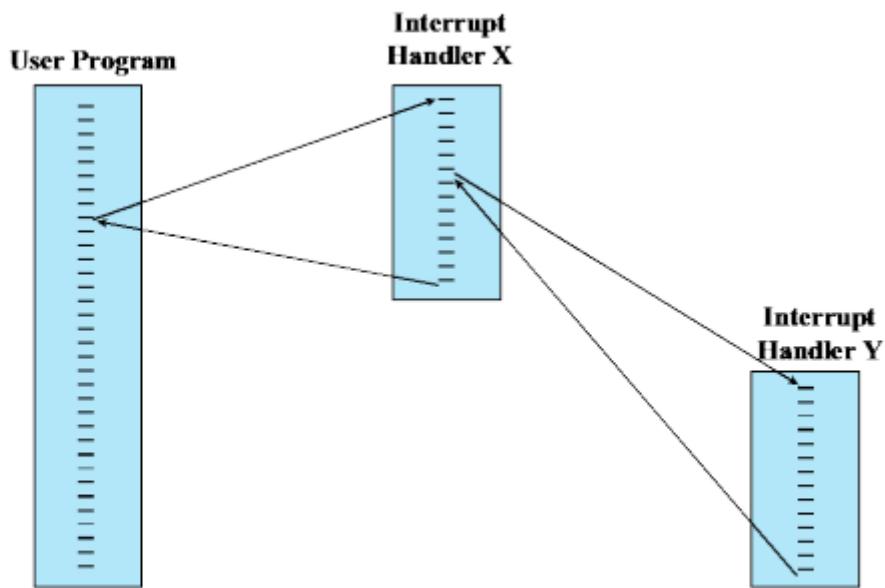


(a) Sequential interrupt processing

Transfer of Control with Multiple Interrupts

This does not take into account relative priority or time-critical needs.

With prioritized interrupts we can interrupt the interrupt handler in order to handle a more important interrupt.



(b) Nested interrupt processing

Figure 1.12 Transfer of Control with Multiple Interrupts

Interrupt Y has a higher priority than interrupt X which results in the execution we see above.

Example

3 I/O devices in order of least to greatest priority:

- printer - priority 2
- disk - priority 4
- comms - priority 5

timeline $t=x$:

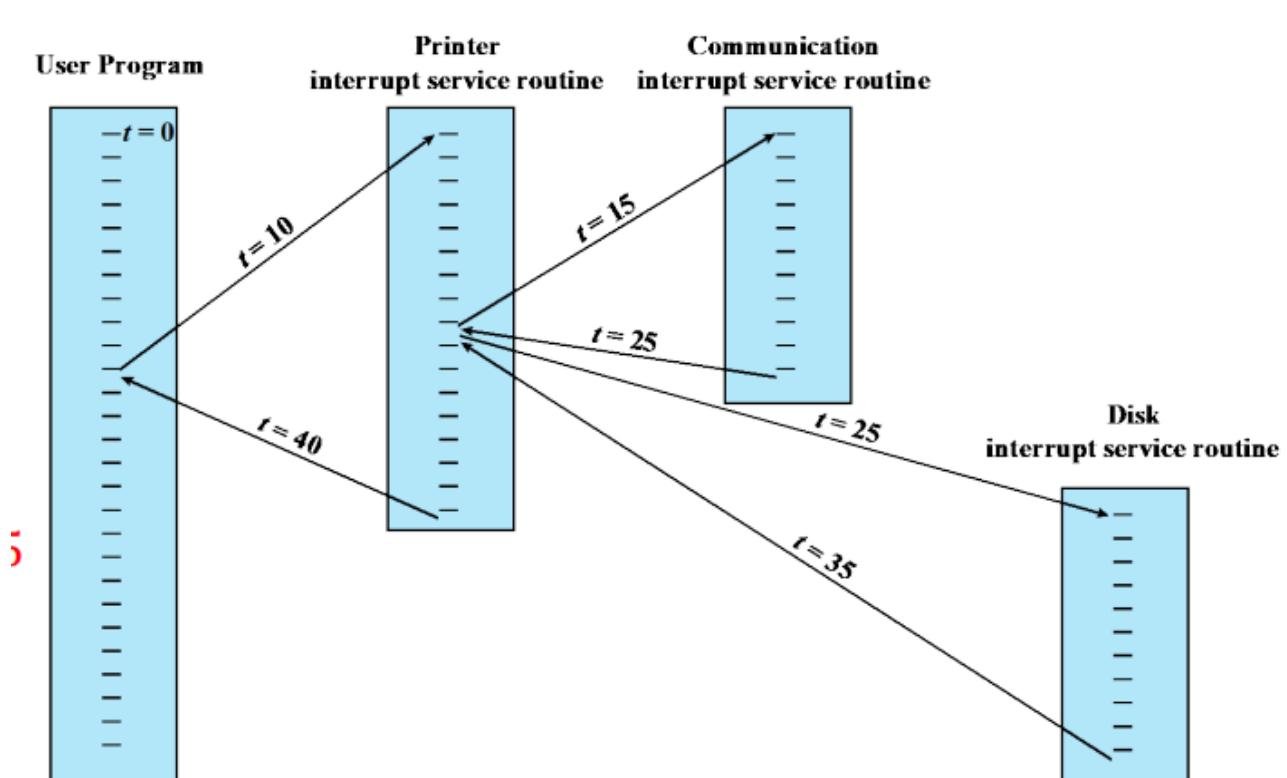
- 00: begin
- 10: printer interrupt
- 15: comms interrupt
- 20: disk interrupt

each interrupt lasts 10 time units

how will these interrupts be executed?

my answer:

- 10: begin printer interrupt
 - 15: begin comms interrupt
 - 25: end comms interrupt
 - 25: begin disk interrupt
 - 35: end disk interrupt
- 40: end printer interrupt



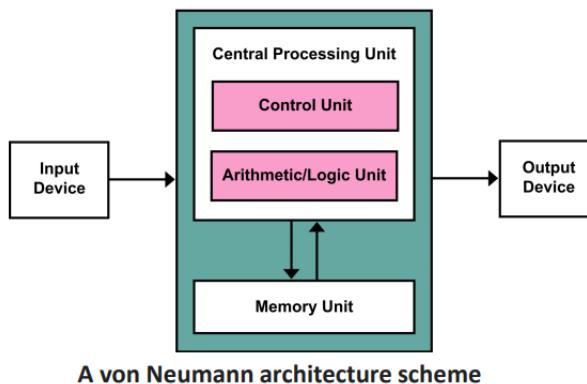
Example Time Sequence of Multiple Interrupts

Storage Structure

Von Neumann Architecture

Von Neumann architecture

- The structure of most computers, in which both process instructions and data are stored in the same main memory.



Source: https://en.wikipedia.org/wiki/Von_Neumann_architecture

	
von Neumann in the 1940s	
Member of the United States Atomic Energy Commission	
In office	
President	Dwight D. Eisenhower
Preceded by	Eugene M. Zuckert
Succeeded by	John S. Graham
Personal details	
Born	Neumann János Lajos December 28, 1903 Budapest, Kingdom of Hungary, Austria-Hungary
Died	February 8, 1957 (aged 53) Washington, D.C., U.S.
Resting place	Princeton Cemetery
Citizenship	Hungary United States

Main Memory

CPU can only load instructions from memory so programs must be loaded into memory in order to run.

Computers run most of their programs from rewritable memory, called main memory (aka random access memory(aka RAM))

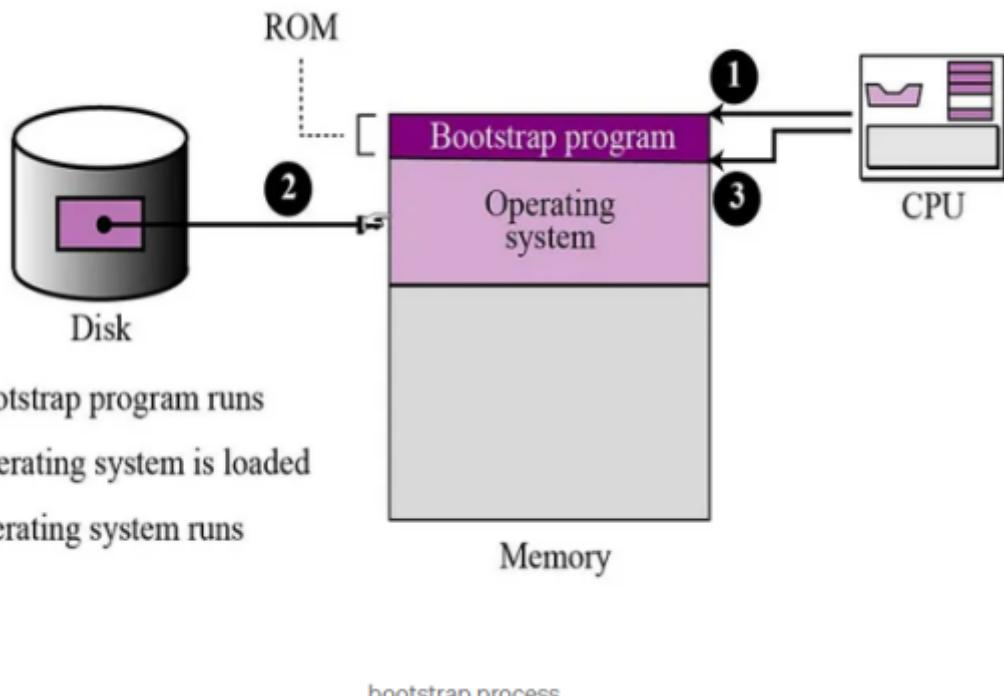
Main memory is volatile (loses everything when powered off) and is commonly implemented in a semiconductor technology called dynamic random-access memory (DRAM).

ROM, EPROM, and EEPROM

bootstrap program is loaded at power-up or reboot:

- the computer is checking itself
 - this is also known as post
- Initializes all aspects of the system
 - all your bios shit
- Loads operating system kernel and starts execution
- We can't store it to the RAM, instead it is typically stored in ROM or EPROM, and is generally known as firmware

- the rom, eprom, or eeprom are non-volatile so we put the bootstrap there



ROM

read only memory.

non-volatile

data is written there permanently during the manufacturing.

Used for storing firmware and software that does not need to be changed frequently

ex:

- BIOS
- embedded system firmware

EPROM

EPROM (Erasable Programmable Read-Only Memory): only reprogrammable after being erased by exposing it to ultraviolet (UV) light.

Commonly used in early embedded systems and microcontrollers.

EEPROM

EEPROM (Electrically Erasable Programmable Read-Only Memory): A type of memory that can be reprogrammed electrically, without the need for UV light.

Can be reprogrammed in place without the need for removal from the circuit.

Used in applications where data needs to be updated or modified periodically, such as in configuration settings.

Secondary Storage

we want programs and data to reside in main memory permanently but main memory is too small (to get a reasonably large size to store data permanently) and it's volatile.

We use non-volatile secondary storage as an extension of the main memory.

Hard-disk drives (HDDs) and nonvolatile memory (NVM) hold onto things before we load them into memory.

HDD

- metal or glass platters covered with magnetic recording material
- head that moves and disk that rotates to read off of the disk
- physical head moves

Solid-state disks (SSD)

- faster than HDD
- is not limited by the disk head movement

It's more accurate to think of it as storage in addition to being an extension of the main memory as opposed to just being an extension of the main memory.

Non-volatile memory (NVM)

- harder than hard disks
- becoming more popular

prof uses external ssd but doesn't know about nvme ssd and sata ssd??

Tertiary Storage

Slow and large.

Used for special purposes, typically as a backup of material stored on other devices.

CD-ROM or blu-ray

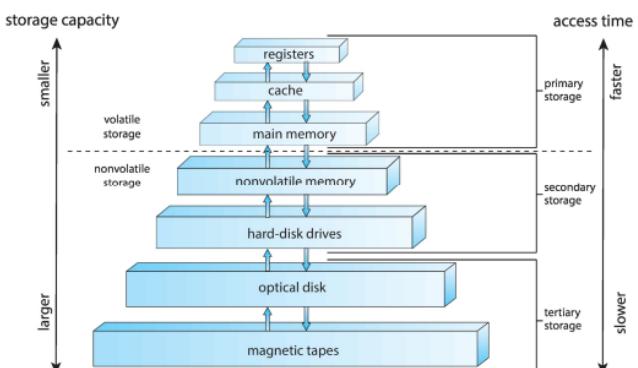
magnetic tapes.

We still use magnetic tapes even though prof doesn't think we do since last she used them was in '97. What a fucking world we live in. Still using magnetic tape.

Storage Hierarchy

■ Going down the hierarchy:

- a) Decreasing cost per bit (**cheaper**)
- b) Increasing capacity (**larger**)
- c) Increasing access time (**slower**)
- d) Decreasing frequency of access to the memory by the processor



Storage-device hierarchy.

going down the hierarchy we get cheaper, larger, and slower memory that is accessed less and less frequently by the processor.

Caching

Used at many levels in a computer from the most basest of hardware to the most frivolous of software.

If we want to use information from a slower piece of storage then we will copy that information to a faster piece of storage.

In the above instance we would describe the faster piece of memory as our cache.

We copy things from the secondary storage (the drives) to the main memory (ram) in order to have programs execute faster. The main memory is the cache.

When we want to grab information we check the cache to see if it's there

- if it is, we use it directly from there
- if not, we copy the data to the cache from the other storage and then use it there

Hit Ratio:

- fraction of all memory accesses that are found in the faster memory
- ex
 - processor has access to two levels of memory
 - lvl 1 - 1,000 bytes and access time of $T_1 = 0.1 \mu s$
 - lvl 2 - 100,000 bytes and access time of $T_2 = 1 \mu s$
 - Suppose 95% of the memory accesses are found in the cache
 - For $H = 0.95$, the average time to access a byte = $(0.95)(0.1 \mu s) + (0.05)(0.1 \mu s + 1 \mu s) = 0.095 + 0.055 = 0.15 \mu s$
 - We want to keep the hit ratio high and the access time(s) low

I/O Structure

A large portion of operating system code is dedicated to managing I/O:

- important to keeping system reliable
- the varying nature of devices makes this a complex task that requires a lot of code
 - cameras, mice, keyboards, printers, fax machines, scanners, and all manner of devices are all classified under I/O device but will all have different needs

When the processor encounters an instruction relating to I/O, it executes that instruction by issuing a command to the appropriate I/O module.

Possible techniques for I/O operations (2 of which covered previously [here](#)):

1. programmed I/O
2. interrupt-driven I/O
3. Direct Memory Access (DMA)

We previously mentioned that programmed I/O is very inefficient as it would require the processor to wait on the I/O devices as no device can keep up with the processor.

Interrupt-driven is fine for moving small amounts of data but can produce high overhead when used for bulk data movement. Bulk memory movement is pretty good for all of this. Remember that all operating systems nowadays are interrupt-driven.

But we run into the problem of I/O requests running into other I/O requests.

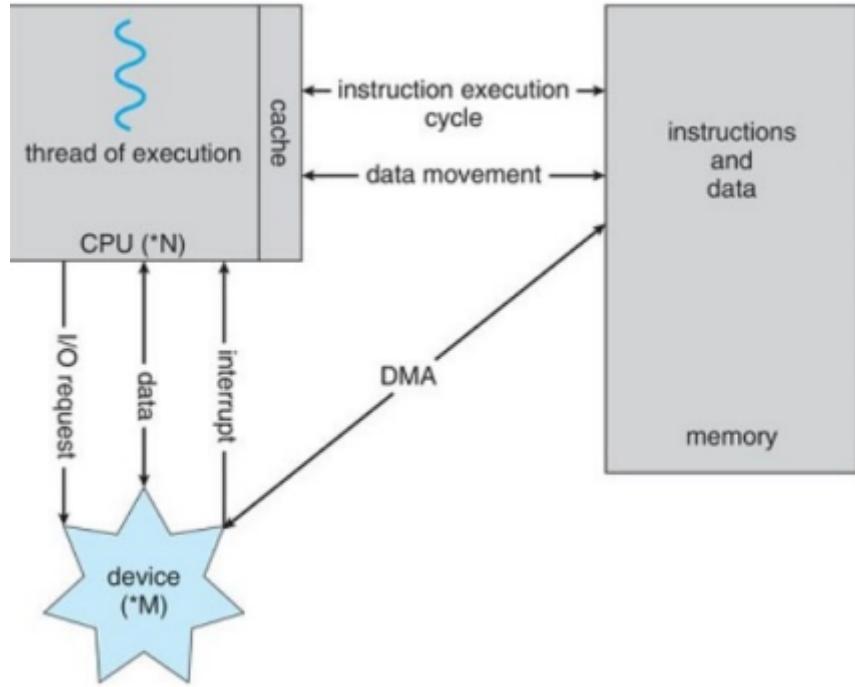
We need more efficiency which is where DMA comes in.

Direct Memory Access Structure

DMA is used for high speed I/O devices capable of transmitting info at close to memory speeds.

Blocks of data is transferred from the buffer storage to the main memory by the device controller w/o the CPU intervening.

Only one interrupt is generated per block, rather than one interrupt per byte as is the case in interrupt-driven I/O.



A von Neumann architecture:

DMA is performed by a separate module on the system bus or incorporated into an I/O module.

When the processor wishes to read or write data it issues a command to the DMA module containing:

- whether a read or write is requested
- address of the involved I/O device
- starting memory location for the read/write
- number of words to read/write

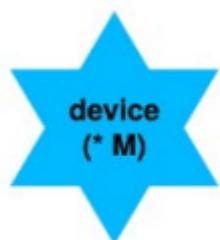
The processor returns to its own work while the I/O operation is continued on by the DMA module.

DMA transfers the block one word at a time, directly to/from memory w/o processor.

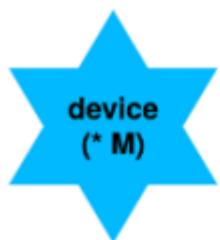
Once complete, DMA sends interrupt signal to the processor.

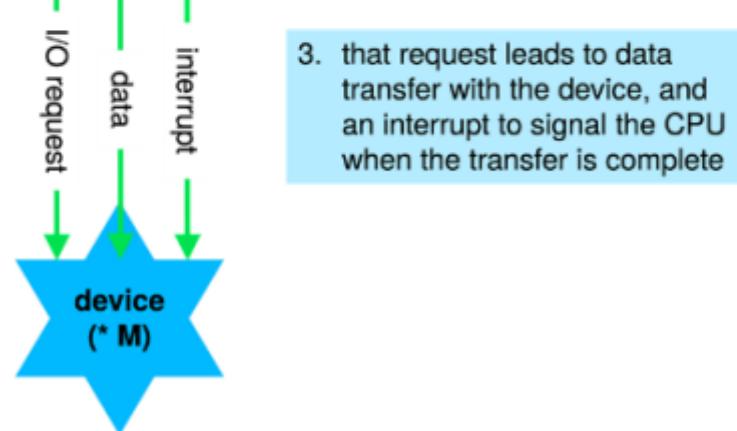
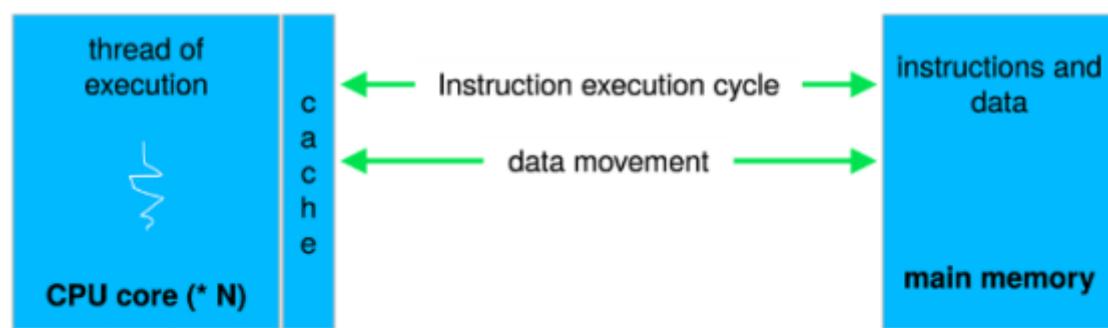
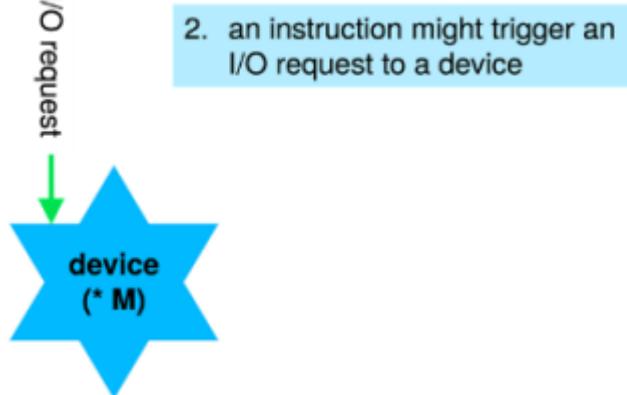
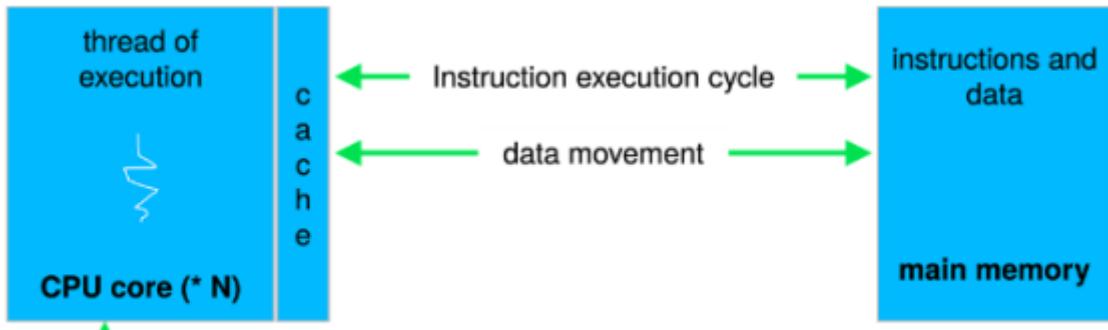
Processor is only involved at the beginning and end of the transfer making it far more efficient than interrupt driven or programmed I/O.

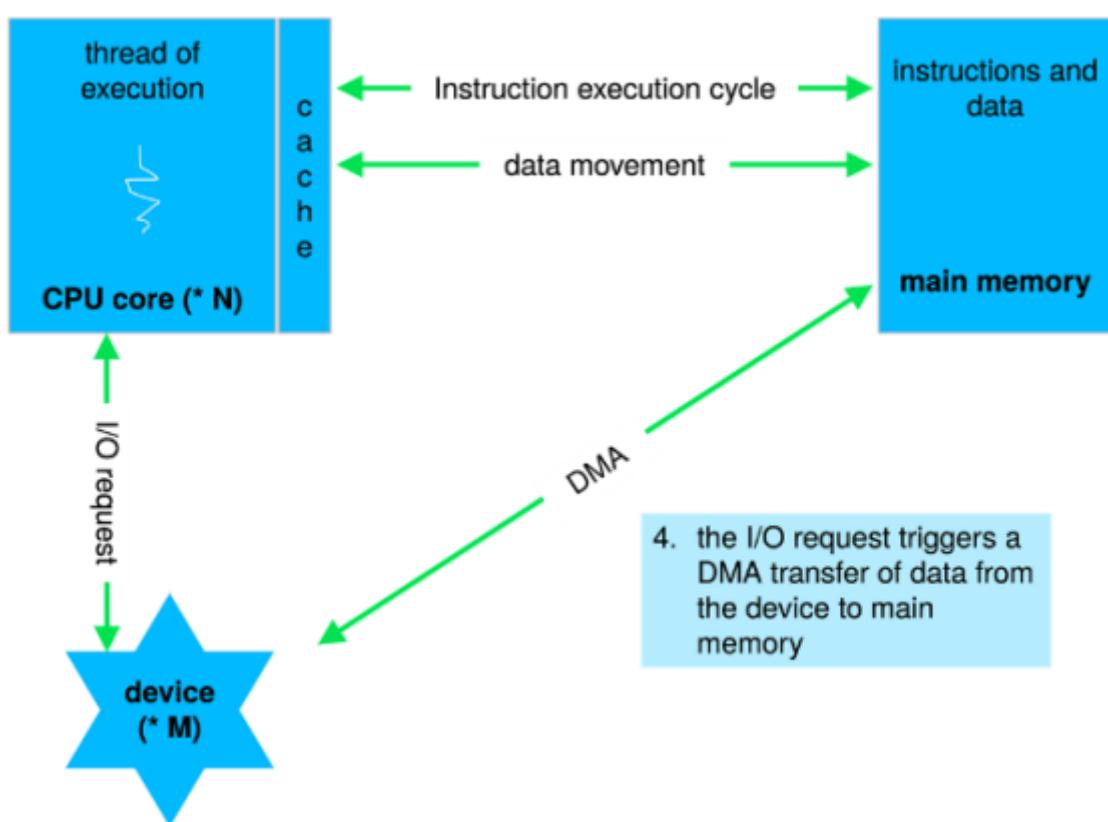
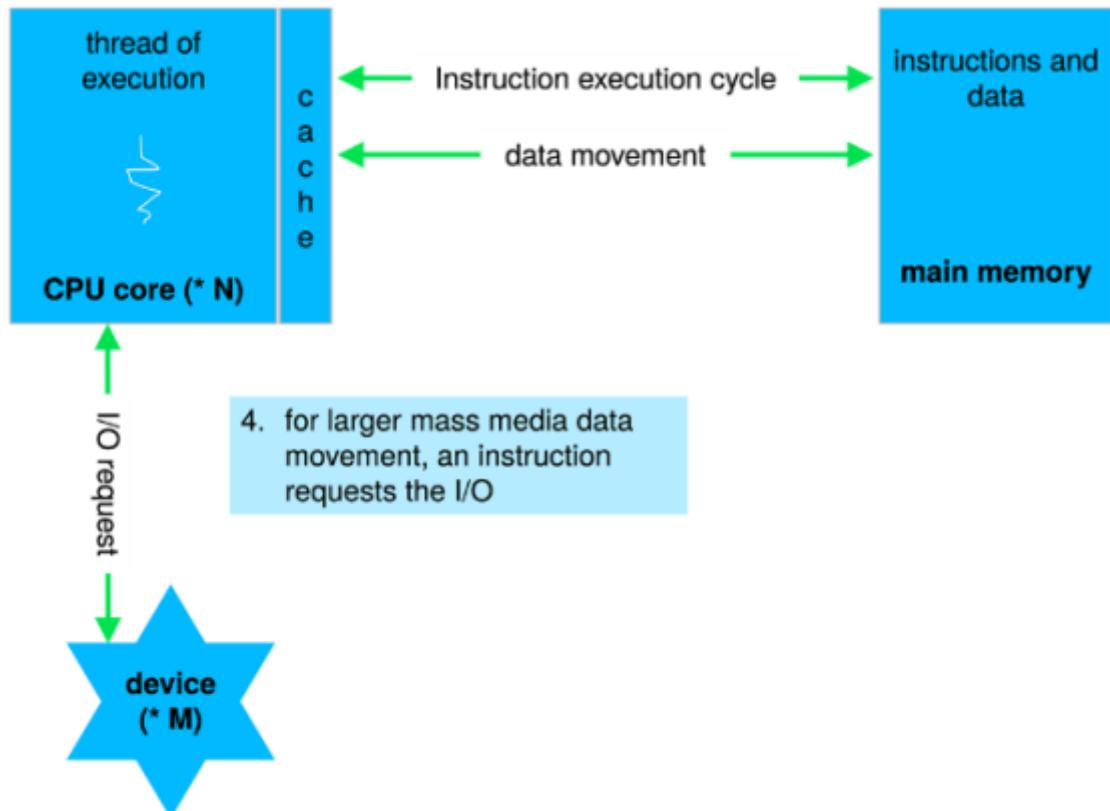
How a Modern Computer Works

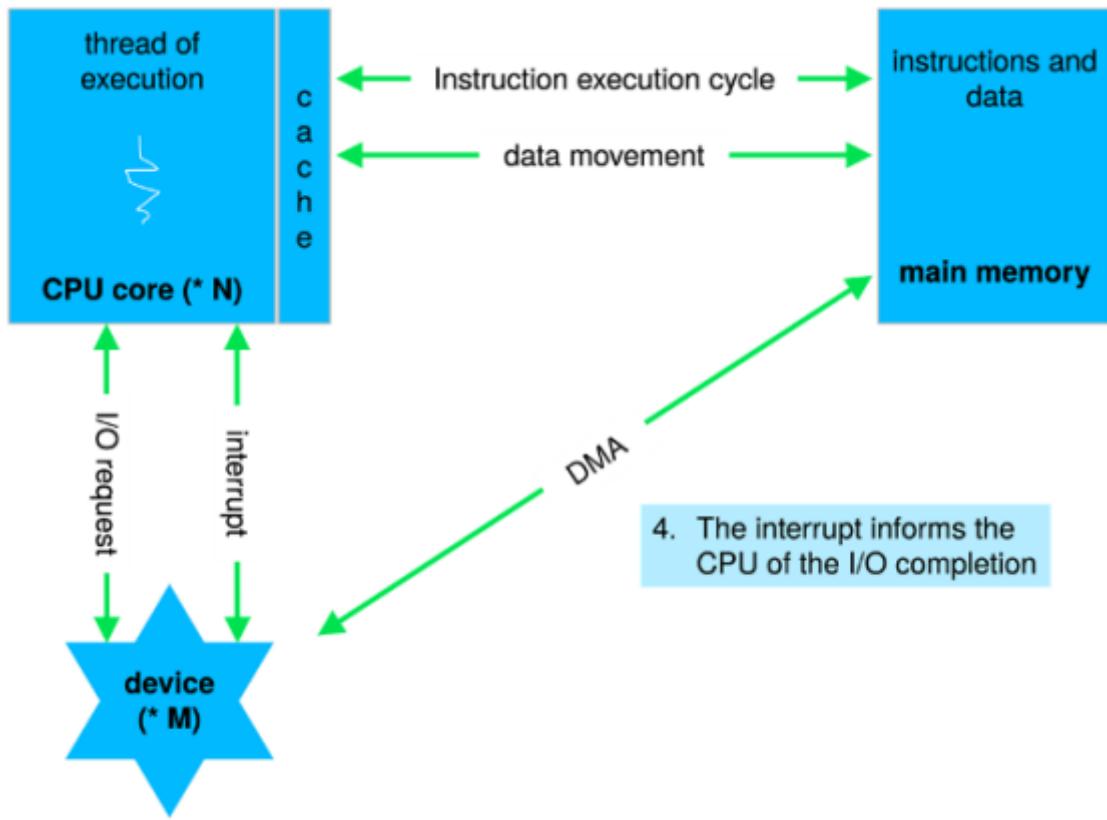


1. the CPU core instruction cycle involves reading instructions and reading and writing data to main memory









Operating System Structures

We look at generic operating system design in this course.

Different designers have different paradigms and ideas.

POSIX is used for linux and macOS

Android uses different things.

The fundamentals are all the same but we may have different names for things and different styles of schedules.

For specific issues we may look at examples of how different operating systems do things

Computer System Architecture

we can categorize a computer system according to the # of general-purpose processors used:

- in contrast to application specific processors
 - found less often or even not at all
- single-processor systems
- multiprocessor systems
- clustered systems

- different machines that work together

general purpose processor:

- type of microprocessor designed to handle a wide range of tasks and applications
- the typical processor

single processor system:

- old days
- most systems used to use only one cpu
- a core is a basic computation unit of the cpu
 - a unit that can execute instructions

prof's dad used to work at msx

that was her first computer

it was popular in the middle east and asia

there was a huge pyramid scheme surrounding msx

she remembered playing hangman on it with her brother

hunt and peck typing to code up hangman.

they had to reprogram hangman everytime they wanted to play since they didn't know how to save shit or I think they couldn't idk.

she saw the rise of the era of the 1.44 mb 3.5 inch floppies

if we had a lab this would've been so cool to listen to but rn idk

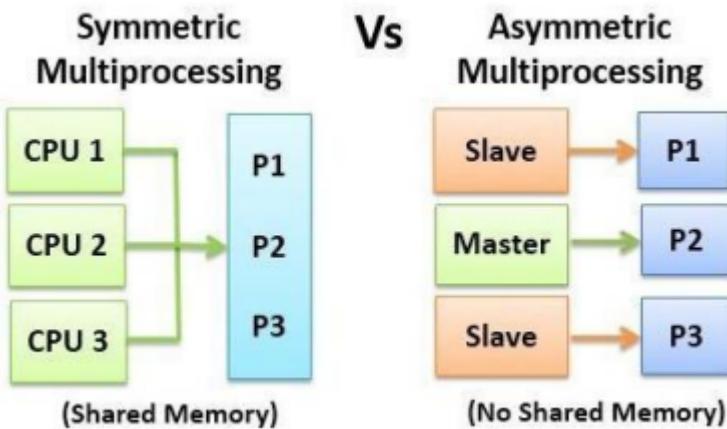
Most systems have special-purpose processors as well.

We might have graphics processors and all that

multiprocessor:

- these now dominate the landscape of computing
- aka parallel systems, tightly-coupled systems
- advantages:
 - increased throughput
 - increased reliability - fault tolerance
- 2 types of systems
 - symmetric multiprocessing
 - each processor performs all tasks
 - processors share memory

- asymmetric multiprocessing
 - each processor performs a specific assigned task



Symmetric:

- most common and what we're working with
- each processor has a CPU w/ their own registers
- they also have their own private/local cache
- All processors share physical memory over the system bus
- **N** processes can run if there are **N** CPUs

multicore systems can be more efficient than multiple chips with single cores because on-chip communication is faster than between-chip communications

each core effectively acts as a separate CPU in terms of execution (there are some caveats) and can more quickly communicate with one another.

One chip w/ multiple cores uses significantly less power than multiple single-core chips.

There's always a lot of overhead that comes in the form of inter-process and inter-program communication as well as resource allocation. These problems feed back into one another as they need to communicate about the resource allocation. It's a big fucking mess

Dual-Core design

each core has its own register set and its own local cache often referred to as L1 cache

The level 2 or L2 cache is local to the chip and is shared by the two processing cores.

multicore processor with **N** cores appears to the operating systems as **N** standard CPUs

Clustered Systems:

- 2 or more individual systems (known as nodes) joined together
- each node is typically a multicore system
- no shared memory
- loosely coupled systems
- usually there is shared storage via a Storage-Area Network (SAN)
 - usually for very large networks

- prof used to use a lot of SANs back in her old job
- communication at a hardware and software level
 - there is clustering software that shows the status of every node and all that
 - at the hardware level there is a sort of mini-operating system
 - maybe not too important but cool
- all cloud servers are clustered so that any server going down is invisible to the user
 - if you ping a server that goes down then you might not even notice as the ping just gets redirected
- all linked together to communicate via a Local-Area Network (LAN)
- provides high availability and fault tolerance service which survives failures and provides increased reliability which is crucial in many applications
 - financial applications

Asymmetric clustering:

- one machine in hot-standby mode while the other is running the applications
 - if server fails then standby becomes the active

Symmetric clustering:

- multiple nodes running the applications and monitoring each other
- more efficient
- all active at the same time

Some clusters are for high-performance computing (HPC)

- applications must be specially written for use parallelization

Prof told us about PS/2, hdmi, dvi, and other ports and all that

Going over the anatomy of the motherboard.

The old university monitors still use vga and all of that.

Prof tells us how to install cpu and ram.

Operating System Operations

Multiprogramming and Multitasking

Allows the operating system to run multiple programs

Multiprogramming (batch system) needed for efficiency:

- single user cannot keep CPU and I/O devices busy at all times

- multiprogramming organizes jobs (code and data) so CPU always has one to execute
- a subset of total jobs in a system is kept in memory
- one job selected and run via job scheduling
- when it has to wait (for I/O for ex), OS switches to another job

multitasking is a logical extension in which CPU executes multiple processes by switching among them very quickly.

Interactive computing.

multitasking

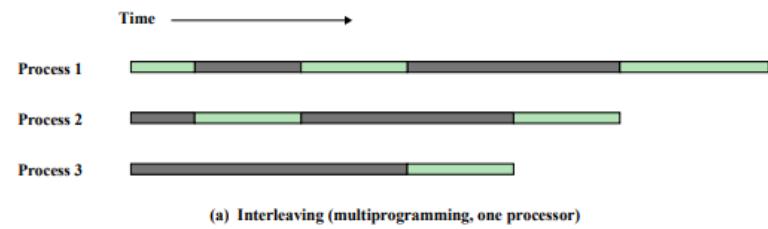
- each user has at least one program executing in memory, a process
- if several jobs are ready to run at the same time then we schedule jobs for the CPU
- if processes don't fit in memory then we swap them in and out of memory in order to run

multiprogramming only can only execute one program at a time meanwhile all other processes are waiting for the processor.

With multiprocessing, more than one process can be running simultaneously, each on a different processor.

Multiprogramming vs. Multiprocessing

- With multiprogramming, only one process can execute at a time; meanwhile all other processes are waiting for the processor.



- With multiprocessing, more than one process can be running simultaneously, each on a different processor.



Both of these let us run programs “at the same time”.

Dual Mode

Dual-mode operation:

- OS able to protect itself and other system components
- 2 modes

- user mode
 - kernel mode
- mode bit is provided by hardware
 - distinguish between system running user or kernel code
 - some instructions are only executable in kernel mode
 - privileged instructions
 - system call changes mode to kernel mode and return from call resets it to user mode

CPUs increasingly support multi-mode operations

- Virtual Machine Manager (VMM) mode for guest VMs

I'm thinking most of the operating system shit lives in kernel mode.

Transition from User to Kernel Mode

Operating system must always be in control of the CPU.

The user program can't get stuck in an infinite loop or fail to call system services.

We use a timer to stop that from happening. When the timer runs out it interrupts the computer. The time length can either be fixed or variable. No matter what we are able to regain control or terminate the program that goes over time.

There are different timers.

For example on linux systems we can observe different types of timers based on how the kernel is configured, what kind of machine it's on, and what the architecture of the components are like. We can configure a certain number of interrupts per second for example and it's a HZ value.

Crucial takeaways:

- a timer can be used to prevent a user program from never returning control to the operating system
- kernel mode has many different names
 - supervisor mode
 - system mode
 - privileged mode

Resource Management

Remember that the operating system is a resource manager.

File space, i/o, processes, memory, cache, etc. all managed by the OS

Process Management

Process management:

- a process is a program in execution
 - program is passive entity
 - process is active entity
 - unit of work within the system
 - once we run a program we have a runtime environment
 - allocated area in the memory where we will load in the code and text and stack and other values/variables that we need for the process
- process needs resources to accomplish its task
 - CPU cycles
 - memory space
 - I/O
 - files
 - initialization data
- upon termination, reusable resources are reclaimed

Threads and Processes:

- threads can be single-threaded or multithreaded
- single-threaded
 - one program counter specifying the location of the next instruction to execute
 - process executes instructions sequentially
- multi-threaded
 - one program counter per thread
 - requires a multicore processor, each core works on a different thread
 - most processes aren't multithreaded

Issues to Consider with Process Management

- synchronization
 - coordinate processes that need to share data
- mutual exclusion
 - how to make processes or users share a resource at the same time
- deadlocks
 - >=2 processes can't continue b/c they are waiting for the others to do something
- livelock
 - >=2 processes continuously change their states in response to changes in other other process(es) without doing anything useful
 - "you go first" "no you go first" and then we do nothing forever
- starvation
 - process is overlooked indefinitely by the scheduler
 - it can proceed but is never chosen

We're going to be looking at the pseudocode and underlying theory for programs that do this stuff correctly.

After that we can look at some real, runnable code to see what's really going on.

There's a lot of groundwork that needs to be done until we can even understand what the demos we're going to get are doing when we run them.

Activities connected to process management:

- creating/deleting user and system processes
- suspending/resuming processes
- providing mechanisms for
 - process
 - synchronization
 - communication
 - handling deadlocks

Memory Management

All or part of the instructions and data needed for a program in the memory in order to execute the process.

Memory management optimizes cpu utilization and response to users.

What's in the memory and when.

Memory Management Activities:

- keep track of what memory is being used by who
- decide which processes (and their parts) and data to move in/out of memory

File-System Management

most visible component

uniform, logical view of information storage

Abstracts from the physical properties of its storage devices (stores 1s and 0s and engravings and shit) to define a logical storage unit - file (`.c` , `.java` , `.py` , `.txt` , `.exe` files)

Each storage medium is controlled by a device with various processes

- access speed
- capacity
- data-transfer rate
- access method
 - sequential or random

Activities:

- organize files into dirs
- manage access perms
- os activities
 - creating/deleting files/dirs
 - primitives to manip files and dirs
 - mapping files onto secondary storage
 - backup files onto stable(non-volatile) storage media

Mass-Storage Management

storing things like code and data when not in memory

Long term storage.

Very important to manage properly

Speed of computer operation hinges on dis subsystem and its algos. How do we read and write in an efficient way

SSDs have semiconductor tech that's really good but still very expensive so we still have to take other forms of storage into account.

Some storage doesn't even have to be fast, just very long lived, large, and cheap.

Optical drives and magnetic tapes fall under the above category but they still need to be managed by OS or applications.

Q: When I move a file to another computer, how does the other computer know what kind of file it is?

A: Every operating system has their own system to manage files. But there are many universally used ways of interpreting the files. The physical files themselves are the same so we just have to worry about the interpretation. Some files are platform specific (`.msi` = microsoft installer, ` `.deb` = debian package file) and they can't be interpreted properly by machines with the wrong operating system

OS Activities:

- (un)mounting
 - (dis)connect to storage medium
- free-space management
 - if you want to write, where do you go to write?
 - do we even have enough space to write?
- storage allocation
- disk scheduling
- partitioning
- protection

Caching Management

Recall that cache is a faster piece of storage that we place frequently accessed data that we would ordinarily store in slower storage.

Caches have limited size so we need to manage what we store in them and when.

The size and replacement policy can greatly impact performance.

Characteristics of Various Types of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid-state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	25,000-50,000	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5,000-10,000	1,000-5,000	500	20-150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Movement between levels of storage hierarchy can be explicit or implicit

CPUs have cache that the OS doesn't really mess with but there's other forms of cache that the OS is concerned with

Access time:

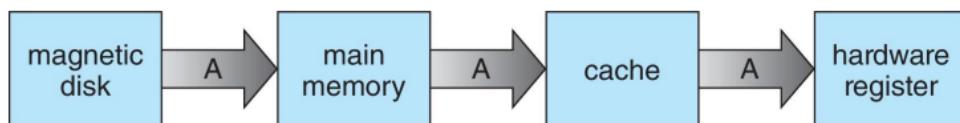
- time needed to arrive at the byte we want to read/write
- takes time to do the calculation and find the exact address
- the time the cache tech needs to arrive at the right location

Bandwidth:

- speed of the read/write
- a unit of data per second, typically megabytes

Migration of data “A” from Disk to Register

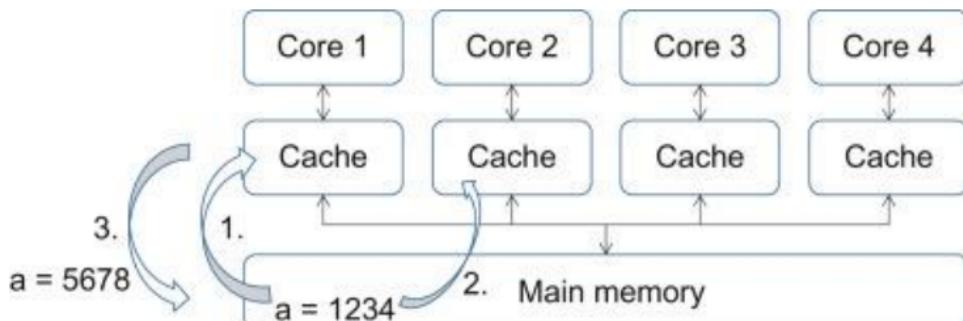
- Multitasking environments must be careful to use the most recent value, no matter where it is stored in the storage hierarchy
- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache



Cache Coherence

Cache Coherence

- **Cache coherency** is a situation where multiple processor cores share the same memory hierarchy, but have their own L1 data and instruction caches.
- Incorrect execution could occur if two or more copies of a given cache block exist, in two processors' caches, and one of these blocks is modified.



There might be multiple processors with their own caches but there could be some incorrect execution that results in 2 or more copies of a cache block to exist in 2 different processors' caches

This is a problem when we try to write back that cache.

I/O Subsystem Management

OS is meant to hide particularities of hardware devices from the user

responsible for:

- memory management of I/O
 - buffering - temp storage of data while being transferred
 - caching - storing parts of data in faster storage for performance
 - spooling - overlapping of output of one job with input of other jobs
- general device-driver interface
- drivers for specific hardware devices

Everything not under processor or memory management can be thought of as being I/O in general but in this course we think of things more granularly

Q: Two important design issues for cache memory are

A: Size and replacement policy

Q: What is the unit of work in a system?

A: Process

Operating System Services

OSes provide an env to execute programs and services for users and other programs

2 types of OS Services with different goals:

- helpful to user
- efficient operation of the system itself

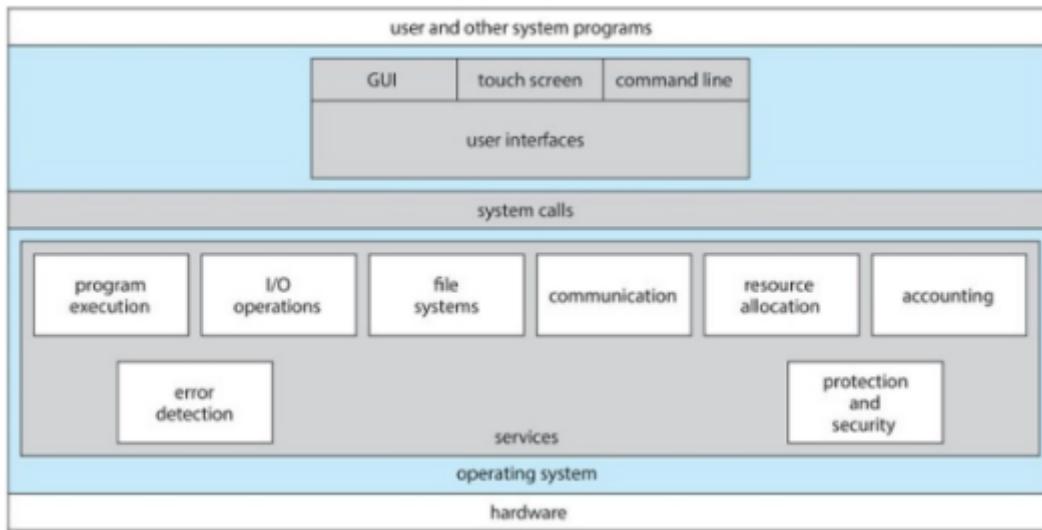
Services for the system and services for the user

For the User

The user does not have to be aware of how all of this works

User Interface:

- almost all operating systems have a User Interface (UI)
- types:
 - Command-Line (CLI)
 - Graphical User Interface (GUI)
 - touch-screen



A view of operating system services

Program Execution:

- load program into memory
- run program
- end execution
 - normally or abnormally and give error

I/O ops

File System Manip:

- read/write files
- create/delete dirs
- search for files/dirs
- list file/dir info
- manage perms

Communications:

- exchange info within or without the computer using a network
- shared memory or message passing via packets moved by the OS

Error Detection:

- constantly aware of all errors and what to do
- errors in CPU, memory, I/O devices, and programs
- give debugging facilities

Q: What service is it when we force quit a program?

A: Things can go wrong on the software level, hardware level, or any number of other levels (within those levels). The best thing to do is dive into the debugging hole. If you're at a low level then you can see each instruction execute one at a time. EAX register, RAX register available to you. If you're at a higher level then you'll

likely just get an error code to work it out.

It depends on what is happening.

For the System

Resource Allocation:

- same as before
- make sure no one can access what they shouldn't be able to

Logging:

- keep track of what's being done and by who with what resources

Protection and Security:

- control information that is stored and used
- concurrent processes should not interfere with each other
- protection ensures that all access to system resources are controlled
- security
 - outside users require authentication
 - passwords
 - biometrics
 - physical key - yubikey
 - 2FA
 - defend against external I/O devices from invalid access attempts
- once users are allowed in the OS looks on the access control list (ACL) to see what the user is allowed to do

System call interface is the boundary between user programs and operating system services.

Yes

Is the GUI the most common UI?

It depends. CLI might be more common depending on what you're doing.

There's probably more CLIs kicking around in total but people mostly use GUIs

Touch screen is used on mobile a lot

yes

User Operating System Interfaces

CLI or command interpreter allows direct command entry

- implemented by kernel or some system programs
- sometimes multiple different kinds of interpreters implemented
 - known as **shells**
 - shell is loaded into run time env and we execute from there
- sometimes commands built-in, sometimes just names of programs
 - in the latter case then we don't have to modify the shell itself to add new features
 - the shell is just a powerful program that can run other programs directly and access the resources
 - we have the latter case a lot in our linux systems, even if we think things are built in they're really just programs in the **/bin/** directory
- shell is merely the interpreter that you put in commands, the OS looks at what to do with that command, and runs the program
- access to the shell is access to the whole computer

prof loves security, very passionate

GUI:

- user-friendly desktop metaphor interface
 - usually mouse, keyboard, and monitor
 - icons represent files, programs, actions, etc
 - invented at Xerox PARC
- Big developments in GUI designs come from open-source projects
 - KDE - K Desktop Environment
 - GNOME desktop by GNU project
 - both run on linux and unix systems
 - source code readily available for reading, modification, and even contribution under specific license terms

Touchscreen:

- impossible/undesired mouse
- gesture based
- virtual keyboard
- uses a lot of voice commands

Remember that there are many different kinds of shells to choose from

System call

Provides a programming interface to the services provided by the OS

System calls are typically written in a high-level language (C or C++).

Mostly accessed by programs via a high-level Application Programming Interface (API) rather than direct system call use

API specifies a set of functions that are available to an application programmer

params passed to functions to return expected values

3 most common:

- Win32 API for Windows
- POSIX API
 - UNIX, Linux, Mac OS X
- Java API for Java Virtual Machine

We might use system calls for:

- copying contents of one file to another file

Types of System Calls:

- Process control
- File management
- Device Management
- Information Maintenance
- Communications
- Protection

System calls are run in kernel mode

Process Description and Control

What is a Process

Computers used to only execute 1 program at a time. This program had complete control of the system and access to all its resources.

Now they can have multiple programs loaded into memory and executed at the same time.

With that evolution we needed more control and compartmentalization of running programs.

Thus came the birth of the process.

A process is:

- a program in execution
- the unit of work in a modern computing system
- an entity that can be assigned to a processor and executed on a processor

- a unit of activity characterized by
 - the execution of a sequence of instructions
 - having a current state
 - having an associated set of system resources

2 essential components of processes:

- program code
 - text section
- set of data associated with that code
 - data section

We may be more precise later

A process can be uniquely characterized by a number of elements:

- identifier
 - unique identifier associated w/ process to distinguish from all other processes
 - pid - process ID unique to a process
 - no 2 processes have the same pid
- state
 - if a process is currently executing it is in the **running** state
- priority
 - priority relative to other processes
 - what processes can it interrupt and what processes can interrupt it
- program counter
 - the address of the next instruction in the program to be executed
- memory pointers
 - includes pointers to the program code and data associated with the process plus any memory blocks shared with other processes
- context data
 - data that are present in the registers during execution
- i/o status information
 - outstanding i/o requests
 - assigned devices
 - list of files used in the process
- accounting information
 - may include
 - amount of processor time
 - clock time used
 - time limits
 - helps us log activities
 - find out who/what accessed things, when and why

This is all very general, every OS has their own implementations but most every OS is going to have these things just under different names.

All of the above is stored in the process control block (PCB):

- created and managed by the operating system
- makes it possible to interrupt a running process and later resume execution as if the interruption had not occurred (context-switching)
 - context-switching is what this course is all about
 - trying to keep the processor busy and feed it new processes in order to increase the performance and allat
 - the speed of the processor is the speed limit of the computer, no one is faster than the processor. less of a road speed limit and more of a physical speed limit, like the speed of light

Each process has only one program counter and thread for now. We talk about multithreading and parallelism later on.

Process States

Process Trace

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A

(b) Trace of Process B

(c) Trace of Process C

5000 = Starting address of program of Process A

8000 = Starting address of program of Process B

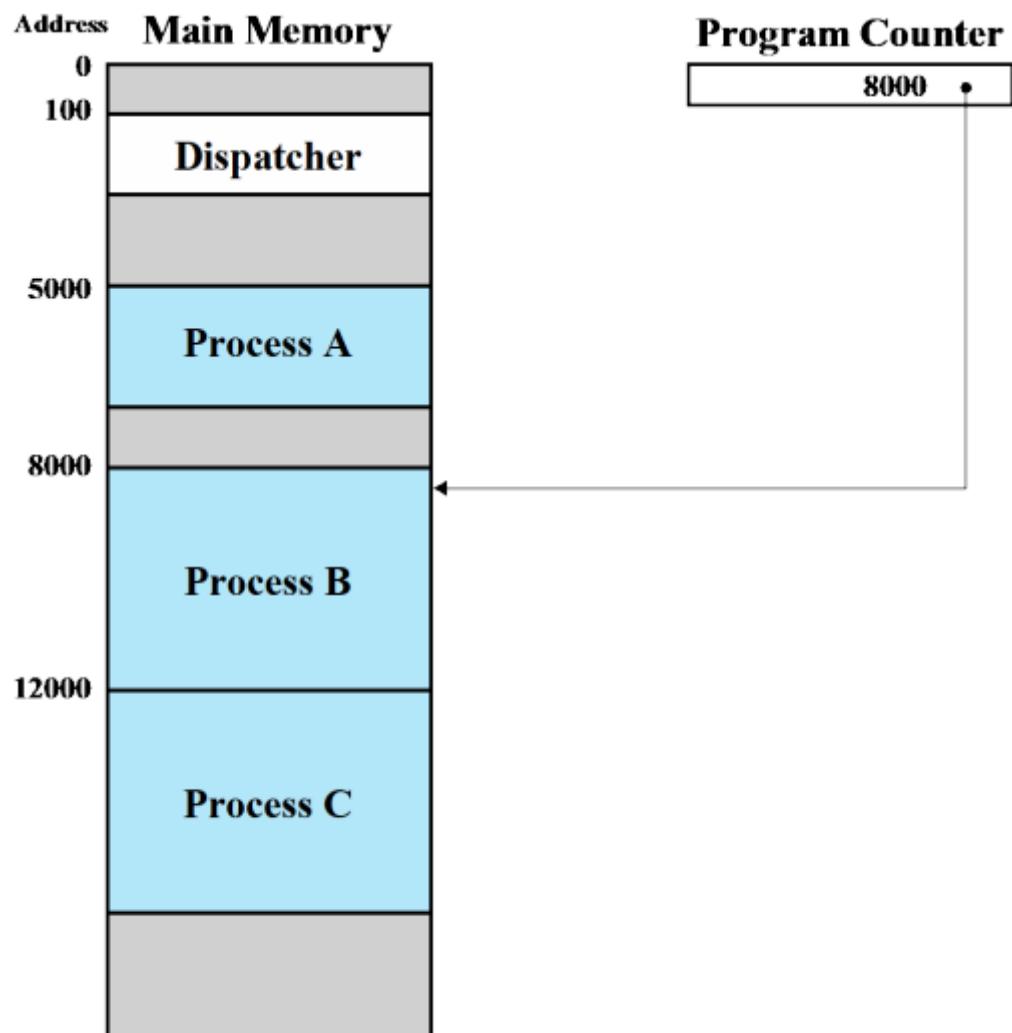
12000 = Starting address of program of Process C

Traces of Processes

Trace:

- Characterize the behavior of an individual process by listing the sequence of instructions that execute for that process
- *set of instructions for a process*
- show how the traces of the various processes are interleaved

Example



Snapshot of Execution at Instruction Cycle - Example

We have a small dispatcher process that switches the process from one process to another.

Process Trace Execution

- Assume the OS only allows **six** instruction cycles per process before timing out.
- 100 = Starting address of dispatcher program.
- Shaded areas indicate execution of dispatcher process.
- How do we design the OS to perform this extremely simple interleaving of processes execution by the processor!

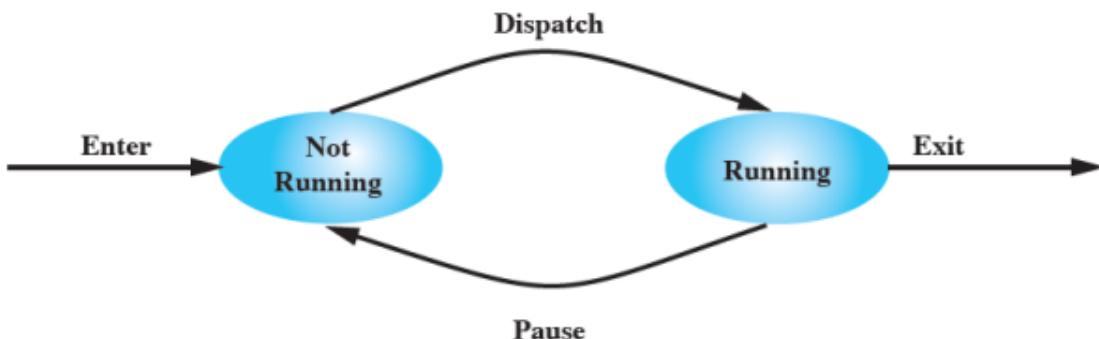
1	5000	27	12004
2	5001	28	12005
3	5002	----- Timeout	
4	5003	29	100
5	5004	30	101
6	5005	31	102
7	100	32	103
8	101	33	104
9	102	34	105
10	103	35	5006
11	104	36	5007
12	105	37	5008
13	8000	38	5009
14	8001	39	5010
15	8002	40	5011
16	8003	----- Timeout	
17	100	41	100
18	101	42	101
19	102	43	102
20	103	44	103
21	104	45	104
22	105	46	105
23	12000	47	12006
24	12001	48	12007
25	12002	49	12008
26	12003	50	12009
		51	12010
		52	12011
		----- Timeout	

Combined Traces of Processes

There's a limit to how long a process can stay in the processor.

The above image shows how the dispatcher switches between the different processes in the previous image to the one above.

Simple Two-State Process Model



(a) State transition diagram

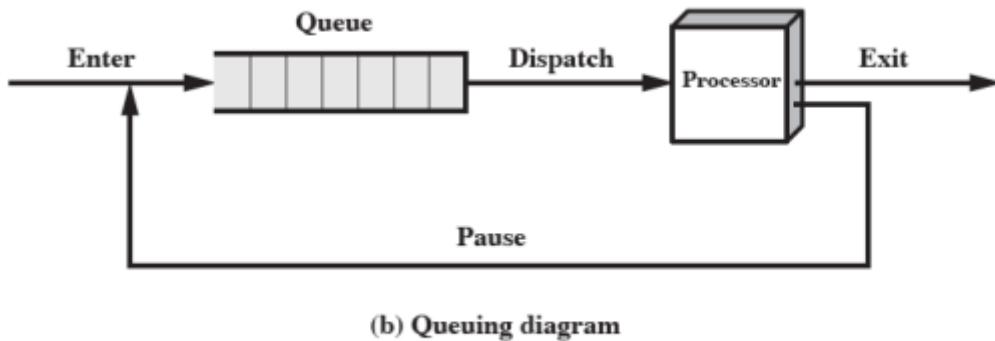
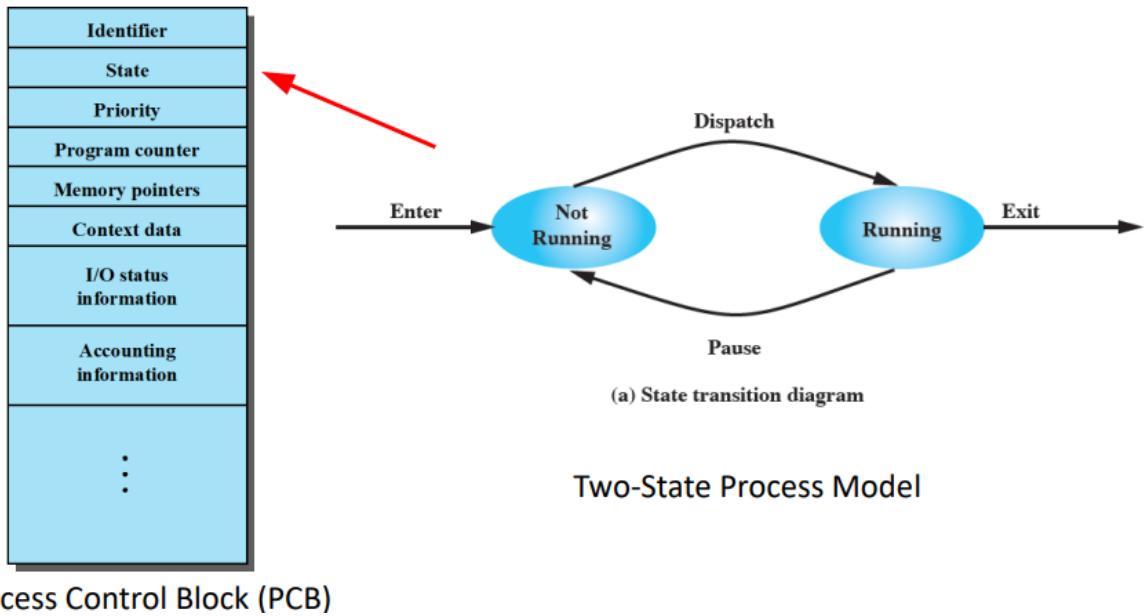
Two-State Process Model

Processes change state as they execute.

With a 2 state model we have 2 states:

- running
 - the process is being executed

- not running
 - not being executed by processor



Two-State Process Model

We can have the processes in a queue stored in the ram and we cycle through them, kicking a process out of the processor and back into the queue. When we do that we set it's state to the **Not Running** state.

When we start executing the new process we get to set that process' state to **Running**.

Note: the entries in the queue are links to the process' PCB as opposed to the whole process itself.

The queue is a first-in-first-out list and the processor operates in round-robin fashion on the available processes.

Round-Robin: each process in the queue is given a certain amount of time, in turn, to execute and then returned to the queue, unless blocked.

I think every prof seems to have a different definition of "round robin"

Operations on Processes

We need to be able to _____ processes:

- create
- terminate

Creation

4 reasons to create a new process:

- new batch job
 - OS is provided with a batch job control system
 - operating system is given a ton of jobs and has to execute them all in the proper priority
- interactive logon
 - system d, inet, etc
 - tons of processes start up to serve the user
- created by os to provide a service
 - when the user does stuff they will often start processes, without the user having to wait
 - send information to printers
- spawned by existing process
 - opportunity to exploit parallelism or for the purposes of modularity
 - modular design will allow any user to take advantage of the work that the process is doing
 - processes made to support other processes

On Spawning:

- parent process creates child processes forming a tree of processes
- they may or may not be communicating with one another and/or sharing resources
- all, some, or none of the resources are shared between the parent and child process
- the parent and child may execute concurrently or the parent may wait until the child terminates
- it all depends on what the purpose for the spawning was and what the processes were designed to do

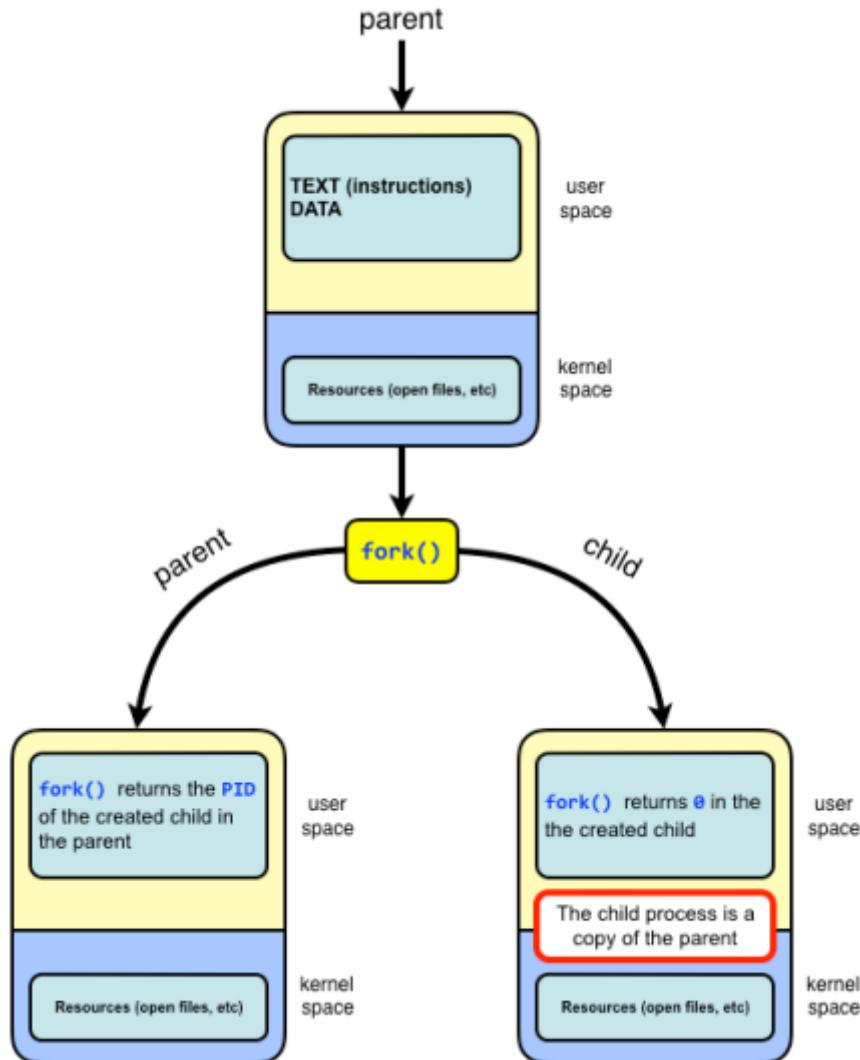
fork

In unix-like OSes, the creation of a child process is established through a system called `fork`

`fork` system call creates a copy by duplicating the parent process

- returns the child process ID to the parent and returns 0 to the child process

- the child and parent processes know each others PIDs



child process can execute a different program using `exec` system call

child can't terminate until the parent collects the output of the child. parent gives a `wait` command to the child and does the collection.

While the child waits then the child is known as a zombie.

If the parent dies then the child becomes an orphan.

fork can be used in a program written by the user.

There are a variety of different reasons why you would want to fork so there's many things that the child and parent processes can do together or separate. We could see that they duplicate and continue to execute the same thing or they'll duplicate and start doing different things and receiving commands from the parent program.

Mircosoft doesn't like `fork()`

- fork is insecure
- fork gives the child more access than it needs since it copies everything by default
- defies the least privilege principle

- give the least privilege possible to get a job done in order to protect the user and system
- programs that fork but don't exec render address-space layout randomisation ineffective, since each process has the same memory layout

fork is fast and has 0 parameters while windows' CreateProcess has many parameters.

We also don't run into the error of running out of memory.

`prof` really likes this paper and found it very interesting.

CreateProcess

Windows doesn't have a direct equivalent to fork.

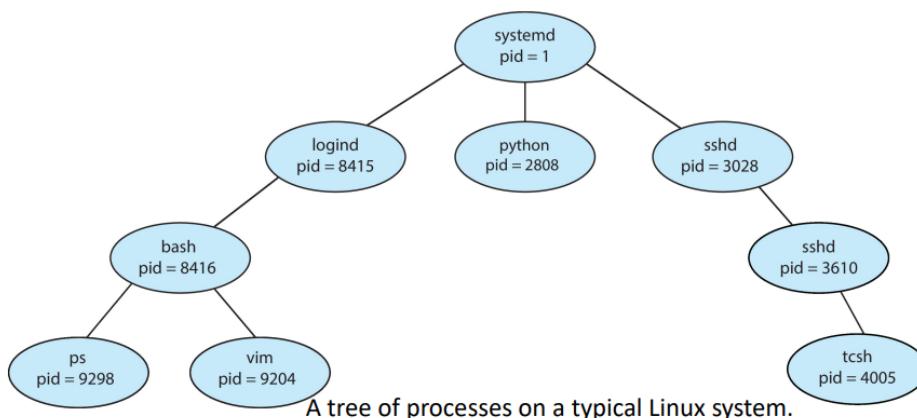
`CreateProcess` is used by Windows API

Spawn vs Fork

- fork
 - existing process creates identical copy of itself
- spawn
 - more general term for creating new process
 - implemented in various different ways depending on OS and programming env
 - doesn't necessarily involve duplicating the existing process
 - spawned process may or may not have a relationship with the parent process
 - spawned process can be a different program altogether

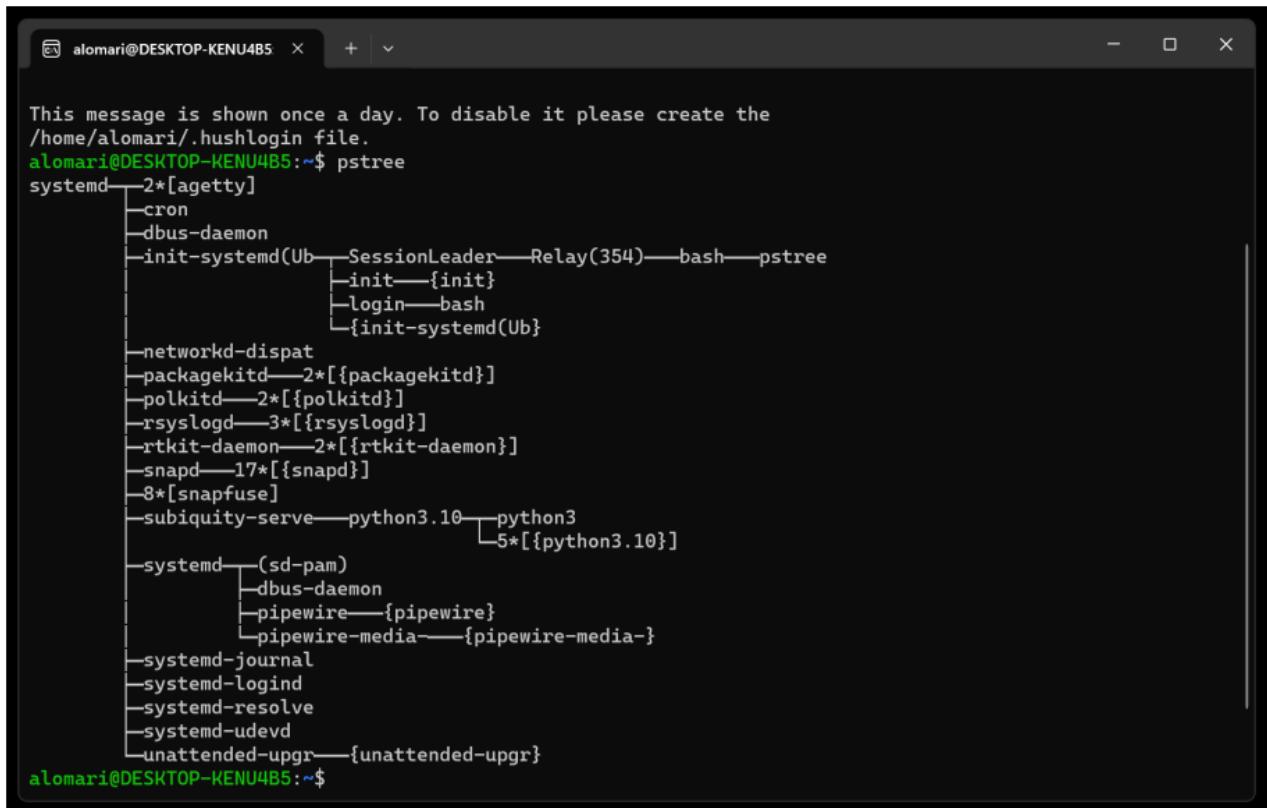
Tree of Processes in Linux

- **systemd** serves as the system's initial process and is identified as the **root** of all child processes. It is assigned a **pid** of 1, and is the first process created when the system is booted.



`pstree`

prof:



This message is shown once a day. To disable it please create the /home/alomari/.hushlogin file.

```
alomari@DESKTOP-KENU4B5:~$ pstree
systemd—2*[agetty]
|   cron
|   dbus-daemon
|   init-systemd(Ub)—SessionLeader—Relay(354)—bash—pstree
|       |   init—{init}
|       |   login—bash
|       |   {init-systemd(Ub)}
|   networkd-dispat
|   packagekitd—2*[{packagekitd}]
|   polkitd—2*[{polkitd}]
|   rsyslogd—3*[{rsyslogd}]
|   rtkit-daemon—2*[{rtkit-daemon}]
|   snapd—17*[{snapd}]
|   8*[snapfuse]
|   subiquity-serve—python3.10—python3
|       |   5*[{python3.10}]
|   systemd—(sd-pam)
|       |   dbus-daemon
|       |   pipewire—{pipewire}
|       |   pipewire-media—{pipewire-media-}
|   systemd-journal
|   systemd-logind
|   systemd-resolve
|   systemd-udevd
|   unattended-upgr—{unattended-upgr}
alomari@DESKTOP-KENU4B5:~$
```

mine:

```
[tt@www ~]$ pstree
systemd—ModemManager—3*[{ModemManager}]
      └─NetworkManager—3*[{NetworkManager}]
        ├─accounts-daemon—3*[{accounts-daemon}]
        └─avahi-daemon—avahi-daemon
      └─bluetoothd
      └─3*[chrome_crashpad—2*[{chrome_crashpad}]]
        └─chrome_crashpad—{chrome_crashpad}
      └─code—code—code—18*[{code}]
        └─code—code—code—17*[{code}]
          └─code—7*[{code}]
          └─code—code—7*[{code}]
            └─14*[{code}]
          └─code—15*[{code}]
          └─code—17*[{code}]
            └─38*[{code}]
      └─conky—{conky}
      └─crond
      └─cupsd
      └─dbus-broker-lau—dbus-broker
      └─lightdm—Xorg—7*[{Xorg}]
        └─lightdm—i3—Discord—Discord—19*[{Discord}]
          └─Discord—Discord
          └─2*[Discord—4*[{Discord}]]
          └─Discord—42*[{Discord}]
            └─28*[{Discord}]
          └─alacritty—bash—pstree
            └─8*[{alacritty}]
          └─chrome—2*[cat]
            └─chrome—chrome—18*[{chrome}]
            └─chrome—chrome—7*[{chrome}]
              └─6*[chrome—12*[{chrome}]]
              └─16*[chrome—11*[{chrome}]]
              └─chrome—9*[{chrome}]
              └─2*[chrome—15*[{chrome}]]
              └─chrome—13*[{chrome}]
              └─chrome—8*[{chrome}]
            └─nacl_helper
          └─chrome—chrome
            └─12*[{chrome}]
          └─chrome—7*[{chrome}]
            └─35*[{chrome}]
          └─clipit—3*[{clipit}]
          └─i3bar—i3status
          └─nm-applet—4*[{nm-applet}]
          └─pa-applet—3*[{pa-applet}]
          └─pamac-tray—3*[{pamac-tray}]
```

Termination

A process terminates when it finishes executing its final statement and askss the operating system to delete it by using the `exit()` system call

all the resources of the process (physical and virtual) are deallocated and reclaimed by the operating system.

Batch job should include a halt instruction or an explicit os service call for termination. Halt will generate an interrupt to alert the os that a process has completed.

Interactive applications (browsers and shit used by user) are terminated when the user decides typically.

parents can terminate children

WOW:

Normal completion	The process executes an OS service call to indicate that it has completed running.
Time limit exceeded	The process has run longer than the specified total time limit. There are a number of possibilities for the type of time that is measured. These include total elapsed time ("wall clock time"), amount of time spent executing, and, in the case of an interactive process, the amount of time since the user last provided any input.
Memory unavailable	The process requires more memory than the system can provide.
Bounds violation	The process tries to access a memory location that it is not allowed to access.
Protection error	The process attempts to use a resource such as a file that it is not allowed to use, or it tries to use it in an improper fashion, such as writing to a read-only file.
Arithmetic error	The process tries a prohibited computation, such as division by zero, or tries to store numbers larger than the hardware can accommodate.
Time overrun	The process has waited longer than a specified maximum for a certain event to occur.
I/O failure	An error occurs during input or output, such as inability to find a file, failure to read or write after a specified maximum number of tries (when, for example, a defective area is encountered on a tape), or invalid operation (such as reading from the line printer).
Invalid instruction	The process attempts to execute a nonexistent instruction (often a result of branching into a data area and attempting to execute the data).
Privileged instruction	The process attempts to use an instruction reserved for the operating system.
Data misuse	A piece of data is of the wrong type or is not initialized.
Operator or OS intervention	For some reason, the operator or the operating system has terminated the process (e.g., if a deadlock exists).
Parent termination	When a parent terminates, the operating system may automatically terminate all of the offspring of that parent.
Parent request	A parent process typically has the authority to terminate any of its offspring.

not reading allat but pretty cool :thumbsup:

Five-State Process Model

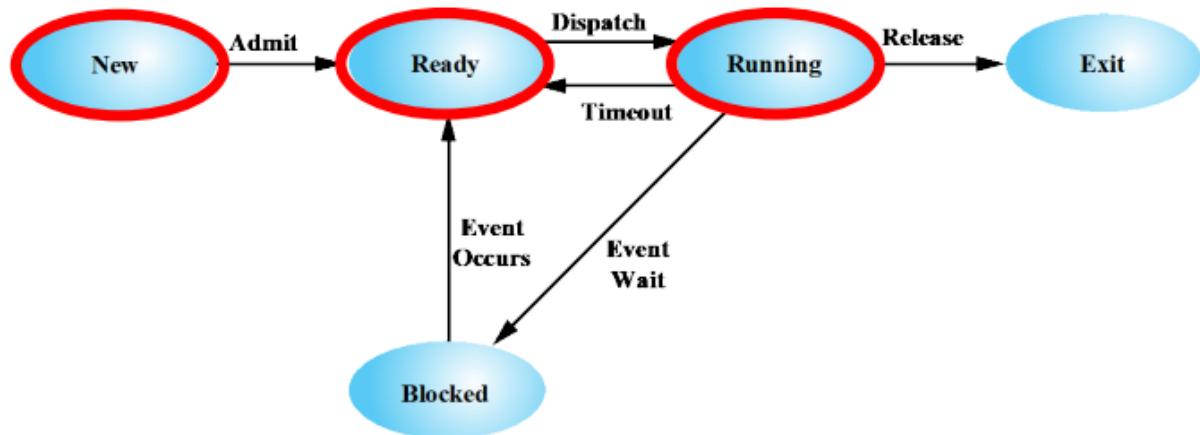
Note: THis is on the assignemnt (Thanks Cate)

USing a single queue the dispatcher could not just select the process at the oldest end.

Dispatcher needs to find a process that is not blocked and that has been in the list for the queue for the longest.

Turn not **Not** Running into **Ready** and **Blocked**

State process diagram



Five-State Process Model

States:

- new
 - newly created process not yet admitted to the pool of executable processes by the OS
- ready
 - prepared to execute when given opportunity
- running
 - process that is currently being executed
- blocked
 - unix-based: waiting
 - process cannot execute until some event occurs such as completeion of an I/O operation
- exit
 - unix-based: terminated
 - process released from pool of executable processes by the OS
 - process has been halted or aborted for whatever reason

Transitions:

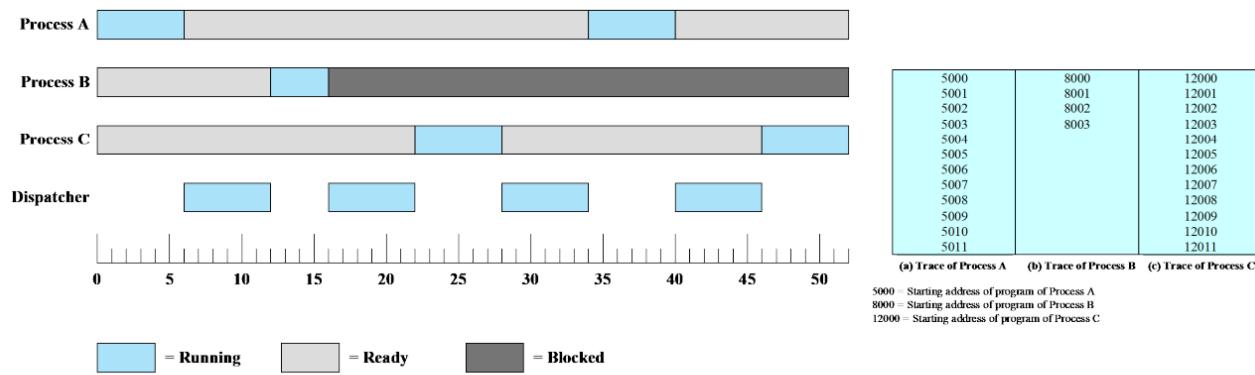
- null -> new
 - new process is created to exec program
- new -> ready
 - os is prepared to take on additional process
- ready -> running

- dispatcher dispatches the process
- OS chooses ready process to run
- running -> exit
 - process is terminated by the OS if the process indicates it has completed or if it aborts
- running → ready
 - running process has reached the maximum allowable time for uninterrupted execution
 - no longer your turn on the xbox
- running → blocked
 - process requests something for which it must wait
- blocked → ready
 - process in blocked state is moved to ready state when the event for which it has been waiting occurs

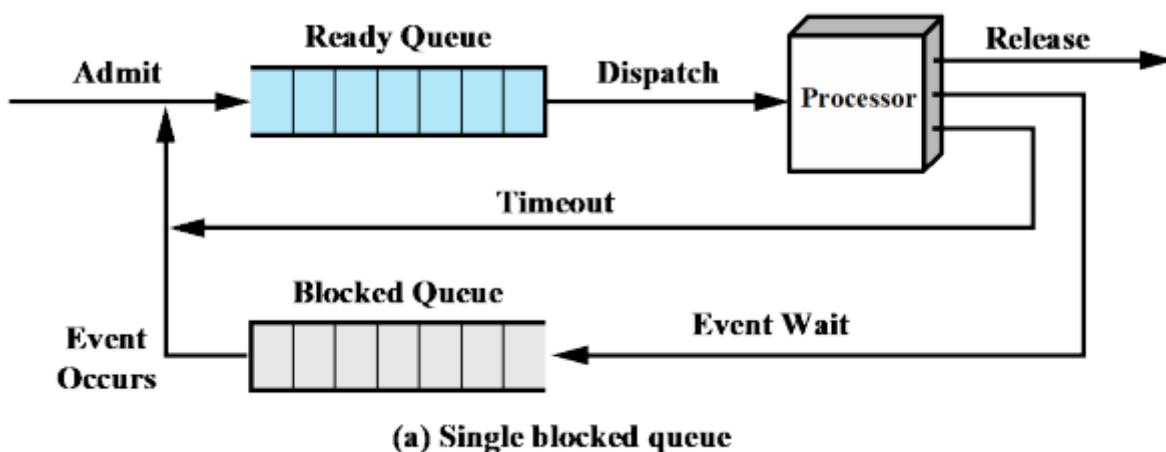
remember that we never have blocked to running.

still not perfect but more later

Example



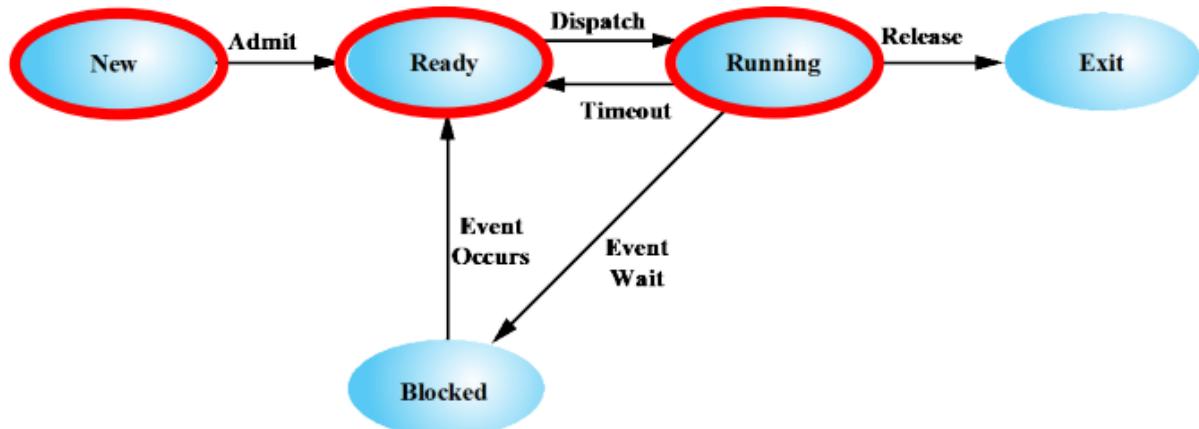
while the processes are simply ready but not running, the dispatcher is running in order to go to dispatch them



There are 2 queues:

1. ready queue

2. blocked queue



Five-State Process Model

So when we ready something up it actually goes into a queue and when something gets hit with an event wait it gets put into a blocked queue.

When a process

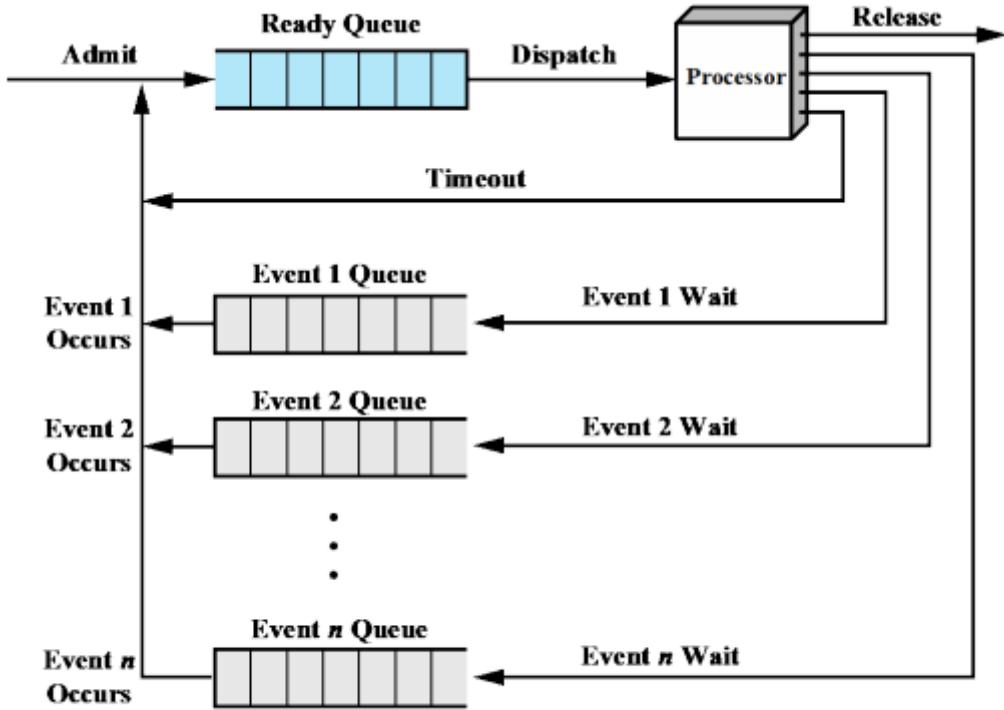
- is admitted to the system → admit to ready queue
- is chosen to run → select from ready queue
- has to wait → send to blocked queue
- times out → send to ready queue again
- finishes → release

In the absence of any priority scheme this can be a simple FIFO queue (round-robin). Otherwise we could use something like a priority queue.

When an event occurs, any process in the blocked queue that has been waiting on that event only is moved to the ready queue.

The problem is that we have to scan the entire blocked queue which could have 100s or 1000s of blocked processes.

We can remedy this with the use of multiple blocked queues



(b) Multiple blocked queues

For each event that will cause an interrupt we will make a queue for it that way we don't have to scan through the whole queue to find a process that corresponds to the event.

Suspended Processes

Problem - it will become common for all of the processes in memory to be waiting for I/O, and for the processor to become idle.

- there's only so much memory space to store so many processes
- we can pretty easily fill up the memory with all the processes and leave the processor idle since it only really works with shit in the memory.

Solution - swapping

- move part or all of a process from main memory to disk
- it's a very specific part of the disk that's allocated for the memory to work with
- "there is a specific page file dedicated by the operating system"
- this is like the linux swap

The only problem is how do we do that?

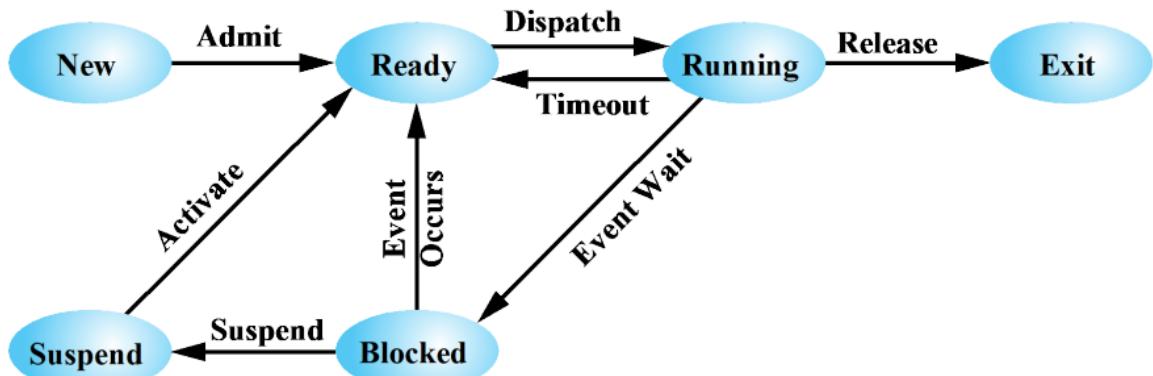
We add another state to our process model

When the OS swaps one of the **blocked** processes out onto the disk we put it into a **suspended** queue:

- a queue of existing processes that have been temporarily kicked out of the main memory or suspended

the OS then brings in another process from the suspend queue or it honors a new-process request

execution then continues with the newly arrived process



(a) With One Suspend State

So now when we block a process from there it can either go back to the ready queue or get sent to the suspend queue.

When it gets sent to the suspend queue then it gets sent to the disk to make more space in the memory.

Once it's in the suspend queue, we can activate it again.

So there are now 4 ways to get a process into the ready queue

1. a newly admitted process
2. a process times out
3. a blocked process is unblocked thanks to the event it was waiting on occurring
4. a suspended process is selected to become active once more

If there is nothing in the ready queue and we're waiting on blocked processes then we will start suspending processes to swap out of main memory so that we can admit new processes or activate one from the suspend.

It's always preferable to activate a suspended process.

The swap is always available - it is built into the architecture - and there's always labels for each state.

When the OS performs a swapping-out operation, it has 2 choices for selecting a process to bring into main memory:

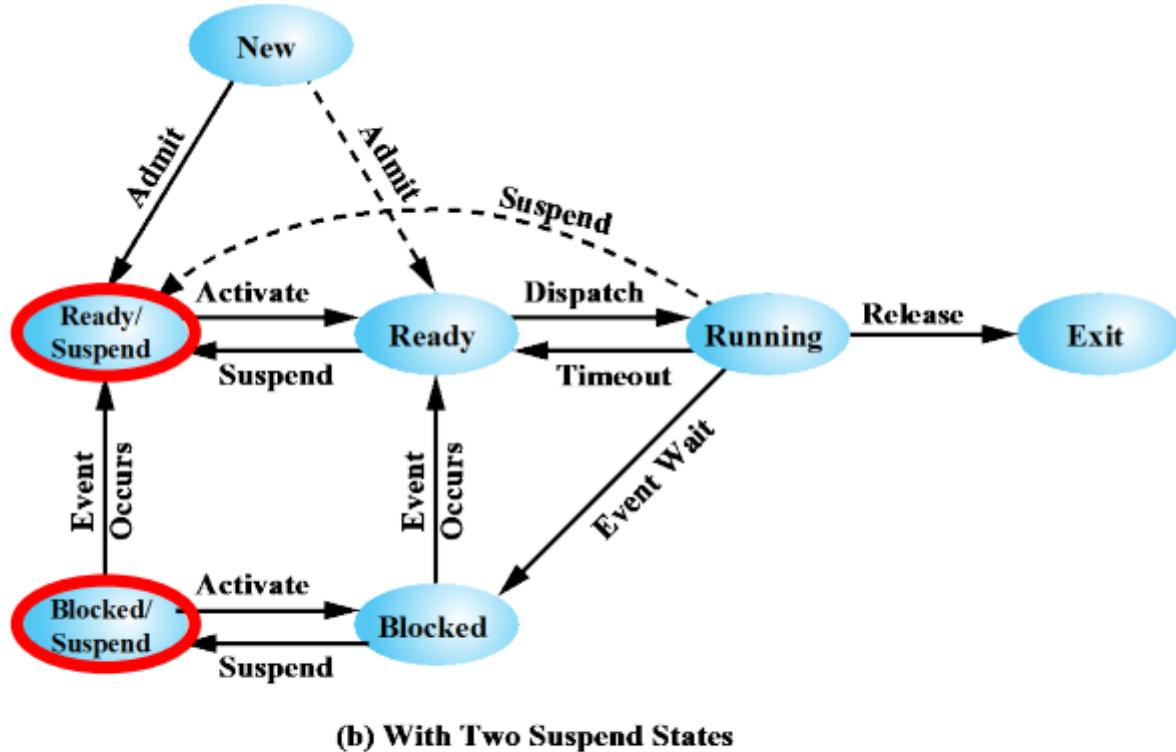
1. admit a newly created process
2. bring a previously suspended process (preferred)

problem:

- all suspended processes were blocked when they were suspended
- which one do we bring back?

Each process in the suspend state was originally blocked on a particular event, something that accounted for with the use of multiple queues. We also know that once that event is over we know that the processes in its corresponding queue are no longer blocked.

We want to make another state for suspended processes



Dashed lines in the figure indicate possible but not necessary transitions.

For the processes in the secondary memory we're able to have a ready and blocked queue for the suspended.

When we hear/read “suspend” then we think of the disk.

Blocked → Blocked+Suspend:

- if there are no Ready processes then at least one blocked process is swapped out to make room for another process that is not blocked.

Blocked+Suspend → Ready+Suspend:

- a process was waiting for an event

- the event finishes
- the process is no longer blocked but it's still suspended so it remains in the disk space but gets sent to a suspended ready queue

Ready+Suspend → Ready:

- when there are no **Ready** processes the OS will need to bring one in to continue execution

Ready → Ready+Suspend:

- to free up a sufficiently large block of main memory
- we have a ton of processes that are ready.

Running → Ready+Suspend:

- the OS is preempting the process because a higher-priority process on the **Blocked+Suspend** queue has just become unblocked
- something higher priority just got unblocked and we need to execute on that one now instead
- the OS could move the running process directly to the **Read+Suspend** queue and free some main memory
- process preemption occurs when an execution process is interrupted by the processor so that another process can be executed

New → Ready+Suspend:

New → Ready:

We could add more states but then there's more work involved in changing between states. There are more states in different designs however.

Q: How do we decide whether a **New** process goes straight to **Ready+Suspend** or **Ready**?

A: we prefer to send a new process to the suspended queue. The OS will prefer grabbing from ready suspend. Very rare that we will send something straight to ready since that would mean that the suspend is empty as well so there would be no point in sending it there first and wasting compute time on switching

A suspended process

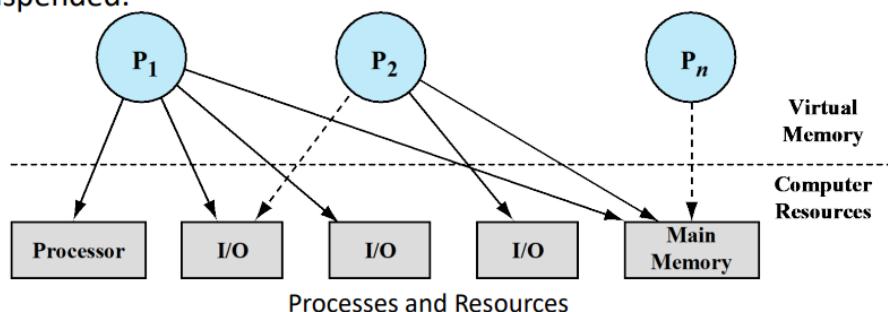
- non immediately available for execution
 - all info sitting on secondary memory instead of main
- process was placed in a suspended state by
 - itself
 - a parent process
 - OS
- suspended in order to prevent execution
- may or may not be waiting on an event
- process may not be removed from this state until the agent orders the removal
- reasons for suspensions

- swapping
- other OS reason
 - illegal memory access
 - usually terminates in that case
 - buffer overflow
- interactive user request
 - debugging
 - want to use resource
 - minimizing applications or switching off of tabs so we don't really need the program working
- timing
- parent process request

Process Description and Control

In a multiprogramming env there are several processes that have been created and exist in virtual memory.

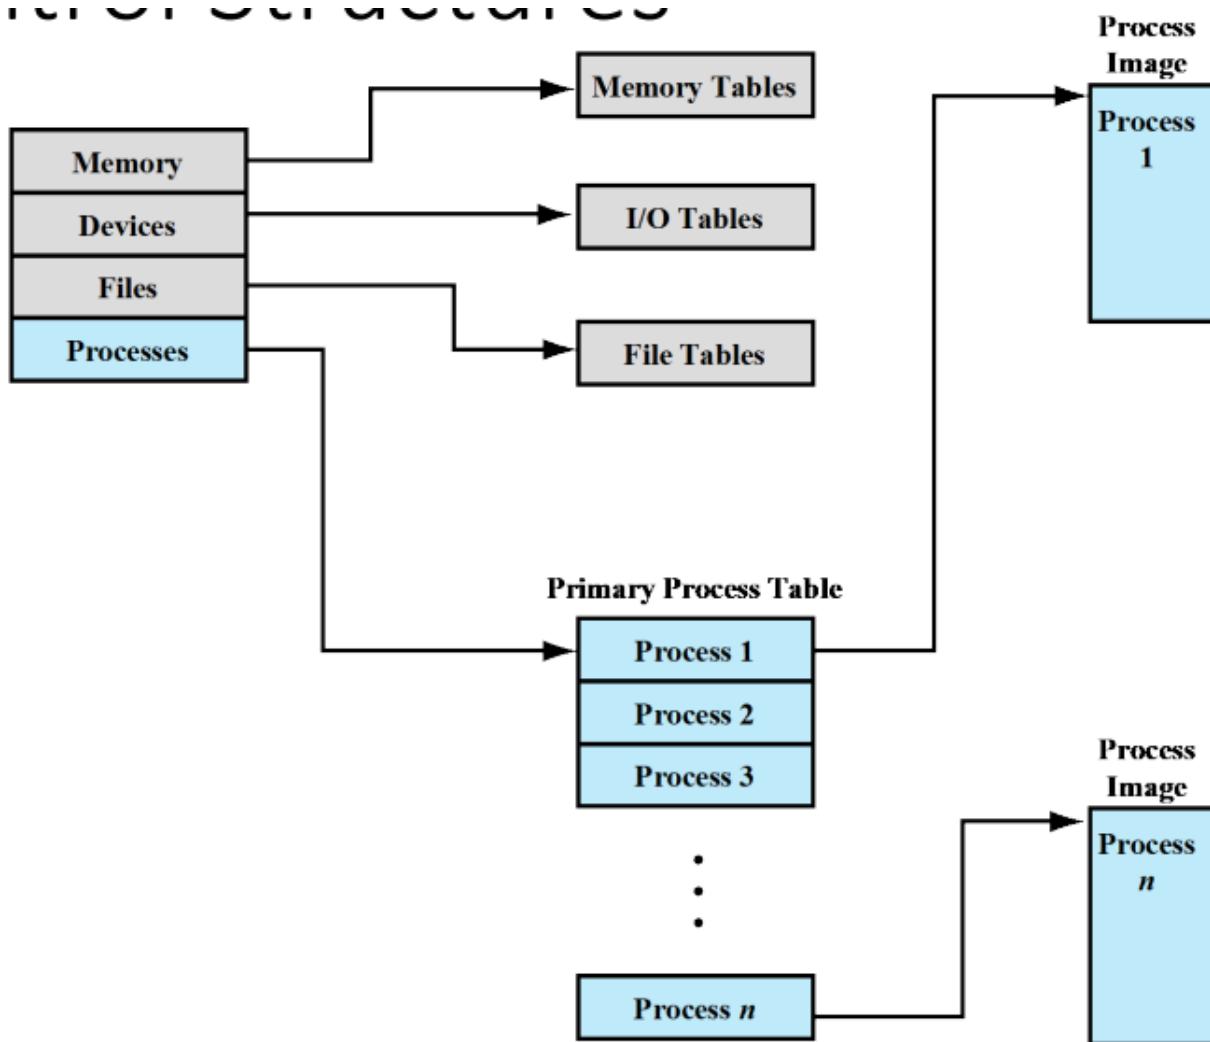
- In a multiprogramming environment, there are several processes (P_1, \dots, P_n) that have been created and exist in virtual memory.
- Each process needs access to certain system resources:
 - P_1 is running; part of the process is in main memory, and it has control of two I/O devices.
 - P_2 is also in main memory but is blocked waiting for an I/O device allocated to P_1 .
 - P_n is suspended.



What info does the OS need to control processes and manage resources for them?

OS constructs and maintains tables of information about each entity that it is managing:

- memory tables
- I/O tables
- file tables
- process tables



General Structure of Operating System Control Tables

Memory Tables

main = real

secondary = virtual

we use memory tables to do:

- allocation of main memory to processes
- allocation of secondary memory to processes
- any protection attributes of blocks of main or virtual memory
 - which processes may access certain shared memory regions
 - prevent programs from accessing main or virtual memory that it's not allowed to
- any information needed to manage virtual memory

I/O Tables

used by the OS to manage the I/O devices and channels of the computer system

at any given time, an I/O device may be available or assigned to a particular process

OS needs to know the status of the I/O operation and the location in main memory being used as the source or destination for the I/O transfer

File Tables

Provide information about the:

- existence of files
- location on secondary memory
- current status
- other attrs

may be maintained by a file management system (windows NTFS, linux/unix FAT) or the OS based on different operating system

Process Tables

Must be maintained to manage processes

References memory, I/O, and files directly or indirectly in the process tables to show what the process needs/uses

when the OS is initialized it must have some access to some configuration data that define the basic environment

- all the info required to build out the tables
- everything you can see and view in the bios
- I think anyways

Process Control

to manage and control a process, the OS must know

- where the process is located
- the attrs of the process that are necessary for its management
- basically everything in the PCB

What is the physical manifestation of a process?

a process must include:

- a program or set of programs to be executed
- a set of data location for local and global variables

- stack that is used to keep track of
 - procedure calls
 - parameter passing between procedures
- a number of attributes that are used by the OS for process control (Process control block)

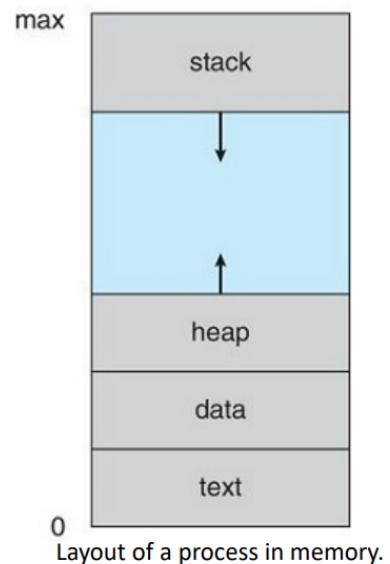
We refer to this as the `process image`

Process Layout in Memory

Process Layout in Memory

The memory layout of a process is typically divided into:

- **Text section**—the executable `code`
- **Data section**—global variables
- **Heap section**—memory that is dynamically allocated during program run time
- **Stack section**—temporary data storage when **invoking functions** (such as function parameters, return addresses, and local variables)



The stack goes from higher to lower addresses while the heap grows from lower addresses to higher addresses.

Process Location

location of a process image will depend on the memory management scheme being used

At any given time parts of the process image can be both in the main and secondary memory

so we need process tables maintained by the OS to show the location of each page of each process image

Process Attributes

info resides in a process control block

there are 3 general categories

1. process identification
2. processor state information
3. process control information

process identification:

- each processor gets a unique numeric identifier
 - PID
- many of the bales controlled by the OS may use process identifiers to cross-reference process tables
- interprocess communication uses PIDs to determine who is who and where
- PIDs let us see who are parent and descendent processes
 - the parents have the child PIDs and the children have the parent PID

Process State Information

- what is the current state of the process
- consists of the contents of processor registers
 - registers on the cpu itself, very fast
- nature of registers involved depend on the design of the processor but typically
 - user-visible registers
 - control and status registers
 - stack pointers
- all processor designs include a register or set of registers often known as the program status word (PSW) that contains condition codes plus status information
 - EFLAGS register is an example of a PSW used by any OS running on an x86 processor



X ID = Identification flag
X VIP = Virtual interrupt pending
X VIF = Virtual interrupt flag
X AC = Alignment check
X VM = Virtual 8086 mode
X RF = Resume flag
X NT = Nested task flag
X IOPL = I/O privilege level
S OF = Overflow flag

C DF = Direction flag
X IF = Interrupt enable flag
X TF = Trap flag
S SF = Sign flag
S ZF = Zero flag
S AF = Auxiliary carry flag
S PF = Parity flag
S CF = Carry flag

S Indicates a Status Flag
C Indicates a Control Flag
X Indicates a System Flag
Shaded bits are reserved

○

Example of Processor Status Word: X86 EFLAGS Register

x86

EFLAGS

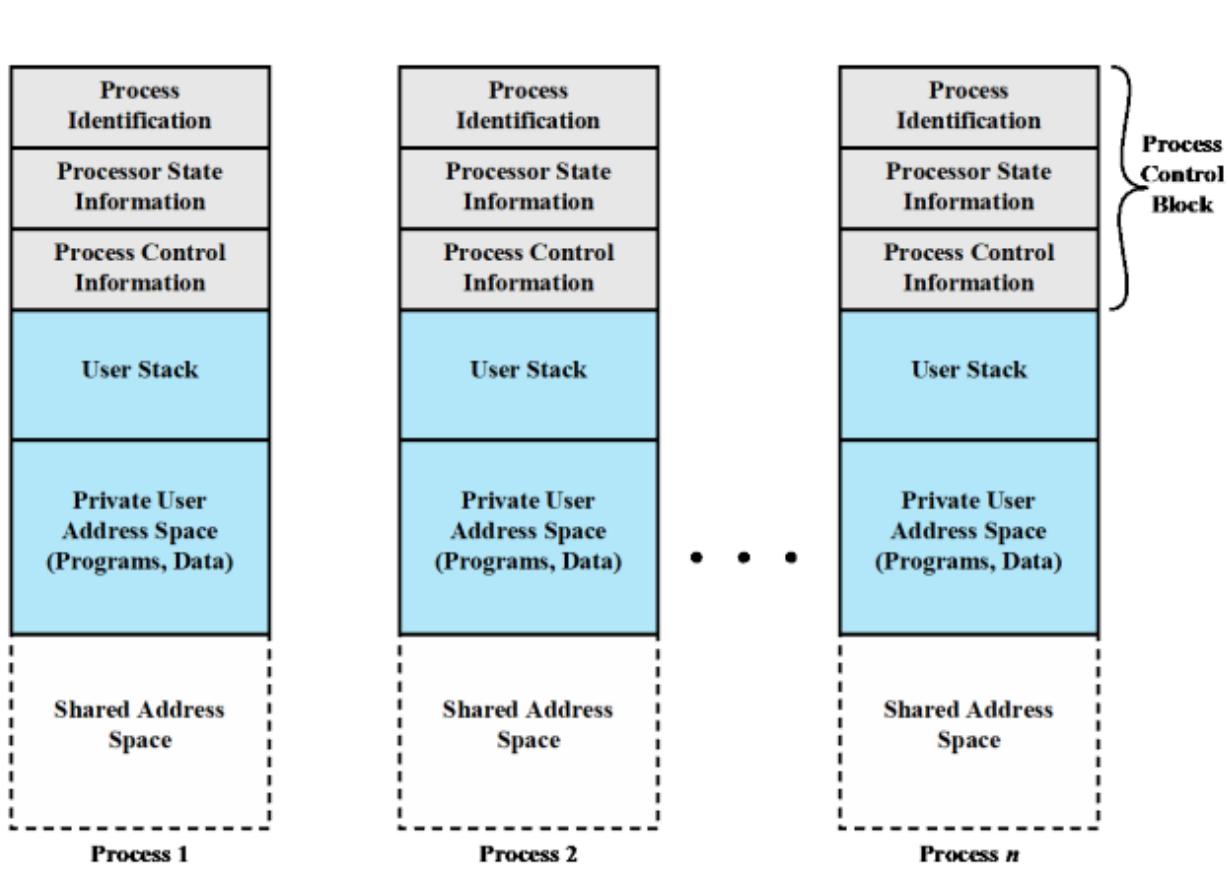
Register Bits

○

Status Flags (condition codes)	
AF (Auxiliary carry flag)	Represents carrying or borrowing between half-bytes of an 8-bit arithmetic or logic operation using the AL register.
CF (Carry flag)	Indicates carrying out or borrowing into the leftmost bit position following an arithmetic operation. Also modified by some of the shift and rotate operations.
OF (Overflow flag)	Indicates an arithmetic overflow after an addition or subtraction.
PF (Parity flag)	Parity of the result of an arithmetic or logic operation. 1 indicates even parity; 0 indicates odd parity.
SF (Sign flag)	Indicates the sign of the result of an arithmetic or logic operation.
ZF (Zero flag)	Indicates that the result of an arithmetic or logic operation is 0.
Control Flag	
DF (Direction flag)	Determines whether string processing instructions increment or decrement the 16-bit half-registers SI and DI (for 16-bit operations) or the 32-bit registers ESI and EDI (for 32-bit operations).
System Flags (should not be modified by application programs)	
AC (Alignment check)	Set if a word or doubleword is addressed on a nonword or nondoubleword boundary.
ID (Identification flag)	If this bit can be set and cleared, this processor supports the CPUID instruction. This instruction provides information about the vendor, family, and model.
RF (Resume flag)	Allows the programmer to disable debug exceptions so that the instruction can be restarted after a debug exception without immediately causing another debug exception.
IOPL (I/O privilege level)	When set, causes the processor to generate an exception on all accesses to I/O devices during protected mode operation.
IF (Interrupt enable flag)	When set, the processor will recognize external interrupts.
TF (Trap flag)	When set, causes an interrupt after the execution of each instruction. This is used for debugging.
NT (Nested task flag)	Indicates that the current task is nested within another task in protected mode operation.
VM (Virtual 8086 mode)	Allows the programmer to enable or disable virtual 8086 mode, which determines whether the processor runs as an 8086 machine.
VIP (Virtual interrupt pending)	Used in virtual 8086 mode to indicate that one or more interrupts are awaiting service.
VIF (Virtual interrupt flag)	Used in virtual 8086 mode instead of IF.

Process Control Information

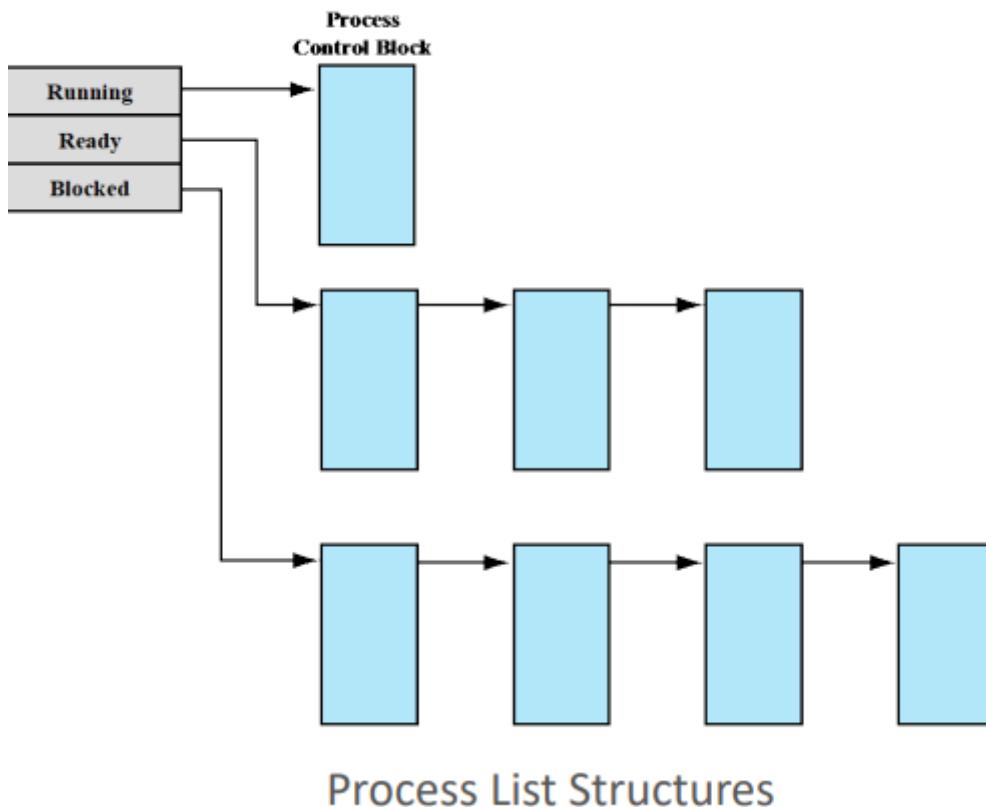
- the additional information needed by the OS to control and coordinate the various active processes



User Processes in Virtual Memory

Role of the Process Control Block

- most important data structure in an OS
 - contains all info needed by OS about process
 - blocks are read and/or modified by virtually every module in the OS
- the set of PCBs define that state of the OS



Process Control

Typical management functions:

- creation and termination
- scheduling and dispatching
- switching
- synchronization and support for interprocess communication
- management of PCBs

we're going to be looking at a lot of these.

Creation

OS decides to create a new process

1. assign a unique PID to the new process
2. allocate space for the process

3. initialize the process control block
4. set the appropriate linkages (put it in the queue)
5. creates or expands other data structures

Lect start

Switching

A process switch may occur at any time that the OS has gained control from the currently running process

Some possible events that may give control to the OS

mechanism	cause	use
interrupt	something external to the execution of the current instruction interrupts	lets the OS react to that asynchronous external event
trap	associated with the execution of the current instruction	handling an error or exception condition
supervisor call	explicit request	call to an operating system function

Interrupt:

- an event external to and independent of the current process
- examples
 - clock interrupt
 - timeout
 - I/O interrupt
 - memory fault
 - wow so epic

Trap:

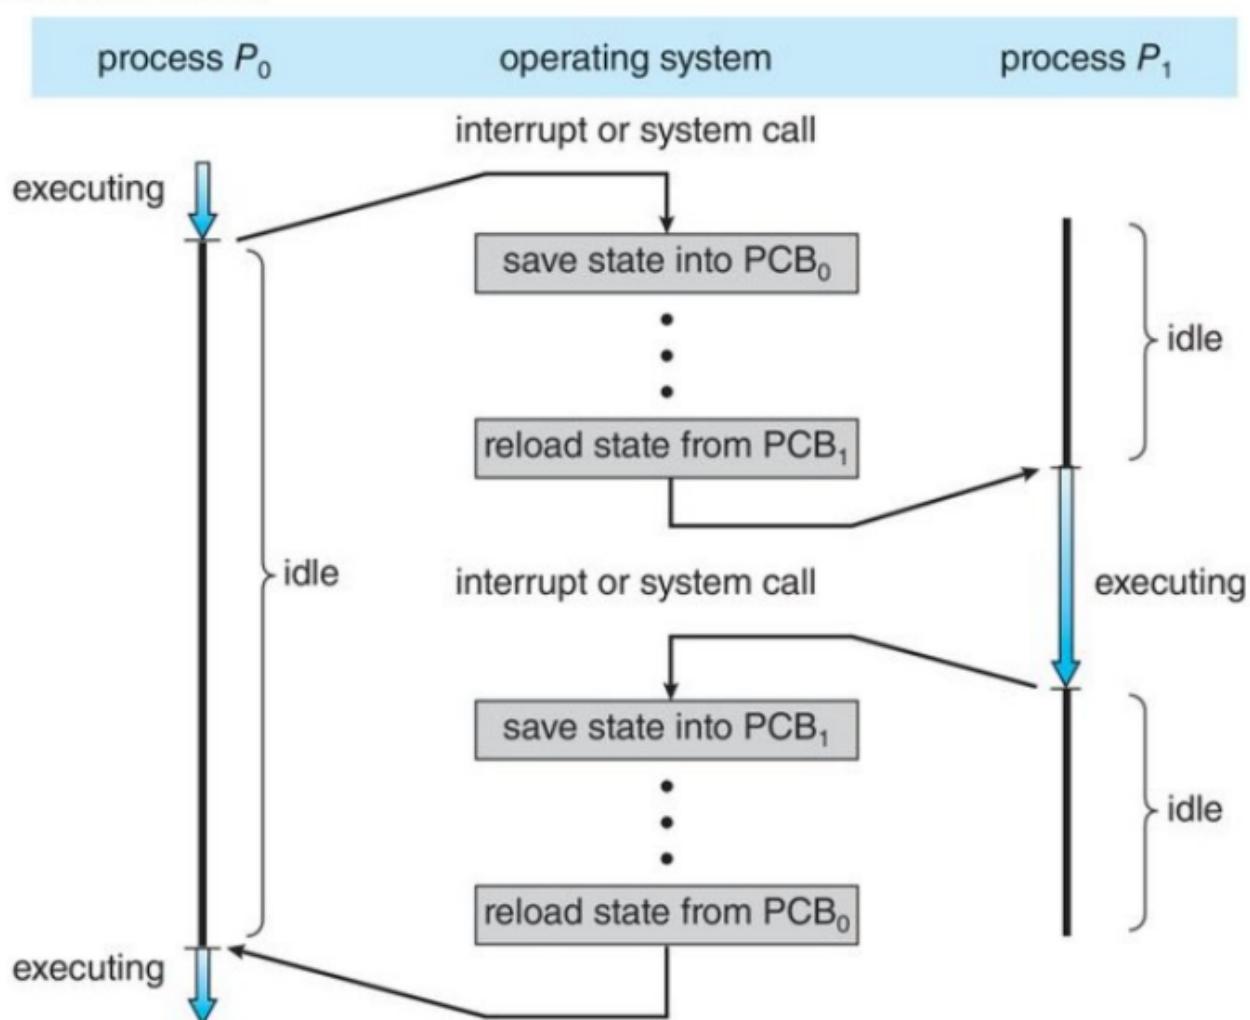
- an error or exception condition generated within the currently running process
- OS determines if the condition is **fatal**
 - fatal
 - move to exit state and a process switch occurs
 - not fatal
 - action will depend on the nature of the error and the design of the OS
 - the OS will try to recover

Supervisory Call:

- OS activated by a supervisor call from the program being executed
- getting the manager involved b/c of smth not on ur level

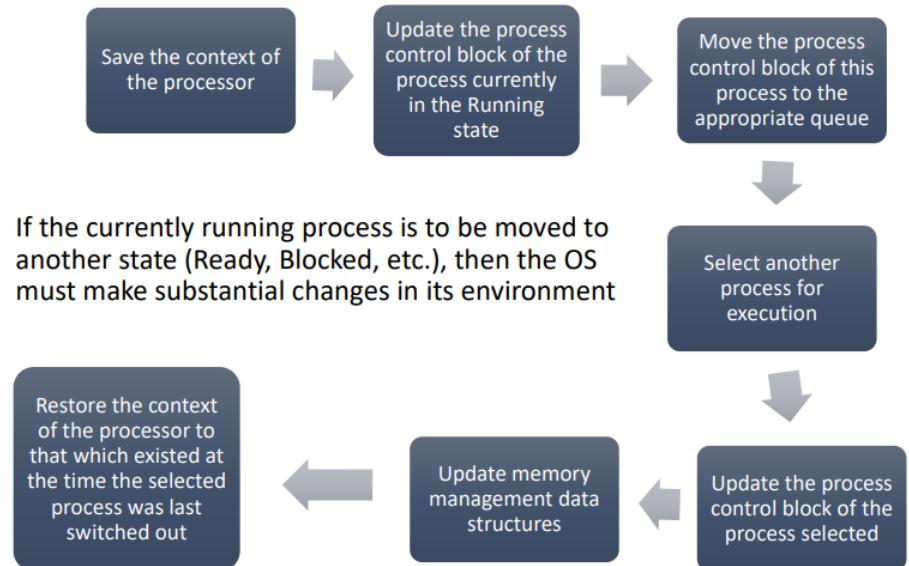
Context Switch

- description:
 - when the CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
 - process' context = PCB
- we have to efficiently save our old PCB and start using the new PCB to execute
- context switch time is overhead
 - no useful work is done while switching
 - more complex OS and PCB = longer switch time
 - we want to minimize context switches so we don't waste too much time
- time dependent on hardware support
 - multiple sets of registers per CPU = multiple loaded contexts
 - not common
 - typical speed is several microseconds



notice that there is idle time for the operating system to perform the saving of PCB_0 and subsequent loading of PCB_1

- The steps in a full process context switch are:



"blah blah blah" - professor

Execution of the Operating System

recall:

- OS functions in the same way as ordinary computer software
 - OS is a set of programs executed by the processor
- OS frequently relinquishes control and depends on the processor to restore control to the OS

Question

- OS is just a collection of programs
- it's executed by the processor like any other program
- is the OS a process?
- how is it controlled?

OS just throws control back to itself.

Whenever it gives away control there's always the contingency that it'll come back.

the collection of processes control one another.

3 approaches:

1. separate kernel
2. execution within user processes
3. process-based operating system

Separate Kernel

Execute the kernel of the OS outside of any process

- the concept of a the proess is considered to apply only to user programs
- OS executed as separate entity operated in privileged mode
- common on older OSes
 - not often used today

User Processes

All OS software is executed in the context of a user process:

- OS a collection of routines that the user calls to perform variousfunctions, executed within the environment of the user's process
- common on OSes on smaller computers (PCs)
 - commonly used in today's laptops

Process-Based OS

OS as a collection of system processes

- software that is part of the kernel executes in a kernel mode
- major kernel functgions are organized as separate processes
- imposes a program design discipline that encourages the use of a modular OS with minimal, clean interfaces between the modules.
- used in
 - servers in more modern computing
 - modular design

Summary

The most fundamental concept in a modern OS is the process.

The principal function of the OS is to create, manage, and terminate processes.

While processes are active, the OS must see that each is allocated time for execution by the processor, coordinate their activities, manage conflicting demands, and allocate system resources to processes.

Threads

Processes and Threads

A Process has 2 Characteristics

1- Resource Ownership

- A process includes a virtual address space to hold the process image.
- A process may be allocated control or ownership of resources, such as main memory, I/O devices, and files.
- The OS performs a protection function to prevent unwanted interference between processes with respect to resources.

2- Scheduling/Execution

- The execution of a process follows an execution path that may be interleaved with other processes.
- A process has an execution state (Running, Ready, etc.) and a dispatching priority, and is the entity that is scheduled and dispatched by the OS.

1- Resource Ownership

Usually referred to as a **process or task**.

2- Scheduling/Execution

Usually referred to as a **thread or lightweight process**.

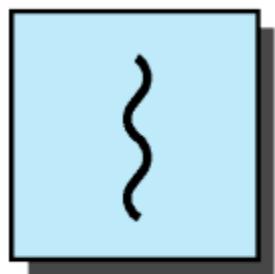
The actual fetching and execution of the instructions is thought of as a thread or lightweight process (LWP).

Most unix based processing will say threads.

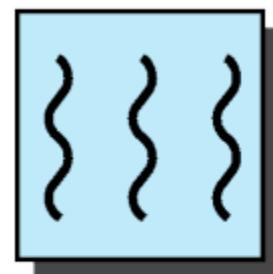
Some OSes don't differentiate between threads and processes, treating them all as processes.

The process model introduced so far assumed that a process as an executing program with a single thread of control.

Virtually all modern operating systems, however, provide features enabling a process to contain multiple threads of control.



**one process
one thread**



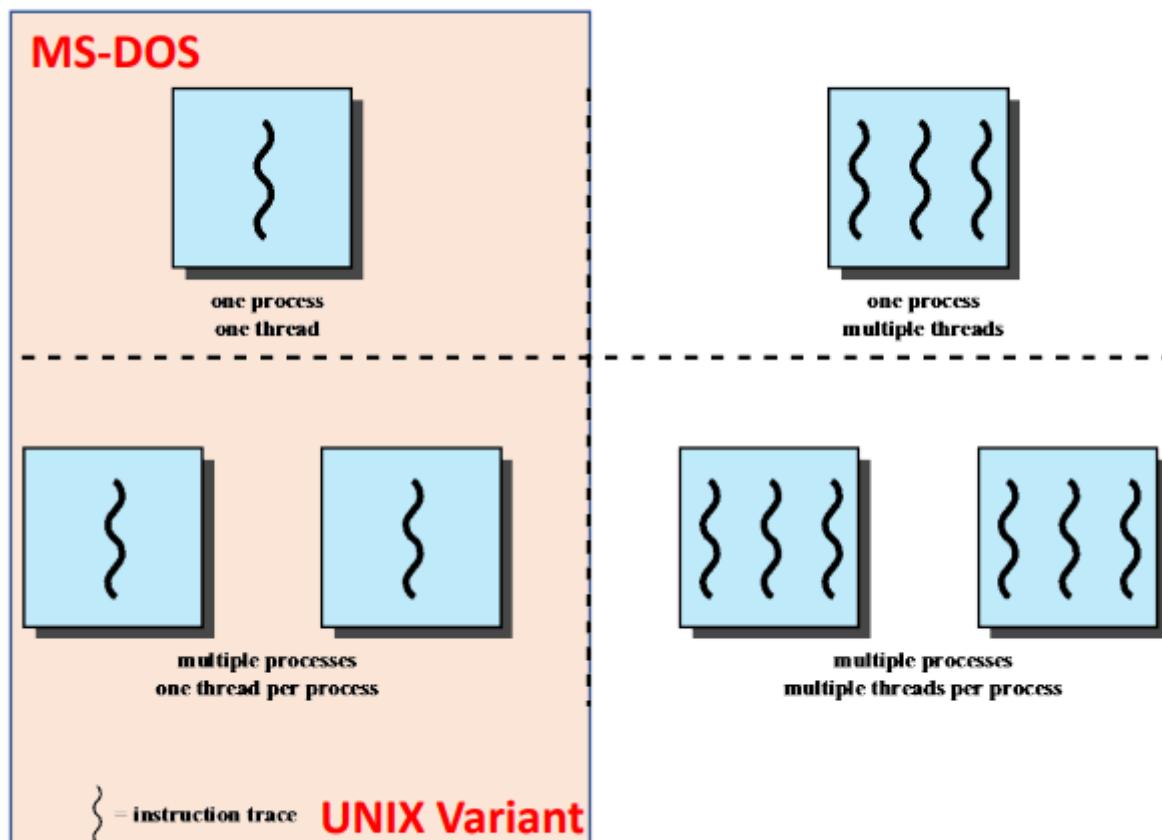
**one process
multiple threads**

Threaded Approaches

Single threaded approaches

a single thread of execution per process

threads are not recognized



Threads and Processes

ex.

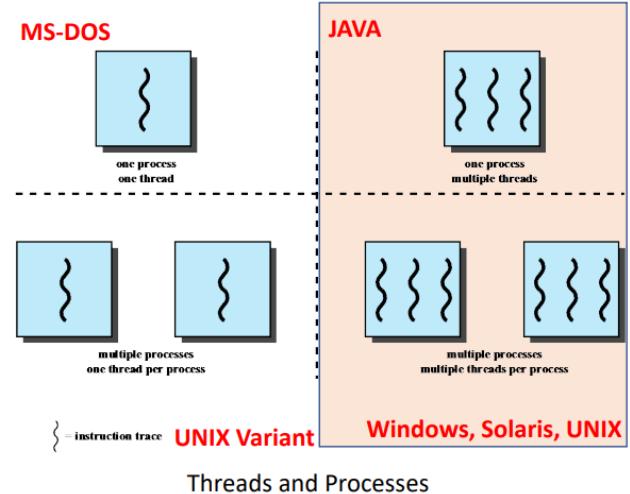
- ms-dos
- old linux

Multithreaded Approaches

- The ability of an OS to support multiple, concurrent paths of execution within a single process.

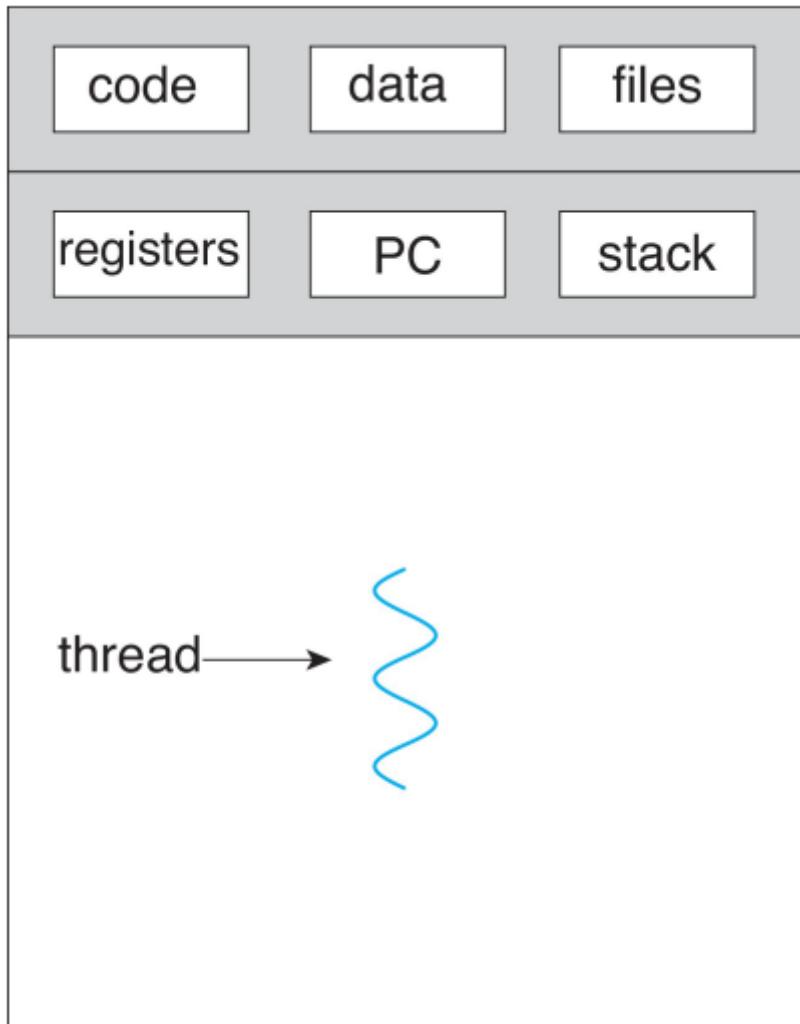
- Examples:

- A Java run-time environment is an example of a system of one process with multiple threads.
- Windows and many modern versions of UNIX are examples of multiple processes, each of which supports multiple threads.



we can make it seem as though the processes are all working simultaneously but we're actually just switching around very quickly

What is a Thread

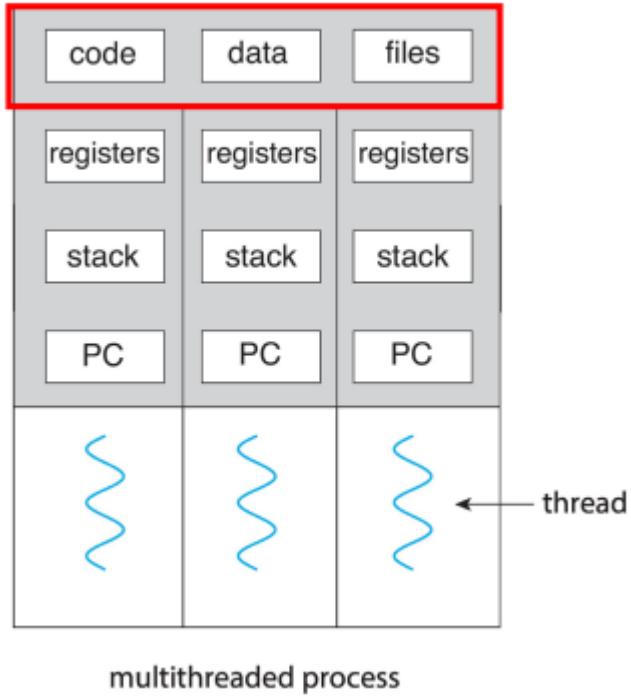


single-threaded process

thread is a basic unit of cpu utilization

- each thread knows its parent process
- thread ID
- program counter
- a register set
- a stack

thread = instruction trace



multithreaded process

If a process has multiple threads of control then it can perform more than one task at a time.

Threads belonging to the same process

- share
 - code section
 - data section
 - other operating system resources
 - open files
 - signals
- have unique
 - registers
 - stack
 - PC
 - ID

Why Multithread?

Most modern applications are multithreaded

an application is typically implemented as a separate process with several threads of control

multiple tasks with the application can be implemented by separate threads for different reasons

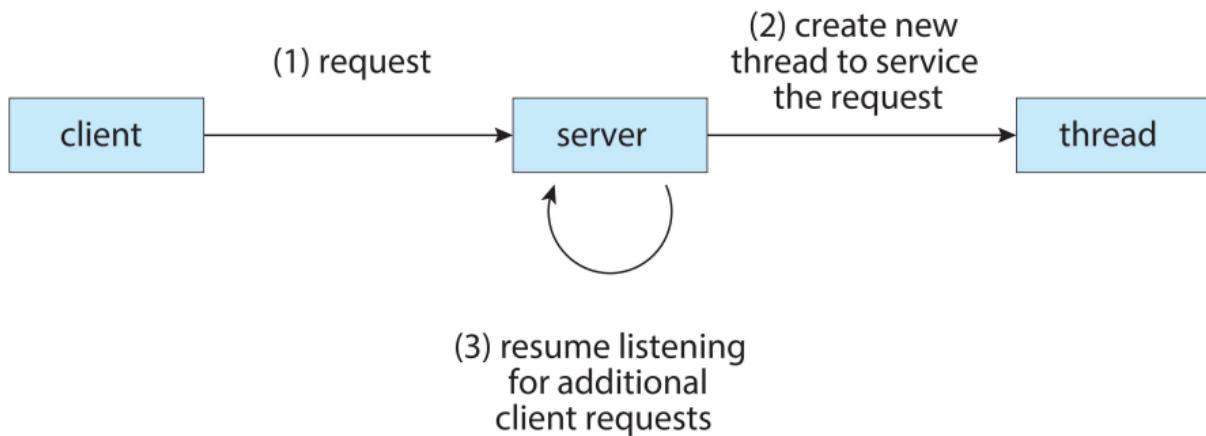
ex:

- making thumbnails of tons of different images at once
 - 1 thread per image

- loading a webpage
 - 1 thread for text
 - 1 thread for images
- word processors
 - 1 thread for responding to keystrokes
 - 1 thread for displaying graphics
 - 1 thread for spell checking

we feed each thread a different task and they don't interfere with one another

multithreading is also good for client/server applications



Multithreaded Server Architecture

Server will have

- have a port open and a thread to listen for connection requests
- make a thread to do what the request wants
- the open port
 - very vulnerable
 - an open door
 - necessary as otherwise no one can connect

Q: why multithread and not multi proces

A: It's much faster

- process creation is heavyweight while thread creation is lightweight
- creation, termination, and switching of threads is faster than with processes
- threads are ten times faster to create than processes in UNIX

kernels are generally multithreaded

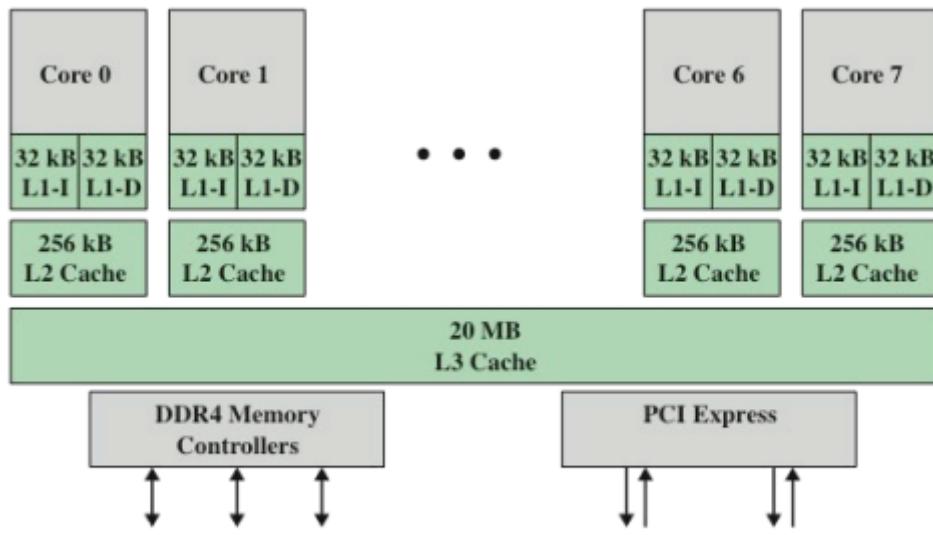
- during system boot time on linux system, several kernel threads are created
- each thread performs a specific task
 - device management
 - memory management
 - interrupt handling

- etc

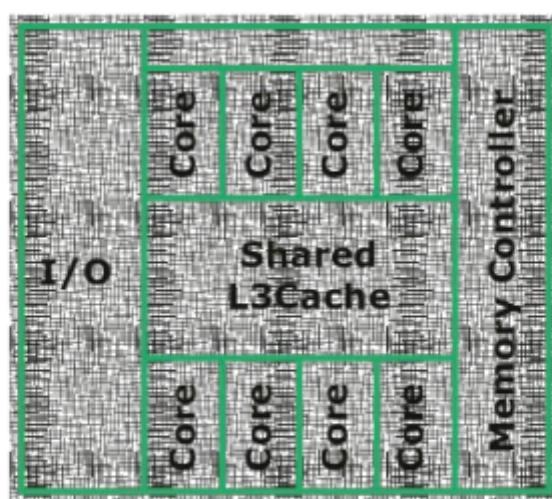
benefits of multithreading

- responsiveness
 - allow continued execution if part of the process is blocked
 - good for user interfaces
- resource sharing
 - threads share resources of process
 - easier than shared memory or message passing b/w processes
- economy
 - cheaper than process creation
 - thread switching lower overhead than context switching
- scalability
 - process can take advantage of multicore architectures
 - one process can use at most one core
 - multi-threads could use all available cores or CPUs
 - speed of execution

Multicore Programming



(a) Block diagram



(b) Physical layout on chip

Figure 1.20 Intel Core i7-5960X Block Diagram

multicore:

- multiple computing cores on a single chip
- each core appears as a separate CPU to the operating system
- each core can execute 1 trace at a time

word on caches, specifically in this intel architecture

L1 cache has all of the instructions and stuff.

All the highest value things

L2 can allow communication between 2 cores physically next to one another. They look separate in the diagram but that's just the allocation we do, they can look at each other's work and talk to one another basically.

L3 is on the chip and is shared by all of the cores

There are different kinds of architectures out there so the interesting L2 cache sharing won't be true of every architecture.

multithreaded programming

- provides a mechanism for more efficient use of these multiple computing cores and improved concurrency

we need to program properly in order to take advantage of the presence of multithreading

Concurrency vs Parallelism

Parallelism

- implies a system can perform more than one task simultaneously
- 2 cores doing 2 threads at the same time
- working in parallel

Concurrency

- supports more than one task by allowing all the tasks to make progress

Q: Is it possible to have concurrency without parallelism?

A: yes. utilize thread switching. parallelism requires parallel execution. We can have all the tasks progress together without parallel execution by way of thread switching.

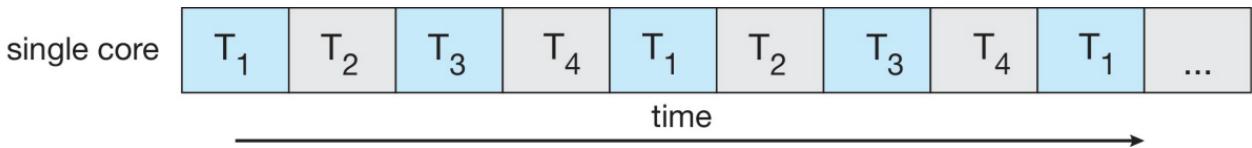
parallelism cannot happen on a single core

parallelism implies concurrency

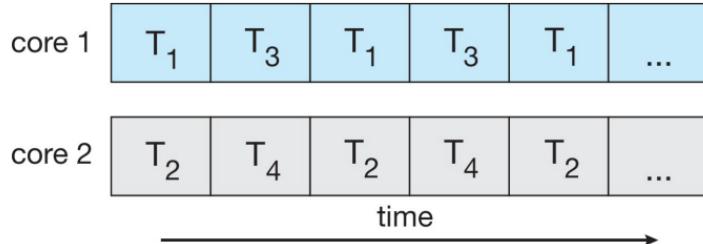
there can be concurrency on a single core

parallelism is faster than concurrency alone

- Concurrent execution on single-core system:



- Parallel execution (parallelism) on a multi-core system:



We are either parallelising data or tasks

data parallelism:

- distribute distinctive chunk of data across multiple cores, and perform the same operation on each
- ex
 - each core gets a different image
 - each core performs the same operation of creating a thumbnail

task parallelism

- distributing threads across cores
- each thread performs a unique operation
- not necessarily all working on the same data
 - 1 core is summing rows
 - 1 core is averaging rows
 - 1 core is finding the median
 - all may or may not be working on the same rows

There are a number of challenges that come with programming for multicore or multiprocessor systems

- dividing activities into separate concurrent tasks
 - what tasks are independent of each other
- balance
 - pick tasks that perform equal work of equal value
- data splitting
 - data accessed and manip'd by the tasks must be divided to run on separate cores
- data dependency
 - see what tasks require what data and see which pieces of data are shared by tasks
- testing and debugging
 - testing and debugging concurrent programs is inherently more difficult than testing and debugging single-threaded applications

Many developers argue that the advent of multicore systems will require an entirely new approach to designing software systems in the future.

Many educators believe that there should be an increased emphasis on parallel programming in software development.

Amdahl's Law

Identifies performance gains from adding additional cores to an application that has both serial and parallel components

$$\text{speedup} \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- S - serial portion
- N - processing cores

example:

- application is
 - 75% parallel
 - 25% serial
- 1 core → 2 cores

$$\begin{aligned}\text{speedup} &\leq \frac{1}{0.25 + \frac{(1-0.25)}{2}} \\ &= \frac{1}{0.25 + \frac{0.75}{2}} \\ &= \frac{1}{0.25 + 0.375} \\ &= \frac{1}{0.625} \\ &= 1.6\end{aligned}$$

- 1 core → 4 cores

$$\begin{aligned}\text{speedup} &\leq \frac{1}{1/4 + \frac{(1-1/4)}{4}} \\ &= \frac{1}{1/4 + \frac{3/4}{4}} \\ &= \frac{1}{1/4 + 3/16} \\ &= \frac{1}{7/16} \\ &= 16/7 \\ &\approx 2.28\end{aligned}$$

as N approaches infinity

- speedup approaches $1/S$
- 50% of application is serial, maximum speedup = 2 regardless of N

serial portion of an application has a disproportionate effect on performance gained by adding cores.

Knowledge Check:

- parallelism involves distributing tasks across multiple computing cores
 - use Amdahl's law to find the speedup gain for an application that is 60% parallel component for 2 processing cores
 - 1.43
 - it is possible to have concurrency without parallelism
-

Thread States and Operations

Key states for a thread

- running
- ready
- blocked

thread operations associated with a change in thread state are:

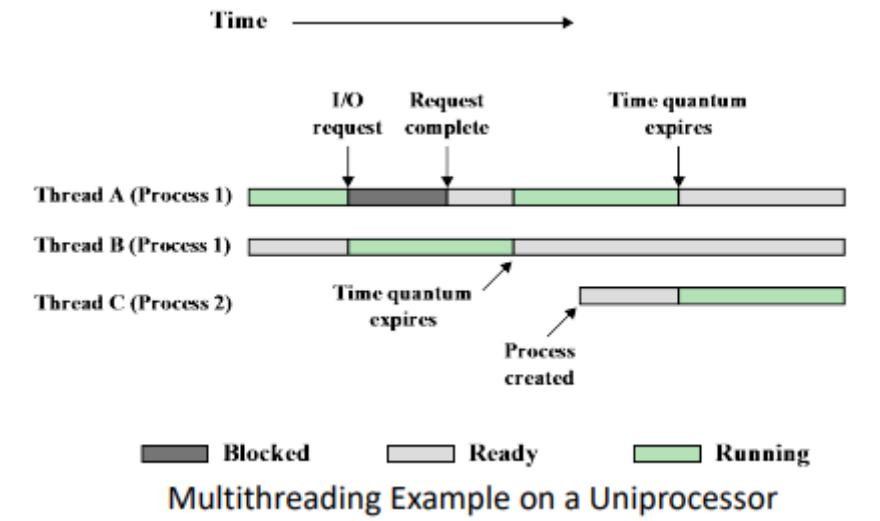
- spawn
 - new process is spawned, a thread for that process is also spawned
- block
 - a thread needs to wait for an event and is blocked
- unblock
 - the event the thread was waiting for has occurred, the thread is moved to the ready queue
- finish
 - thread completes. deallocate register, context, and stacks

Thread Synchronization

We need to sync up the activities on the threads because:

- all threads share the same resources
 - namely address space

- one thread altering the shared resource will affect the other threads in the same process



Multithreading Models

Types of threads:

- user threads
 - supported above the kernel
 - managed by user-level thread library w/o kernel support
- kernel threads
 - supported and managed directly by the kernel

virtually all contemporary operating systems support kernel threads

thread libraries contain code for:

- creating and destroying threads
- passing messages and data between threads
- scheduling thread execution
- saving and restoring thread contexts

3 primary thread libraries

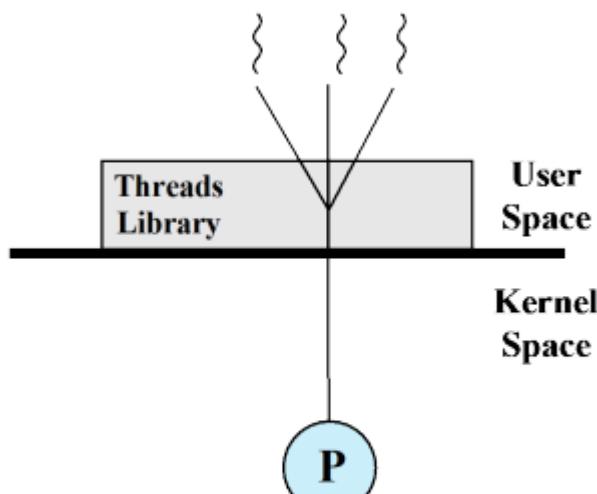
- POSIX Pthreads
- Windows threads
- Java threads

User Threads

User Threads

- many are mapped to a single kernel thread
- kernel is not aware of the existence of threads

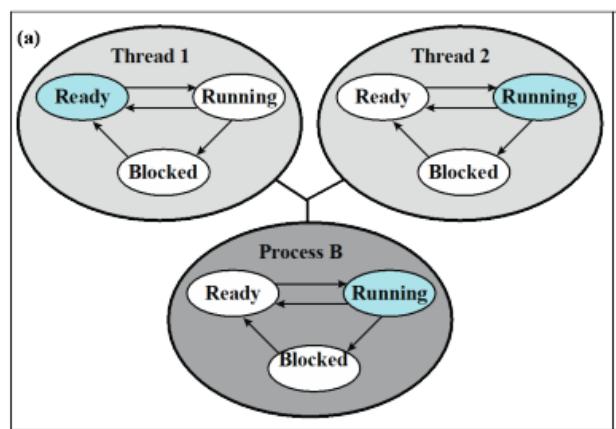
- thread management is done by threda library at the user level



examples

Example – User Threads /1

- Process B (pure user level threading) is executing in its thread 2; the states of the process and two User Threads that are part of the process are shown below.
- Colored state is current state.

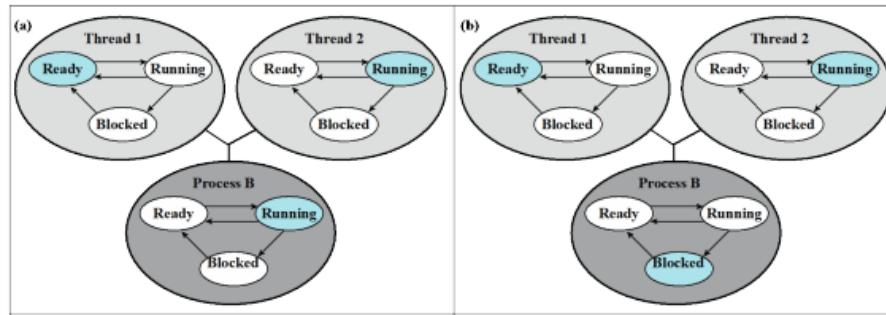


Examples of the Relationships Between User Thread States and Process States

Process B has 2 threads. OS only sees process B , not its threads

Example – User Threads /2

- **Scenario (b):** The application executing in **thread 2** makes a system call that **blocks Process B**. Control is transferred to the kernel.
- The kernel invokes the I/O action, places process B in the Blocked state, and switches to another process.



Examples of the Relationships Between User Thread States and Process States

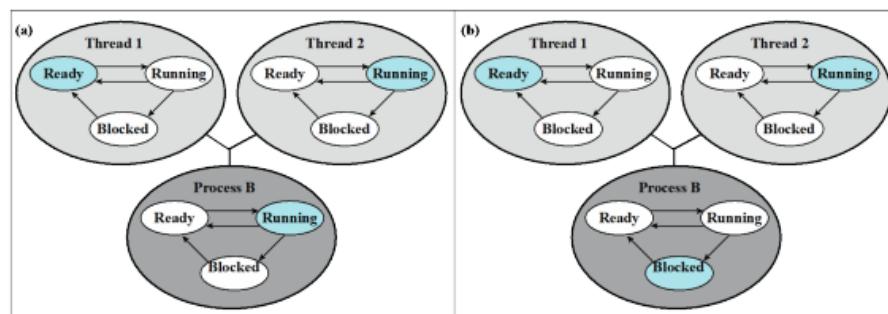
We're blocking the whole process.

The threads remain in their states but they aren't actually doing anything.

We can't just switch to thread 1 as the kernel isn't aware of the threads, only the process.

Example – User Threads /3

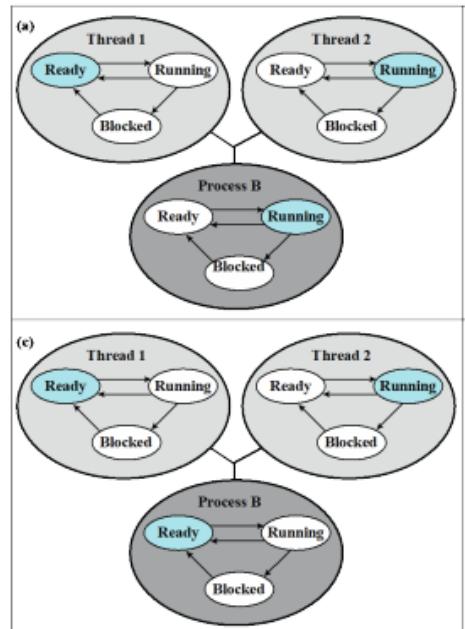
- **Scenario (b) – continued:** According to the threads library, thread 2 of process B is still in the Running state.
 - Thread 2 is not actually running (as in being executed on a processor); but it is perceived as being in the Running state by the threads library.



Examples of the Relationships Between User Thread States and Process States

Example – User Threads /4

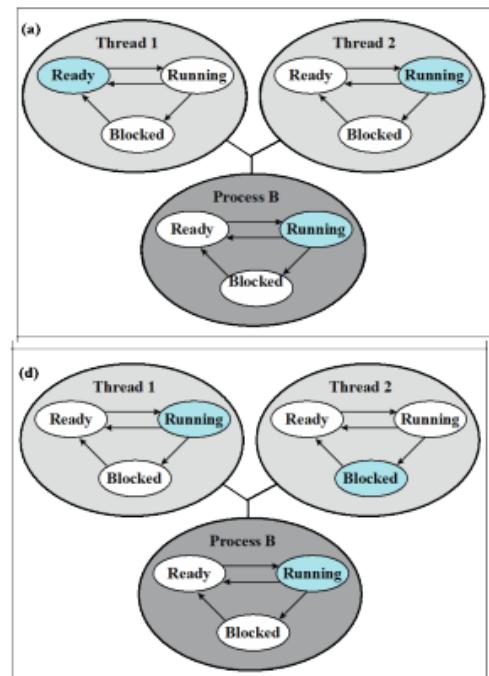
- **Scenario (c):** Process (B) has exhausted its time slice.
- The kernel places process B in the Ready state and switches to another process.
- The threads library sees thread 2 of process B in the Running state.



Examples of the Relationships Between User Thread States and Process States

Example – User Threads /5

- **Scenario (d):** Thread 2 has reached a point where it needs some action performed by thread 1 of process B.
- Thread 2 enters a Blocked state and thread 1 transitions from Ready to Running.
- The process itself remains in the Running state.



Examples of the Relationships Between User Thread States and Process States

pros:

- thread switching does not require kernel mode privileges
 - saves overhead of two mode switches
- scheduling can be application specific
 - scheduling algo can be tailored to the application
- ULTs can run on any OS

cons:

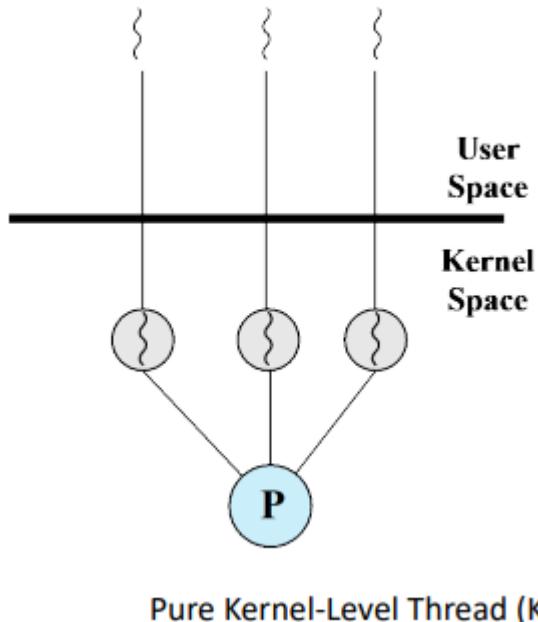
- entire process will block if a single thread makes a blocking system call
- pure ULT strat, multithreaded application cannot take advantage of multiprocessing/parallelism
- very few systems use this model because of its inability to utilize multiple cores
 - solaris green threads

Kernel Threads

thread management is all done by the kernel

creating a user thread creates a kernel thread (one-to-one)

kernel maintains context information for the process as a whole and for individual threads within the process



pros:

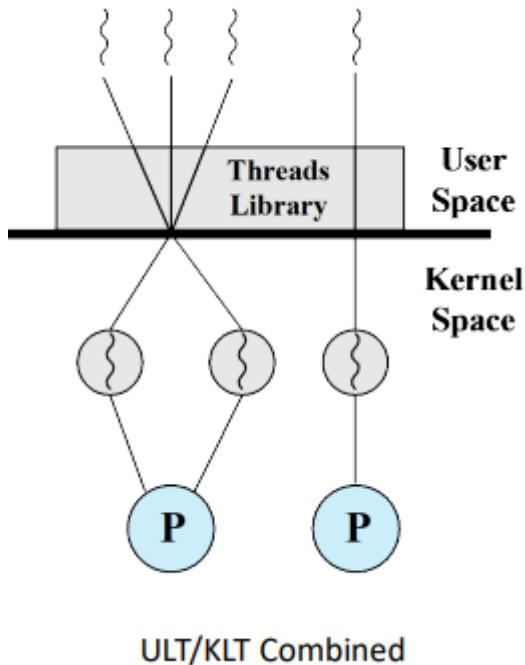
- kernel can simultaneously schedule multiple threads from the same process on multiple processes
- if one thread in a process is blocked, the kernel can schedule another thread of the same process
- linux and windows implement this

cons:

- 1-1 correspondence b/w user and kernel threads will burden the performance of a system after creating a large number of threads
 - so many threads to have to switch between

Combined User and Kernel level threading

- thread creation done in user space
- most of scheduling and synchronization of threads within an application
- multiple user-level threads from a single application are mapped onto smaller or equal number of kernel threads
 - many-to-many
- number of kernel threads may be specific to either a particular application or a particular machine



ULT/KLT Combined

Note:

- we can't map 3 threads to 4 cores.
- a kernel thread can only belong to one processor
- a user thread can only map to one kernel thread

pros:

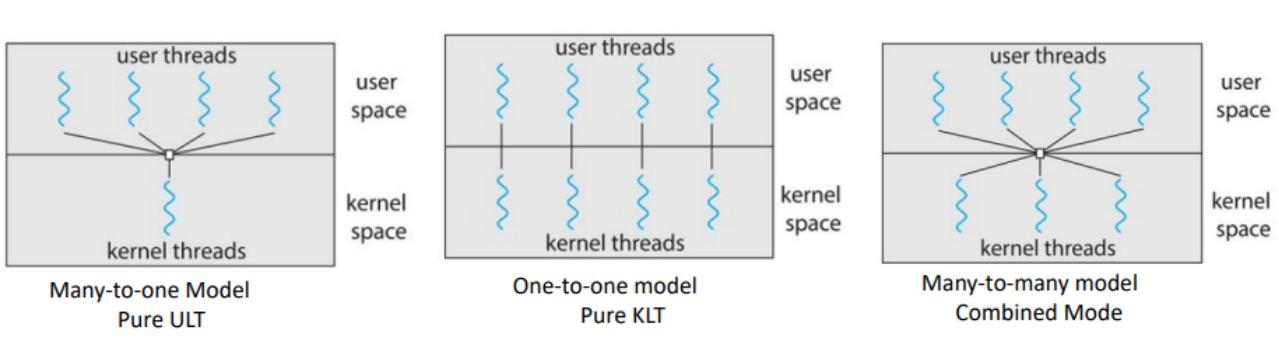
- appears flexible
- number of threads and parallelism
 - many-to-one (pure ULT)
 - create many threads but no parallelism
 - one-to-one (pure KLT)
 - better concurrency but high overhead
 - many-to-many (combined approach)
 - as many threads as needed running parallel on a multiprocessor
- blocking
 - when a thread performs a blocking system call, the kernel can schedule another thread for execution

cons:

- difficult to implement
 - windows with the `ThreadFiber` package
 - not very common outside of the above

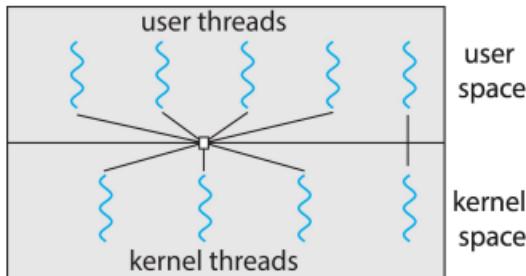
Whether KLT or ULT remember that

a relationship must exist between user threads and kernel threads no matter the cardinality



Two-Level Model

- Similar to Many-to-Many, except that it allows a user thread to be **bound** to kernel thread.



2-level model is like many-to-many but we're able to bind a user thread to a kernel thread.

we can basically dedicate a core to a thread.

Q: can we interleave threads from different processes?

A: This is just normal process switching.

knowledge check

- The _____ model allows a user-level thread to be bound to one kernel thread.
 - a) many-to-many
 - b) two-level
 - c) one-to-one
 - d) many-to-one
-

Question?

- Why must a user thread always be mapped to a specific kernel thread?
 - Users' programs make their own "threads" and simulate context-switches to switch between them. However, these threads aren't kernel threads. Each user thread can't actually run on its own, and the only way for a user thread to run is if a kernel thread is actually told to execute the code contained in a user thread.
 - In other words, user threads have to be associated with kernel threads because a user thread is just a bunch of data in a user program. Kernel threads are the real threads in the system that can be scheduled on the CPU, so for a user thread to progress; the user program has to have its scheduler take a user thread and then run it on a kernel thread.

prof:

"the threads need to have a connection to the processor in order for execution"

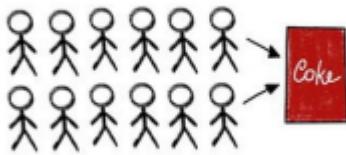
Concurrency: Mutual Exclusion and Synchronization

Background for Concurrency

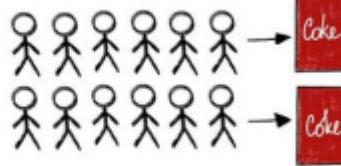
processes can execute concurrently or in parallel

may be interrupted at any time, partially completing execution

Concurrent = 2 queues → 1 coke



Parallel = 2 queues → 2 coke



we have a shared memory where we store variables and everything we want to work with.

If we don't manage that shared memory properly we end up with data inconsistency

- Maintaining data consistency requires mechanisms to ensure the orderly execution of **cooperating** processes that share a logical address space.
 - A process that can affect or be affected by other processes executing in the system.

make sure the processes execute in an orderly fashion and don't fuck with each other when they share a resources

Race Condition

- several processes access and manip the same data at the same time
- outcome depends on who got there first

Thread 1	Thread 2	Integer value
		0
read value	←	0
increase value		0
write back	→	1
	read value	1
	increase value	1
	write back	2

Thread 1	Thread 2	Integer value
		0
read value	←	0
	read value	0
increase value		0
	increase value	0
write back	→	1
	write back	1

Critical-Section Problem

critical segment of code:

- process may be accessing and updating data that is shared with at least one other process
- only one process in critical section at a time

when a process wants to enter the critical section

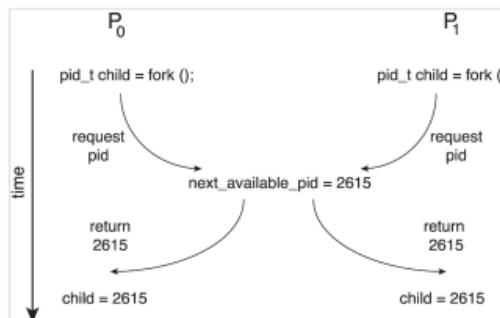
- entry section - process asks permission to enter
- it does the critical section
- exit section - let's everyone know it's done
- remainder section is the rest of the noncritical stuff

we need to ensure the 3 following conditions

1. mutual exclusion
 - only one process in critical section at a time
2. progress
 - if no one else is in critical and someone else wants to go critical then we have to select the next process
 - we can't postpone this selection indefinitely
3. bounded waiting
 - there must be a limit on the number of times that other processes can enter the critical sections after a process has made a request to enter
 - assume same speed of all processes

Kernel-Mode Race Condition Example

- Processes P_0 and P_1 are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (`pid`)



- Unless there is mutual exclusion, the same pid could be assigned to two different processes!

Single Core

very shrimples

we could prevent interrupts from occurring while a shared variable was being modified

useless when we try to move to multiprocessor since it would just start disabling everything.

OS

2 approaches depending on preemption

- preemptive
 - allows preemption of process when running in kernel mode
 - processor can tell the process to pack all their bags and get ready to move
- non-preemptive
 - runs until kernel mode, blocks, or voluntarily yields CPU
 - essentially free of race conditions as only one process is active in the kernel at a time

preemptive are difficult to design.

it allows two kernel-mode processes to run simultaneously on diff cores

it also makes things more responsive

knowledge check

- A race condition ____.
 - a) results when several threads try to access the same data concurrently
 - b) results when several threads try to access and modify the same data concurrently
 - c) will result only if the outcome of execution does not depend on the order in which instructions are executed
 - d) None of the above
- Instructions from different processes can be interleaved when interrupts are allowed.
 - True
 - False
- A(n) _____ refers to where a process is accessing/updating shared data.
 - a) critical section
 - b) entry section
 - c) mutex
 - d) test-and-set

- A non-preemptive kernel is safe from race conditions on kernel data structures.
 - True
 - False
 - A solution to the critical section problem does not have to satisfy which of the following requirements?
 - a) mutual exclusion
 - b) progress
 - c) atomicity
 - d) bounded waiting
-

Atomic Operation

- A function or action implemented as a sequence of one or more instructions that appears to be **indivisible**:
 - No other process can see an intermediate state or interrupt the operation.
 - The sequence of instruction is guaranteed to execute as a group, or not execute at all.



atomic operations

Mutual Exclusion

Software Approaches

Dekker's Algo

Dijkstra went to go publish it for him so that's why there's a jpeg of him on the slide

first known sol'n to the mutual exclusion problem in concurrent programming

allows 2 threads to share a single-use resource without conflict, using only shared memory for communication

only one access to a memory location can be made at a time

we have a global memory location labelled `turn` that is shared between the two processes

Processes check to see if it's their turn and execute accordingly. Once they're done they give the turn to the other process.

Mutual exclusion is guaranteed.

What is the problem?

prof left off here for 05/02/24 lect

midterm is all question formats and on paper

Processes must strictly alternate in their use of their critical section

- the pace of execution is dictated by the slower of the two processes

One process being slow makes the whole thing slower since everyone has to wait on that slower process.

If one process fails then the other process is permanently blocked

- this happens because the failed process can't give the turn over to the other process by exiting and setting the flag
- this happens whether or not the process is in the critical section or not

2nd attempt - a revision:

- each process has its own key (**flag**)
- if one fails, the other can still access its critical section
- processes can look at others flags
- one process can't change the flag of another process
- processes share a boolean vector that holds the flags
 - Boolean vector `flag[2]`

```

/* PROCESS 0 *          /* PROCESS 1 *
.
.
while (flag[1])        while (flag[0])
    /* do nothing */;   /* do nothing */;
flag[0] = true;         flag[1] = true;
/*critical section*/;  /* critical section*/;
flag[0] = false;       flag[1] = false;
.

```

drawbacks:

- permanent blocking still possible
 - process fails inside critical section or after setting flag to true w/o entering
- no guaranteed mutual exclusion
 - P0 executes the while statement and finds flag[1] set to false
 - P1 executes the while statement and finds flag[0] set to false
 - P0 sets flag[0] to true and enters its critical section
 - P1 sets flag[1] to true and enters its critical section
 - now both processes are now executing their critical sections so the program is incorrect

3rd attempt - another revision

- the problem from before was that a process can change its state after the other process has checked it but before the other process can enter its critical section
- swap around the two statements to guarantee mutual exclusion

```

/* PROCESS 0 *          /* PROCESS 1 *
.
.
flag[0] = true;         flag[1] = true;
while (flag[1])        while (flag[0])
    /* do nothing */;   /* do nothing */;
/* critical section*/;  /* critical section*/;
flag[0] = false;        flag[1] = false;
.

```

we set the flag's value before we execute the while loop

drawbacks:

- this results in a deadlock

- both processes set their flags to true
- later on they will check each other's flags and decide they need to wait for the other
 - the flag value makes them think the other process is in their critical section

4th attempt - yet another revision:

- Each process sets its flag to indicate its desire to enter its critical section but is prepared to reset the flag to defer to the other process
- mutual exclusion is guaranteed under this model

```

/* PROCESS 0 *          /* PROCESS 1 *
.
.
flag[0] = true;        flag[1] = true;
while (flag[1]) {      while (flag[0]) {
    flag[0] = false;   flag[1] = false;
    /*delay */;       /*delay */;
    flag[0] = true;   flag[1] = true;
}
/*critical section*/;  /* critical section*/;
flag[0] = false;       flag[1] = false;
.
.
```

drawbacks:

- ~~flag[0] = true;~~ ~~flag[1] = true;~~
 - each process sets their flag to true
- **`while (flag[1]) {`** **`while (flag[0]) {`**
 - each process checks the other's flag before proceeding with the while loop
- **`while (flag[1]) {`** **`while (flag[0]) {`**
 - each flag flips from false and back to true
- the loop ends up repeating in what's known as a *livelock*
- this won't sustain for a long time
 - the processes are different and we can presume that the while loops won't always take the same amount of time and that one will finish their while loop before the other can so they'll enter the critical section first

5th iteration - the correct solution:

- we introduce ***the right to insist***

```

boolean flag [2];
int turn;
void P0()
{
    while (true) {
        flag [0] = true;
        while (flag [1]) {
            if (turn == 1) {
                flag [0] = false;
                while (turn == 1) /* do nothing
*/;
                    flag [0] = true;
            }
        }
        /* critical section */;
        turn = 1;
        flag [0] = false;
        /* remainder */;
    }
}
void P1( )
{
    while (true) {
        flag [1] = true;
        while (flag [0]) {
            if (turn == 0) {
                flag [1] = false;
                while (turn == 0) /* do nothing
*/;
                    flag [1] = true;
            }
        }
        /* critical section */;
        turn = 0;
        flag [1] = false;
        /* remainder */;
    }
}
void main ()
{
    flag [0] = false;
    flag [1] = false;
    turn = 1;
    parbegin (P0, P1);
}

```

- when P0 wants to enter the critical section
 - it sets its flag to true
 - checks P1's flag
 - if P1 flag is false then P0 gets to go critical
 - else, P0 consults turn
 - turn==0
 - it's P0's turn to insist
 - P0 will just keep on checking P1's flag in order to try and enter its critical section
- P1 will at some point note that it is its turn to defer and set its flag to false, allowing P0 to proceed
- after P0 has used its critical section
 - P0 flag set to false
 - turn set to 1 to transfer the right to P1

Peterson's Solution

Dekker's algo solves things but it's hard to follow

Peterson has a simpler sol'n that allows two processes to share a single-use resource w/o conflict

only shared memory is used for communication

it is kind of rewriting Dekker's algorithm

```
bool flag[2] = {false, false};  
int turn;
```

```
P0:      flag[0] = true;  
P0_gate: turn = 1;  
        while (flag[1] && turn == 1)  
        {  
            // busy wait  
        }  
        // critical section  
        ...  
        // end of critical section  
        flag[0] = false;
```

```
P1:      flag[1] = true;  
P1_gate: turn = 0;  
        while (flag[0] && turn == 0)  
        {  
            // busy wait  
        }  
        // critical section  
        ...  
        // end of critical section  
        flag[1] = false;
```

`int turn` - indicates whose turn it is to enter the critical section

`bool flag[2]` - indicate if a process is ready to enter the critical section

`flag[i]==true` implies that process P_i is ready

```

while (true) {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;
        /* critical section */
    flag[i] = false;
    /* remainder section */
}

```

3 critical-section requirements are met(provably):

1. mutual exclusion is preserved

- P_i enters critical section only if: either `flag[j]==false` or `turn == i`
 - the other process is not ready or
 - it is P_i 's turn

2. Progress requirements is satisfied

3. Bounded-waiting requirement is met

Useful for demonstrating an algo but not guaranteed to work on modern architectures.

to improve performance, processors and/or compilers may reorder operations that have no dependencies

- fine for single-threaded as result is always the same
- may produce inconsistent or unexpected results for multithreading

two threads share `_flag_` and `_x_` variables

```

boolean flag = false;
int x = 0;

```

Thread 1 performs

```

while (!flag)
    print x;

```

Thread 2 performs

```
x = 100;
flag = true;
```

what is the expected output

100 is the expected output however the operations for thread 2 may be reordered so the output might be 0

Original

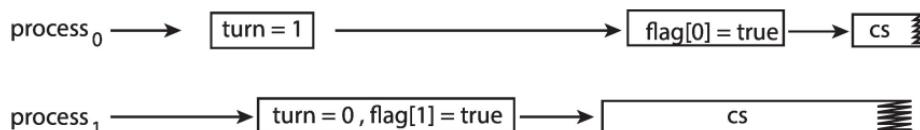
```
while (true){
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;
    /* critical section */

    flag[i] = false;
    /* remainder section */
}
```

Reordered

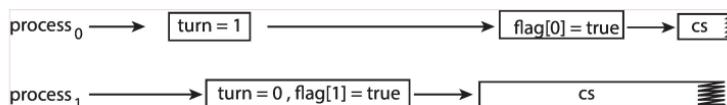
```
while (true){
    turn = j;
    flag[i] = true;
    while (flag[j] && turn == j)
        ;
    /* critical section */

    flag[i] = false;
    /* remainder section */
}
```



- Process 0: turn = 1;
- Process 1: turn = 0;
- Process 1: flag[1] = true;
- Process 1: while(flag[0] && turn==0) // terminates, since flag[0]==false
- Process 1: enter critical section 1
- Process 0: flag[0] = true;
- Process 0: while(flag[1] && turn==1) // terminates, since turn==0
- Process 0: enter critical section 0

```
while (true){
    turn = j;
    flag[i] = true;
    while (flag[j] && turn == j)
        ;
    /* critical section */
    flag[i] = false;
    /* remainder section */
}
```



We have to use a **Memory Barrier** to make sure that Peterson's solution will work correctly on modern computer architecture.

So this software based support will need the aid of some hardware.

Knowledge check:

- In Peterson's solution, the ___ variable indicates if a process is ready to enter its critical section
 - flag[i]

Hardware Support Approaches

“There is no bullet solution... all solutions build on each other”

Many systems provide hardware support for implementing the critical section code

uniprocessors - could disable interrupts

- currently running code would execute w/o preemption
- generally, too inefficient on multiprocessor systems
 - OS using this are not broadly scalable

looking at 3 forms of hardware support

1. memory barriers (aka memory fences)
2. hardware instructions
3. atomic variables

Memory Barrier

How a computer architecture determines what memory guarantees it will provide to an application program is known as its memory model.

models may be either

- strongly ordered
 - memory modification of one processor is immediately visible to all other processors
- weakly ordered
 - not always immediately visible to all other processors

developers cannot make an assumption about whether or not the memory model is strongly or weakly ordered.

in lieu of this guarantee, we use memory barriers

memory barriers

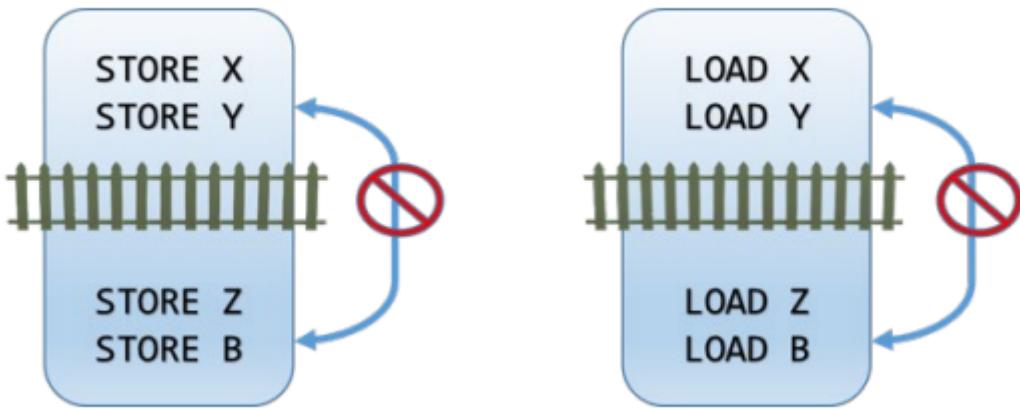
- computer instructions that force any changes in memory to be propagated to all other processors in the system

`Sfence` is a memory barrier instruction for example with regards to storing

`Lfence` is another type of instruction with regards to loading

`Mfence` prevents reordering of reading and writing before or after the fence

do not reorder anything before the barrier with anything that comes after the barrier



5/07/10/memory-harriers-in-dot-net.aspx/

we're not able to move the loads to either side of the fence

- We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100:
- **Thread 1** now performs


```
while (!flag)
    memory_barrier();
print x
```
- **Thread 2** now performs


```
x = 100;
memory_barrier();
flag = true
```
- For Thread 1 we are guaranteed that the value of `flag` is loaded before the value of `x`.
- For Thread 2 we ensure that the assignment to `x` occurs before the assignment `flag`.

Hardware Instructions

```
boolean test_and_set (boolean *target){
    boolean rv = *target;
    *target = true;
    return rv;
}
```

`test_and_set()`

- executed atomically
- read a value in memory
- set it to true
- return the value that was in memory

This let's us make a lock out of the flag we want us to make sure we don't enter the critical section while another program is in its critical section.

```
do {
    while (test_and_set(&lock));
    /*critical section*/
    lock = false;
    /*remainder section*/
} while (true);
```

This is supported by nearly every architecture.

`compare_and_swap()`

```
int compare_and_swap(int *value, int expected, int new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

- executed atomically
- returns the original value of `value`
- set the `value` to the value of `new_value`
 - but only if `*value == expected`

We can also use this to make a shared lock for mutual exclusion. Have it initialized to 0.

```
while (true) {
    //lock is the lock, 0 is the expected value, 1 is the prospective new value
    while (compared_and_swap(&lock, 0, 1) != 0) {/*do nothing, someone else is in
critical section*/}
    /*critical section, lock has been set to by us*/
    lock = 0;      //we left the critical section so we're able to set the lock back to
0
    /*remainder section*/
}
```

Atomic Variables

Typically instructions such as compare-and-swap are used as building blocks for other synchronization tools.

One tool is an atomic variable that provides *atomic* (uninterruptible) updates on basic data types such as integers and booleans.

Ex:

- `sequence` - our atomic variable
- `increment()` - an operation on `sequence`
- command
 - `increment(&sequence)`
 - ensures the sequence is incremented without interruption

we make atomic updates so that we're guaranteed to update a variable.

`increment()` can be implemented like so:

```
void increment (atomic_int *v) {  
    int temp;  
    do {  
        temp = *v;  
    }  
    while (temp != compare_and_swap(v, temp, temp+1));  
    // the above keeps looping until we are able to successfully swap, effectively  
    incrementing the value of `v'  
}
```

Knowledge Check:

- race conditions are prevented by requiring that critical regions be protected by locks
 - true
- `test_and_set()` instruction is executed atomically
 - true
- an instruction that executes atomically
 - executes as a single, uninterruptible unit

OS & Programming Languages Approaches

Mutex Locks

Previous solutions are complicated and generally inaccessible to application programmers.

OS designers build software tools to solve critical section problem.

The simplest solution is the mutex lock.

fun fact: mutex is a portmanteau of `mutual exclusion`, as it was created to solve that problem

Common Concurrency Mechanisms

Semaphore	An integer value used for signaling among processes. Only three operations may be performed on a semaphore, all of which are atomic: initialize, decrement, and increment. The decrement operation may result in the blocking of a process, and the increment operation may result in the unblocking of a process. Also known as a counting semaphore or a general semaphore
Binary Semaphore	A semaphore that takes on only the values 0 and 1.
Mutex	Similar to a binary semaphore. A key difference between the two is that the process that locks the mutex (sets the value to zero) must be the one to unlock it (sets the value to 1).
Condition Variable	A data type that is used to block a process or thread until a particular condition is true.
Monitor	A programming language construct that encapsulates variables, access procedures and initialization code within an abstract data type. The monitor's variable may only be accessed via its access procedures and only one process may be actively accessing the monitor at any one time. The access procedures are <i>critical sections</i> . A monitor may have a queue of processes that are waiting to access it.
Event Flags	A memory word used as a synchronization mechanism. Application code may associate a different event with each bit in a flag. A thread can wait for either a single event or a combination of events by checking one or multiple bits in the corresponding flag. The thread is blocked until all of the required bits are set (AND) or until at least one of the bits is set (OR).
Mailboxes/Messages	A means for two processes to exchange information and that may be used for synchronization.
Spinlocks	Mutual exclusion mechanism in which a process executes in an infinite loop waiting for the value of a lock variable to indicate availability.

note that the mutex is an integer that alternates between 0 for locked and 1 to unlocked, this is effectively a boolean.

Protect a critical section by first `acquire()` a lock then `release()` the lock. Boolean variable indicating if `lock` is available or not.

`acquire()` and `release()` must be atomic. Usually implemented via hardware atomic instructions such as compare-and-swap.

```

while (true) {
    acquire lock

    critical section

    release lock

    remainder section
}

```

You `acquire()` the lock and no one else can touch it while you're in the critical section.

Afterwards you can `release()` since you're done with the critical section.

The instructions look something like this in implementation.

```
acquire() {  
    while (!available) { /* busy wait */}  
    available = false;  
}  
release () {  
    available = true;  
}
```

- these two functions must be implemented atomically, typically through the use of something like `test_and_set()` and `compare_and_swap()`

cons:

- sol'n requires busy waiting
 - a process trying to enter its critical section will loop call `acquire()` over and over again to try and grab hold of it
- this lock is called a *spinlock*
 - if a lock is to be held for a `short duration`, one thread can “spin” on one processing core while another thread performs its critical section on another core.

short duration:

- how short is short?
- waiting on a lock requires 2 context switches
 - move the thread to the waiting state
 - storing away all of the process info (pcb, psw, etc.) introduces overhead
 - restore the waiting thread once the lock becomes available
- the general rule is to use a spinlock if the lock will be held for a duration of less than two context switches
- we don't know how long it will take for a critical section to take but we can get some good estimates

Knowledge check

- a mutex lock
 - is essentially a boolean variable
- what is the correct order of operations for protecting a critical section using mutex locks?
 - `acquire()` followed by `release()`
- busy waiting refers to the phenomenon that while a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the call to acquire the mutex lock
 - yes

Prof very briefly touched upon a notion of 2 different kinds of implementation of semaphore and mutex.

Something about a classic implementation versus some other implementation.

Semaphores

One problem we had with the mutex was the busy waiting.

basic idea:

- ≥ 2 processes can work together using simple signals
- a process can be forced to stop at a specified place until it has received a specific signal
- any complex coordination can be done with an appropriate amount of signals
- our previously mentioned mutex is akin to a single flag that we can raise up and down.

Sempahores are a more sophisticated synchronization tool than mutex locks.

A semaphore `s` is an integer variable, that can only be accessed through 2 atomic operations:

1. `wait()` - decrements the semaphore value
2. `signal()` - increments the semaphore value

these each have other names.

`wait()` originally P() from Dutch proberen, "to test", `semWait()`

`signal()` originally V() from verhogen, "to increment", `semSignal()`

we may find these in other literature

Implementation

```
wait(s) {  
    while (s <= 0) /* busy wait */  
        s--;  
}  
signal(s) {  
    s++;  
}
```

There are 2 different kinds of semaphores:

1. counting sempahore
2. binary semaphore

Counting semaphores (aka general semaphore) is an integer value that can range over an unrestricted domain

Binary semaphores are an integer value that can only range between 0 and 1. This is a mutex lock.

We can implement a counting semaphore as a binary semaphore.

Good tool to solve synch problems.

- Solution to the CS Problem

- Create a semaphore “**mutex**” initialized to 1
- ```
wait(mutex);
CS
signal(mutex);
```

- Consider  $P_1$  and  $P_2$  processes that require  $S_1$  to happen before  $S_2$

Create a semaphore “**synch**” initialized to 0

**P1 :**

```
S1;
signal(synch);
```

**P2 :**

```
wait(synch);
S2;
```

Initializing the semaphore to 0 will allow us to only allow  $P_2$  to execute  $S_2$  after the execution of  $S_1$ .

A semaphore must guarantee that no 2 processes can execute the `wait()` and `signal()` on the same semaphore at the same time

Implementation becomes the critical section problem where the `wait` and `signal` code are placed in critical section.

Could now have busy waiting in critical section implementation.

Process can just suspend instead of waiting for the semaphore. Each semaphore has a waiting queue w/ 2 items

- value (of type integer)
- pointer to next record in the list

2 operations:

- block (aka sleep) - place the process invoking the operation on the appropriate waiting queue
- wakeup - remove one of the processes from the waiting queue and place in ready queue

waiting queue:

```
typedef struct {
 int value;
 struct process *list;
} semaphore
```

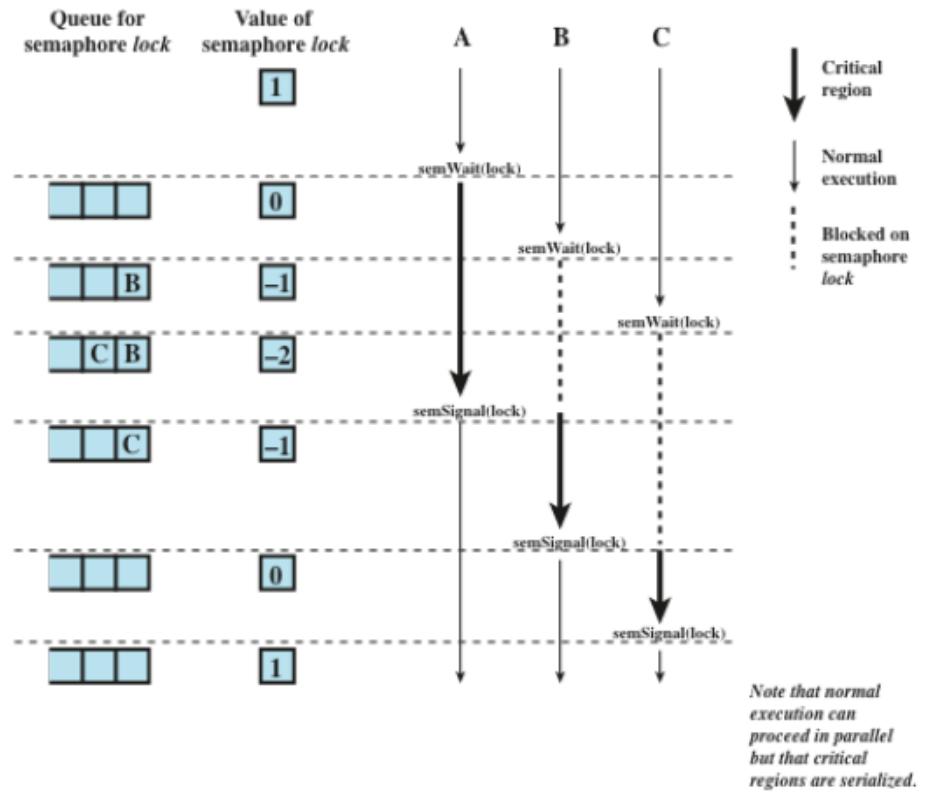
No busy waiting implementation:

```
wait (semaphore *S) {
 S->value--;

 if (S->value < 0) {
 add this process to S->list;
 block;
 }
}

signal(semaphore *S) {
 S->value++;
 if (S->value <= 0) {
 remove a process P from S->list;
 wakeup(P);
 }
}
```

# Semaphores Usage Example



Processes Accessing Shared Resource Protected by a Semaphore

We're using the incrementing and decrementing to iterate through the queue.

A strong semaphore uses a queue of blocked processes, the process unblocked is the one in the queue for the longest time.

A weak semaphore uses a set of blocked processes. The processes unblocked is unpredictable. This can lead to process starvation as a process may remain blocked forever

# Problems with Semaphores

- Incorrect use of semaphore operations:

```
 signal(mutex) ;
 ...
 critical section
 ...
 wait(mutex) ;
or
 wait(mutex) ;
 ...
 critical section
 ...
 wait(mutex) ;
```

- Or omitting of `wait(mutex)` and/or `signal(mutex)`
- In these cases, either mutual exclusion is violated, or the process will permanently block.

---

knowledge check

- a counting semaphore
  - is essentially an integer variable
- \_\_\_ can be used to prevent busy waiting when implementing a semaphore
  - waiting queues
- mutex locks and binary semaphores are essentially the same thing
  - true

## Monitors

# Deadlock and Starvation

## Principles of Deadlocks

Systems have different resources

Threads must:

- request a resource before using it
- release the resource after using it

The number of resources requested cannot exceed the total number of resources available in the system

- a thread cannot request 2 net interface if there is only 1
- can't ask for what we don't have

Request:

- Thread requests the resource
- if you can give it, then give it immediately
- If it can't be given right now then the requesting thread must wait until it can acquire the resource

Use:

- the thread can operate on the resource
- ex: use the mutex lock to access a process' critical section

Release:

- resource is released and made usable again

Reusable Resources:

- can be used safely by one process at a time
- not depleted by that use
- not consumable
- ex: processors, I/O channels, memory (main and secondary), I/O devices, data structures (files, databases, and semaphores)

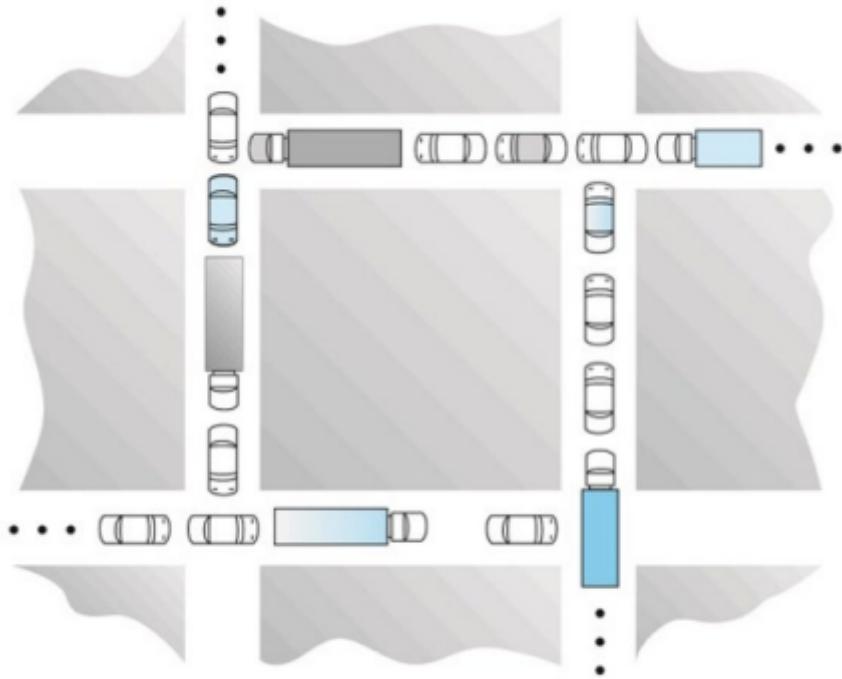
Consumable Resources:

- created/produced and destroyed/consumed
- ceases to exist after being acquired by the consuming process
- ex: interrupts, signals, messages, info in I/O buffers

## Deadlock

A deadlock is the **permanent** block.

Two or more processes are waiting for the other process to release a shared resource.



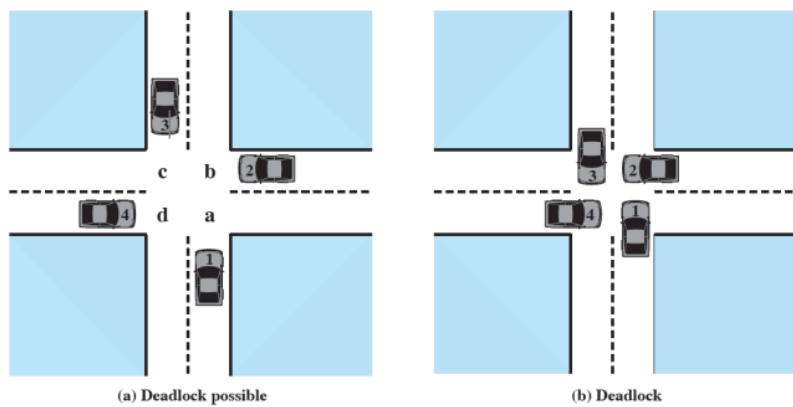
## Traffic Deadlock

A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set.

The block is permanent because none of the events ever get triggered.

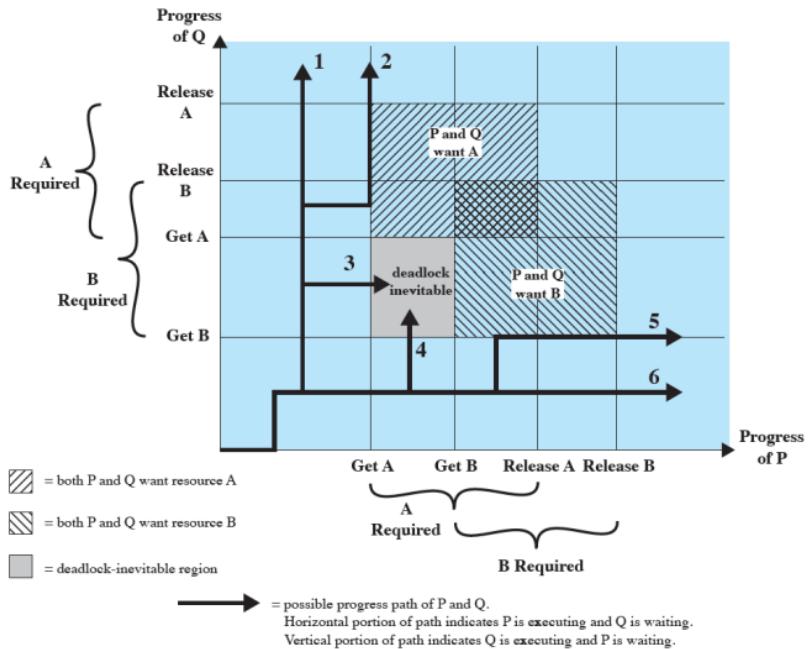
## Illustration of Deadlock

- All deadlocks involve conflicting needs for resources by two or more processes.



# Joint Progress Diagram – Deadlock

| Process P | Process Q |
|-----------|-----------|
| •••       | •••       |
| Get A     | Get B     |
| •••       | •••       |
| Get B     | Get A     |
| •••       | •••       |
| Release A | Release B |
| •••       | •••       |
| Release B | Release A |
| •••       | •••       |



Trace 1 sees process  $Q$  execute all of its desire operations before process  $P$  can perform any single operation

Trace 2 sees process  $Q$  get A then make process  $P$  wait until it was done.

Trace 3 sees  $Q$  get B and  $P$  get A. 4 does the same but in reverse order

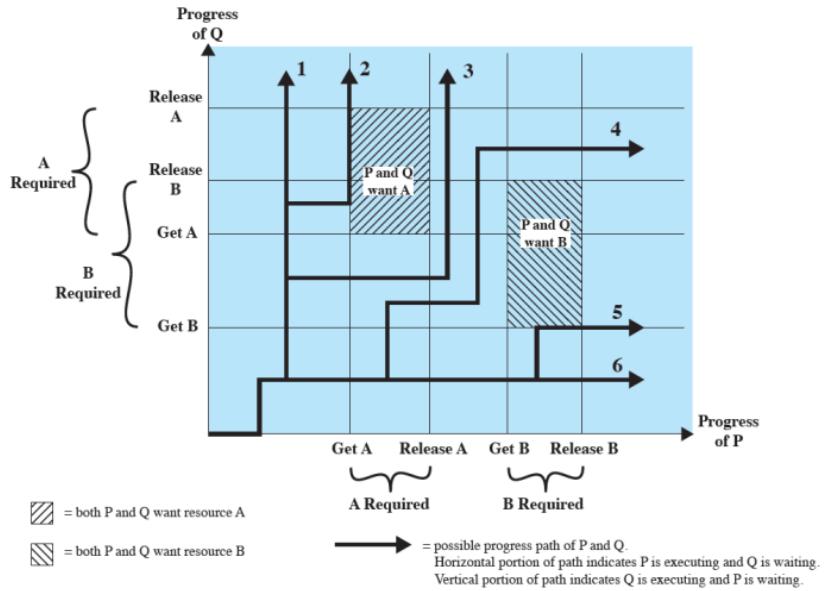
Traces 3 & 4 eventually leads to a deadlock as each process will want the other's acquired resource that they can't release until they get the resource they want.

Trace 5 sees process  $P$  get B then make process  $Q$  wait until it's done.

Trace 6 sees  $P$  execute everything before  $Q$  gets make some meaningful progress.

# Joint Progress Diagram – No Deadlock

| Process P | Process Q |
|-----------|-----------|
| •••       | •••       |
| Get A     | Get B     |
| •••       | •••       |
| Release A | Get A     |
| •••       | •••       |
| Get B     | Release B |
| •••       | •••       |
| Release B | Release A |
| •••       | •••       |



Joint Progress Diagram

“I would like to leave that for the banking algorithm” - Prof, no one knows what she meant by that

This rearrangement is not always a possible solution since we don't know what the order of the instructions/operations will be beforehand.

Traces 1, 2, 5, and 6 are unaffected but now traces 3 and 4 don't lead to deadlock.

instead with Trace 3

- $Q$  gets B
- $P$  gets A then releases A
- $Q$  gets A, releases B, then releases A
- $P$  gets B then releases B

Similarly with trace 4

- 2 processes that compete for exclusive access to a disk file **D** and a tape drive **T**.
- **Questions:** Are these reusable or consumable resources? Will there be a deadlock? What sequence of execution will result in a deadlock?

| Step           | Process P Action | Step           | Process Q Action |
|----------------|------------------|----------------|------------------|
| P <sub>0</sub> | Request (D)      | q <sub>0</sub> | Request (T)      |
| P <sub>1</sub> | Lock (D)         | q <sub>1</sub> | Lock (T)         |
| P <sub>2</sub> | Request (T)      | q <sub>2</sub> | Request (D)      |
| P <sub>3</sub> | Lock (T)         | q <sub>3</sub> | Lock (D)         |
| P <sub>4</sub> | Perform function | q <sub>4</sub> | Perform function |
| P <sub>5</sub> | Unlock (D)       | q <sub>5</sub> | Unlock (T)       |
| P <sub>6</sub> | Unlock (T)       | q <sub>6</sub> | Unlock (D)       |

These are reusable resources.

There will be a deadlock.

$p_0 q_0 p_1 q_1$  introduces the possibility of a deadlock

$p_0 q_0 p_1 q_1 p_2 q_2$  is where the deadlock happens.

- Memory Space is available for allocation of **200Kbytes**, and the following sequence of requests occur.
- **Question:** Are these reusable or consumable resources? When will there be a deadlock?

| P1                 | P2                 |
|--------------------|--------------------|
| ...                | ...                |
| Request 80 Kbytes; | Request 70 Kbytes; |
| ...                | ...                |
| Request 60 Kbytes; | Request 80 Kbytes; |

These are reusable resources.

There will be a deadlock.

P1 and P2 can make their requests then the next requests will cause a deadlock.

P1 or P2 can make both requests but then upon the next request from the other process we will run into deadlock. That is unless there is some kind of release condition that we don't see beyond the processes' second requests.

The cause of the deadlock being that we don't have enough resources to give.

- Each process is attempting to receive a message from the other process and **then** send a message to the other process:
- **Question:** Are these reusable or consumable resources? Will there be a deadlock?

| P1             | P2             |
|----------------|----------------|
| ...            | ...            |
| Receive (P2);  | Receive (P1);  |
| ...            | ...            |
| Send (P2, M1); | Send (P1, M2); |

Consumable resource.

There is a deadlock as they are both waiting on each other.

If receive is not blocking - i.e. the processes just make themselves open to receiving and don't wait - then there is no deadlock.

It depends on the nature of the receive operation.

Prof will tell us if it's blocking or not blocking.

## Deadlock in Multithreaded Applications

2 mutex locks are created and initialized:

```
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;

pthread_mutex_init(&first_mutex,NULL);
pthread_mutex_init(&second_mutex,NULL);
```

2 threads are created and both threads have access to both mutex locks.

```

/* thread_one runs in this function */
void *do_work_one(void *param)
{
 pthread_mutex_lock(&first_mutex);
 pthread_mutex_lock(&second_mutex);
 /**
 * Do some work
 */
 pthread_mutex_unlock(&second_mutex);
 pthread_mutex_unlock(&first_mutex);

 pthread_exit(0);
}

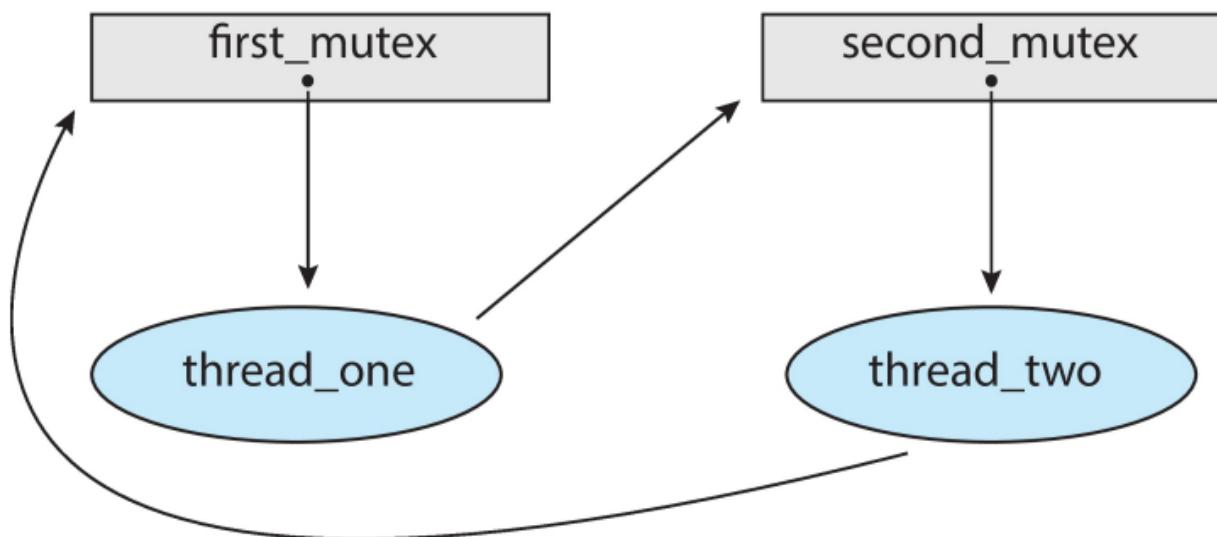
/* thread_two runs in this function */
void *do_work_two(void *param)
{
 pthread_mutex_lock(&second_mutex);
 pthread_mutex_lock(&first_mutex);
 /**
 * Do some work
 */
 pthread_mutex_unlock(&first_mutex);
 pthread_mutex_unlock(&second_mutex);

 pthread_exit(0);
}

```

Deadlock is possible if thread 1 acquires `first_mutex` and thread 2 acquires `second_mutex`

Each thread then waits for the other's mutex.



The order depends on how the threads are scheduled by the CPU scheduler. Hard to test for deadlocks as they may only occur under certain scheduling circumstances.

This is the livelock of waiting continually trying to acquire

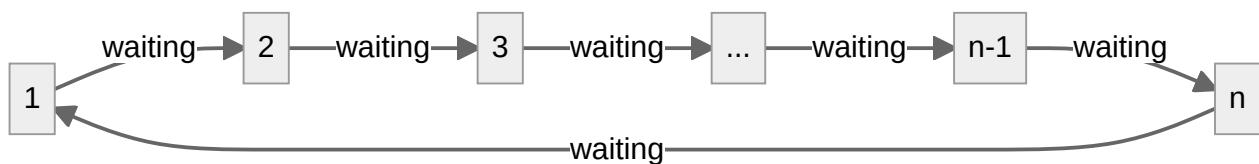
knowledge check:

- a deadlocked state occurs whenever
  - every process in a set is waiting for an event that can only be caused by another process in the set
- Deadlock occurs when every thread in a set is blocked waiting for an event that can be caused only by another thread in the set, while livelock occurs when a thread continuously attempts an action that fails.
  - true
- in the dining philosophers problem, there is a possibility of deadlock but not livelock
  - false
  - everyone picks a fork up then there is deadlock
  - if they keep picking it up and putting it down then there is livelock

## Deadlock Characterization

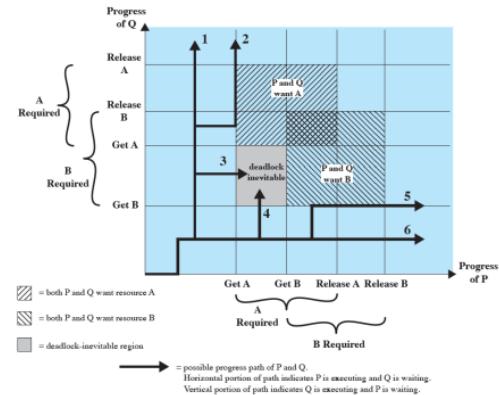
deadlock can arise if 4 conditions hold at the same time:

- mutual exclusion
  - only one process at a time can use a resource
- hold and wait
  - a process holding at least one resource is waiting to acquire additional resources held by other processes
- no preemption (for resources no the process)
  - resource only released by the process holding it after it's done its task
  - resource cannot be preempted
    - processor can't take resources from process and give it to other processes arbitrarily
- circular wait
  - process 1 is waiting on process 2
  - process 2 is waiting on process 3
  - ...
  - process n-1 is waiting on process n
  - process n is waiting on process 1



- All four conditions must hold for a deadlock to occur.

| Possibility of Deadlock | Existence of Deadlock |
|-------------------------|-----------------------|
| 1.Mutual exclusion      | 1.Mutual exclusion    |
| 2.No preemption         | 2.No preemption       |
| 3.Hold and wait         | 3.Hold and wait       |
|                         | 4.Circular wait       |



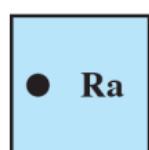
## Resource Allocation Graph

- Deadlocks can be described more precisely in terms of a directed graph called a system **resource-allocation graph**.
  - A set of vertices  $V$  and a set of edges  $E$ .
  - $V$  is partitioned into two types:
    - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the **processes** in the system
    - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all **resource types** in the system

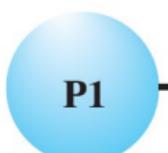
A digraph with vertices  $V$  partitioned into  $P$  processes and  $R$  resources

**request edge** – directed edge  $P_i \rightarrow R_j$

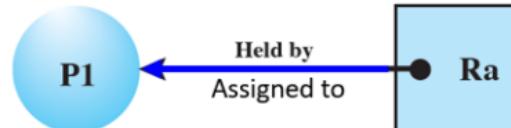
**assignment edge** – directed edge  $R_j \rightarrow P_i$



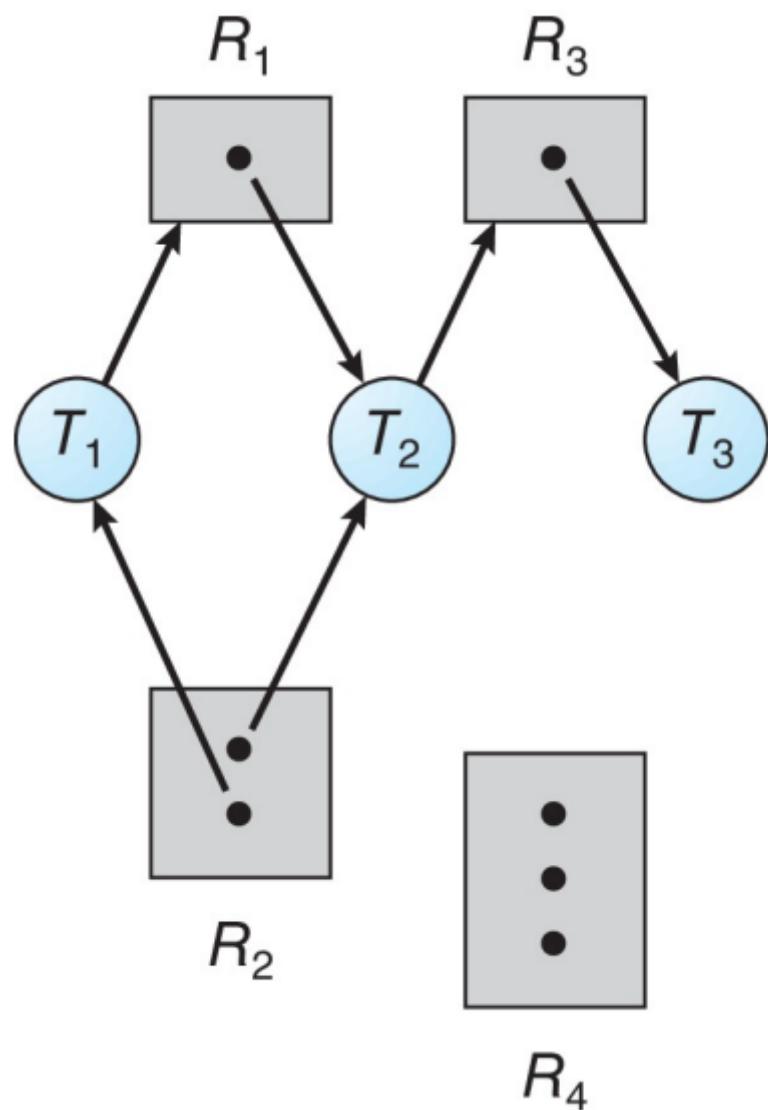
Resource



Process/Thread



● 1 Instance of the resource



There are 3 threads and 4 resources.

| .     | $T_1$    | $T_2$    | $T_3$    |
|-------|----------|----------|----------|
| $R_1$ | req      | assigned |          |
| $R_2$ | assigned | assigned |          |
| $R_3$ |          | req      | assigned |
| $R_4$ |          |          |          |

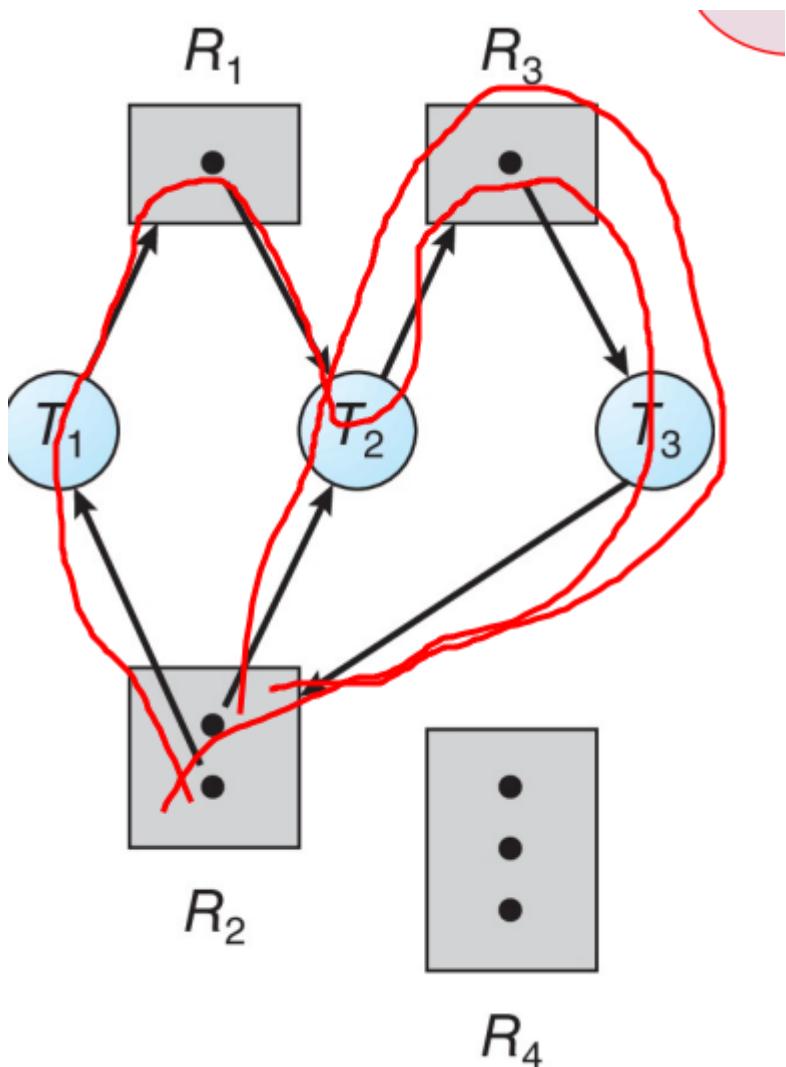
There are 1 instance of  $R_1$  and  $R_2$

2 of  $R_2$  and 3 of  $R_4$

There are no cycles in the graph.

Therefore there are no deadlocks

If there is no cycle then there is no deadlock.



There are 2 cycles.

$T_1$  is assigned  $R_2$  and wants  $R_1$

$T_2$  is assigned  $R_1+R_2$  and wants  $R_3$

$T_3$  is assigned  $R_3$  and wants  $R_2$  which it can't get since there's no more instances of  $R_2$

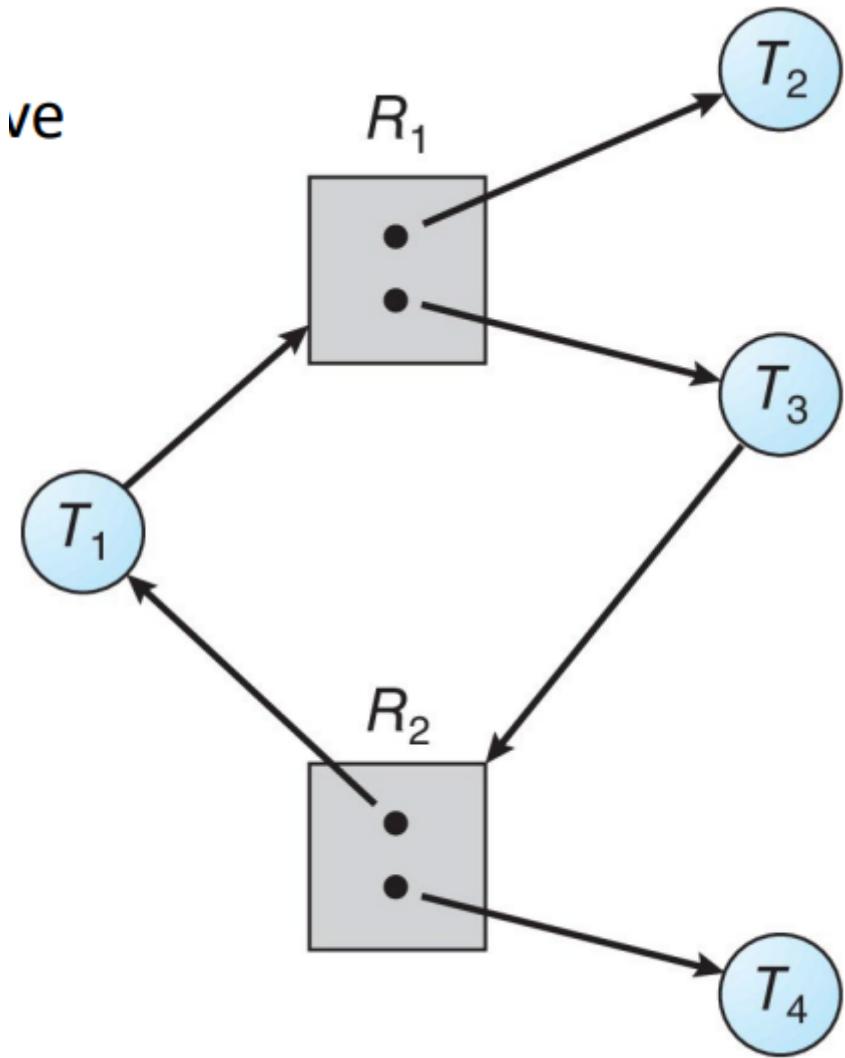
It might seem like there's no deadlock if we look at the smaller cycle.

$T_2$  is assigned  $R_2$  and wants  $R_3$

$T_3$  is assigned  $R_3$  and wants  $R_2$  which it can get since there's another instance of  $R_2$

$T_3$  finishes with  $R_3$  and gives it up for  $T_2$ .

This is a static situation so we have a deadlock



There is a cycle with the middle 4 nodes.

There is no deadlock.

The outer threads of  $T_2$  and  $T_4$  can finish and release their instances of  $R_1$  and  $R_2$  respectively.

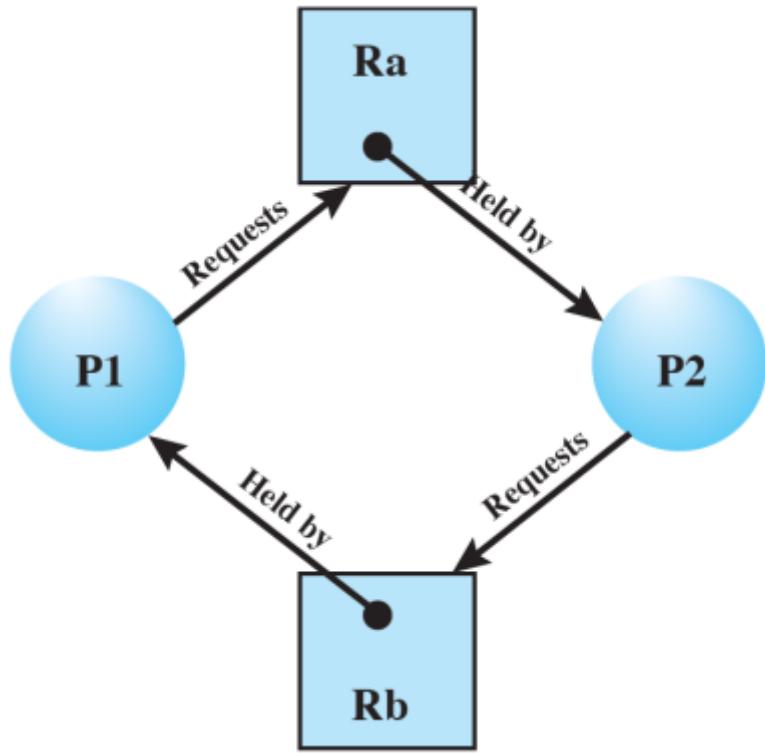
Then  $T_1$  and  $T_3$  can get hold of the resources and work on them.

If there is no cycle, then there is no deadlock

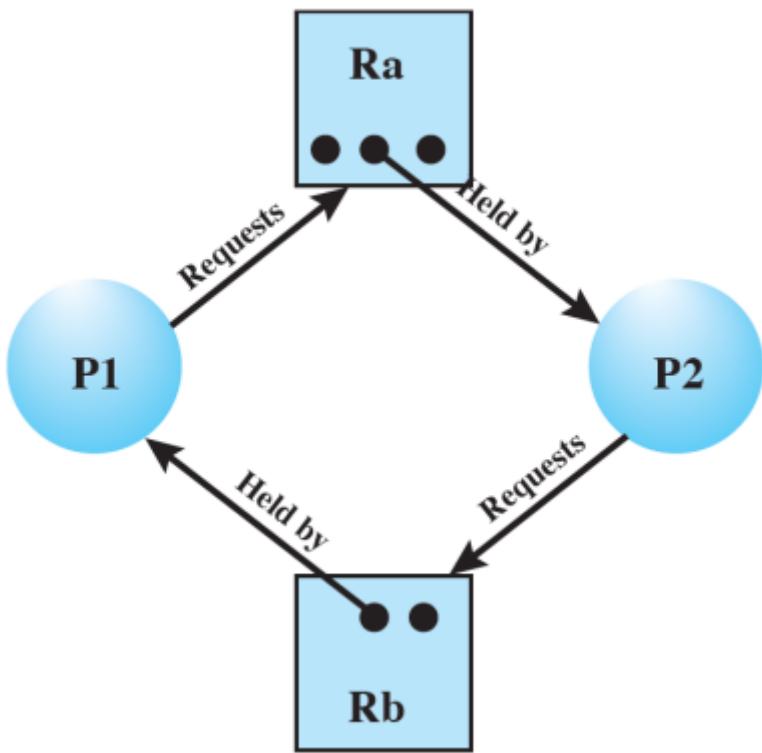
If there is a cycle then there may or may not be a deadlock

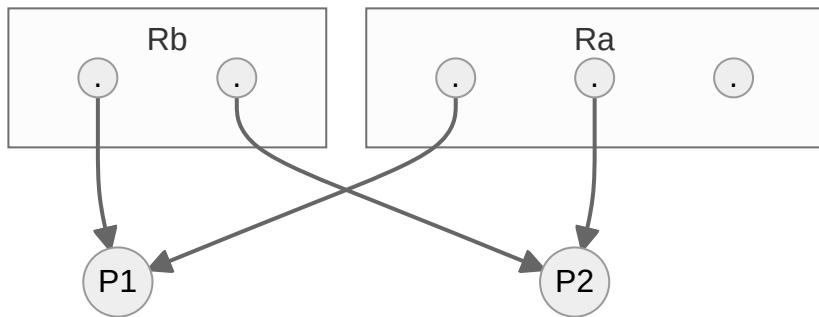
- if only one instance per resource type, then deadlock
- if there are several instances per resource type, possibility of deadlock

deadlock:

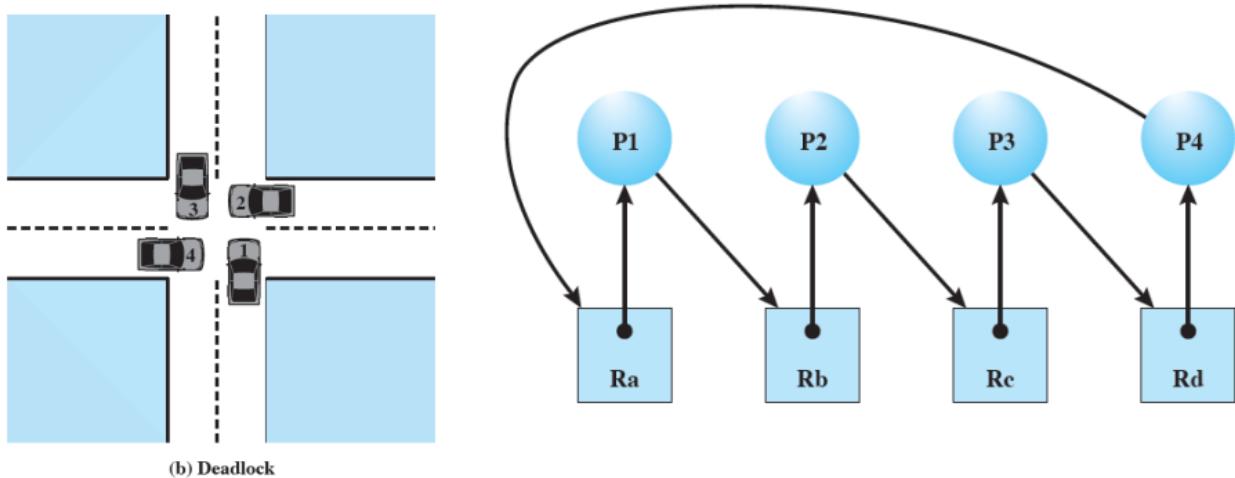


no deadlock:





Resource Allocation graph for the car deadlock example showing circular wait



We can take the numbers on the cars to be the numbers of the processes.

Car 1 is sitting lane **Ra** and wants to cross lane **Rb**

Car 2 is sitting lane **Rb** and wants to cross lane **Rc**

Car 3 is sitting lane **Rc** and wants to cross lane **Rd**

Car 4 is sitting lane **Rd** and wants to cross lane **Ra**

#### Knowledge Check

- one necessary condition for deadlock is \_\_\_, which states that a process must be holding one resource and waiting to acquire additional resources
  - a: hold and wait
- a cycle in a resource-allocation graph is \_\_\_
  - d. a necessary and sufficient condition for a deadlock in the case that each resource has exactly one instance
- if a resource-allocation graph has a cycle, the system must be in a deadlocked state.
  - false

## Methods for Handling Deadlocks

In general, there are 3 ways to deal with deadlocks:

1. prevention and avoidance
2. detection and recovery
3. ignoring it

most OSes including windows and linux just pretend that deadlocks don't happen ???

Prof later explained that this is just left up to the programmer/user and it's not up to the OS.

pulled a pontius pilate and washed their hands of the whole situ

## Prevention

In order to prevent deadlocks we have to invalidate one of the 4 necessary conditions for deadlock

as a refresher

deadlock can arise if 4 conditions hold at the same time:

- mutual exclusion
  - only one process at a time can use a resource
- hold and wait
  - a process holding at least one resource is waiting to acquire additional resources held by other processes
- no preemption (for resources no the process)
  - resource only released by the process holding it after it's done its task
  - resource cannot be preempted
    - processor can't take resources from process and give it to other processes arbitrarily
- circular wait
  - process 1 is waiting on process 2
  - process 2 is waiting on process 3
  - ...
  - process n-1 is waiting on process n
  - process n is waiting on process 1

Though there are shareable resources, we cannot invalidate mutual exclusion as a whole since some resources just cannot be shared and used at the same time.

## Deny Hold and Wait

We can invalidate hold and wait by ensuring that whenever a requests a resource it won't be holding any other resources. Put another way, If the process is holding a resource then they won't be making any requests.

how:

- require process to request and allocate all resources before execution
  - **dynamic nature of resource requesting makes this impractical**
- only allow resource requests when process has no resource allocated
  - must release resources before requesting for more

cons:

- low resource utilization
  - resources allocated but not used
- starvation possible
  - a thread may have to wait indefinitely if it needs several popular resources
  - those resources will always end up allocated to processes that will only need those resources
  - *I imagine that even with bounded wait time considerations, priority will cause this to happen as well*

## Deny No Preemption

reminder: non-preemptive is when a process enters into the processor and only leaves when it's finished or voluntarily leaves when waiting on something

if

- a process is holding some resources
- the process requests another resources
- that resource cannot be immediately allocated to it

then **all resources currently being held are preempted**, released by the processor.

Preempted resources are added to the list of resources that the process is waiting for.

The process starts again once it regains all of the resources it's requesting (the old preempted ones and the new ones).

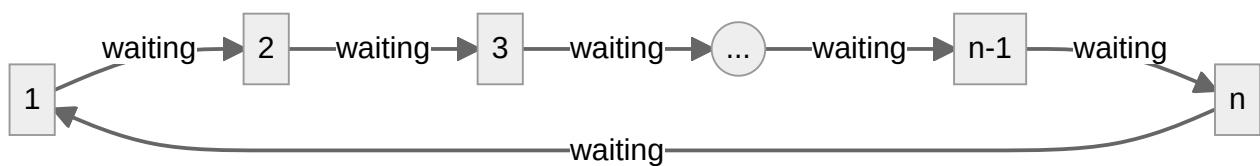
This is often **selectively** applied to resources whose state can be easily saved and restored later, such as cpu registers and database transactions.

cons:

- cannot be applied to resources like mutex locks and semaphores which are the type of resources where deadlock occurs most commonly
  - it's hard to save the condition of a mutex when a process wants to lock for a critical section
  - not easy to store and retrieve → not easy to preempt
  - *we can't do this for the shit we actually care about bruhhhhhh*

# Deny Circular Wait

This is the most common approach



how:

- assign each resource a unique number
- resources must be acquired in the order of their unique number

Knowledge check:

- to handle deadlocks, operating systems most often \_\_\_\_
  - a. pretend that deadlocks never occur
- both deadlock prevention and deadlock avoidance techniques ensure that the system will never enter a deadlocked state
  - true
- most operating systems choose to ignore deadlocks, because
  - d. all of the above
  - handling is expensive, they occur infrequently, livelock recovery methods can be used on deadlocks

# Deadlock Avoidance

Dynamically made decision to see if an allocation request will potentially lead to a deadlock if granted.

System considers

- currently available resources
- currently allocated resources
- and the future requests and releases of each thread

requires knowledge of future process requests.

basics:

- state
  - reflects current allocation of resources to processes
- safe state
  - state where there is at least one execution path where all the processes will finish
  - there will be no deadlock
- unsafe state

- state resulting in possibility of deadlock
- avoidance
  - ensuring system will never enter an unsafe state

## Safe State

When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.

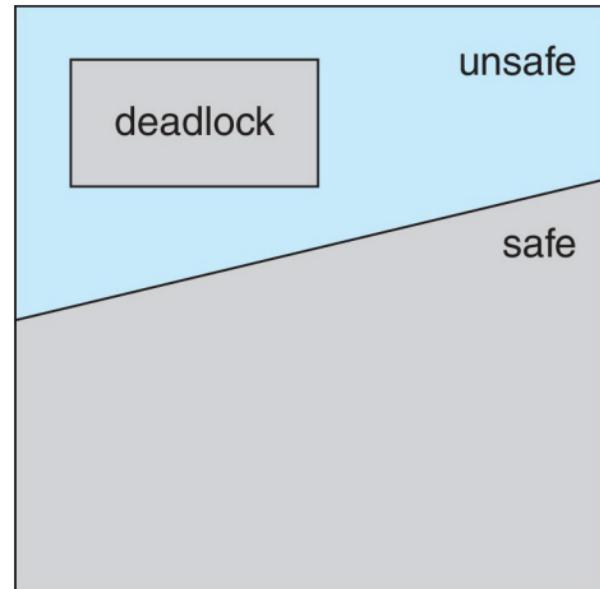
System is in **safe state** if:

- There exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$

That is:

- If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
- When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
- When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on..

- If a system is in **safe state**  $\Rightarrow$  no deadlocks
- If a system is in **unsafe state**  $\Rightarrow$  possibility of deadlock
- **Avoidance**  $\Rightarrow$  ensure that a system will never enter an unsafe state.



- A system has 12 resources and 3 threads  $T_0$ ,  $T_1$ , and  $T_2$  shown below.
- At  $t_0$ ,  $T_0$  is holding 5 resources,  $T_1$  holding two, and  $T_2$  holding two.
- **Question:** Does the sequence  $\langle T_1, T_0, T_2 \rangle$  satisfies the safety condition?
  - $T_1: 2 + 2 \rightarrow$  done, release back 4  $\rightarrow$  5 available.
  - $T_0: 5 + 5 \rightarrow$  done, release back 10  $\rightarrow$  10 available.
  - $T_2: 2 + 9 \rightarrow$  done, release back 9  $\rightarrow$  12 available.

|       | <u>Maximum Needs</u> | <u>Current Needs</u> |
|-------|----------------------|----------------------|
| $T_0$ | 10                   | 5                    |
| $T_1$ | 4                    | 2                    |
| $T_2$ | 9                    | 2                    |

In the above example we only have 3 resources available.

For  $T_1$  we give it 2 then release back 4 once it's done.

Then the above steps just go as shown.

So we are in a safe state

- A system has 12 resources and 3 threads  $T_0$ ,  $T_1$ , and  $T_2$  shown below.
- At  $t_1$ , thread  $T_2$  requests and is allocated one more resource.
- **Question:** Is the system now in a safe state?
  - $T_2: 3$ , may request 6  $\rightarrow$  at  $t_1$ : 2 available.
  - $T_0: 5$ , may request 5  $\rightarrow$  at  $t_1$ : 2 available.
  - $T_1: 2 + 2 \rightarrow$  done, release back 4  $\rightarrow$  at  $t_2$ : 4 available.

|       | <u>Maximum Needs</u> | <u>Current Needs</u> |
|-------|----------------------|----------------------|
| $T_0$ | 10                   | 5                    |
| $T_1$ | 4                    | 2                    |
| $T_2$ | 9                    | 3                    |

This is the same example as seen above but now  $T_2$  needs 1 more resource, requests the resource and receives it, resulting in only 2 resources available at the beginning.

We can start the sequence with  $T_1$  and allocate 2 then get back 2 once done.

We end up with only 4 available resources which isn't enough to satisfy the maximum needs of  $T_0$  and  $T_2$ .

# Avoidance Algos

There are 2 algos:

1. using a resource-allocation when there's a single instance of a resource type
2. use the Banker's Algo when there's multiple instances of a resource type

## Resource-Allocation Graph Scheme

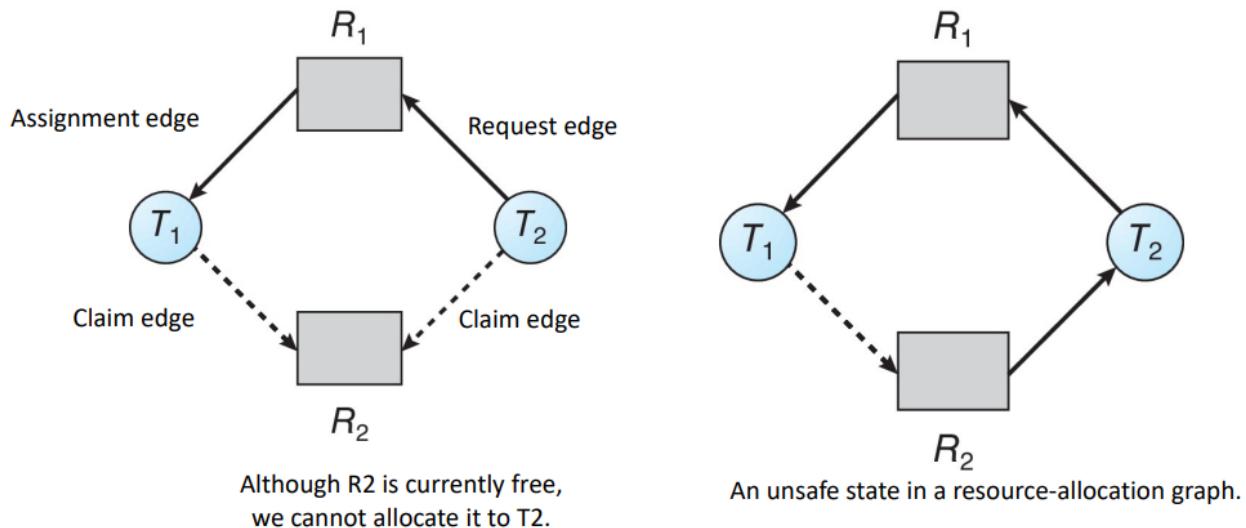
**Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_i$  **may request** resource  $R_j$  represented by a dashed line

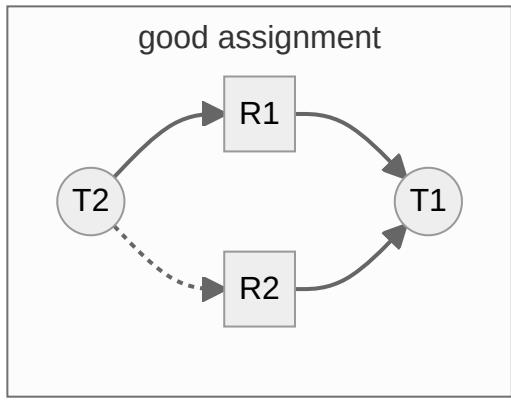
- Claim edge converts to **request edge** when a process requests a resource
- Request edge converted to an **assignment edge** when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge

Resources must be claimed *a priori* in the system.

*a priori* = ahead of time

claims turn into requests which turn into assignments before turning back into claim edges





Suppose that process  $P_i$  requests a resource  $R_j$

The request can be granted **only** if converting the request edge to an assignment edge **does not result in the formation of a cycle** in the resource allocation graph.

make sure there is no cycles forming if you convert a request edge to an assignment edge

## Banker's Algorithm

prof misspells this is as Baker at some point. It's Banker.

conditions:

- multiple instances of resources
- each process must have a prior claim maximum use
- when a process requests a resource, it may have to wait
- when a process gets all its resources it must return them in a finite amount of time

Let  $n = \text{number of processes}$ , and  $m = \text{number of resources types}$ .

- **Available:** Vector of length  $m$ . If  $\text{available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- **Max:**  $n \times m$  matrix. If  $\text{Max}[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- **Allocation:**  $n \times m$  matrix. If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- **Need:**  $n \times m$  matrix. If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

# Safety Algorithm

1. Let  $\text{Work}$  and  $\text{Finish}$  be vectors of length  $m$  and  $n$ , respectively. Initialize:  
 $\text{Work} = \text{Available}$   
 $\text{Finish}[i] = \text{false}$  for  $i = 0, 1, \dots, n-1$
2. Find an index  $i$  such that both:
  - (a)  $\text{Finish}[i] = \text{false}$
  - (b)  $\text{Need}_i \leq \text{Work}$If no such  $i$  exists, go to step 4
3.  $\text{Work} = \text{Work} + \text{Allocation}_i$   
 $\text{Finish}[i] = \text{true}$   
go to step 2
4. If  $\text{Finish}[i] == \text{true}$  for all  $i$ , then the system is in a **safe** state otherwise we are in **unsafe** state.

$\text{Request}_i$  = request vector for process  $P_i$ . If  $\text{Request}_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $\text{Request}_i \leq \text{Need}_i$ , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $\text{Request}_i \leq \text{Available}$ , go to step 3. Otherwise,  $P_i$  must wait, since resources are not available.
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:  
 $\text{Available} = \text{Available} - \text{Request}_i$   
 $\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$   
 $\text{Need}_i = \text{Need}_i - \text{Request}_i$ 
  - If **safe**  $\Rightarrow$  the resources are allocated to  $P_i$  (*use Safety alg. in the prev slide*).
  - If **unsafe**  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored.

# Example of Banker's Algorithm /1

- 5 processes  $P_0$  through  $P_4$ ;
- 3 resource types:  
 $A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)
- Snapshot at time  $T_0$ :

|       | <u>Allocation</u> | <u>Max</u>  | <u>Available</u> |
|-------|-------------------|-------------|------------------|
|       | $A \ B \ C$       | $A \ B \ C$ | $A \ B \ C$      |
| $P_0$ | 0 1 0             | 7 5 3       | 3 3 2            |
| $P_1$ | 2 0 0             | 3 2 2       |                  |
| $P_2$ | 3 0 2             | 9 0 2       |                  |
| $P_3$ | 2 1 1             | 2 2 2       |                  |
| $P_4$ | 0 0 2             | 4 3 3       |                  |

## Need

|       | $A \ B \ C$ |
|-------|-------------|
| $P_0$ | 7 4 3       |
| $P_1$ | 1 2 2       |
| $P_2$ | 6 0 0       |
| $P_3$ | 0 1 1       |
| $P_4$ | 4 3 1       |

Q: Does the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfy safety criteria?

A: yes

we start with  $\langle 3, 3, 2 \rangle$

$P_1$  can be allocated and give back  $\langle 5, 3, 2 \rangle$

$P_3$  can be allocated and give back  $\langle 7, 4, 3 \rangle$

$P_4$  can be allocated and give back  $\langle 7, 4, 5 \rangle$

$P_2$  can be allocated and give back  $\langle 10, 4, 7 \rangle$ , making all instances of  $A$  and  $C$  available

$P_0$  can be allocated and give back  $< 10, 5, 7 >$ , making all instances of all resources available.

---

## Example of Banker's Algorithm /3

- Does the sequence  $< P_1, P_3, P_4, P_2, P_0 >$  satisfy safety criteria?

|       | <u>Allocation</u> | <u>Need</u> | <u>Available</u> |
|-------|-------------------|-------------|------------------|
|       | A B C             | A B C       | A B C            |
| $P_0$ | 0 1 0             | 7 4 3       | 3 3 2            |
| $P_1$ | 2 0 0             | 1 2 2       |                  |
| $P_2$ | 3 0 2             | 6 0 0       |                  |
| $P_3$ | 2 1 1             | 0 1 1       |                  |
| $P_4$ | 0 0 2             | 4 3 1       |                  |

- Executing safety algorithm shows that sequence  $< P_1, P_3, P_4, P_2, P_0 >$  **satisfies** safety requirement.

you do the same process with this.

A: the sequence satisfies safety requirement

## Example: P1 Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true

|       | <u>Allocation</u> | <u>Need</u> | <u>Available</u> |
|-------|-------------------|-------------|------------------|
|       | A B C             | A B C       | A B C            |
| $P_0$ | 0 1 0             | 7 4 3       | 2 3 0            |
| $P_1$ | 3 0 2             | 0 2 0       |                  |
| $P_2$ | 3 0 2             | 6 0 0       |                  |
| $P_3$ | 2 1 1             | 0 1 1       |                  |
| $P_4$ | 0 0 2             | 4 3 1       |                  |

- Is the new state safe?
  - Executing safety algorithm shows that sequence  $< P_1, P_3, P_4, P_0, P_2 >$  **satisfies** safety requirement.
  - Can request for  $(3,3,0)$  by  $P_4$  be granted?
  - Can request for  $(0,2,0)$  by  $P_0$  be granted?
- 

In essence:

- we are in a safe set if there is at least one sequence of resource allocations to processes that does not result in a deadlock
- for  $i \in [0..n]$ 
  - compare the available with the need vector
  - if the available is lesser than the need vector then unsafe

- else, it can be allocated
  - add allocation vector to available vector as the resources are now returned from process
  - at no point should the available exceed the maximum for each entry in the vector
  - if at any point in the vector of processes does the above **for-loop** exit then it does not satisfy the safety requirement
- 

Knowledge Check:

- which of the following statements is true:
  - d. An unsafe state may lead to a deadlocked state
  - we may avoid the deadlocked state
- Suppose that there are **10 resources** available to **3 processes**. At  $t_0$ , the following data is collected.

| Process        | Maximum Needs | Currently Owned |
|----------------|---------------|-----------------|
| P <sub>0</sub> | 10            | 4               |
| P <sub>1</sub> | 3             | 1               |
| P <sub>2</sub> | 6             | 4               |

- ■ Which of the following correctly characterizes this state?
  - b. It is not safe

Suppose that there are **12 resources** available to **3 processes**. At  $t_0$ , the following data is collected.

| Process        | Maximum Needs | Currently Owned |
|----------------|---------------|-----------------|
| P <sub>0</sub> | 10            | 4               |
| P <sub>1</sub> | 3             | 2               |
| P <sub>2</sub> | 7             | 4               |

- Which of the following correctly characterizes this state?
  - It is safe

## Deadlock Detection

If we don't use deadlock prevention or avoidance then a deadlock situation may occur

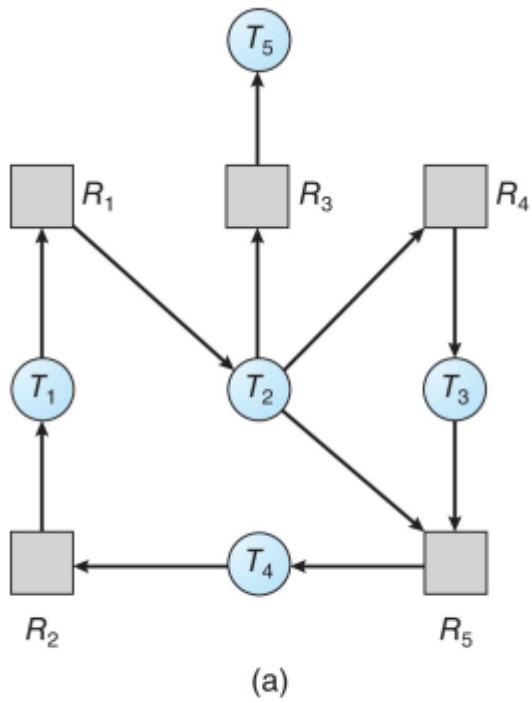
in lieu of prevention and avoidance we can use detection and recovery

A detection algorithm to look at the system state see if a deadlock has happened.

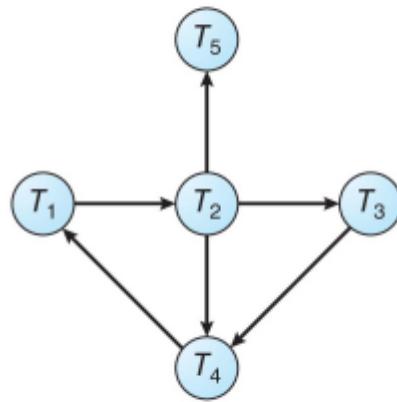
A recovery scheme to recovery from the deadlock.

For situations where we have a single instance of each resource type:

- maintain a **wait-for** graph
  - nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically check for a cycle in the graph
  - if there is a cycle then there is a deadlock
- takes  $O(n^2)$  where  $n$  is the number of vertices in the graph



(a)



(b)

## Resource-Allocation Graph

## Corresponding wait-for graph

In the above, we see that there is a deadlock as there are cycles so we are in deadlock.

For situations where we have multiple instances for each resource type:

- we use something similar to the banker's algo
- all the matrices and all of that
- "it's exactly similar to the Banker's algorithm" - prof

How often we invoke this depends on how likely it is to occur and how many threads would be affected by a deadlock if it happens.

The more often and the more impactful then the more frequently we should invoke.

Invoking everytime there is a resource request is expensive (except in the case of Banker's algo since we use it anyways).

A less expensive alternative is to invoke the algo at defined intervals.

### Knowledge Check

- the wait-for graph scheme is not applicable to a resource allocation system with multiple instances of each resource type
  - True
- in a wait-for graph
  - d. None of the above
  - b and c are each others' contrapositive so they're both false
- The decision of when to invoke a detection algorithm depends on the likely frequency of deadlocks and the number of threads likely to be affected by a deadlock when it happens

- True

# Recovery from Deadlock

We have 2 options in a deadlock:

- break circular wait by aborting one or more involved threads
- preempt some resources from one or more of the deadlocked threads
  - take some resources away

for abortion we have some options

- abort all the deadlocked processes
- abort one at a time until the deadlock cycle is eliminated

Let's not kill everything.

What affects the process abortion order:

1. priority of process
2. elapsed and remaining computation time
3. allocated resources
4. required resources
5. number of processes to be terminated

copying slides now

selecting a victim that minimizes cost

rollback to some safe state, restart process for that state once the resources are available.

We may have starvation if we keep selecting the same victim over and over again

we have to include the number of rollbacks as part of the cost factor

---

## Knowledge Check

- a system can recover from a deadlock by
  - d. All of the above
- to recover from a deadlock using resource preemption,
  - d. All of the above

This concludes the **Deadlock and Starvation** slide deck.

# CPU Scheduling

toc:

- Recap
- Basics
- Scheduling Criteria
- Scheduling Algos
- Multi-Processor Scheduling
  - Multicore
  - Multithreaded Cores
  - Load Balancing
  - Processor Affinity
  - NUMA Systems
  - Heterogeneous Multiprocessing
- Real-Time CPU Scheduling
  - Priority Based Scheduling
- Algo Eval
  - Deterministic
  - Queueing Models
  - Simulations
  - Real Implementation

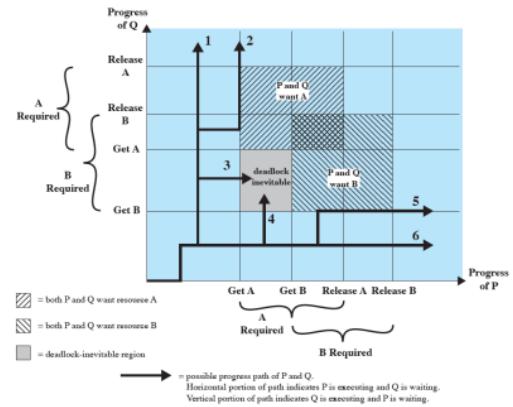
## Where we are...

- Computer System Overview
- Operating System Structures
- Process Description and Control
- Threads
- Synchronization and Mutual Exclusion
- Deadlock and Starvation
- **CPU Scheduling**

## Recap for CPU Scheduling

- All four conditions must hold for a deadlock to occur.

| Possibility of Deadlock | Existence of Deadlock |
|-------------------------|-----------------------|
| 1.Mutual exclusion      | 1.Mutual exclusion    |
| 2.No preemption         | 2.No preemption       |
| 3.Hold and wait         | 3.Hold and wait       |
|                         | 4.Circular wait       |



Generally speaking, we can deal with the deadlock problem in one of **three** ways:

- Ensure that the system will **never** enter a deadlock state:
  - Deadlock prevention.
  - Deadlock avoidance.
- Allow the system to enter a deadlock state, **detect** it, and then recover.
- **Ignore** the problem and pretend that deadlocks never occur in the system.
  - Used by most OSes including Linux and Windows.

## Basics for CPU Scheduling

Kernel threads are the ones being scheduled.

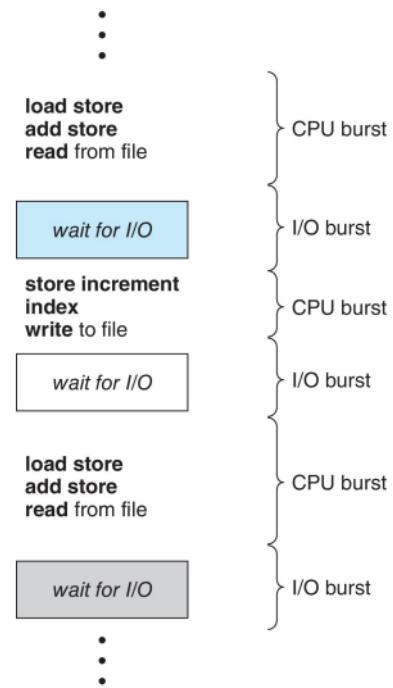
“Process scheduling” and “thread scheduling” are often used interchangeably.

We say CPU to mean CPU core as a computation unit

Maximum CPU utilization obtained with multiprogramming

# CPU-I/O Burst Cycle

- Process execution consists of a **cycle** of CPU execution and I/O wait.
- Process execution begins with a **CPU burst** followed by an **I/O burst**, and so on.
  - CPU burst: Scheduling process state in which the process executes on CPU.
  - I/O burst: Scheduling process state in which the CPU performs I/O.

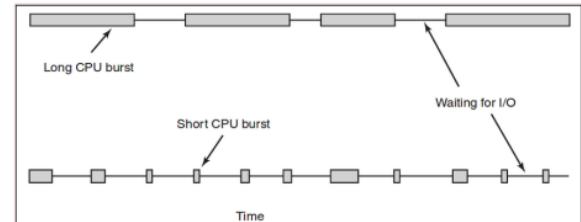
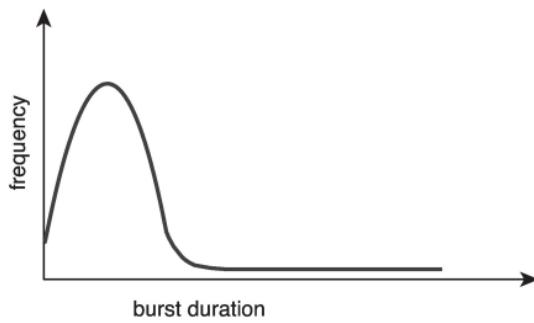


We alternate working (cpu) and wait (i/o) bursts.

We end on an I/O burst as the process is ending

## Histogram of CPU-burst Times

- CPU burst distribution is of main concern.
  - Large number of short bursts: An I/O-bound program typically has many short CPU bursts.
  - Small number of longer bursts: A CPU-bound program might have a few long CPU bursts.



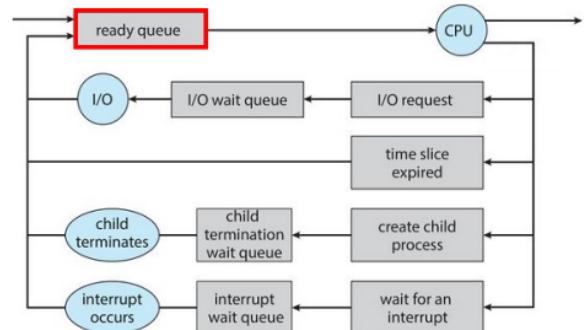
I/O and CPU bound comparison.

On the histogram:

- above is cpu-bound
- below is i/o bound

# CPU Scheduler

- The **CPU scheduler** selects from among the processes in ready queue, and allocates a CPU core to one of them.
  - Queue may be **ordered in various ways**.
- The records in the queues are generally process control blocks (PCBs) of the processes.



Queueing-diagram representation  
of process scheduling.

Reminder: the whole process itself isn't in the queue, just the pcb (which has all the info and points to where the process actually is)

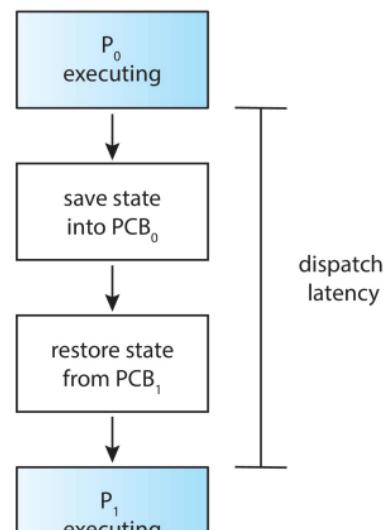
## Preemptive and Nonpreemptive Scheduling

- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state (I/O or child process)
  2. Switches from running to ready state (interrupts)
  3. Switches from waiting to ready (I/O completion or child finish)
  4. Terminates (last instructions)
- If only do **1** and **4**, the scheduling is **nonpreemptive**, or cooperative:
  - Once the CPU has been allocated to a process, the process keeps the CPU until it releases it either by **terminating** or by **switching to the waiting state**.
- All other scheduling is **preemptive** (all modern OSs use this).

Preemptive scheduling makes decisions at all these times.

# Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the CPU scheduler. It involves:
  - Switching context.
  - Switching to user mode.
  - Jumping to the proper location in the user program to resume that program.
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running.



The role of the dispatcher.

---

## Knowledge Check

- which of the following is true of cooperative scheduling
  - b. A process keeps the CPU until it releases the CPU either by terminating or by switching to the waiting state.
- In preemptive scheduling, the sections of code affected by interrupts must be guarded from simultaneous use
  - True
  - we want to ensure mutual exclusion
  - we want to make sure no other process is accessing at the same time

# Scheduling Criteria

- cpu util
  - keep as busy as possible
  - maximize this
- throughput
  - num of processes that complete their execution per time unit
  - maximize this
- turnaround time
  - amount of time to execute a particular process
  - minimize this
- waiting time
  - aggregate amount of time a process has been waiting in the ready queue
  - will be affected by scheduling algo
  - minimize this
- response time
  - time between request submission and first response for interactive systems

- minimize this

All numbers in milliseconds (ms) unless stated otherwise.

For now we assume we have 1 cpu.

\_\_\_\_\_ is the number of processes that are completed per time unit.

- CPU utilization
- Response time
- Turnaround time
- Throughput

## Scheduling Algos

Algorithms to decide which of the processes in the ready queue is to be allocated to the CPU's core:

1. First-Come, First-Served Scheduling (FCFS)
2. Shortest-Job-First Scheduling (SJF)
3. Round-Robin Scheduling (RR)
4. Priority Scheduling
5. Multilevel Queue Scheduling
6. Multilevel Feedback Queue Scheduling

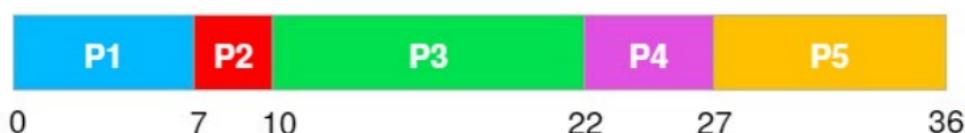
We will be using a Gantt chart.

- A bar chart that illustrates a particular schedule, including the start and finish times of each of the participating processes.
- All times are in ms unless otherwise specified.

### First- Come, First-Served (FCFS) Scheduling

| Process | Burst Time |
|---------|------------|
| P1      | 7          |
| P2      | 3          |
| P3      | 12         |
| P4      | 5          |
| P5      | 9          |

- Suppose processes arrive in the order  $P_1, P_2, P_3, P_4, P_5$ , then the Gantt Chart for the FCFS is:



we're assuming these all arrived at (more or less) the same time since there is no `arrival time` column in the chart. We assume in order of number just because it's convenient here.

### A nonpreemptive algorithm:

- Once the CPU has been allocated to a process, that process keeps the CPU until it releases the CPU, either by **terminating** or by **requesting I/O**.

FCFS can be easily implemented with a FIFO queue.

different orders will result in different wait times but FCFS can't account for that

Suppose processes arrive in the order:  $P_1, P_2, P_3$ .

| Process | Burst Time |
|---------|------------|
| $P_1$   | 24         |
| $P_2$   | 3          |
| $P_3$   | 3          |



Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$

Average waiting time:  $(0 + 24 + 27)/3 = 17$

Suppose that the processes arrive in the order:  $P_2, P_3, P_1$

The Gantt chart for the schedule is:



Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$

Average waiting time:  $(6 + 0 + 3)/3 = 3$

Much better than previous case!

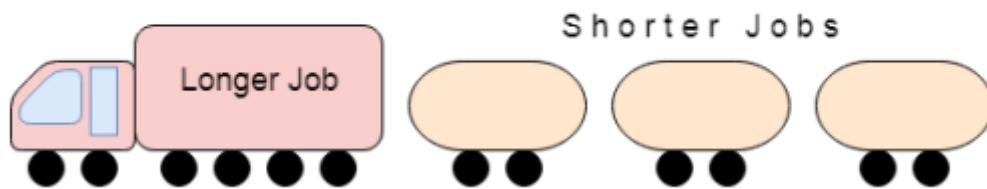
Convoy effect:

- short process behind long process
- consider one cpu-bound and many i/o-bound processes

FCFS is not good for interactive systems

- important that each process get a share of the cpu at regular intervals

### The Convoy Effect, Visualized Starvation



### Shortest-Job-First (SJF) Scheduling

- for each process, associate the length of its next CPU burst
  - use these lengths to schedule the process with the shortest time
  - shortest-next-cpu-burst is a more accurate description
- SJF is optimal
  - gives minimum average waiting time for a given set of processes

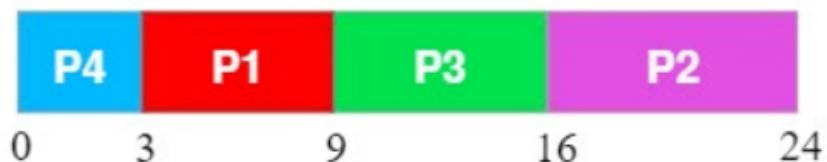
If the next bursts of 2 processes are the same, FCFS scheduling is used to break the tie.

But how do we know how long the next cpu burst is? We can ask the user but there's more to it.

## Example of SJF

| Process | Burst Time |
|---------|------------|
| P1      | 6          |
| P2      | 8          |
| P3      | 7          |
| P4      | 3          |

### SJF scheduling chart



▪ Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$

Remember that waiting time is the time that the process waits to **start** not to finish.

So the above calculation w/ start time  $s_i$  for process  $p_i$  would look like this:

$$\text{average waiting time} = (w_1 + w_2 + w_3 + w_4)/4$$

For FCFS we'd have

$$\text{average waiting time} = (0 + 6 + 14 + 21)/4 = 10.75$$

We can only estimate the length, typically assuming it to be similar to the previous one. Then pick the process with the shortest predicted next CPU burst.

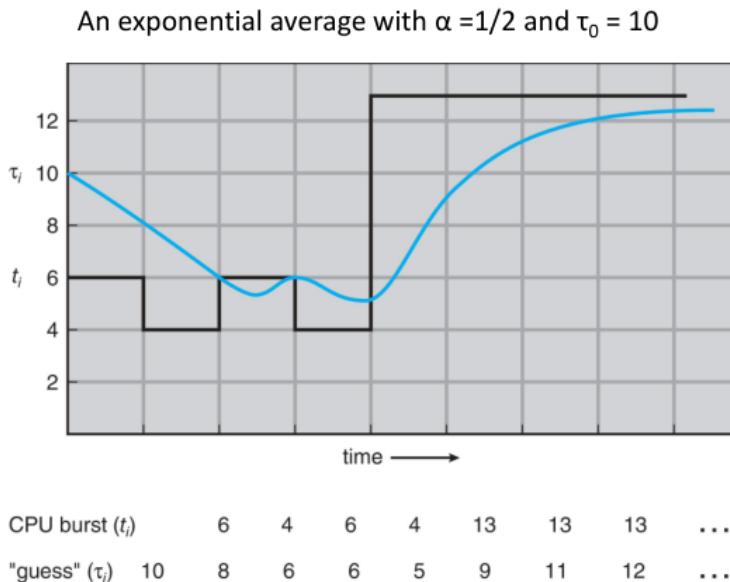
We use an exponential averaging of the length of the previous cpu bursts

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$

- $t_n$  = actual length of  $n^{th}$  CPU burst
- $\tau_{n+1}$  = predicted value of the next CPU burst
- $\alpha \in [0, 1]$ , typically  $\frac{1}{2}$  for our applications

## Prediction of the Length of the Next CPU Burst

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$



We start by predicting 10 ms for the first burst.

It actually uses 6 ms.

We use these values to predict the next burst.

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n = \frac{1}{2}6 + (1 - \frac{1}{2})10 = 8$$

Then we check against the next actual burst length.

So on and so forth.

The predictions aren't the tightest at times but they're not too far off.

Note that we're using the actual previous burst time as well as the previous prediction in our calculations to predict the next burst time.

---

More on exponential averaging:

- $\alpha = 0$ 
  - $\tau_{n+1} = \tau_n$
  - recent history does not count
- $\alpha = 1$ 
  - $\tau_{n+1} = \alpha t_n$
  - only the actual last CPU burst counts

Exponential averaging is a series that we use for prediction in general, not just for processors.

Tweaking the  $\alpha$  value will allow us to put more weight on near or far.

---

Preemptive vs Non-Preemptive

- SJF can be either
- the choice arises when a new process arrives at the ready queue while a previous process is still executing
- the next burst of the newer process may be shorter than what is left of the currently executing process
- in this scenario
  - if preemptive
    - preempt the currently executing process
    - sometimes known as `shortest-remaining-time-first-scheduling`
  - if non-preemptive
    - allow the currently running process to finish its CPU burst

The context switches take about 10 microseconds so preempting here is fine for us as we're typically just working in milliseconds.

# Example of Shortest-remaining-time-first

- Now we add the concepts of **varying arrival times** and **preemption** to the analysis:

| Process | Arrival Time | Burst Time |
|---------|--------------|------------|
| $P_1$   | 0            | 8          |
| $P_2$   | 1            | 4          |
| $P_3$   | 2            | 9          |
| $P_4$   | 3            | 5          |

## Preemptive SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5 \text{ msec}$

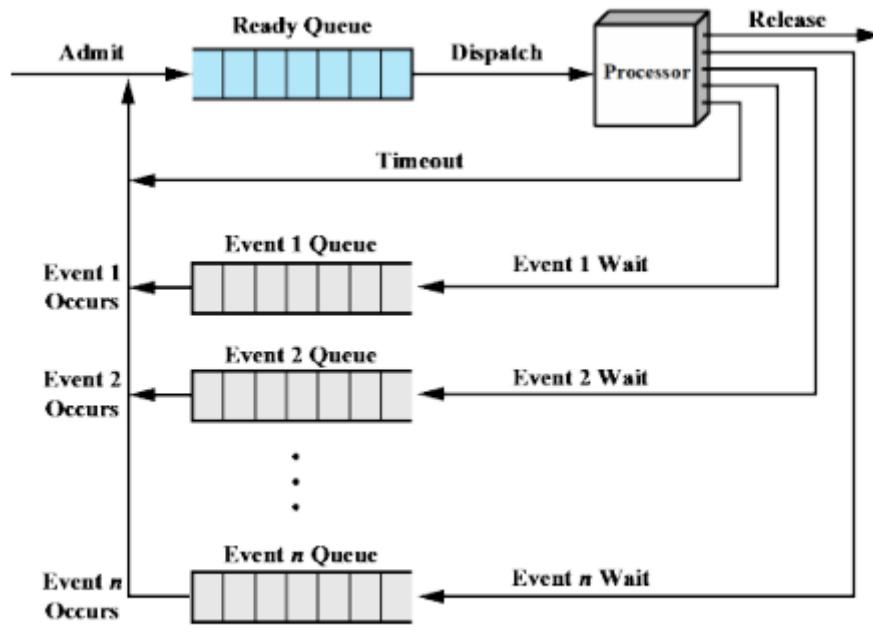
- $P_1$  arrives and works for 1 ms
- $P_2$  arrives and has 4 ms compared to  $P_1$ 's 7 ms
- $P_1$  is preempted and  $P_2$  begins execution
- $P_3$  arrives and has 9 ms compared to  $P_2$ 's 3 ms
- $P_4$  arrives and has 5 ms compared to  $P_2$ 's 2 ms
- @ $t = 5$ ,  $P_2$  finishes
- We choose  $P_4$  as it has the least remaining time to finish
- $P_4$  finishes
- $P_1$  is chosen to execute as it has 7 ms remaining
- $P_3$  is chosen after  $P_1$  finishes as it has 9 ms remaining

$$\text{Average waiting time} = [P_1 + P_2 + P_3 + P_4]/4$$

Remember that waiting time is the time waiting in the queue. So we have to add in the 9 ms that  $P_1$  was sent back to the queue for.

## Round Robin (RR) Scheduling

- each process gets a time quantum (q)
  - a small unit of cpu time
  - usually 10-100 ms
- after the quantum has passed
  - process is preempted
  - process added to end of ready queue



○

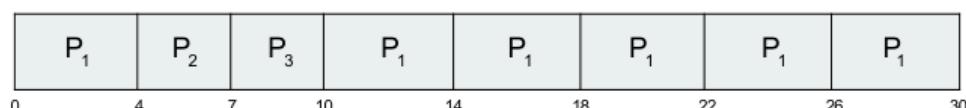
(b) Multiple blocked queues

- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once.
- No process waits more than  $(n-1) \times q$  time units.
- Timer interrupts every quantum to schedule next process.
- Performance:
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow$   $q$  must be large with respect to context switch, otherwise overhead is too high.
- The average waiting time under the RR policy is often long.

## Example of RR with Time Quantum = 4

| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
| $P_1$          | 24                |
| $P_2$          | 3                 |
| $P_3$          | 3                 |

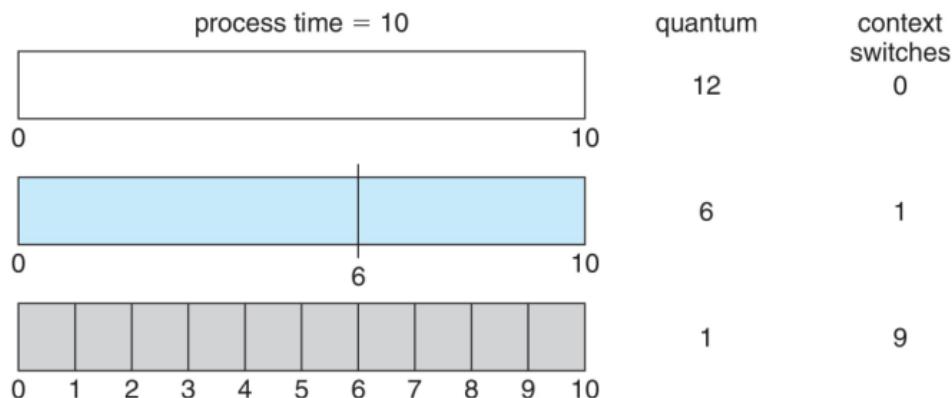
- The Gantt chart is:



- $P_1$  starts executing
- 4 ms pass and the cpu checks the queue
- $P_1$  is preempted and  $P_2$  is put in its place for execution
- $P_2$  finishes execution
- $P_3$  executes and finishes
- $P_1$  begins to execute once more and runs the course of the 4 ms quantum

averaging waiting time =  $\frac{6+4+7}{3} = 5.666\dots$

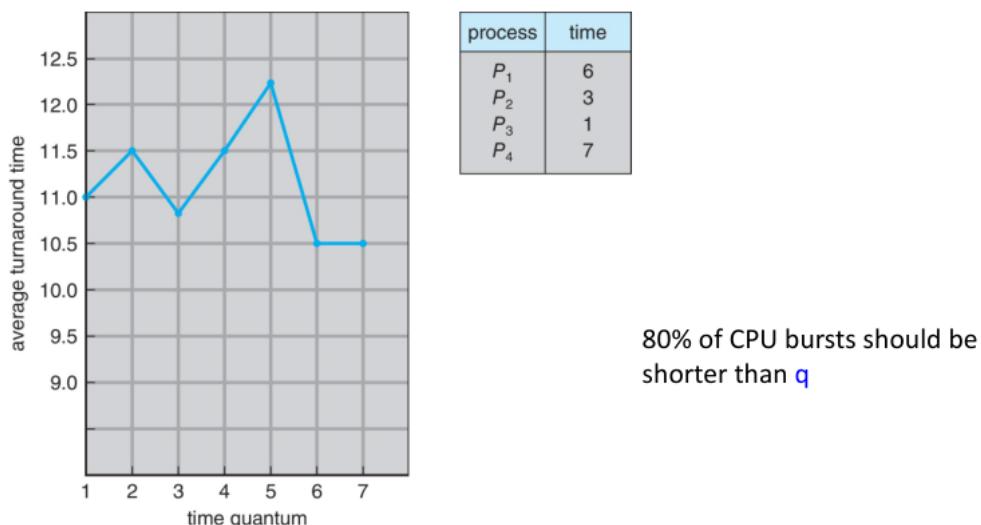
## Time Quantum and Context Switch Time



- $q$  should be large compared to context switch time
- $q$  usually 10ms to 100ms, context switch < 10  $\mu$ sec

## Turnaround Time Varies With The Time Quantum

- Typically, higher average turnaround than SJF, but better **response**:



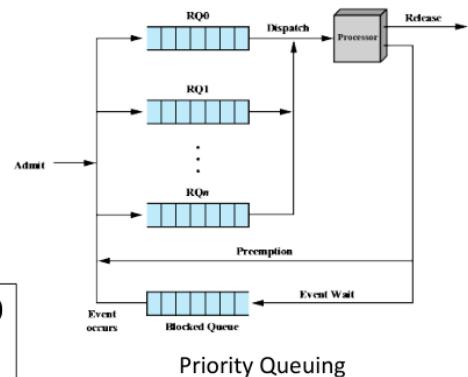
## Priority Scheduling

- priority number (integer) associated w/ each process

- CPU is allocated to the process with the highest priority
  - can be done preemptively or nonpreemptively
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
  - larger burst = lower priority
  - smaller burst = higher priority

## Is 0 the Highest or Lowest Priority?

- Some systems use low numbers to represent low priority; others use low numbers for high priority.
- We assume that low numbers represent high priority.



Linux and Windows use the same numbering system

## Priority Scheduling

- Problem ≡ **Starvation** – low priority processes may never execute.
  - Rumor has it that when they shut down the IBM 7094 at MIT in 1973, they found a low-priority process that had been submitted in 1967 and had not yet been run.
- Solution ≡ **Aging** – as time progresses increase the priority of the process.



IBM 7094 Operator's Console

Starvation is a problem for both preemptive and nonpreemptive.

We up priority for old processes

# Example of Priority Scheduling

| <u>Process</u> | <u>Burst Time</u> | <u>Priority</u> |
|----------------|-------------------|-----------------|
| $P_1$          | 10                | 3               |
| $P_2$          | 1                 | 1               |
| $P_3$          | 2                 | 4               |
| $P_4$          | 1                 | 5               |
| $P_5$          | 5                 | 2               |

- Priority scheduling Gantt Chart



- Average waiting time = 8.2 msec (exercise).

## Priority Scheduling w/ Round-Robin

- Run the process with the highest priority. Processes with the same priority run round-robin.

| <u>Process</u> | <u>Burst Time</u> | <u>Priority</u> |
|----------------|-------------------|-----------------|
| $P_1$          | 4                 | 3               |
| $P_2$          | 5                 | 2               |
| $P_3$          | 8                 | 2               |
| $P_4$          | 7                 | 1               |
| $P_5$          | 3                 | 3               |

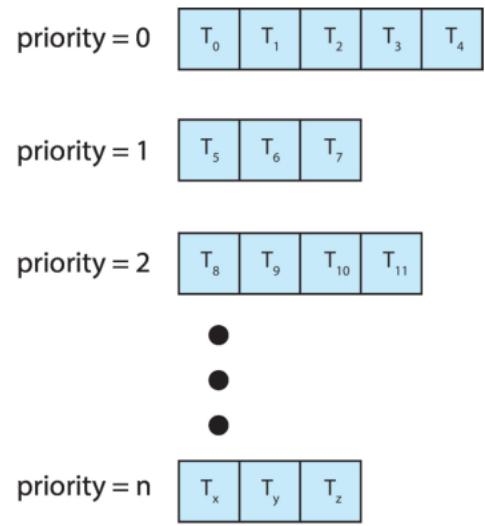
- Gantt Chart with 2 ms time quantum



A process can't be preempted for a process with a lower priority. If a process is the only process within its priority tier and there are no other higher priority processes then the process gets to run with impunity.

# Multilevel Queue Scheduling

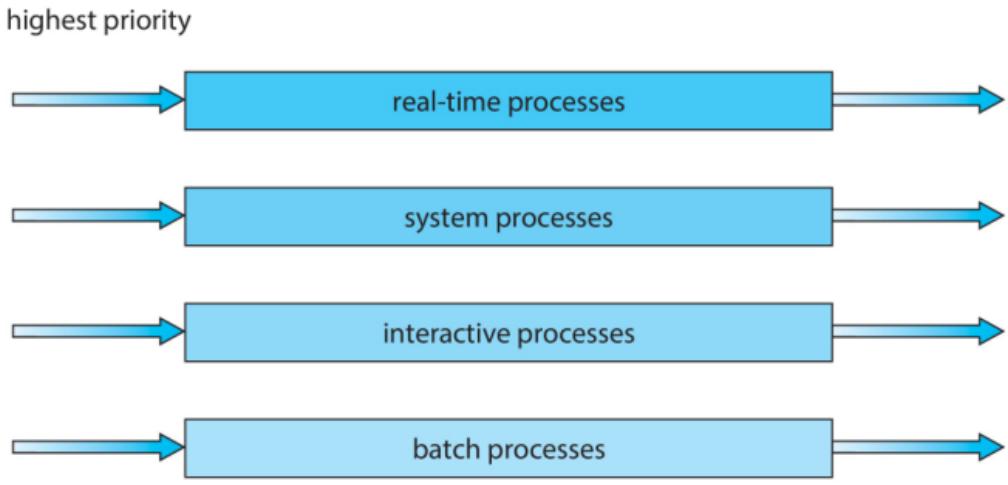
- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue.



Separate queues for each priority.

# Multilevel Queue Scheduling

- Prioritization based upon process type



Multilevel queue scheduling.

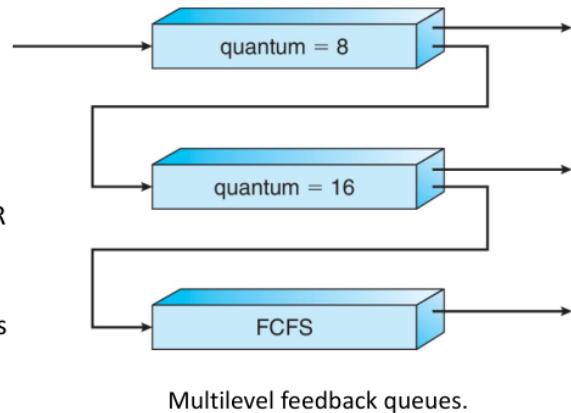
A process can move between the various queues; aging can be implemented this way.

Multilevel-feedback-queue scheduler defined by the following parameters:

- number of queues
- scheduling algorithms for each queue
- method used to determine when to upgrade a process
- method used to determine when to demote a process
- method used to determine which queue a process will enter when that process needs service

## Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS
- Scheduling
  - A new process enters queue  $Q_0$  which is served in RR
    - When it gains CPU, the process receives 8 milliseconds
    - If it does not finish in 8 milliseconds, the process is moved to queue  $Q_1$
  - At  $Q_1$  job is again served in RR and receives 16 additional milliseconds
    - If it still does not complete, it is preempted and moved to queue  $Q_2$



Multilevel feedback queues.

Aging is done via moving them do a different queue.

---

### Knowledge Check

- \_\_\_ scheduling is approximated by predictin the next CPU burst with an exponential average of the measured lengths of previous CPU bursts
  - d. SJF
- The \_\_\_ scheduling algo is designed especially for time-sharing systems
  - c. RR
- which of the following is ture of multilevel queue scheduling?
  - b. each queue has its own scheduling algorithm
  - the most general CPU-scheduling algo is multilevel feedback
  - processes can only move between queues using the feedback variant

## Multi-Processor Scheduling

CPU scheduling is more complex when multiple CPUs available.

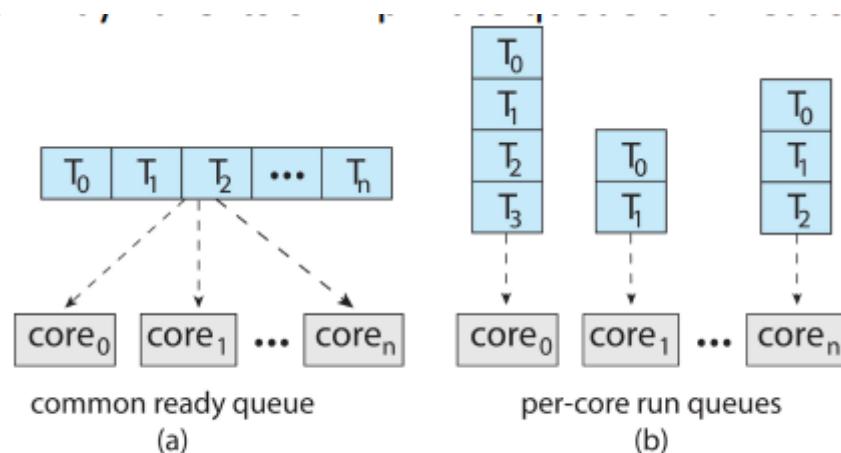
multiprocessor may be any one of the following architectures:

- multicore CPUs
- multithreaded cores
- NUMA systems
- heterogeneous multiprocessing

## Multicore

Symmetric Multi Processing

- each processor is self scheduling
- all threads may be in a common ready queue
- each processor may have its own private queue of threads

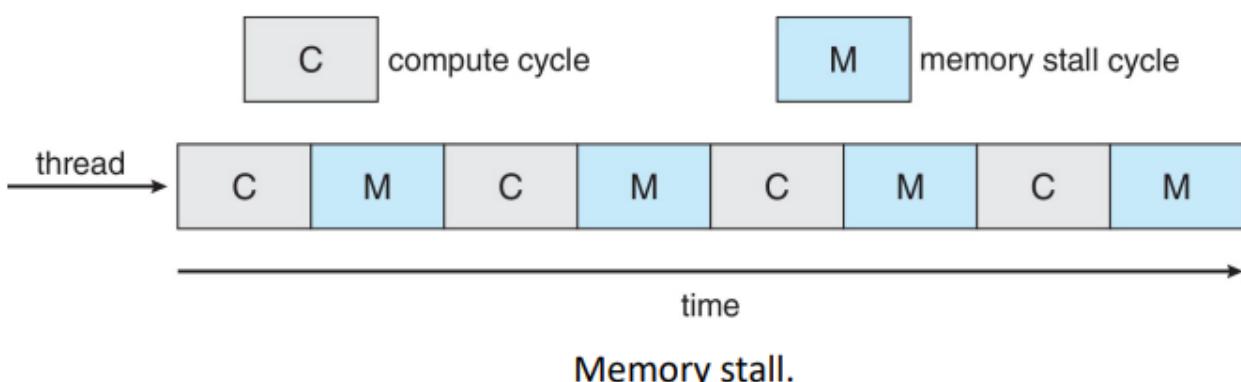


Organization of ready queues.

Lately we've been making multicore chips that are faster and less power hungry.

Multithreaded cores are also a recent advancement we've made.

This takes advantage of memory stall to make progress on another thread while memory retrieve happens.



What is Memory Stall?

When a processor accesses memory it waits around for the data to become available.

The processor is much faster than the memory so it has to wait on the memory.

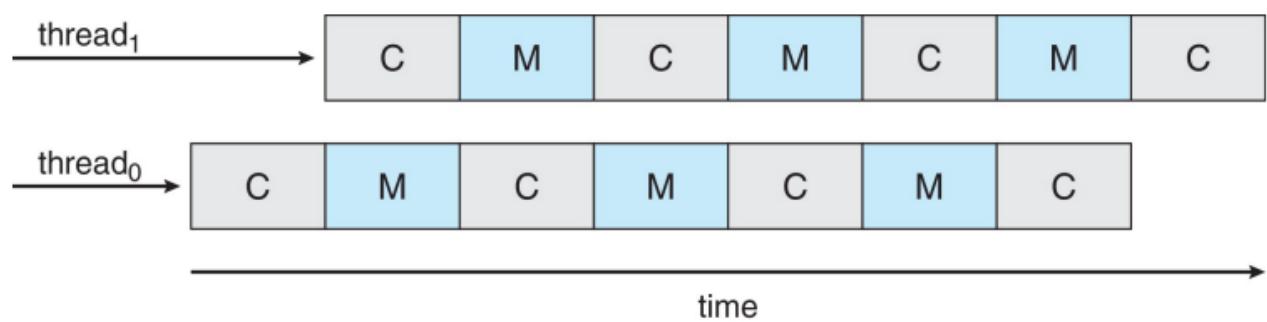
This can also occur during due to a cache miss

---

## Multithreaded Cores

Each core has > 1 hardware thread

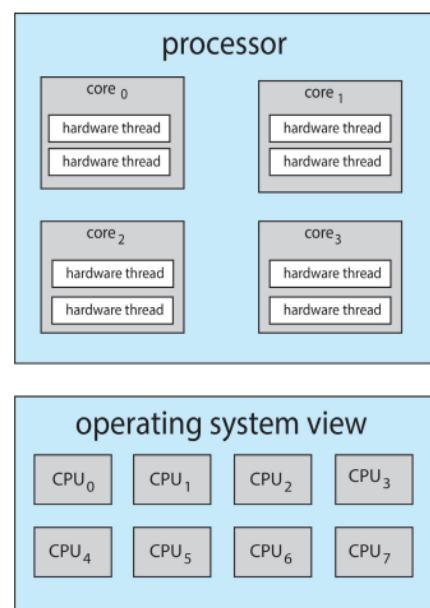
If one thread has a memory stall then we switch to another thread for 0 downtime.



Multithreaded multicore system.

## Multithreaded Multicore System

- **Chip-multithreading (CMT)** assigns each core multiple hardware threads. (Intel refers to this as **hyperthreading**.)
- On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.



Operating System Concepts, 10th Edition, by Silberschatz, Gagne, & Galvin

Chip multithreading.

each thread is basically treated as their own CPU from a logical perspective

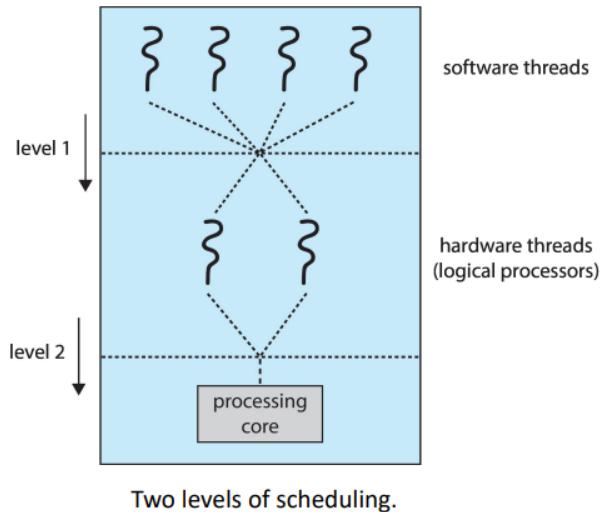
- Two levels of scheduling:

1. The operating system deciding which software thread to run on a logical CPU.

- Scheduling algorithms.

2. How each core decides which hardware thread to run on the physical core.

- RR, Priority, ...



Two levels of scheduling.

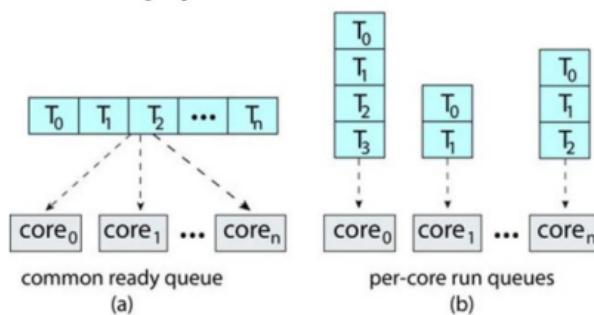
choose a chip then choose a thread

## Load Balancing

a concept for symmetric multiprocessing systems

Necessary only on systems where each processor has its **own private ready queue** of eligible threads to execute.

Unnecessary on systems with a **common run queue**: when a processor becomes idle, it immediately extracts a runnable thread from the common ready queue.



Organization of ready queues.

push migration

- periodic task checks load on each processor
- imbalance found → task pushed from overloaded CPU to other CPUs

pull migration

- idle processors pull a waiting task from a busy processor

## Processor Affinity

When a thread has been running on one processor, the cache contents of that processor stores the memory accesses by that thread.

Successive memory accesses by thread are often satisfied in cache memory (aka “warm cache”).

Prof uses an example of repetitive actions on a data structure, something like a for-loop iterating through an array or spreadsheet.

We start to move things to the cache to speed things up

If load balancing causes a thread to migrate to another processor:

- contents of cache memory must be invalidated for the 1st processor
- cache for 2nd processor must be repopulated
- **high cost**

This high cost is undesirable so OS and SMP try to avoid migrating a thread from one processor to another.

Processor Affinity:

- a process has an affinity for the processor on which it is currently running
- the warm cache is there for the process to take advantage of

Soft Affinity

- OS attempts to keep a thread running on the same processor
- no guarantees

Hard Affinity

- allows a process to specify a set of processor it may run on

initially the cache is empty (usually almost all 16 kilobytes)

when the process starts it misses over and over again but then starts filling out the cache

processor starts prepopulating the cache in order to help it out.  
seeing what to remove and what to get rid of.

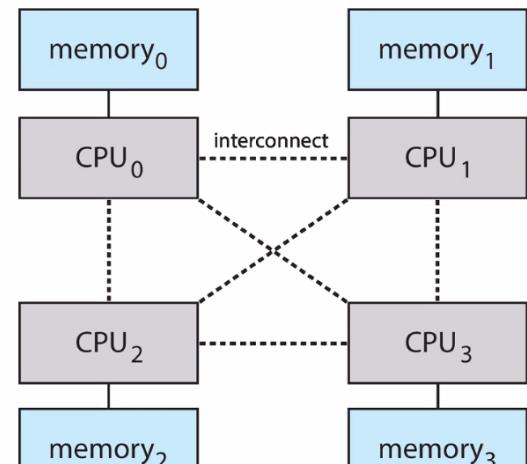
for hard affinity the process decides on a number of CPUs.

It can choose CPUs to benefit from a shared L2 cache

# NUMA Systems

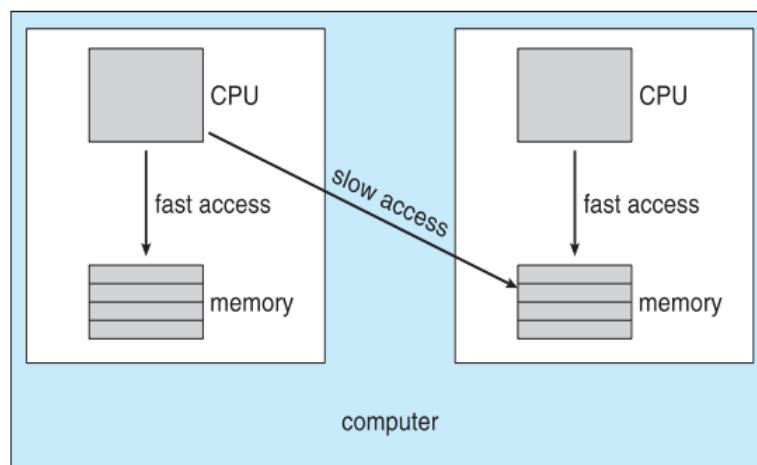
## Recall – Week 1 - Non-Uniform Memory Access System (NUMA)

- Each CPU (or group of CPUs) has its own local memory that is accessed via a small, fast local bus.
- The CPUs are interconnected by a shared system, so that all CPUs share one physical address space.



NUMA multiprocessing architecture.

If the operating system is NUMA-aware, then it will allocate memory closest to the CPU where the thread is running, thus providing the thread the fastest possible memory access -> processor affinity.



NUMA and CPU scheduling.

## Heterogeneous Multiprocessing

so far we have assumed all processors are identical so we can allow any thread to run on any processing core.

We only differed based on load balancing, processor affinity policies, and NUMA systems

Some mobile systems are now designed using cores that run the same instruction set but vary in terms of their clock speed and power management.

They can also adjust the power draw of a core to the point of idling.

These are known as Heterogeneous Multiprocessing (HMP).

This type of architecture is known as **big.LITTLE**:

- **big** cores are higher-performance cores
  - greater power draw → used for short periods of time
  - interactive tasks
- **LITTLE** are energy efficient cores
  - use less energy → used for longer periods of time
  - background tasks

Windows 10 supports HMP scheduling by allowing a thread to select a scheduling policy that best supports its power management

---

knowledge check

- \_\_\_ allows a thread to run on only one processor
  - a. processor affinity
- two general approaches to load balancing are \_\_\_ and \_\_\_
  - d. push migration, pull migration
  - coarse grained is when a process keeps executing on a core until there's a huge stall event
  - fine grained is when we are switching on a very fine level of granularity between processes
    - the boundary of an instruction cycle
      - fetch and execute then now we switch to another thread
    - each comes with their pros and cons with relation to process switching

## Real-Time CPU Scheduling

can present special issues

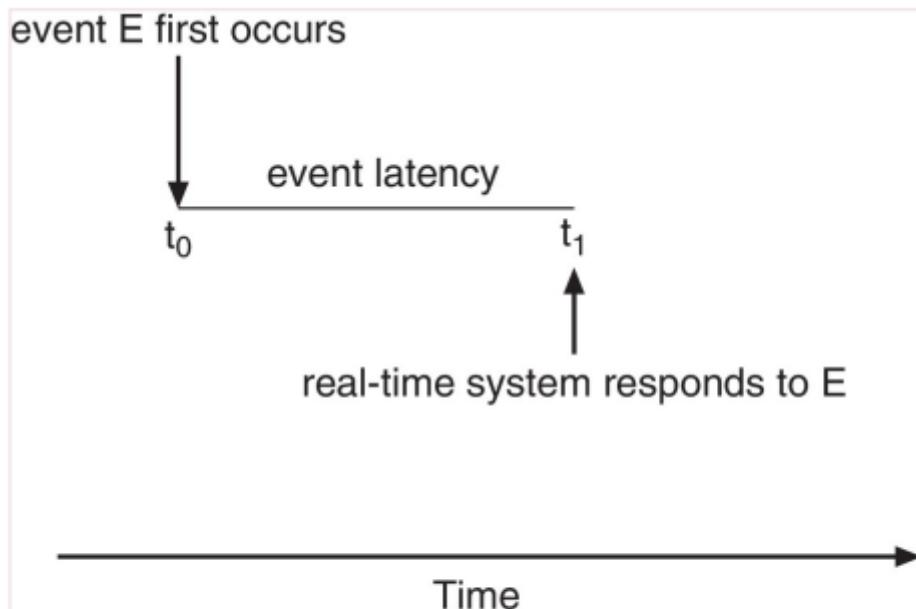
2 types

- soft real-time systems
  - critical real-time tasks have the highest priority
  - no guarantee as to when tasks will be scheduled
- hard real-time systems
  - task must be serviced by its deadline
  - service after deadline has expired
    - same as no service at all

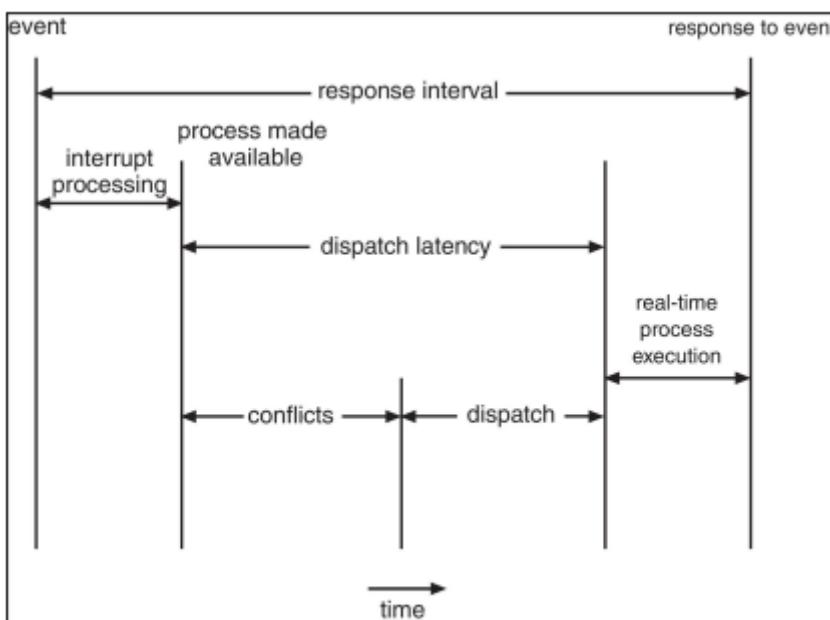
event latency

- amount of time that elapses from when an event occurs to when it is serviced
- 2 types that affect performance
  - interrupt
  - dispatch

- time for schedule to take current process off CPU and switch to another



**Event latency.**



**Dispatch latency.**

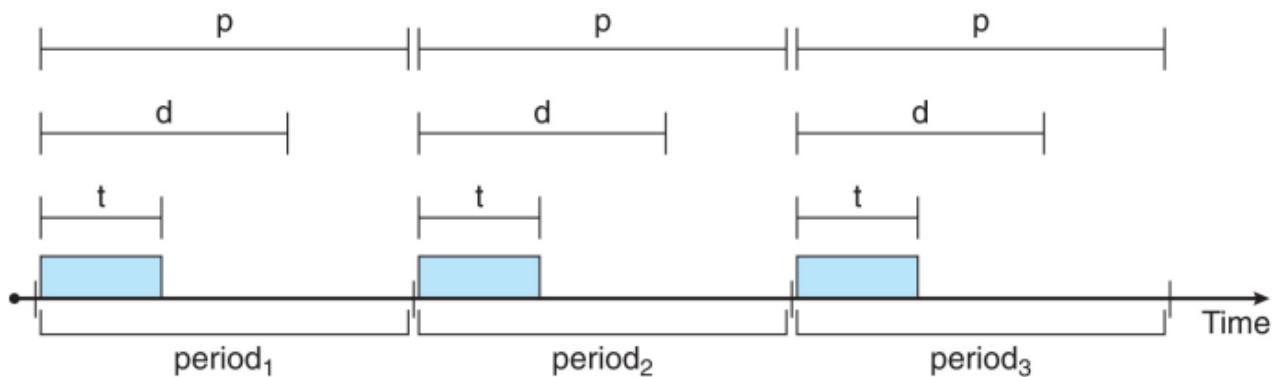
conflict phase of dispatch latency:

1. preemption of any process running in kernel mode
2. release by low-priority process of resources needed by high-priority processes

## Priority Based Scheduling

for soft real-time scheduling, scheduler must support preemptive, priority-based scheduling.

For hard real-time we must also provide the ability to meet deadlines which requires additional scheduling features.



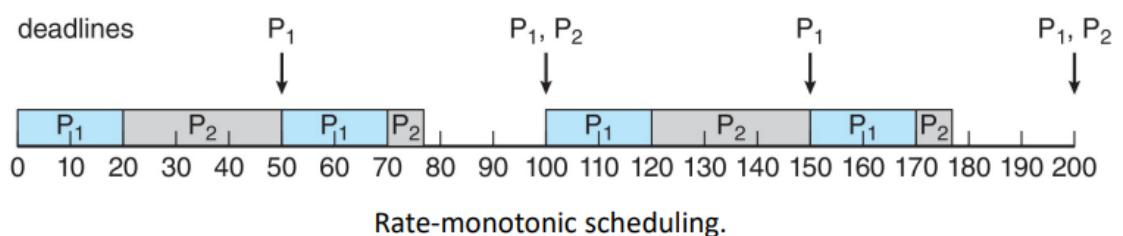
Periodic task.

Processes have new characteristics: **periodic** ones require CPU at constant intervals.

- Has a fixed processing time  $t$ , deadline  $d$  by which it must be serviced by, and a period  $p$ .
- $0 \leq t \leq d \leq p$
- **Rate** (aka **frequency**) of periodic task is  $1/p$
- **Utilization (u)** =  $t/p$

## Rate Monotonic Scheduling (RMS)

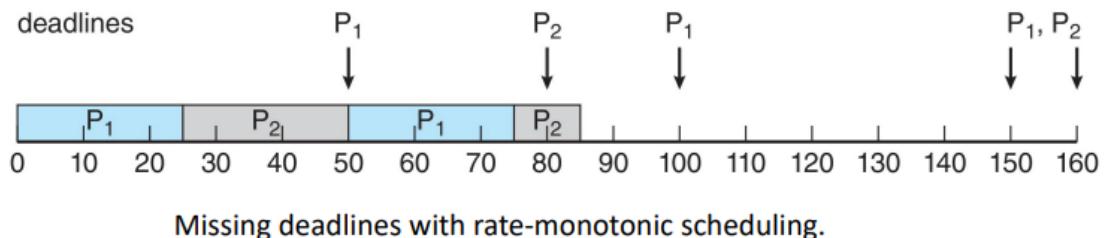
- A priority is assigned based on the inverse of its period.
  - Shorter periods = higher priority;
  - Longer periods = lower priority
- Example:
  - $P_1=50, t_1=20, P_2=100, t_2=35$
  - $P_1$  is assigned a higher priority than  $P_2$  as it has the shorter period.
  - $u_1=20/50 = 0.4, u_2=35/100 = 0.35 \rightarrow u_t = 0.75$



Rate-monotonic scheduling.

# Missed Deadlines with Rate Monotonic Scheduling

- $P_1 = 50, t_1 = 25, P_2 = 80, t_2 = 35$
- $P_1$  is assigned a higher priority than  $P_2$  as it has the shorter period.
- $u_1 = 25/50 = 0.5, u_2 = 35/80 = 0.44 \rightarrow u_t = 0.94$



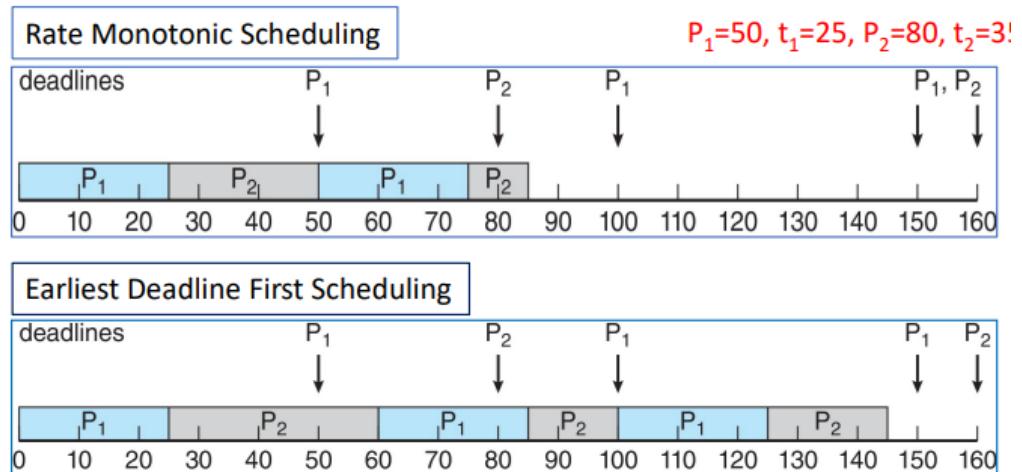
- Process P2 misses finishing its deadline at time 80.

# Missed Deadlines with Rate Monotonic Scheduling

- Max utilization when having N processes in the system:  
 $N(2^{1/N} - 1)$
- With two processes, CPU utilization is bounded at about 83%.
- With one process in the system, CPU utilization is 100%, but it falls to approximately 69% as the number of processes approaches infinity.

# Earliest Deadline First Scheduling (EDF)

- Priorities are assigned dynamically according to deadlines:
  - the earlier the deadline, the higher the priority;
  - the later the deadline, the lower the priority.



knowledge check

- the rate of a periodic tasks in a hard real-time system is \_\_\_, where p is a period and t is the processing time.
  - a. 1/p
- which of the following is true of the rate-monotonic scheduling algo?
  - a is for first scheduling
  - rate-monotonic scheduling does not have a dynamic priority policy
  - it has to be preemptive
  - answer: c. cpu utilization is bounded when using this algo
- which of the following is true of earliest-deadline-first (EDF) scheduling algo?
  - the rule for b is reversed
  - priorities aren't assigned statically
  - priorities are not fixed
  - answer: a. when a process becomes runnable, it must announce its deadline

## Algo Eval

How to select CPU-scheduling algo for an OS?

Determine criteria, then eval algos

example of criteria:

- maximizing utilization under the constraint that the maximum response time is 300 milliseconds
- maximizing throughput such that turnaround time is (on average) linearly proportional to total execution time

## Deterministic

Deterministic Modelling:

- type of analytic evaluation
- takes a particular predetermined workload and defines the performance of each algo for that workload

**Example: Consider 5 processes arriving at time 0:**

| Process | Burst Time |
|---------|------------|
| $P_1$   | 10         |
| $P_2$   | 29         |
| $P_3$   | 3          |
| $P_4$   | 7          |
| $P_5$   | 12         |

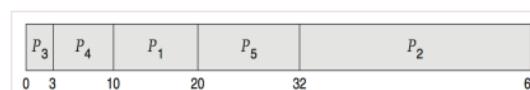
For each algorithm below, calculate minimum average waiting time,  $q=10\text{ms}$ .

Simple and fast, but **requires exact numbers for input**, applies only to those inputs.

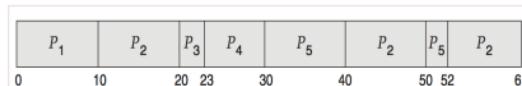
- FCFS is 28ms:



- Non-preemptive SJF is 13ms:



- RR is 23ms:

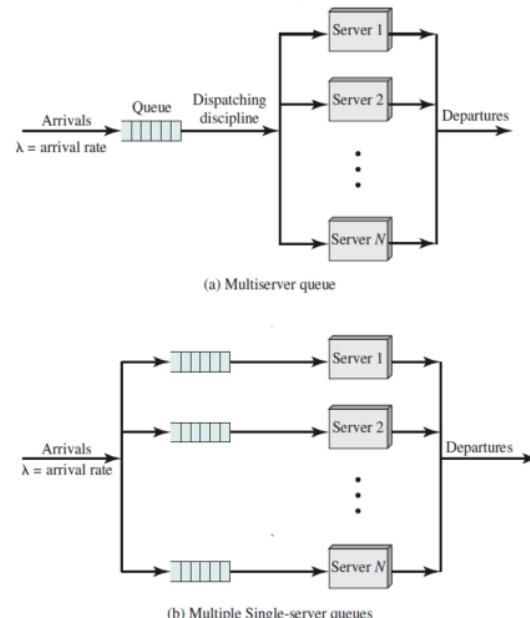


| Process | Burst Time |
|---------|------------|
| $P_1$   | 10         |
| $P_2$   | 29         |
| $P_3$   | 3          |
| $P_4$   | 7          |
| $P_5$   | 12         |

## Queueing Models

# Queueing Models

- Describes the arrival of processes, and CPU and I/O bursts probabilistically.
- Computer system described as network of servers, each with queue of waiting processes
  - Knowing arrival rates and service rates
  - Computes utilization, average queue length, average wait time, etc...



## Little's Law

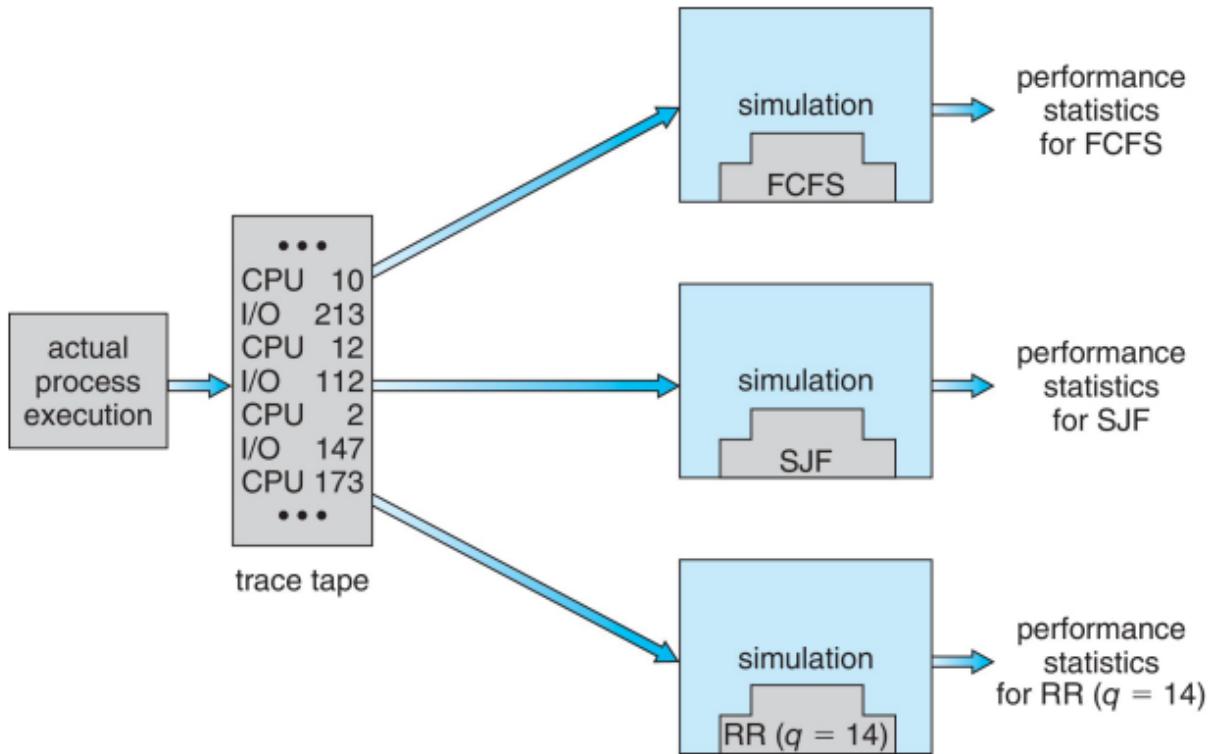
- $n$  = average queue length
- $W$  = average waiting time in queue
- $\lambda$  = average arrival rate into the queue
- Little's law – in steady state, processes leaving the queue must equal processes arriving, thus:  
$$n = \lambda \times W$$
  - Valid for any scheduling algorithm and arrival distribution
  - Non-preemptive scheduling algorithms only.
- For example, if on average 7 processes arrive per second, and normally 14 processes in the queue, then the average wait time per process = 2 seconds.

## Simulations

Queueing models are complicated and limited.

Simulations give more accurate evaluation of scheduling algos:

- programmed model of computer system
- clock is a variable
- gather stats indicating algo performance
- data to drive sim gathered via:
  - rng according to probabilities
  - dists defined mathematically or empirically
  - trace files record sequences of real events in real systems



I did not fucking hear anything she said about this

## Real Implementation

The best way is to just implement this stuff for real.

This is very costly.

In general, most flexible scheduling algos are those that can be altered by the system managers or by the users so that they can be tuned for a specific (set of) app(s)

Environments can vary however.

### Knowledge Check

- in Little's formula,  $\lambda$ , represents the \_\_\_:
  - b. average arrival time for new processes in the queue
- 

## Main Memory

## Background for Memory Management

recap:

- a program must be brought from the disk into the memory and placed within a process for it to be run
- main memory and registers are the only storage the CPU can access directly
  - accessing the non-volatile memory is limited to copying over data to the memory to work with
  - we don't work directly off of disk

a memory unit only sees a stream of either:

- addresses & read requests
- addresses & data and write requests

register access is done in one CPU clock cycle.

main memory can take many cycles, causing a stall where the CPU has to wait on the memory.

cache sits b/w main mem & CPU regs for faster access than mem

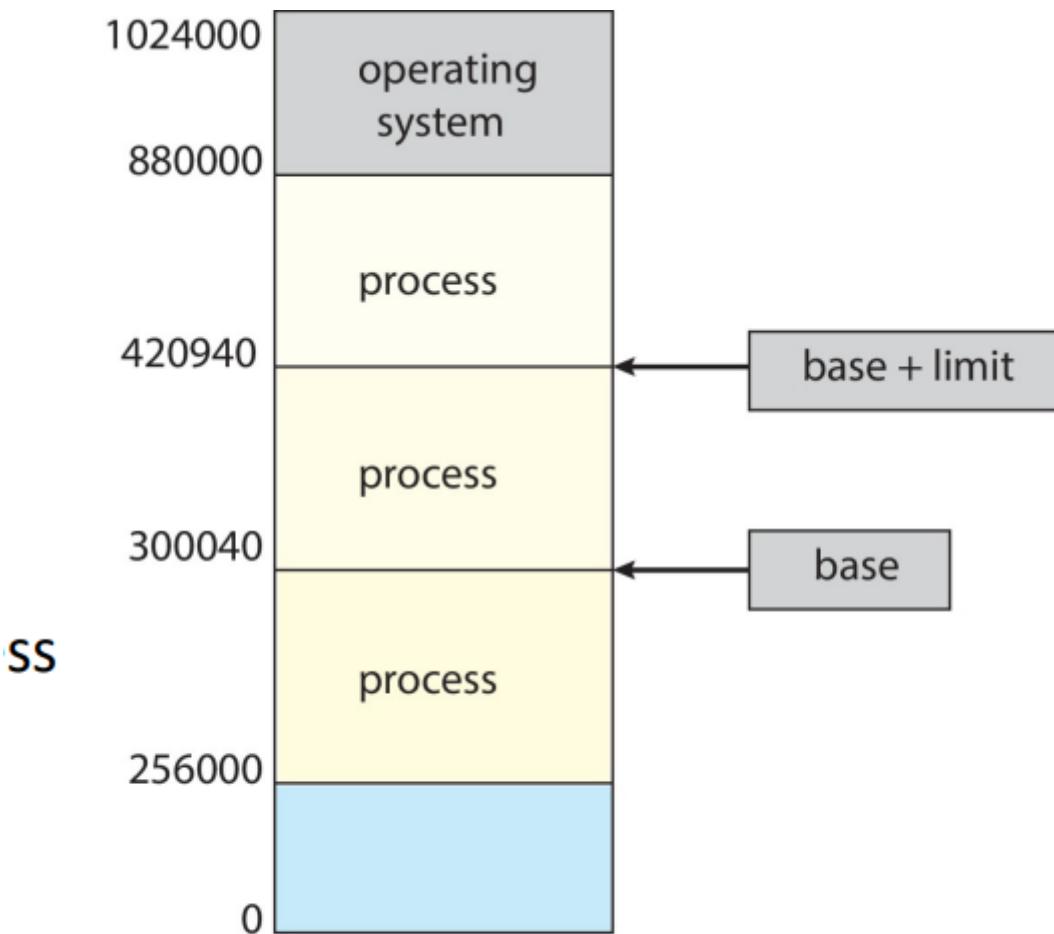
- principle of locality
  - whenever a task is being done on a specific collection of data, it is very likely that we're going to select the next piece of data in the collection so we may fetch it preemptively as we anticipate that it will be required next
  - hit ratio keeps increasing
  - at some point we start doing something else however

## Memory Protection

Protection is required to ensure correct operation. **recall: mutual exclusion**

Make sure that a process can only access its allocated area in the address space.

We can use a pair of **base** and **limit** registers to define the logical address space of a process



A base and a limit register define a logical address space.

---

**base** = smallest legal physical mem address

**limit** = size of range

**base+limit** = highest legal physical mem address

trying to access anything outside of the range is not allowed.

In the above we see a base of 300040 and a limit of 120900 added together to get the highest legal address of 420940.

prof brought up the fact that modern operating systems sit in the highest addresses while older operating systems may do other things like sitting at the oldest

prof will designate which one is which on assessment(s)

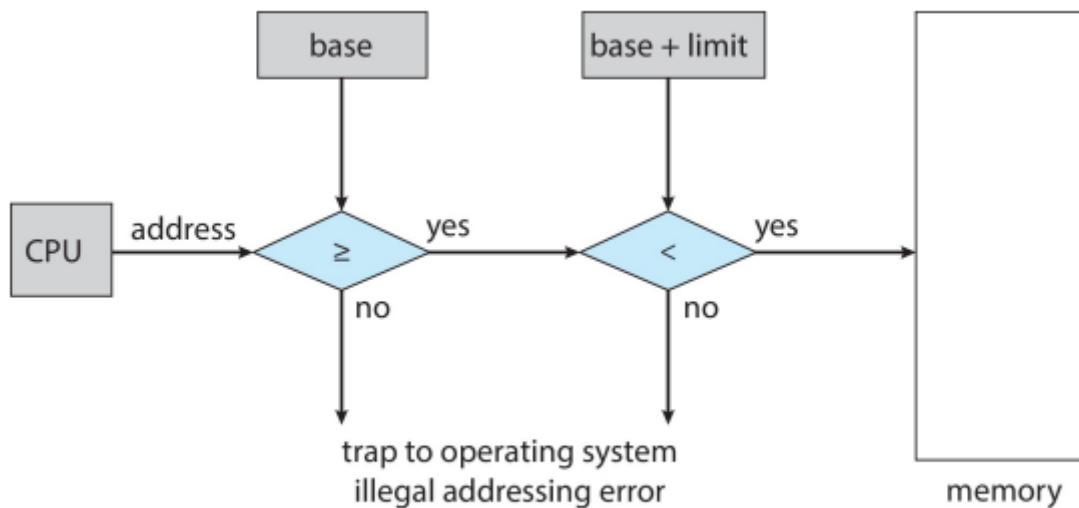
Q: Are all the ranges mutually exclusive?

A: Yes. All the ranges are mutually exclusive

## Hardware Address Protection

- CPU has to check every memory access request generated in user mode to make sure it's inside of the range

- the instructions to loading the base and limit registers are privileged

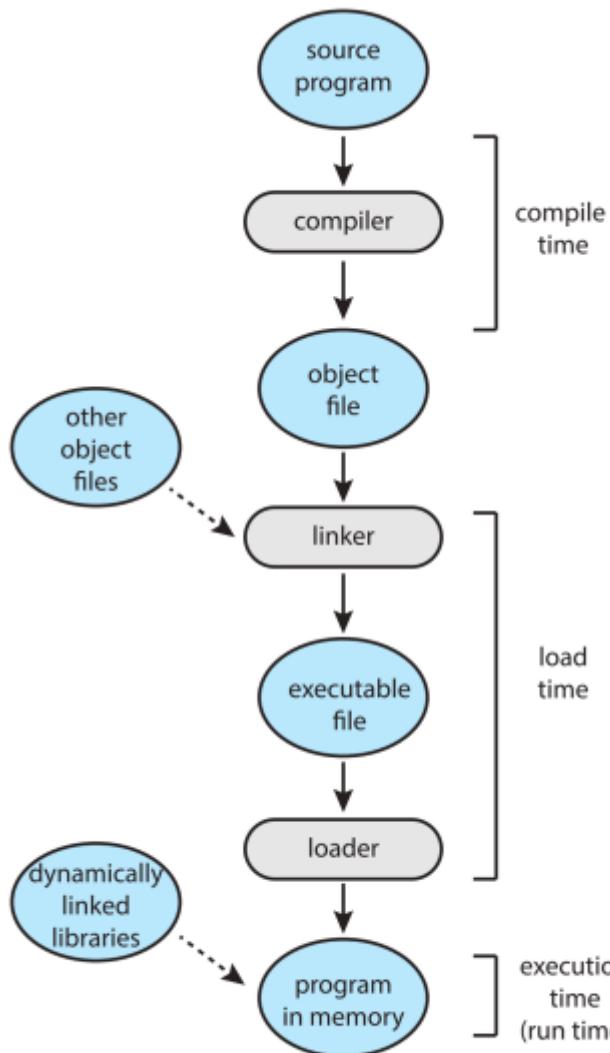


### Hardware address protection with base and limit registers.

malicious programs often try to break out of this limit.

### Address Binding

- programs are stored on disk, ready to be brought into memory to execute from an input queue
  - most systems allow a user process to reside in any part of the physical memory
  - the address space of the computer may start at 00000 but the first address of the user process does not need to be 00000
    - it's inconvenient to enforce it to be 00000
      - why?
      - you can waste memory as you get closer to the OS in mem space
      - we have processes starting and finishing all the time so we have to fit them in wherever we can



## Multistep processing of a user program.

- we start out with the source code
- compile it
- link with the object file
- turn it into an executable
- dynamically linked libraries are provided by the OS and are linked by the loader
- program gets swapped in and swapped out as it needs the libraries

addresses are represented in different ways at different stages of a program's life:

- source code addresses are usually symbolic (`int x`)
- compiled code addresses bind to relocatable addresses
  - “14 bytes from beginning of this module”
- after this point we have to have actual concrete physical addresses
- linker or loader will bind relocatable addresses to absolute addresses
  - ex: 74014
- each binding maps one address space to another

Address binding of instructions and data to memory addresses can happen at three different stages:

1. **Compile time:** If you know at compile time where the process will reside in memory, **absolute code** can be generated.
  - Must **recompile** code if starting location changes.

it's very unlikely that you will know at compile time where the process will reside in memory.

- . **Load time:** If it is not known at compile time where the process will reside in memory, then the compiler must generate **relocatable code**.

Must **reload** code if starting location changes.

- Execution time: If the process can be moved during its execution from one memory segment to another, then binding must be delayed until run time.

Need hardware support for address maps (e.g., base and limit registers)

Most OSs use this method.

Most OSes do this binding during execution time.

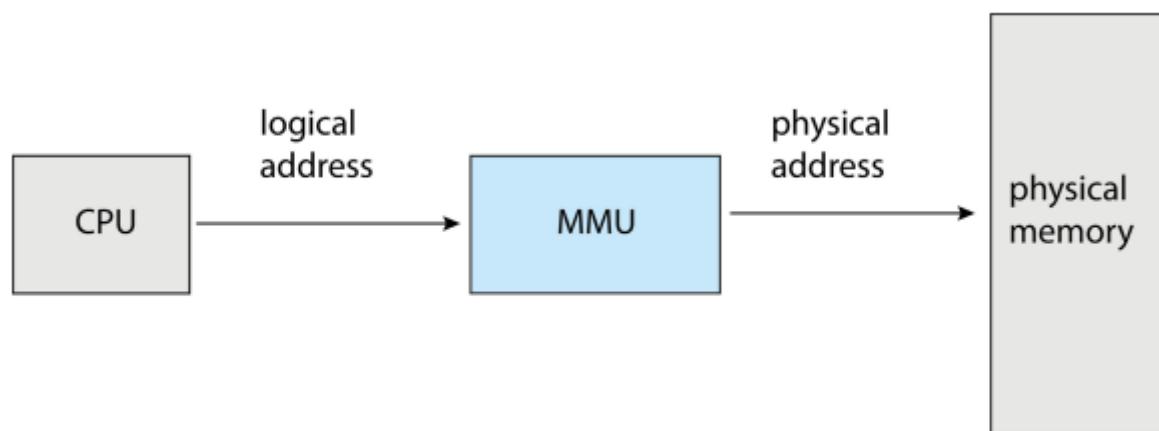
## Logical vs Physical Address Space

Logical Address

- generated by cpu
- aka virtual address

Physical Address

- address seen by the memory unit



logical address space is bound to the physical address space.

The CPU does not directly see/access the physical addresses, only the logical.

Both types of addresses are the same in compile-time and load-time address-binding schemes.

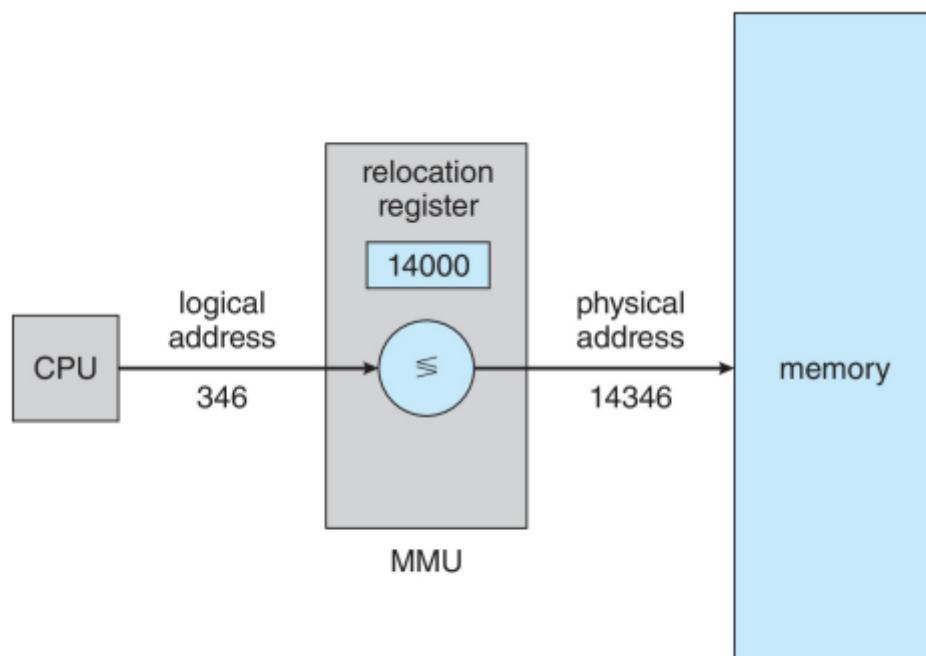
They differ in the execution-time address-binding schemes (so most OSes)

The addresses are all generated by a program by refer to different address spaces (logical or physical)

## Memory Management Unit

The memory management unit (MMU) is a hardware device that maps virtual to physical addresses at run time.

many methods to do this mapping



Dynamic relocation using a relocation register.

simple scheme:

- generalization of the base-register scheme
- base register is now called relocation register
- value of in the relocation register is added to every address generated by a user process at the time it is sent to memory
- user program deals with logical addresses
  - never sees real physical addresses
- execution-time binding occurs when a reference is made to a location in memory
- memory-mapping hardware converts logical addresses into physical addresses

## Dynamic Loading

sometimes the program is larger than the size of memory

with dynamic loading, a routine is not loaded until it is called

- better memory-space util
  - unused routine never loaded
- all routines kept on disk in relocatable load format

this is very useful when large amounts of code are needed to handle infrequently occurring cases. **heavy, memory-hungry edge cases**

There is no special support from the OS required

- implement via program design
- OS *can* help by providing libraries to implement dynamic loading

## Dynamic Linking

Dynamically Linked Libraries (DLLs) are system libraries that are linked to user programs when they run (during execution time)

DLLs are also known as shared libraries

- used extensively in windows and linux systems
- small pieces of code (stub) used to locate the appropriate memory-resident library routine
- stub replaces itself with the address of the routine and executes the routine

Dyn. Linking and DLLs require help from the OS

- if the processes in mem are protected from one another, OS is the only one that can check to see that
  - whether the needed routines are in another process' memory space
  - can allow multiple processes to access the same memory address
- w/o this each program needs to include its own copy of its language library

Static linking

- used only by older computers
- system libraries and program code combined by the loader into the binary program image

---

### Knowledge Check

- absolute code can be generated for \_\_\_\_
  - a. compile-time binding
  - load-time binding can gen relocatable
- In a dynamically linked library, \_\_\_\_
  - d. a stub is included in the image for each library-routine reference
- the mapping of a logical address to a physical address is done in hardware by the \_\_\_\_

- a. memory-management-unit (MMU)

# Contiguous Memory Allocation

Contiguous means continuous. No broken pieces. Only one large piece of memory supports OS and user processes.

Contiguous allocation is an early method of memory management.

Memory is usually divided into 2 partitions:

- OS partition
  - usually held in high memory with interrupt vector
  - used to be in the lowest vector in older systems
- User processes held in low memory
  - each process contained in single contiguous section of memory.

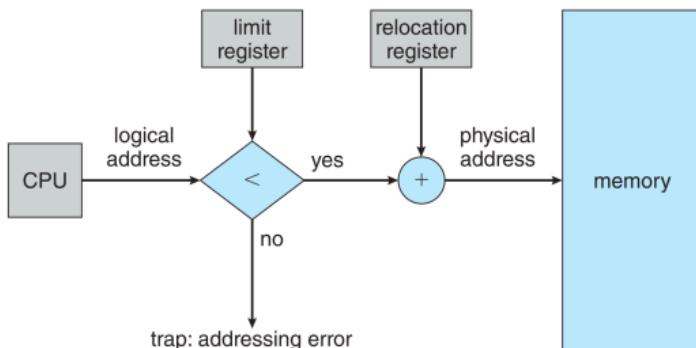
## Memory Protection in Contiguous Memory Allocation

### Relocation Registers

- protect user processes from each other
- prevent user processes from changing operating-system code and data
- base contains value of smallest physical address
- limit contains range of logical addresses
  - example: relocation = 100040 and limit = 74600
- MMU maps logical address dynamically
- relocation-register scheme allows OS' size to change dynamically

When the CPU scheduler selects a process for execution, the dispatcher loads the relocation and limit registers with the correct values as part of the context switch.

We can protect both the operating system and the other users' programs and data from being modified by this running process.

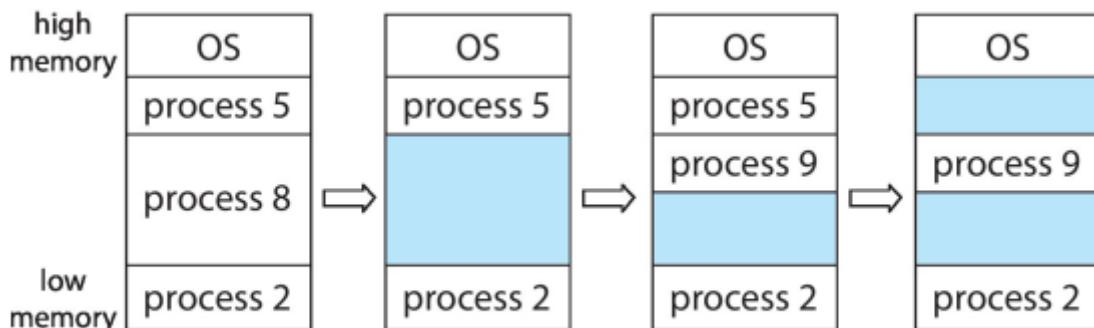


Hardware support for relocation and limit registers.

## Memory Allocation

a simple method to allocate memory is to use variably sized partitions

- used for efficiency
- hole
  - block of available memory
  - holes of various size are scattered throughout memory



Variable partition.

- Process 8 finishes and leaves a hole b/w processes 5 and 2.
- we allocate process 9 and it fits into the hole
- process 5 finishes and makes another hole

Holes are a problem that we need to deal with later on.

When a process arrives, it is allocated memory from a hole large enough to accommodate it.

A process exiting frees its partition.

Adjacent free partitions are combined.

OS maintains info about

- allocated partitions
- free partitions (holes)

How do we satisfy an allocation request of size  $n$  from a list of free holes?

fitting approaches:

- first-fit
  - allocate the first hole that is big enough
- best-fit
  - allocate the smallest hole that is big enough
  - must search entire list unless ordered by size
  - produces smallest leftover hole
- worst-fit
  - allocate the largest hole
  - must search entire list unless ordered by size
  - produces the largest leftover hole
    - may be more useful than the smaller leftover hole from a best-fit approach
    - not so much smaller fragments of memory
    - usually not used

- next-best-fit
  - not used in the course
  - we find the best-fit then take the next one
  - prof thought it was cool so she brought it up

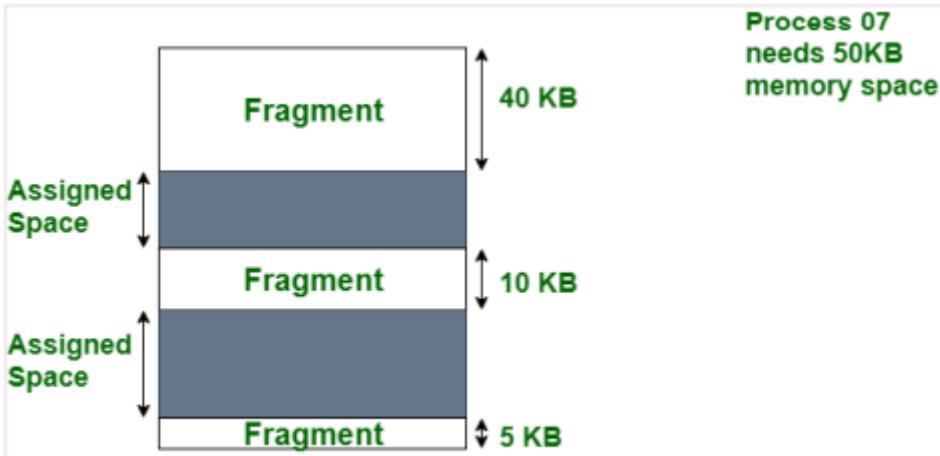
simulations show

- first-fit and best-fit are better than worst-fit in terms of
  - decreasing time util
  - decreasing storage util
- neither is clearly better than the other in terms of storage util
- first-fit is generally faster than best-fit

## Fragmentation

External Fragmentation

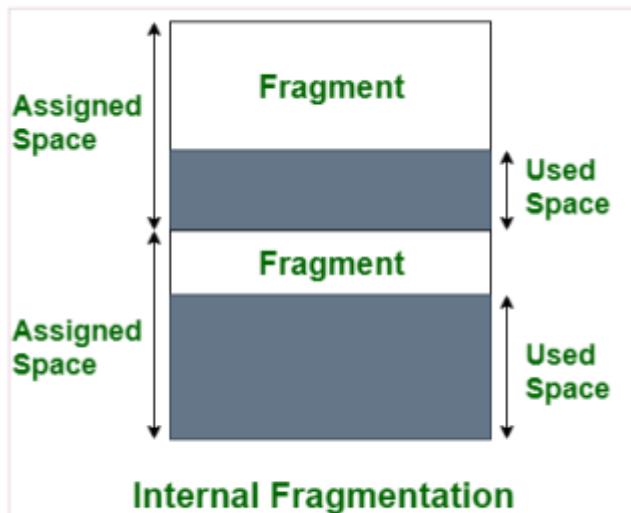
- outside the process
- total memory space exists to satisfy a request but not contiguous
  - the holes are not within an assigned space of a process



### External Fragmentation

Internal Fragmentation

- allocated memory may be slightly larger than requested memory
- size difference is memory interla to a partition but not being used



### Internal Fragmentation

### Internal Fragmentation

The processes in the above don't use up all of their allocated memory

given  $N$  blocks allocated,  $N/2$  blocks will be lost to (internal and external) fragmentation.

- 1/3 ? may be unusable -> 50% rule

Reduce external fragmentation by compaction:

- shuffle memory contents to place all free memory together in one large block
- requires
  - dynamic relocation
  - done at execution time
    - not compile or load time

Fragmentation applies to more than just memory and can occur wherever we manage blocks of data (disk, external storages, etc.)

The simplest compaction is algo is to move all processes toward one end of memory

- all holes move in the other direction to make one large hole
- can be very expensive

We will discuss other strategies in the following 2 sections

---

Knowledge Check:

- \_\_\_ is the dynamic storage-allocation algo which results in the smallest leftover hole in memory.
  - b. best fit
- \_\_\_ is the dynamic storage-allocation algorithm which results in the largest leftover hole in memory
  - c. worst fit
- which of the following is true of compaction
  - c. it is possible only if relocation is dynamic and done at execution time

## Paging

Permitting the physical address space of processes to be non-contiguous in order to solve the external-fragmentation problem.

This still has internal fragmentation.

set up a **page table** to translate logical to physical addresses

- divide physical memory into fixed-sized blocks called **frames**
  - size is a power of 2
    - between 512 bytes and 16 Mbytes
- divide logical memory into fixed-sized blocks called **pages**
  - size is a power of 2
    - typically either 4 kb or 8kb in size

These aren't equivalent to one another so frames can be different sizes as pages.

We can eventually use paging to load processes that are too big for our memory.

- keep track of all free frames

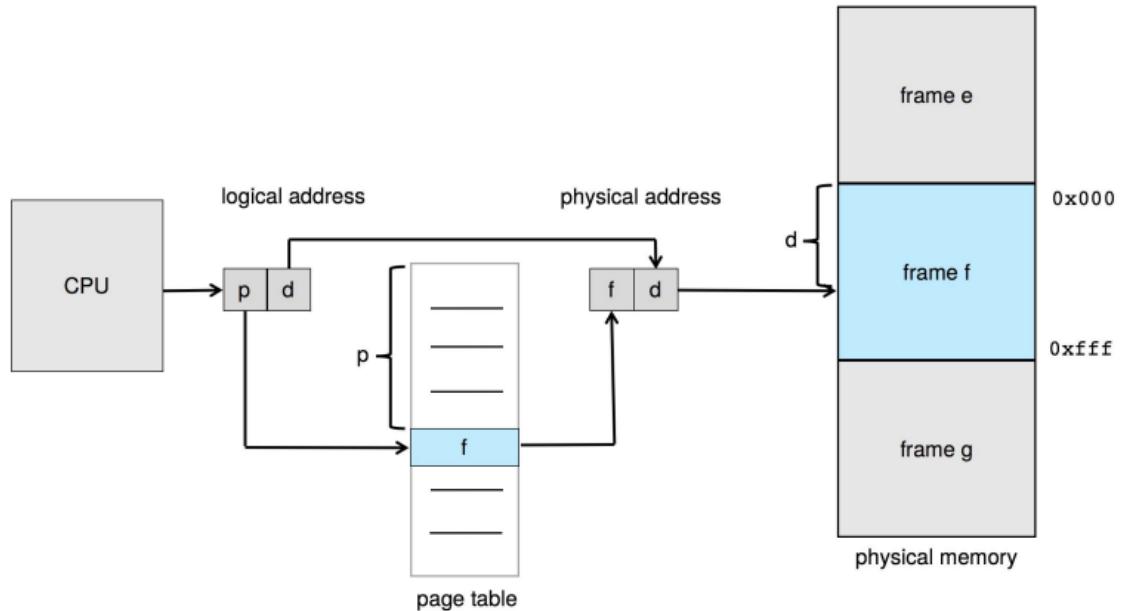
to run a program of size  $N$  pages, need to find  $N$  free frames and the program.

Address generated by CPU divided into:

- p - page number
  - index into page table
  - contains base address of each page in physical memory
- d - page offset

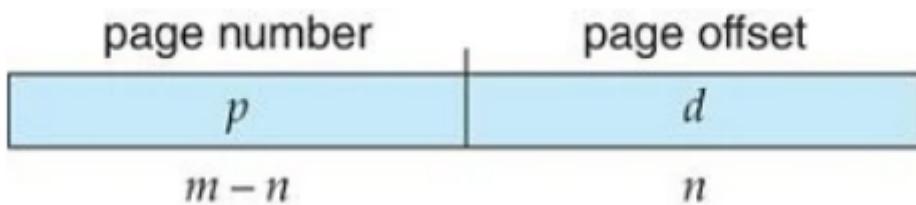
- combined with base address to define the physical memory address that is sent to the memory unit

# Paging Hardware



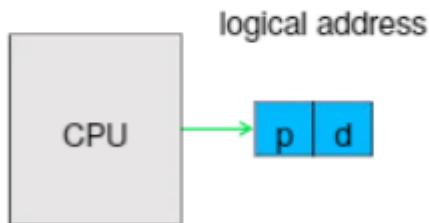
page size:

- defined by hardware
  - like frame size
- size is a power of 2
  - makes translation of logical address to page number and page offset easier this way
- if the size of the logical address space is  $2^m$  and page size is  $2^n$ 
  - the high-order  $m - n$  bits of a logical address designate the page number
  - $n$  low-order bits designate the page offset

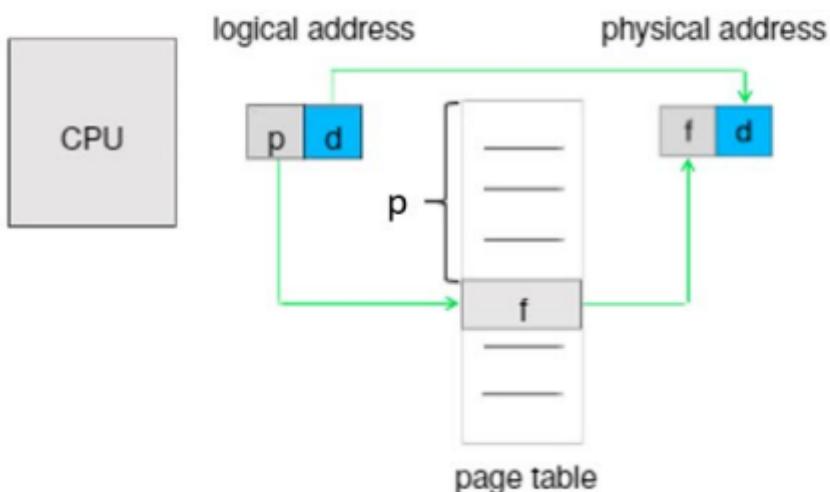


# Paging Hardware

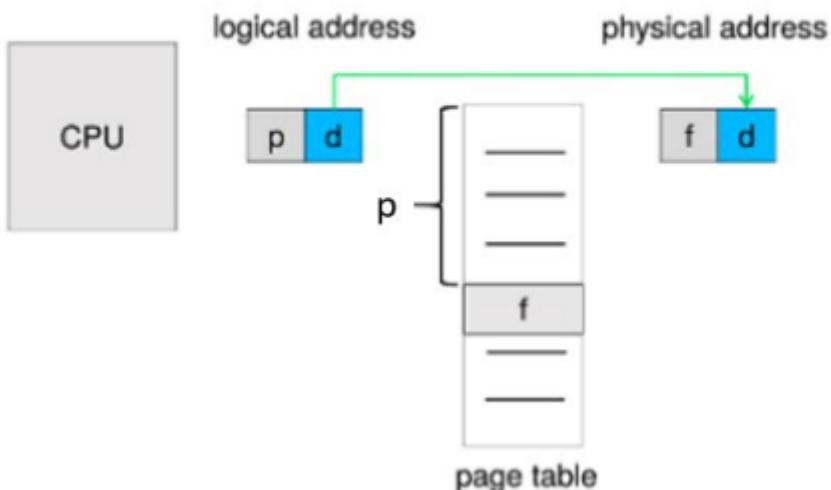
- the logical address generated by the CPU is sent to the MMU where it is divided into a page number (p) and an offset (d)
  - the number of bits in each part depends on the page size



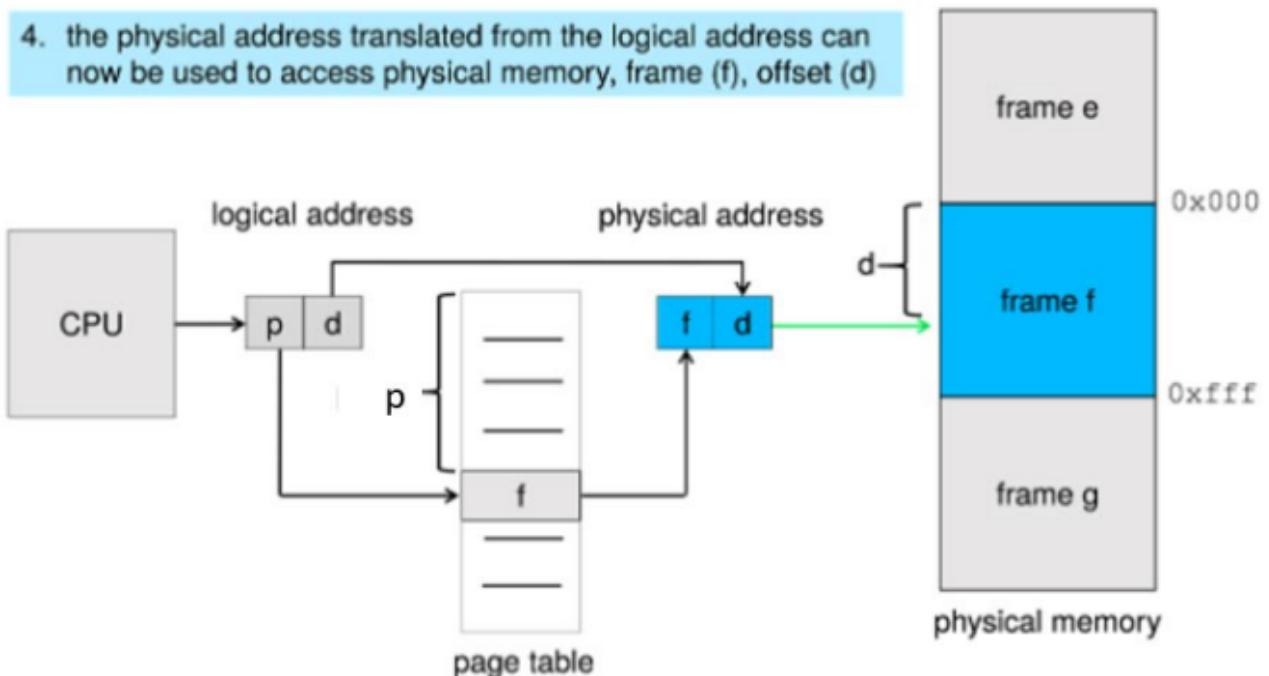
- the page number is used as an index into the page table
  - the entry in the page table at that index is the number of the frame of physical memory containing the page
  - that frame number is the first part of the physical memory address



3. the offset (d) within the page is the same as the offset within the frame, so it is retained, together with the frame number forming the physical address



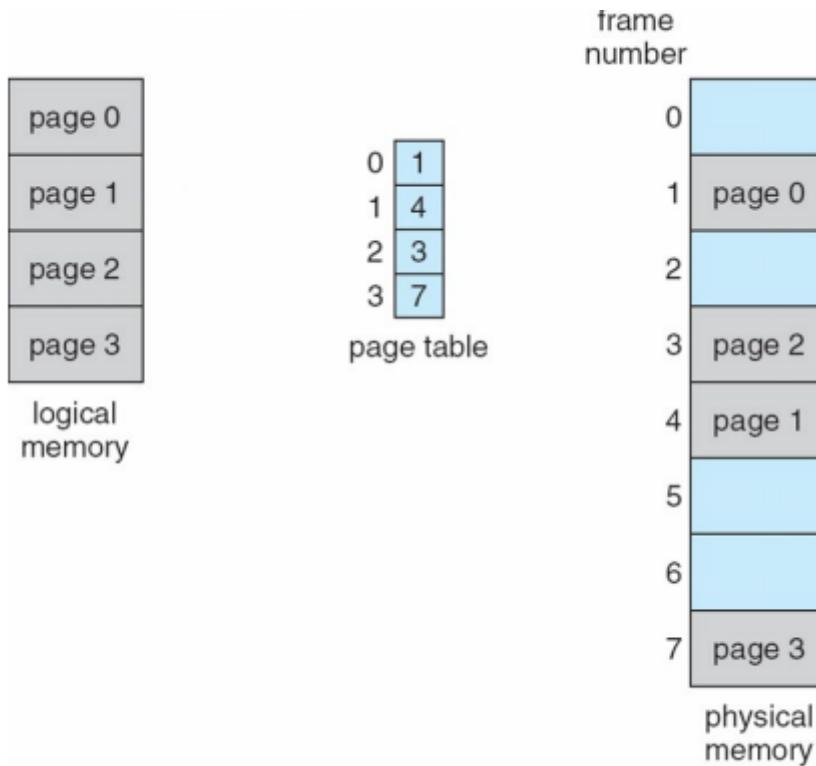
4. the physical address translated from the logical address can now be used to access physical memory, frame (f), offset (d)



steps the MMU takes to translate a logical address gen'd by cpu to a phys address

- extract page number **p** and use it to index into page table
- get corresponding frame number **f** from page table
- replace **p** in the logical address with **f** to get physical address

the offset **d** is unchanged and unreplace, combining with **f** to get the physical address



### Paging model of logical and physical memory.

Logically we see the pages as being contiguous but physically they are all over the place and in a different order.

our paging scheme makes it so that there is no external fragmentation since any free frame can be allocated to a process that needs it.

## Example

given

- $n = 2$
- $m = 4$

map the following logical addresses to physical addresses:

- 0
- 3
- 4
- 13

logical memory

|    |   |
|----|---|
| 0  | a |
| 1  | b |
| 2  | c |
| 3  | d |
| 4  | e |
| 5  | f |
| 6  | g |
| 7  | h |
| 8  | i |
| 9  | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

page table

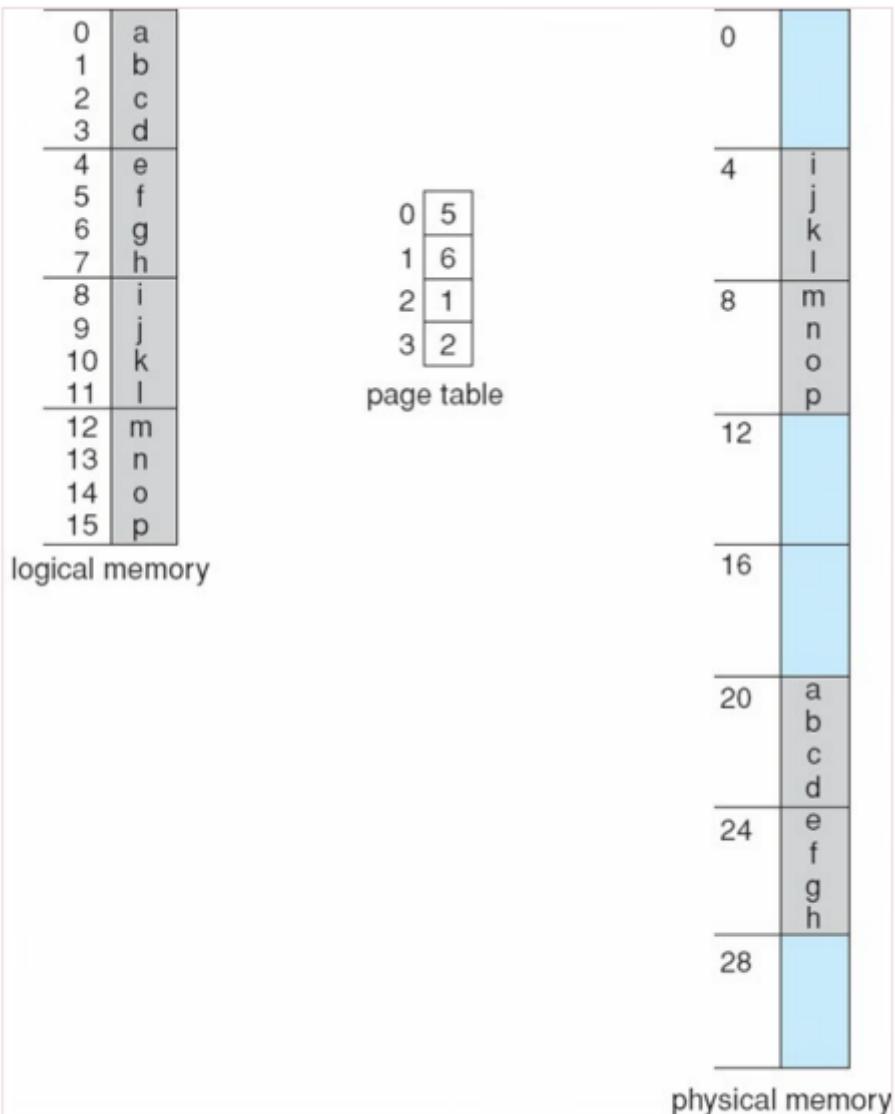
|   |   |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

using a page size of  $2^n = 4$  bytes and a physical memory of  $2^m = 32$  bytes (8 pages)

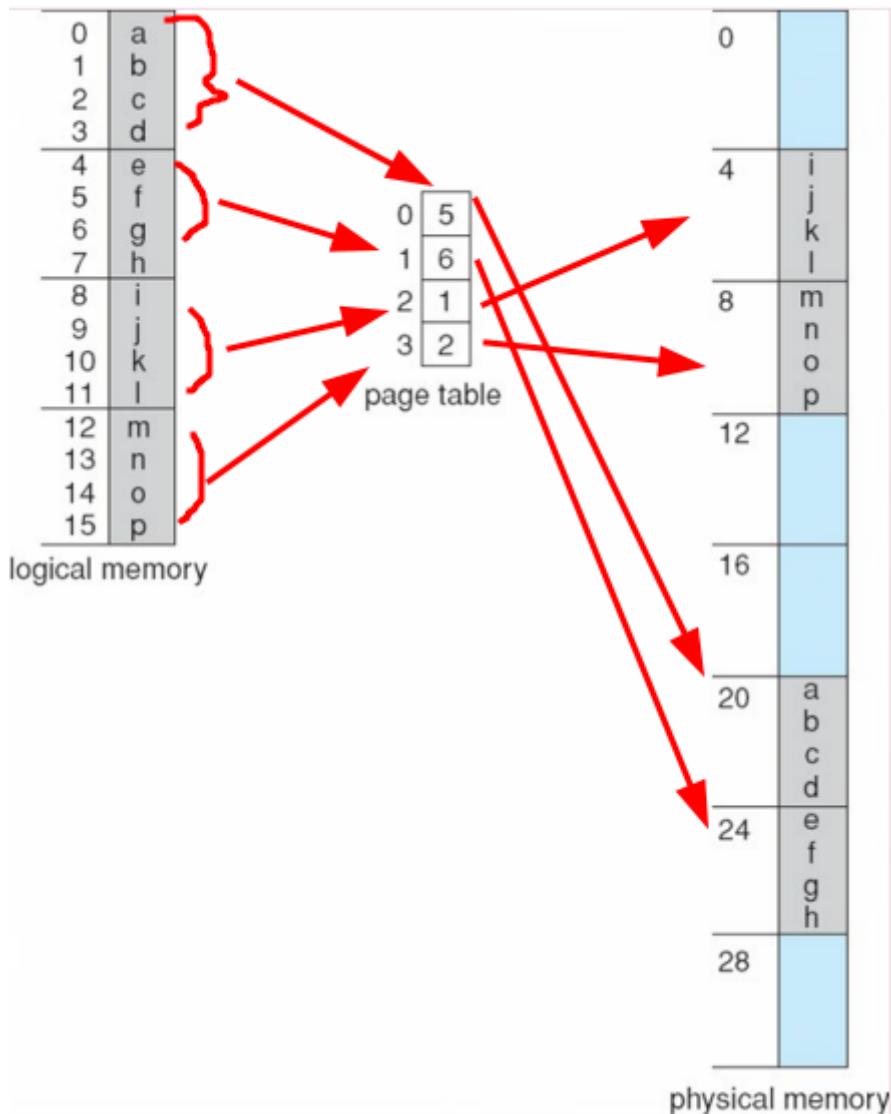
so the logical addresses translate as follows:

- $0 = (5 \times 4) + 0 = 20$
- $3 = (5 \times 4) + 3 = 23$
- $4 = (6 \times 4) + 0 = 24$
- $13 = (2 \times 4) + 1 = 9$

2



Paging example for a 32-byte memory with 4-byte pages.



## Calculating Internal Fragmentation

if page size = 2048 bytes, how many pages will a process of size 72,766 bytes need?

- 35 pages + 1,086 bytes

It will be allocated 36 pages with an internal fragmentation of  $2048 - 1086 = 962$  bytes

can you think of a worst case internal fragmentation scenario

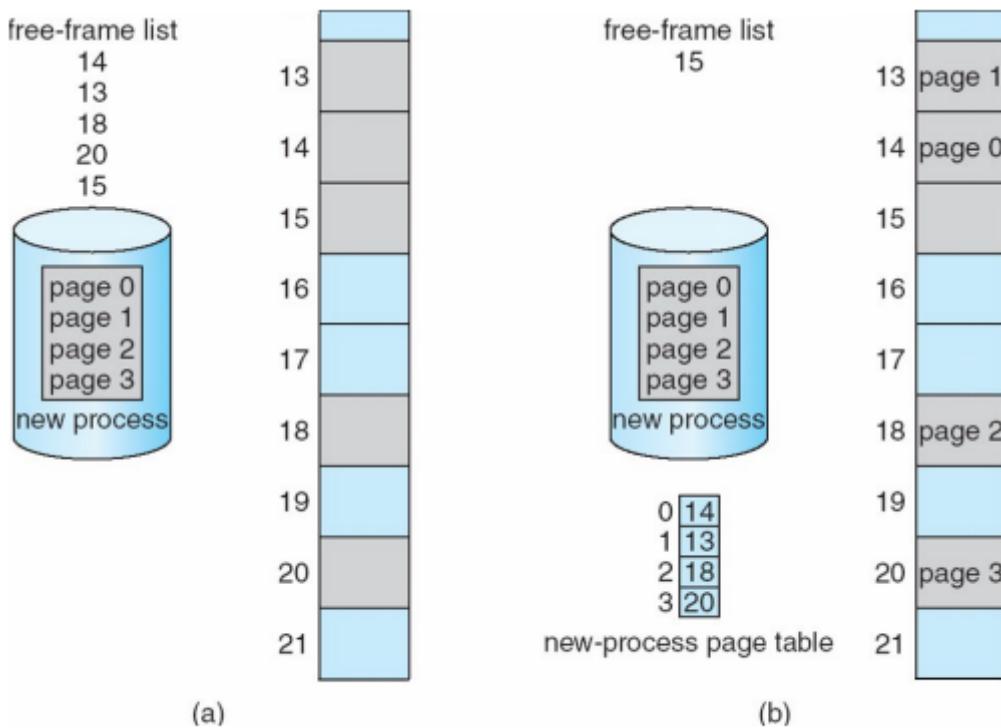
- process needs n pages + 1 byte
- will be allocated n+1 frames with an internal fragmentation of almost a frame

on average fragmentation = 1/2 frame size. so are the small frames more desirable?

- not necessarily as each page table entry takes memory to track.

## Free Frames

The frame table has one entry for each physical page frame indicating whether the latter is free or allocated and, if it is allocated, to which page of which process.



Free frames (a) before allocation and (b) after allocation.

## Implementation of Page Table

OS maintains a copy of the page table for each process just as it maintains a copy of the instruction counter and register contents

page table is kept in main memory

- page-table base register (PTBR) points to the page table
- page-table length register (PLTR) indicates size of the page table

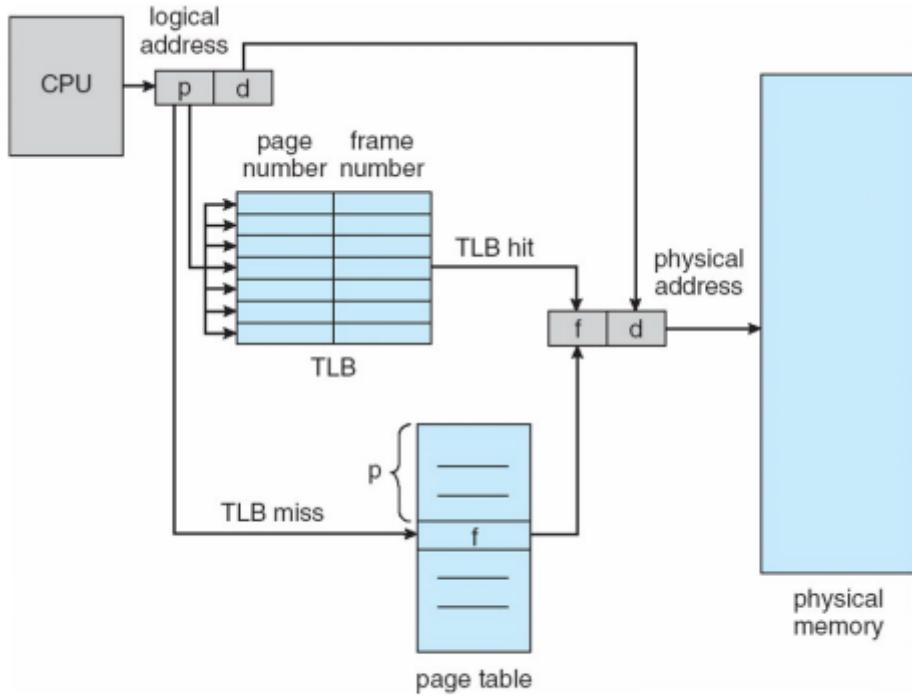
in this scheme every data/instruction access requires two memory accesses

- one for the page table
- one for the actual data or instruction
- memory access is slowed by a factor of 2

sol'n:

- use a special fast-lookup hardware cache called translation look-aside buffers (TLBs) (aka associative memory)

## TLB



## Paging hardware with TLB.

each entry in the TLB has 2 parts:

1. key/tag
2. value

TLB is associative, hihg-speed memory

- when the associative memory is presented with an item
- the item is compared with all keys simultaneously

typically small

- 32-1024 entries

tlb is a hardware feature

This is a lot like a cache

on hit:

- mmu checks if page number is present in the tlb
- if the number is found then its frame number is immediately available and is used to access memory

on miss:

- page number is not in the tlb
- have to reference the page table
- page and frame number are loaded into TLB for faster access next time

if TLB is full:

- an existing entry must be selected for replacement
  - replacement policies must be considered
    - range from least recently used (LRU) through round-robin to random
  - some CPUs allow the OS to participate in LRU entry replacement
    - other handle the matter themselves
  - some entries can be wired down for permanent fast access
    - never replaced
    - entries for key kernel code are wired down, typically

## Effective Access Time

Hit ratio

- % of times that a page number is found in the TLB
- % of the time we find our desired page number

ex:

- suppose it takes 10 nanoseconds to access memory
- if we find the desired page in TLB then a mapped-memory access takes 10 ns
- otherwise we need 2 memory accesses, one to the page table and frame number then one more to access the desired byte in memory, which takes 20 ns
- calculate effective access time

Effective Access Time

- statistical or real measure of how long it takes the CPU to read or write to memory
  - 80% hit ratio
    - $EAT = 0.8 * 10 + 0.2 * 20 = 12 \text{ ns}$
    - 20% slowdown in access time implied
    - $((12-20)/10)*100\%$
  - 99% hit ratio (more realistic)
    - $EAT = 0.99 * 10 + 0.01 * 20 = 10.1 \text{ ns}$
    - 1% slowdown in access time implied
    - $((10.1-10)/10)*100\%$

## Memory Protection in Page Table

memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed

- can also add more bits to indicate page execute-only, and so on

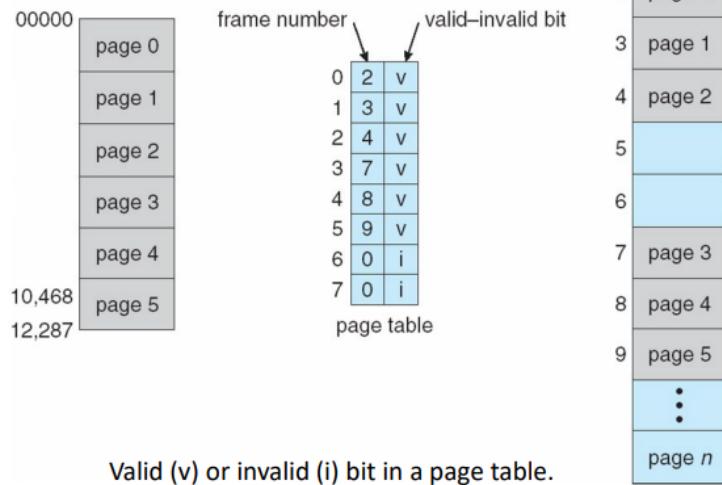
for now we're doing a valid-invalid bit attached to each entry in the page table

- valid
  - the associated page is in the process' logical address space
  - thus a legal page
- invalid
  - indicates the page is not in the process' logical address space
- or we can use the page-table length register (PTLR)

any violations result in a trap to the kernel

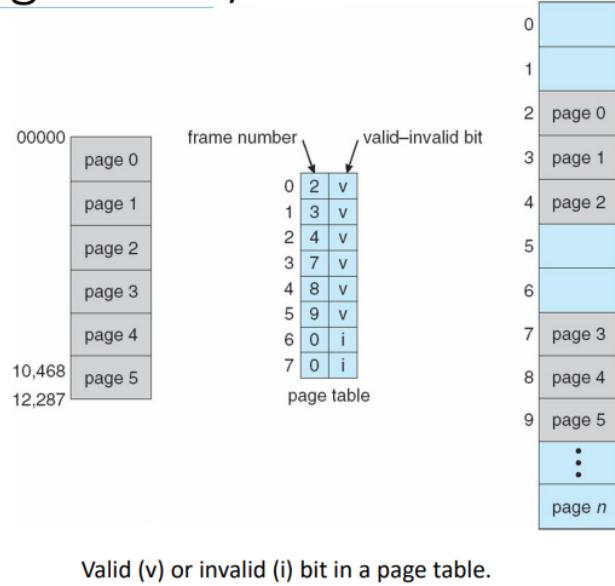
## Memory Protection in Page Table /2

- Given a page size of 2 KB, how many pages will be valid for the process shown?
- If the process tries to generate an address in pages 6 or 7, the computer will trap the OS (invalid page reference).



## Memory Protection in Page Table /2

- In a system with a 14-bit address space (0 to 16383), we have a program that should use only addresses 0 to 10468.
- Given a page size of 2 KB, how many pages will be valid for the process shown?
- If the process tries to generate an address in pages 6 or 7, the computer will trap the OS (invalid page reference).



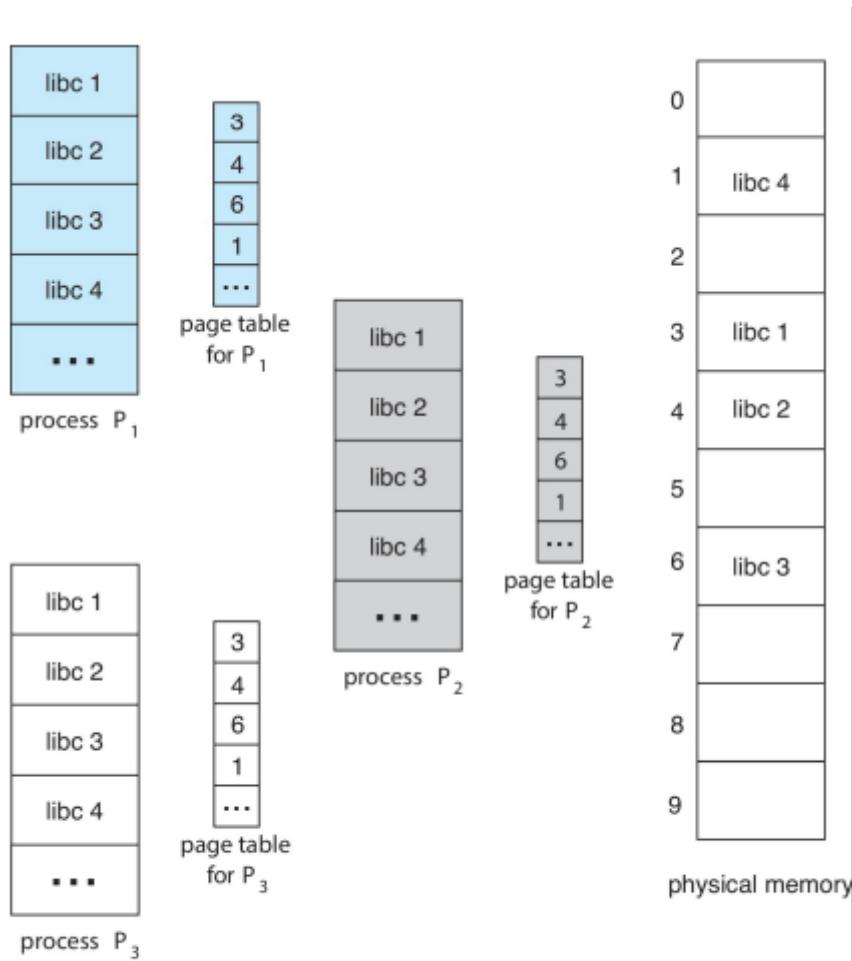
## Shared Pages

## shared code

- one copy of read-only code shared among processes
  - text editors, compilers, window systems
- multiple threads sharing the same process space

## private code and data

- each process keeps a separate copy of the code and data
- the pages for the private code and data can appear anywhere in the logical address space



**Sharing of standard C library in a paging environment.**

## Knowledge Check

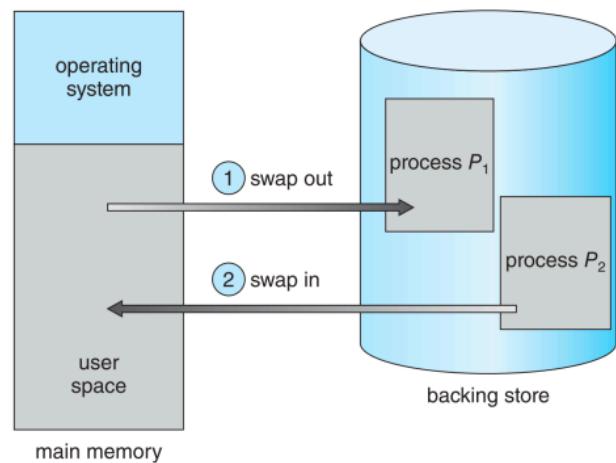
- consider a logical address with a page size of 8 kb. How many bits must be used to represent the page offset in the logical address?
  - the page offset in a logical address is determined by the size of the page
  - to represent the page offset within this 8 kb page we need enough bits to address each byte within the page
    - $(8 \times 1024 = 8192)$
  - number of bits required =  $\log_2(\text{Page Size}) = \log_2(8192) = 13$

prof recommends doing the knowledge checks at home.

# Swapping

A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution.

Makes it possible for the total physical address space of all processes to exceed physical memory of the system.



Standard swapping of two processes using a disk as a backing store.

the backing store can be a harddrive or ssd or whatever

standard swapping

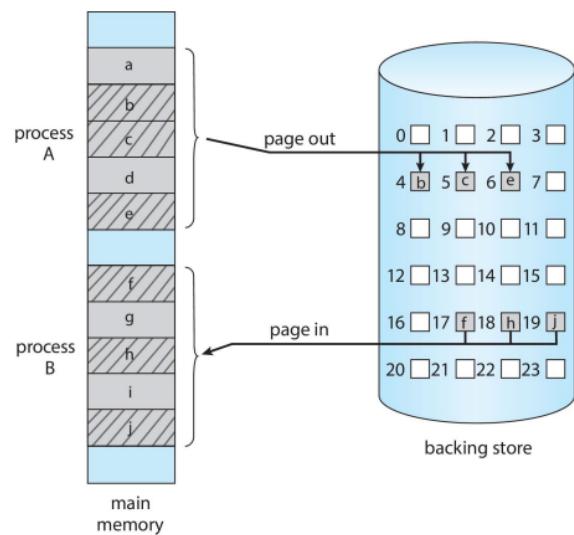
- moving entire processes between main memory and a backing store

Backing store

- fast disk large enough to accomodate copies of all memory images for all users
- must provide direct access to these memory images

## Swapping with Paging

- Most systems, including Linux and Windows, now use a variation of swapping in which pages of a process—rather than an entire process—can be swapped.
- The term **swapping** now generally refers to standard swapping, and **paging** refers to swapping with paging.



Swapping with paging.

page out

- move page from memory to the backing store

page in

- move page from backing store to memory

# Virtual Memory

## Background for Virtual Memory

code needs to be in physical memory to execute but entire program is rarely used

- error code
  - used to handle unusual error conditions
  - almost never executed
- large data structures
  - often allocated more mem than needed
- unusual routines
  - lesser used options and features of a program

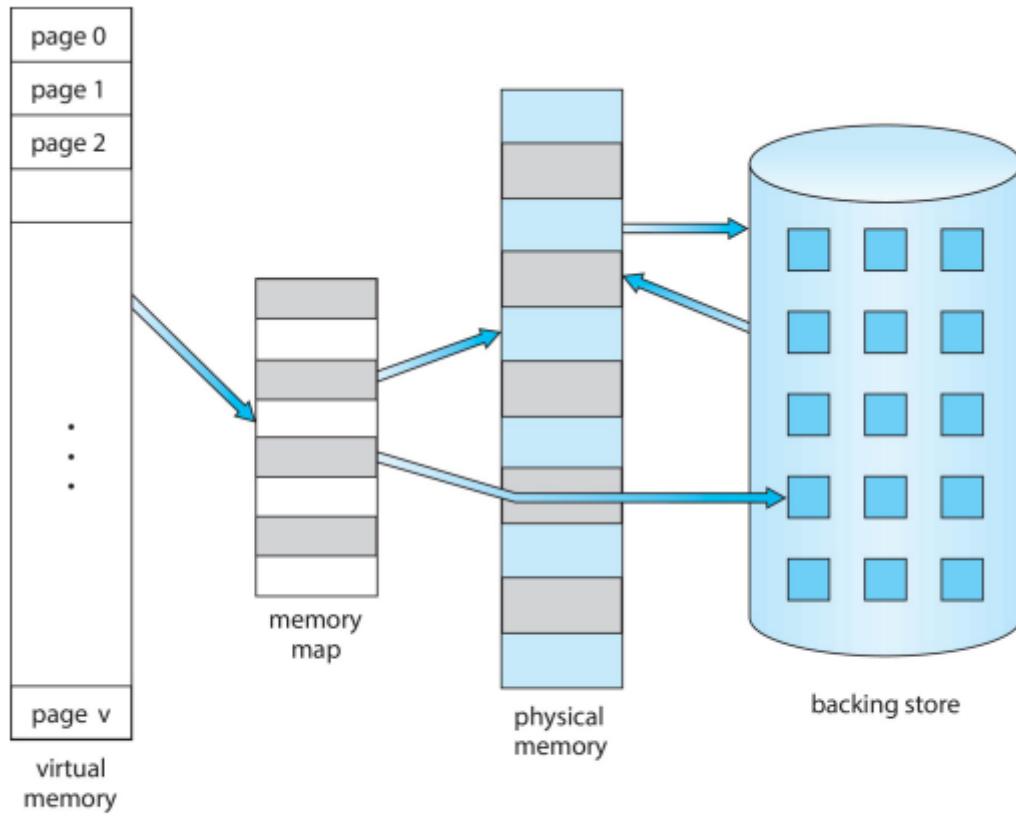
we only have to load what we need for execution

ability to execute partially-loaded program

- no longer limited by physical mem
- each program takes less memory while running
  - more programs can run at the same time
    - cpu util and throughput increase
    - response time and turnaround time the same
- less i/o needed to load or swap programs into mem
  - each user program runs faster

virtual mem

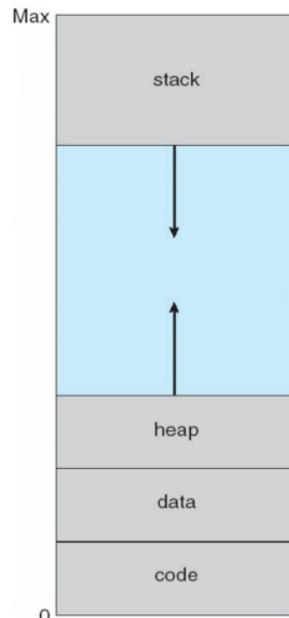
- separation of user logical mem from physical mem
- only part of the program needs to be in mem for execution
- logical address space can be larger than physical address space
- allows address spaces to be shared by several processes
- allows for more efficient process creation
- more programs running concurrently
- less i/o needed to load or swap processes



### Virtual Memory That is Larger Than Physical Memory

## Virtual Address Space

- **Virtual address space** – logical view of how process is stored in memory:
  - Usually start at address 0, contiguous addresses until end of space.
  - Meanwhile, physical memory organized in page frames.
  - MMU must map logical pages to physical page frames in memory.
- Logical memory is also known as virtual address space.



Virtual address space of a process in memory.

The large blank space (or hole) between the heap and the stack is part of the virtual address space.

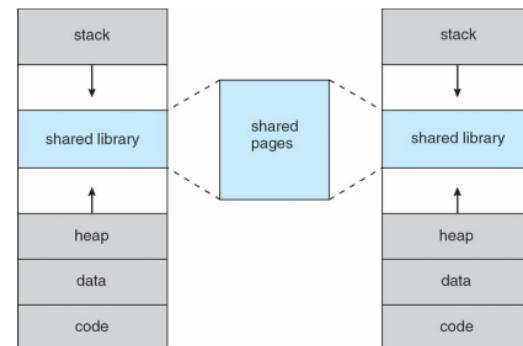
- Will require actual physical pages only if the heap or stack grows.

Virtual address spaces that include holes are known as **sparse address spaces**.

- Beneficial because the holes can be filled as the stack or heap segments grow or if we wish to dynamically link libraries during program execution.

## Shared Library Using Virtual Memory

- Virtual memory allows files and memory to be shared by two or more processes through **page sharing**.
- Examples:
  - Library mapped as read-only.
  - Shared memory for 2 processes to communicate.



Shared library using virtual memory.

## Logical Memory vs. Virtual Memory vs. Real Memory

- Logical memory is the memory space perceived by a process.
- Real memory is the physical RAM installed in the system.
- Virtual memory is a memory management technique used by the operating system to extend available memory beyond physical RAM.
  - logical memory
    - memory space perceived by a process
  - real/physical memory

- physical ram sticks in you're machine
  - virtual memory
    - memory management technique
    - used by operating system
    - extends beyond available physical ram
- 

Knowledge check:

- which of the following is a benefit of allowing a program that is only partially in memory to execute?
  - d. all of the above
  - programs can be written to use more memory than is available in physical memory
  - cpu util and throughput is incdreasde
  - less io needed to load or swap each user program into mem
- in general virtual memory decreases the degree of multiprogramming in a system
  - false
  - why would these even effect each other???

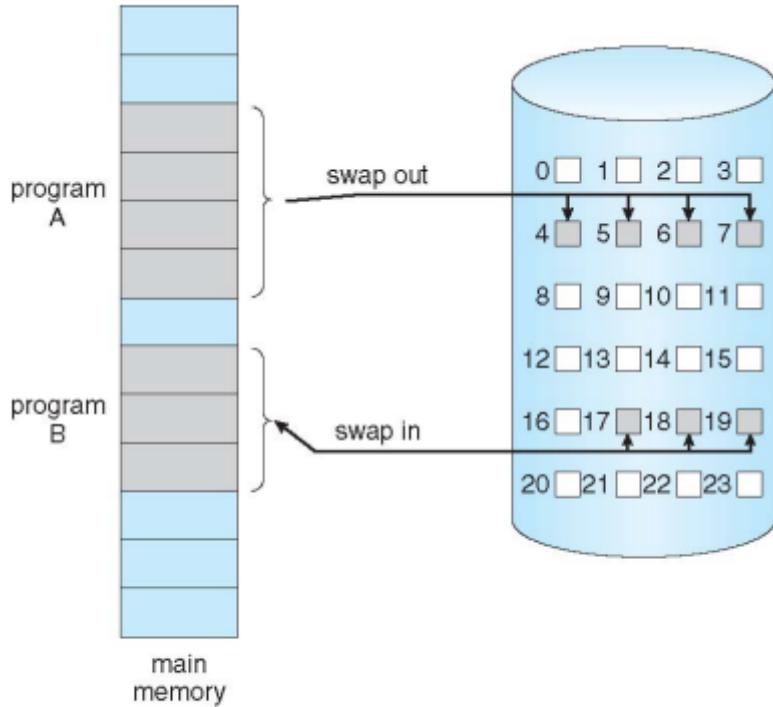
## Demand Paging

virtual memory can be implemented via

- demand paging
- demand segmation

demand paging:

- bring a page into memeory ONLY when it is needed
  - less required
    - I/O
    - memory
  - faster response
  - more users
- similar to a paging system with swapping



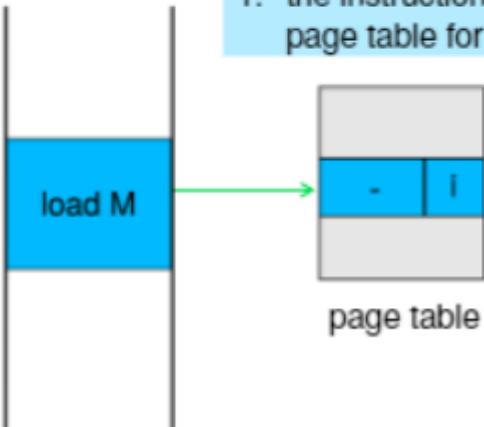
### Swapping with paging.

valid-invalid bit

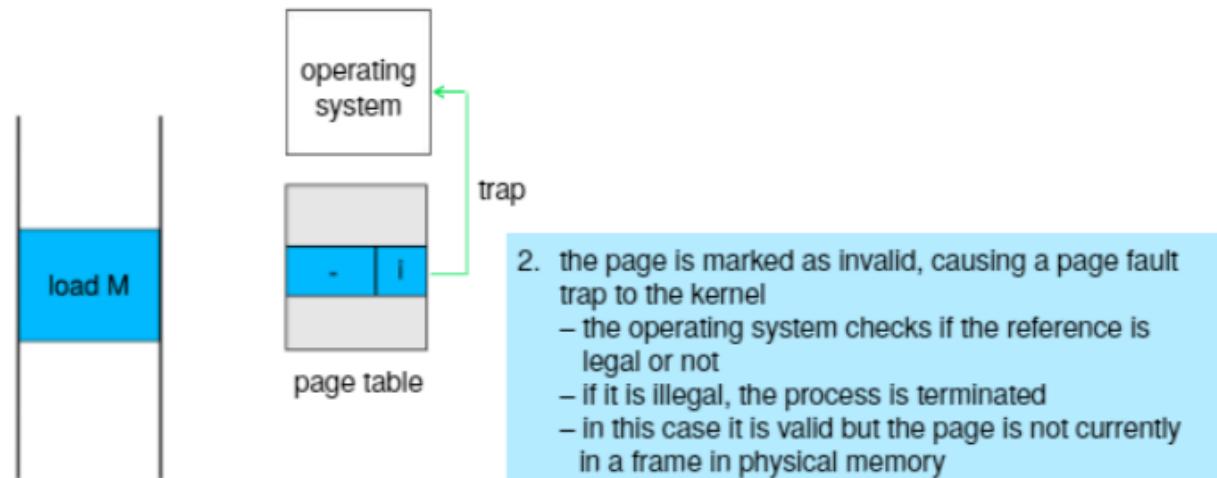
- with each page entry there is a valid-invalid bit
  - v → in-memory - memory resident
  - i → not-in-memory
- initially valid-invalid bit is set to **i** on all entries
  - all invalid
- during MMU address translation
  - if bit is invalid for a page entry then there's a page fault

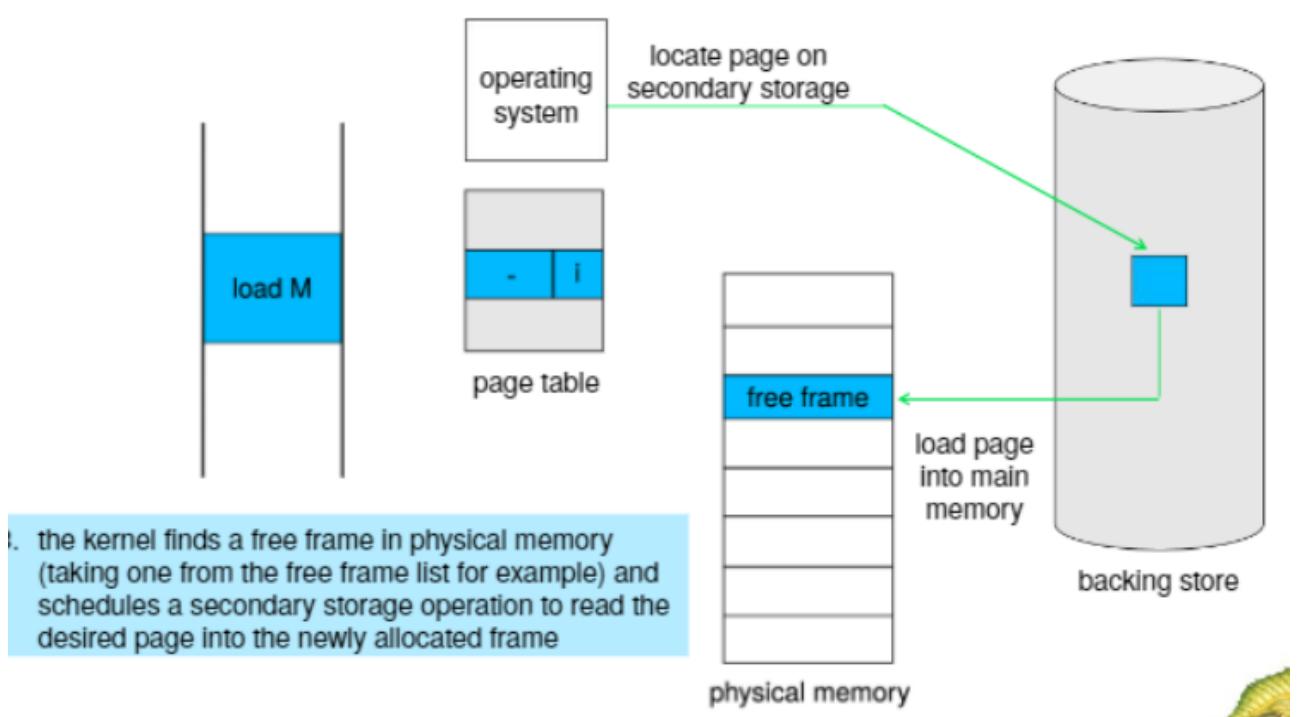
# Steps in Handling a Page Fault

1. the instruction "load M" executes, accessing the page table for the address containing "M"



2. the page is marked as invalid, causing a page fault trap to the kernel
  - the operating system checks if the reference is legal or not
  - if it is illegal, the process is terminated
  - in this case it is valid but the page is not currently in a frame in physical memory

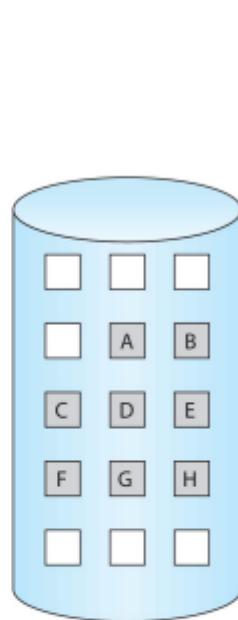
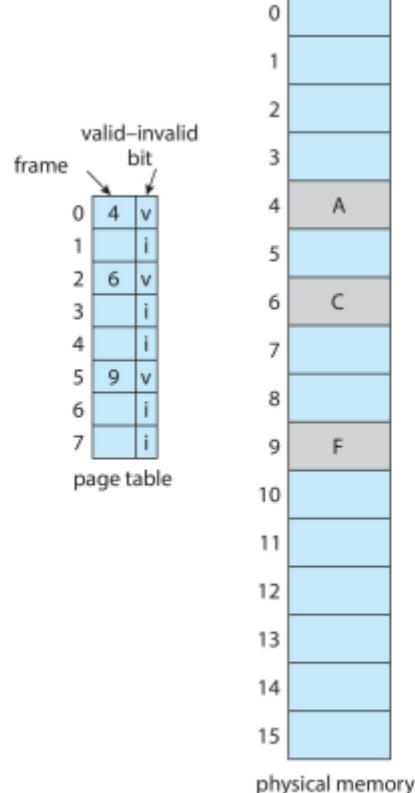




valid

|   |   |
|---|---|
| 0 | A |
| 1 | B |
| 2 | C |
| 3 | D |
| 4 | E |
| 5 | F |
| 6 | G |
| 7 | H |

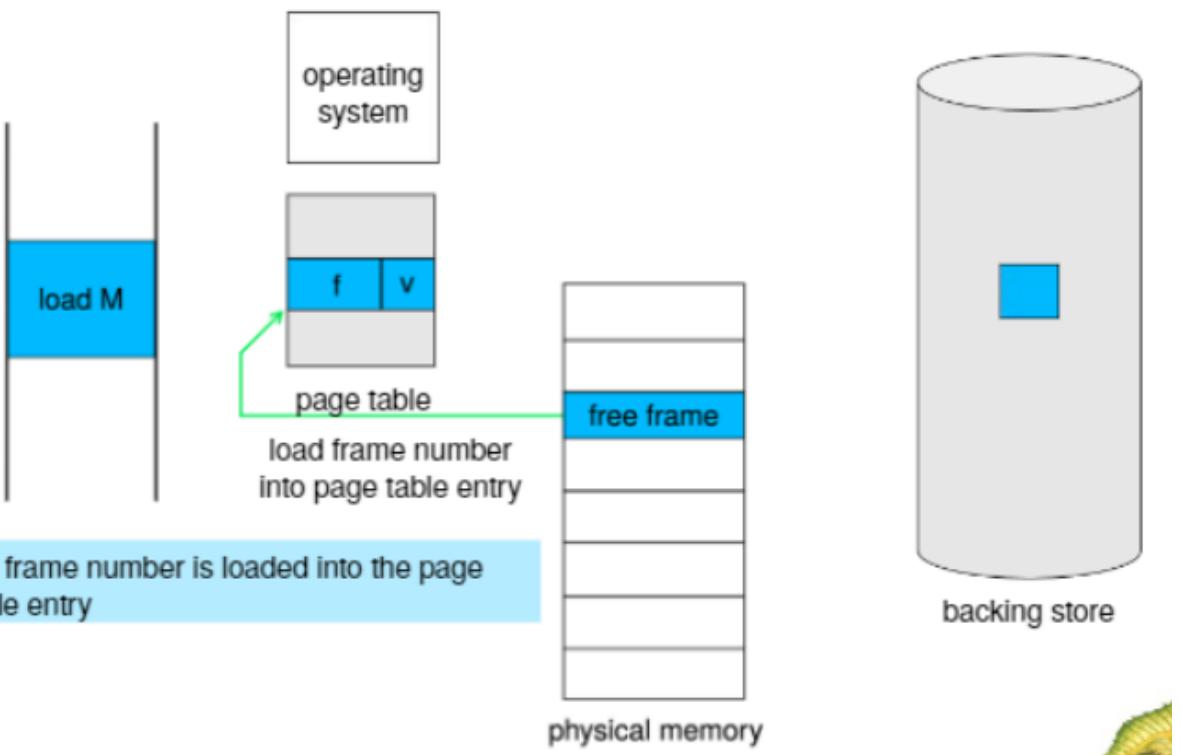
logical memory



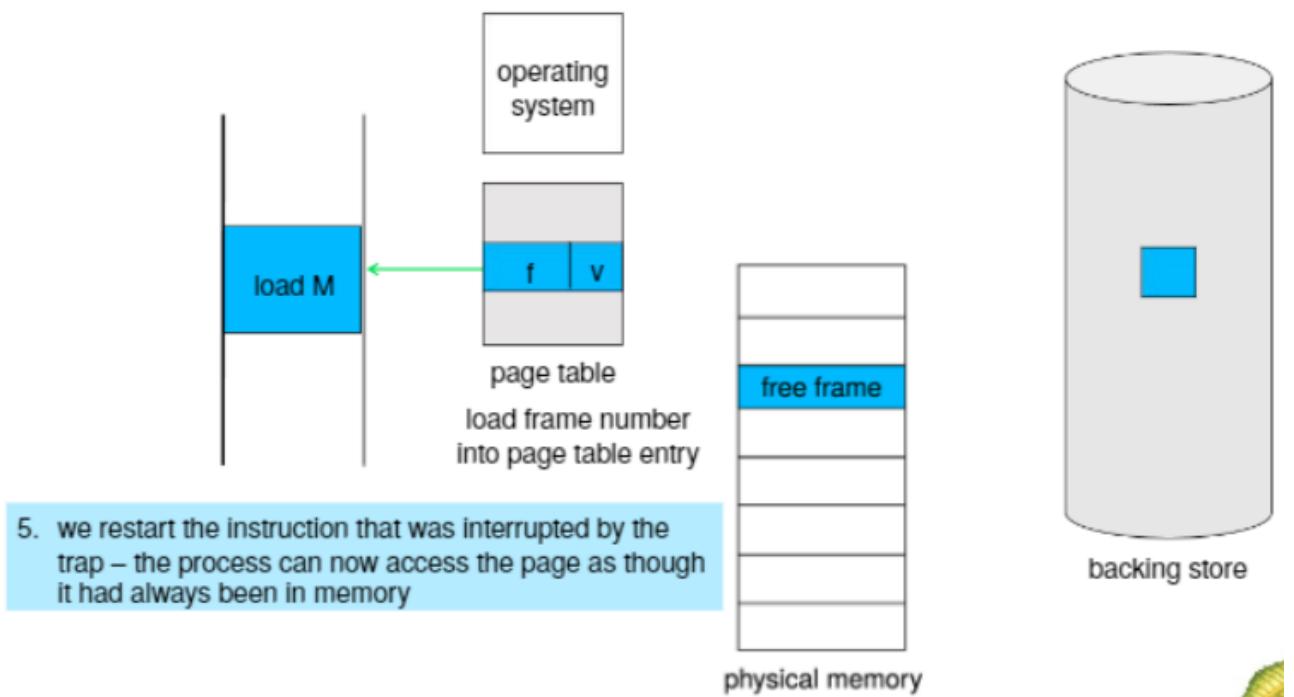
backing store

d–  
e

Page Table When Some Pages Are Not in Main Memory



## Steps in Handling a Page Fault

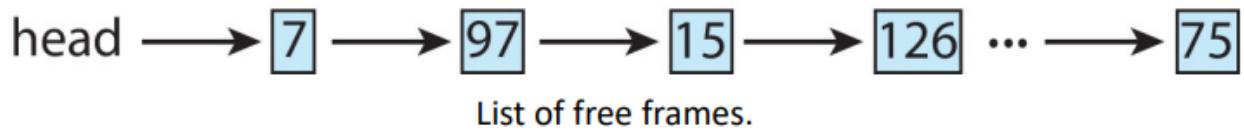


there are some more slides that are the same thing

free-frame list

- a pool of free frames for satisfying such requests
- maintained by most operating systems

- when a page fault occurs, the operating system must bring the desired page from secondary storage into main memory
- OS typically allocate free frames using zero-fill-on-demand
  - we flush out the frame
  - the content of the frames zeroed-out before being allocated
- when a system starts up, all available memory is placed on the free-frame list
  - this lets us start allocating frames as they are demanded



performance of demand paging

- 3 major activities of servicing a page-fault
  - service the page-fault interrupt
  - read in the page
    - this is where most of the time is spent
  - restart the process
- page fault rate  $0 \leq p \leq 1$ 
  - $p = 0 \rightarrow$  no page faults
  - $p = 1 \rightarrow$  every reference is a fault
- Effective Access Time (EAT) =
  - $(1 - p) \times \text{memory access} + p(\text{page fault overhead} + \text{swap page out} + \text{swap page in})$
  - we're basically averaging it out

Example of Demand Paging

memory access time = 200 ns

avg page-fault service time = 8 ms

$$\begin{aligned}\text{EAT} &= (1 - p) \times 200 + p(8\text{ms}) \\ &= ((1 - p) \times 200) + (p \times 8,000,000) \\ &= 200 + (p \times 7,999,800)\end{aligned}$$

if 1 in 1000 accesses causes a page fault ( $p = 0.001$ ) then EAT = 8.2 microseconds. a slowdown by a factor of 40

if we want a performance degradation (slowdown) less than 10 percent

- 10% of 200 (effective access time we want) = 20 ns
- $20 > 200 + 7,999,800 * p$
- $20 > 7,999,800 * p$
- $p < 0.0000025$
- $p <$  one page fault in every 400,000 memory accesses

This needs hardware support, it can't all be done by the OS.

---

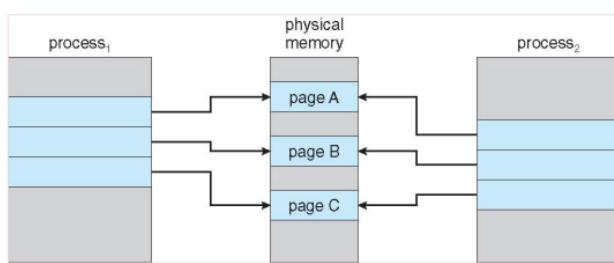
### Knowledge Check

- on a system with demand-paging, a process will experience a high page fault rate when the process begins execution
  - true
  - similar principle to caching where we need to slowly fill things up in the memory that the process asks for
  - since the page table is invalid to start with we end up faulting a ton as we're setting things up for the process
  - as it starts going on it'll start hitting the pages we already loaded for it
- a page fault occurs when
  - c. a process tries to access a page that is not loaded in memory
- if memory access time is 250 ns and average page fault service time 10 ms the probability of page faults must be less \_\_\_ to keep the performance degradation less than 20%
  - options
    - a) 0.0000025
    - b) 0.000005
    - c) 0.0000075
    - d) 0.00001
  - takehome exercise

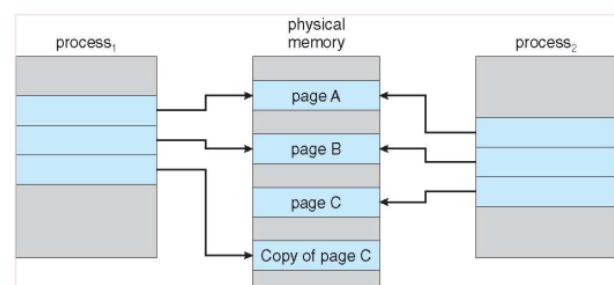
## Copy on Write

copy-on-write (COW):

- allows both parent and child processes to initially share the same pages in memory
- the page is copied once either process modifies a shared page
  - read: write
- allows more efficient process creation as only modified pages are copied



Before process 1 modifies page C.



After process 1 modifies page C.

---

### Knowledge Check

- \_\_\_ allows the parent and child processes to initially share the same pages, but when either process modifies a page, a copy of the shared page is created.

- a. copy-on-write

# Page Replacement

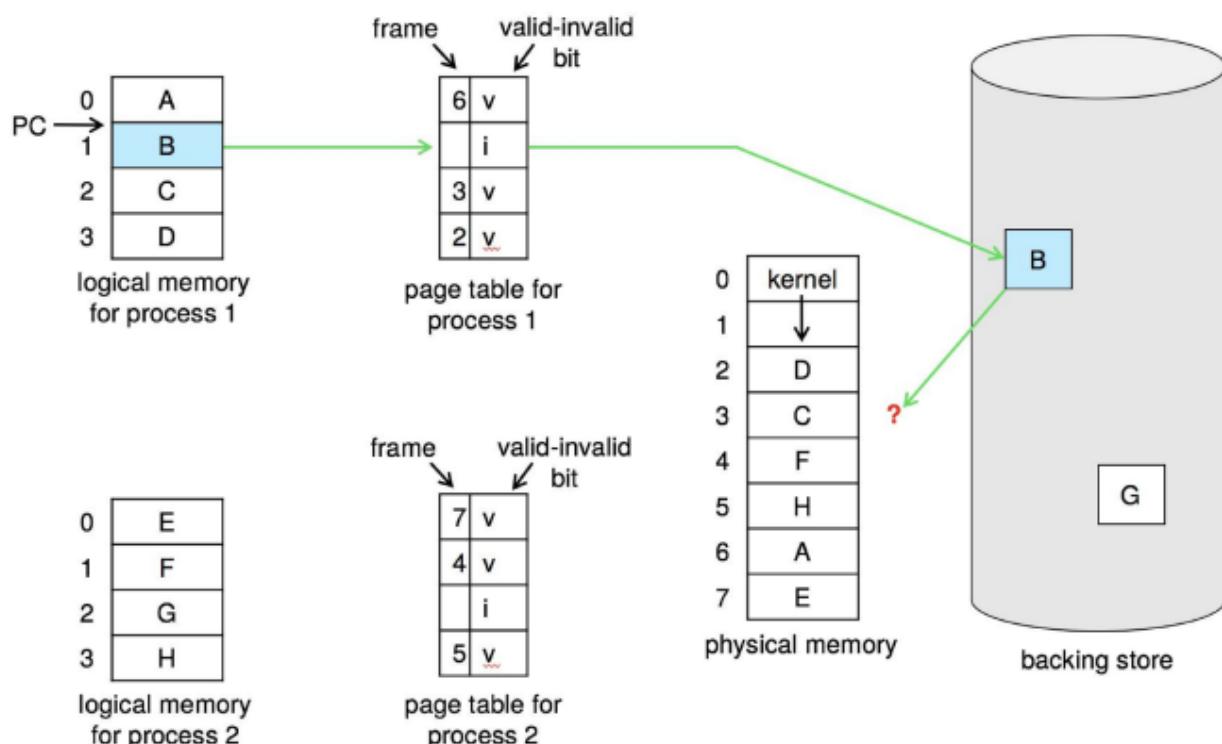
sometimes there is no free frame

how:

- memory is used up by process pages but also in demand from the kernel, I/O buffers, etc
- the page is a hot commodity that everyone needs access to

scenario

- process executing
- page fault occurs
- OS determines where the desired page is residing on secondary storage
- but there are no free frames on the free-frame list
  - all mem is in use



**Need for page replacement.**

need for page replacement

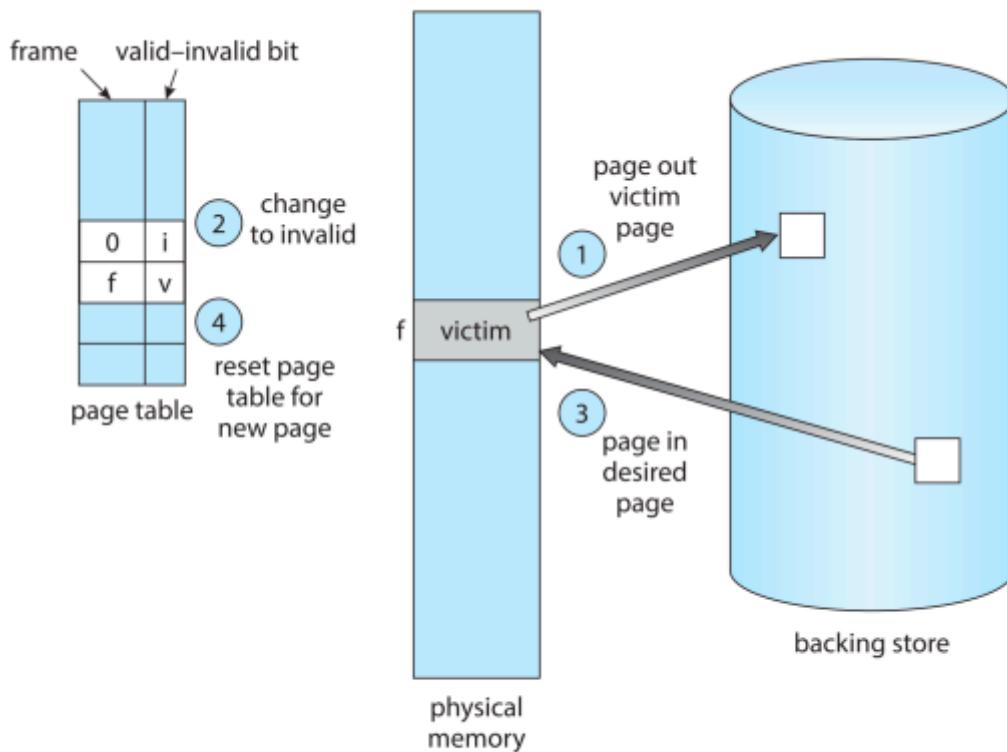
- algorithms
  - do we terminate it
  - do we swap out the whole process till there are more free frames
  - do we replace the page

- we can find some page in memory but not really in use, page it out
- most operating systems now combine swapping pages with page replacement

we want an algo that has the minimum number of page faults

basic page replacement:

- prevent over-allocation of mem
  - modify page-fault service routine to include page replacement
- use modify (dirty) bit to reduce the overhead of page transfers
  - only modified pages are written to disk
  - unmodified pages have no need to be written back to backing store since it wouldn't make a difference
- page replacement completes separation b/w logical mem and phys mem
  - large virt mem can be provided on a smaller phys mem



## Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame:
  - a) If there is a free frame, use it
  - b) If there is no free frame, use a page replacement algorithm to select a **victim frame**
  - c) Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables.
4. Continue the process from where the page occurred.

frame allocation algo determines how many frames given to each process

- some small processes are smaller
- are frames reserved
- etc

page-replacement algo

- which pages to replace
- want lowest page-fault rate on both first access and re-access

## Page Replacement Algorithms

### Page Replacement Algo Evaluation

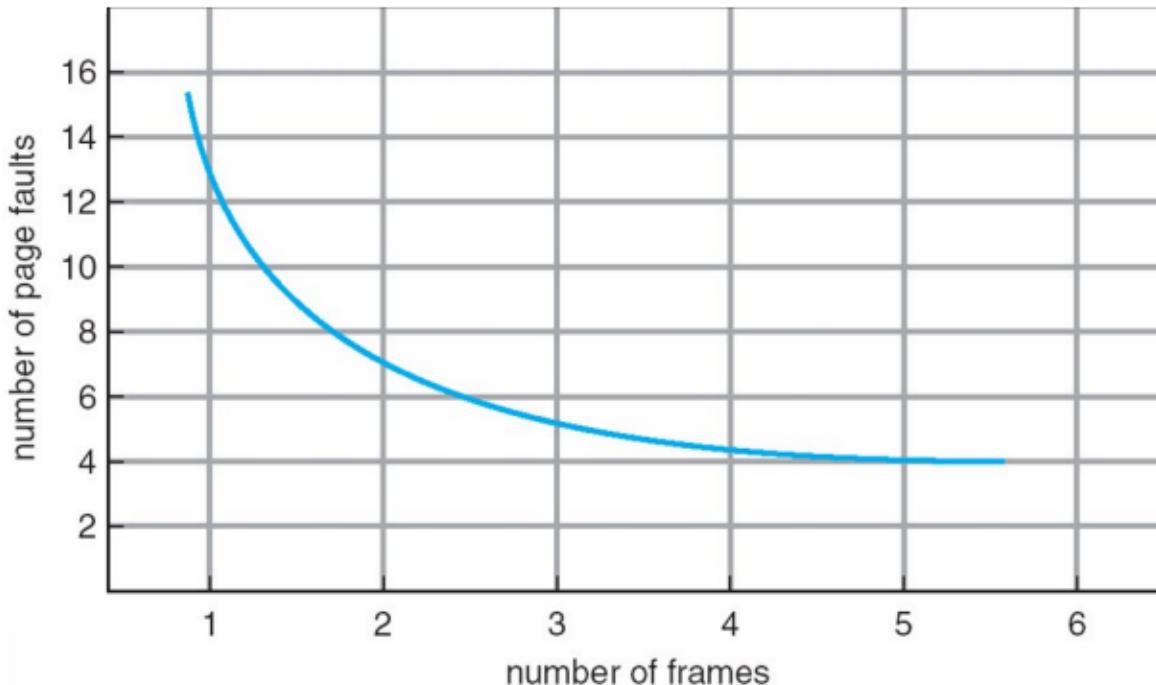
evaluate algo by running it on a particular string of memory reference (reference string) and computing the number of page faults on that string

- string is just page numbers, not full addresses
- if we have a reference to a page  $p$ , then any references to page  $p$  that immediately follow will never cause a page fault
- results depend on the number of frames available

- Example: if we trace a particular process, we may record the address sequence:

**0100, 0432, 0101, 0612, 0102, 0103, 0104, 0101, 0611, 0102, 0103,  
0104, 0101, 0610, 0102, 0103, 0104, 0101, 0609, 0102, 0105**

- This sequence is reduced to the following reference string:  
**1, 4, 1, 6, 1, 6, 1, 6, 1**



Graph of Page Faults Versus The Number of Frames (in general)

## Different Types

algos:

- FIFO
- Optimal
- Least Recently Used (LRU)
- LRU-Approximation
  - Second-Chance
  - Enhanced Second-Chance
- Counting-Based

In all our examples, the reference string of referenced page numbers is: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

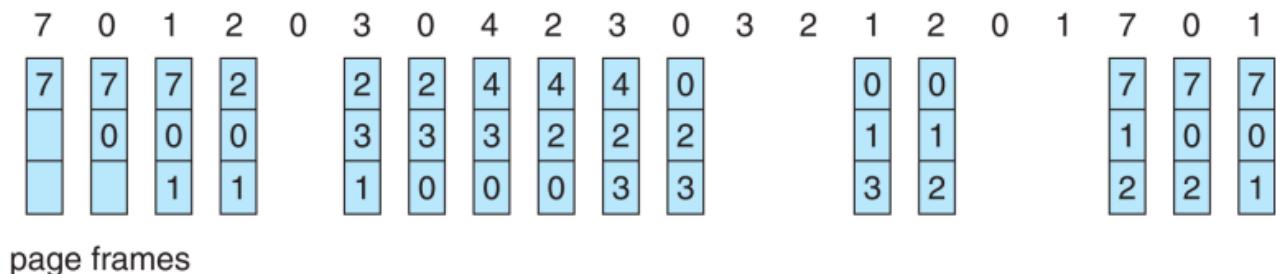
## FIFO

first-in-first-out algo

Reference string: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

3 frames (3 pages can be in memory at a time per process)

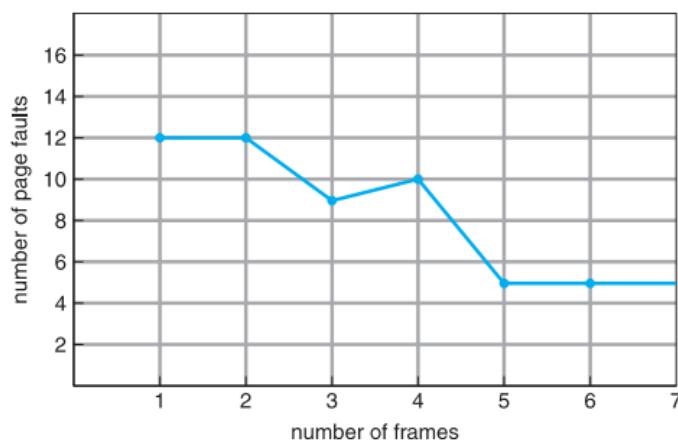
How to track ages of pages? Just use a FIFO queue.



**FIFO: 15 page-faults**

## FIFO Illustrating Belady's Anomaly

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
  - Adding more frames can cause more page faults! → **Belady's Anomaly**



Page-fault curve for FIFO replacement on a reference string.

## Optimal Algo

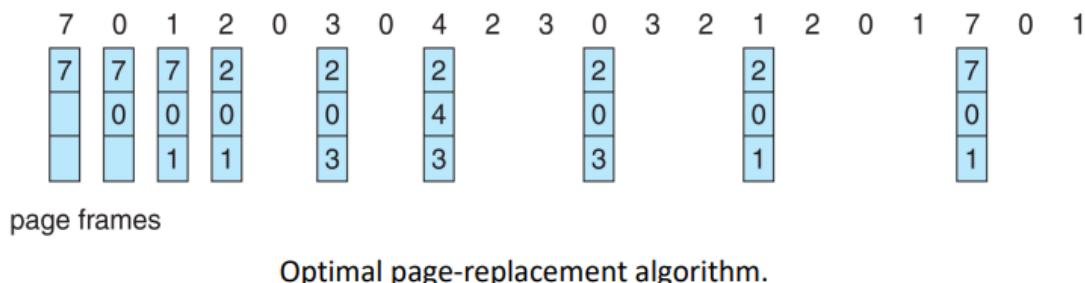
Replace page that **will** not be used for longest period of time.

- **9 page-faults** is optimal for the example.

How do you know this?

- Can't read the future.

Used for measuring how well your algorithm performs.



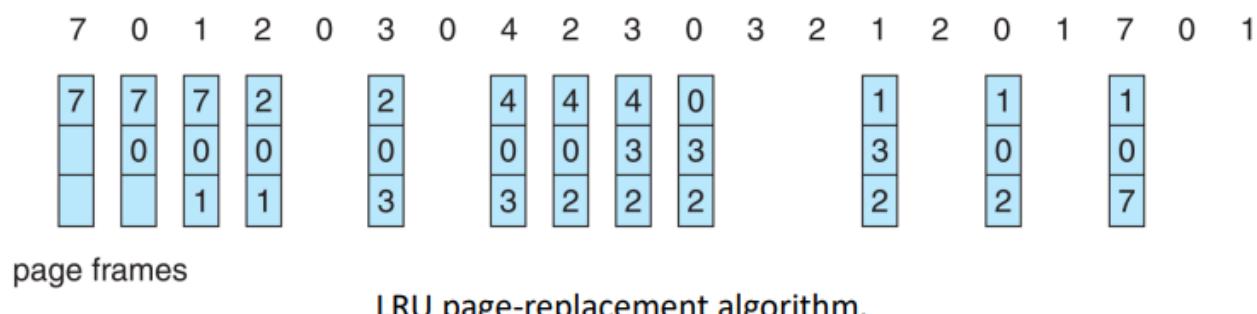
This is an example, not an actual algo.

## Least Recently Used (LRU) Algorithm

Replace the page that has not been used in the most amount of time. Associate time of last use with each page.

Results in 12 faults which is closer to optimal than FIFO

generally good and frequently used.



|                         |
|-------------------------|
| 7 - assign to open page |
| 0 - assign to open page |
| 1 - assign to open page |
| 2 - replace 7 since LRU |
| 0 - 0 already has page  |
| 3 - replace 1 since LRU |
| 0 - 0 already has page  |
| 4 - replace 2 since LRU |
| 2 - replace 3 since LRU |

3 - replace 0 since LRU

...

It's possible that with a different reference string we'll end up with performance similar to FIFO

---

## Implementation

- requires hardware support
- hardware counter implementation
  - every page entry has a counter
    - every time the page is referenced through this entry, copy the clock into the counter
  - when a page needs to be changed
    - look at the counters (clock) of the pages
    - find the one w/ smallest value (least recently used)
    - replace that one
- *prof also mentioned a stack hardware implementation which is interesting*

---

## LRU approximation algos

- LRU needs special hardware
  - the aforementioned counter implementation
- LRU is still slow even with the special hardware
- reference bit
  - with each page associate a bit
    - initially set to 0
  - when a page is referenced, the bit is set to 1
  - if a reference page has reference bit = 0 then we replace it
    - we don't know the order so this isn't perfect
- for this reason, we make many algos that approximate LRU

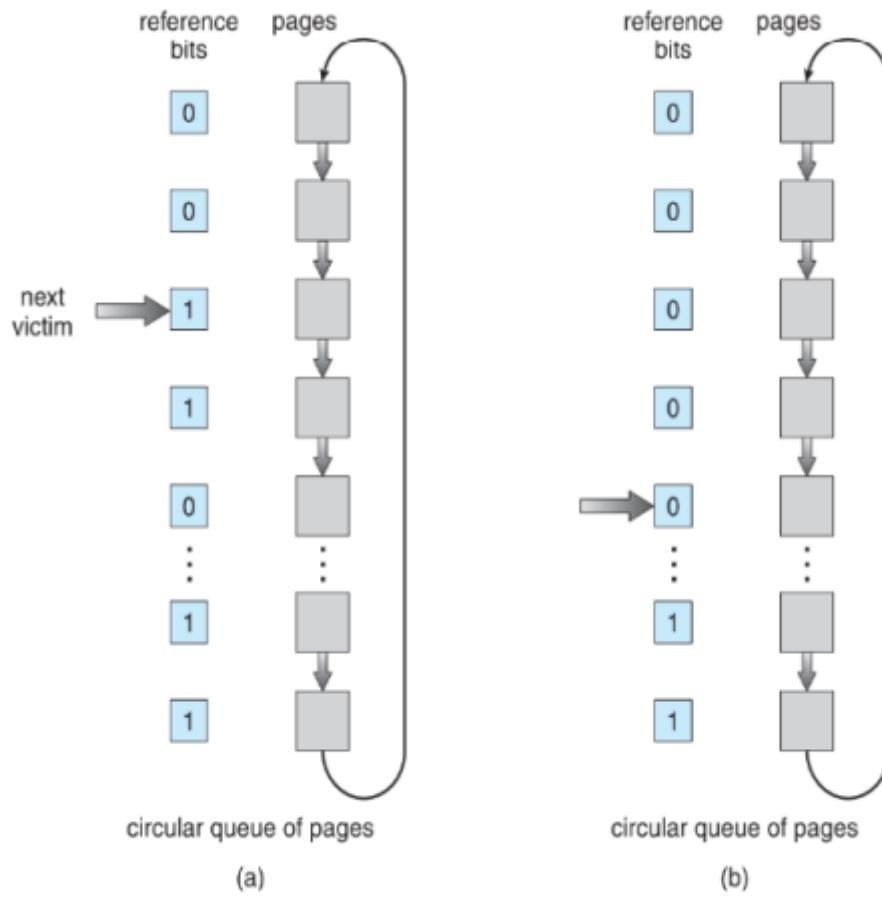
## **Second-Chance Algo (Clock) - LRU Approximation**

Generally implemented using FIFO and a hardware-provided reference bit

if the reference bit of the page to be replaced is equal to:

- 0 → replace it
- 1
  - page is given second chance
  - move on to select the next FIFO page
  - clear reference bit
    - set it to 0
  - arrival time is set to current time

if all the bits are set then this algo degenerates to FIFO replacement.



### Second-chance (clock) page-replacement algorithm.

“The hand is like a clock going tick, tick, tick around the circle... that is why it is called the clock algorithm” - Professor

### Enhanced Second-Chance Algo

Uses 2 bits

- reference bit
- modify bit (if available)

replacing a modified page introduces more overhead

So we can eval pages with the ordered pair (reference, modify)

| pair  | meaning                            | assessment                                       |
|-------|------------------------------------|--------------------------------------------------|
| (0,0) | neither recently used nor modified | best page to replace                             |
| (0,1) | not recently used but modified     | must write out before replacement, more overhead |
| (1,0) | recently used but clean            | will probably be used again soon                 |

| pair  | meaning                    | assessment                                                           |
|-------|----------------------------|----------------------------------------------------------------------|
| (1,1) | recently used and modified | probably will be used again AND need to write out before replacement |

The above table is in order of replacement preference

When page replacement is called for, use the clock scheme but use the four classes & replace page in lowest non-empty class

- this may require searching the circular queue several times

## Counting Algos

Keep a counter of the ## of references that have been made to each page

Least Frequently Used (LFU) algo

- replaces the page with the smallest count

Most Frequently Used (MFU) algo

- based on the argument that the page with the smallest count was probably just brought in and has yet to be used

neither are common

- implementation is expensive
- does not approximate OPT well

### Knowledge Check

- Q:
  - Suppose we have the follow page accesses:
    - 12342341211314
  - 3 frames within our system
  - use LRU replacement algo
  - what is the ## of page faults for the given ref string?
- A:
  - take home exercise
  - The answer is: c. 8
- Q: Belady's anomaly states that \_\_\_\_
- A:
  - d. for some page replacement algorithms, the page-fault rate may increase as the number of allocated frames increases
- Q: in the enhanced second chance algo, which ordered pair represents a page that would be the best choice for replacement?

- A: a. (0,0)

# Allocation of Frames

Much easier design principle compared to page replacement according to prof

The question of how we allocate the fixed amount of free memory among the various processes

ex:

- 128 frames
- os takes 35
- 93 free for user
- we have 2 processes
  - how many frames does each process get?

1. when a user process started execution, it would generate a sequence of page faults

The first requests need to generate the page faults to populate their frames

2. first 93 page faults would all get free frames from the free-frame list

3. when the free-frame list was exhausted, a page-replacement algo would be used to select one of the 93 in-memory pages to replaced by the 94th, and so on

4. when the process terminated, the 93 frames would once again be placed on the free-frame list

2 major allocation schemes that variate this simple strat

- fixed allocation
  - equal allocation
    - each process gets an equal amount of frames
    - ex. 93 frames split among 5 processes result in 18 frames/process and 3 free frames
  - proportional allocation
    - bigger processes get more frames than smaller processes
    - dynamic as degree of multiprogramming
    - process sizes change
- priority allocation
  - proportional allocation scheme using priority rather than size
  - higher priority processes get more frames

# Thrashing

thrashing

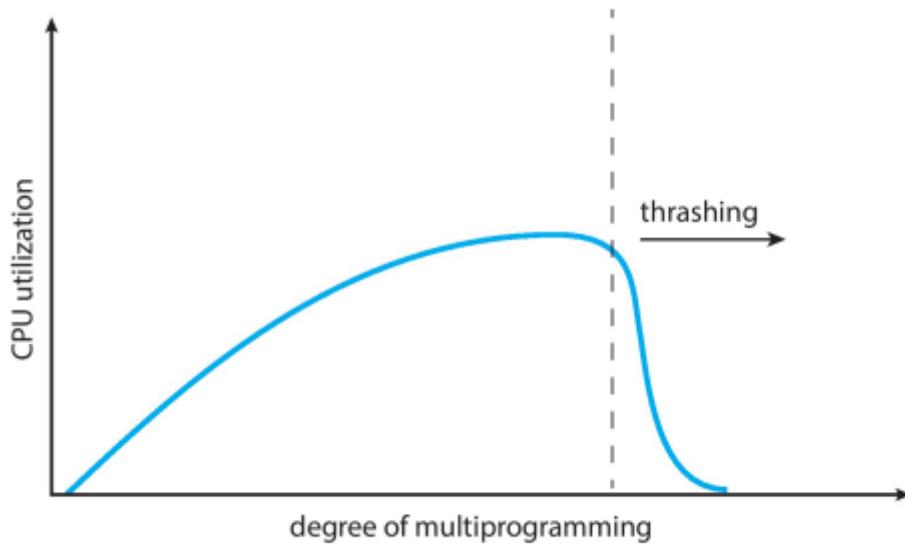
- process is spending more time paging than executing

if there aren't "enough" pages, page fault rate is very high

- page fault to get page
- replace existing frame
- quickly need replaced frame back

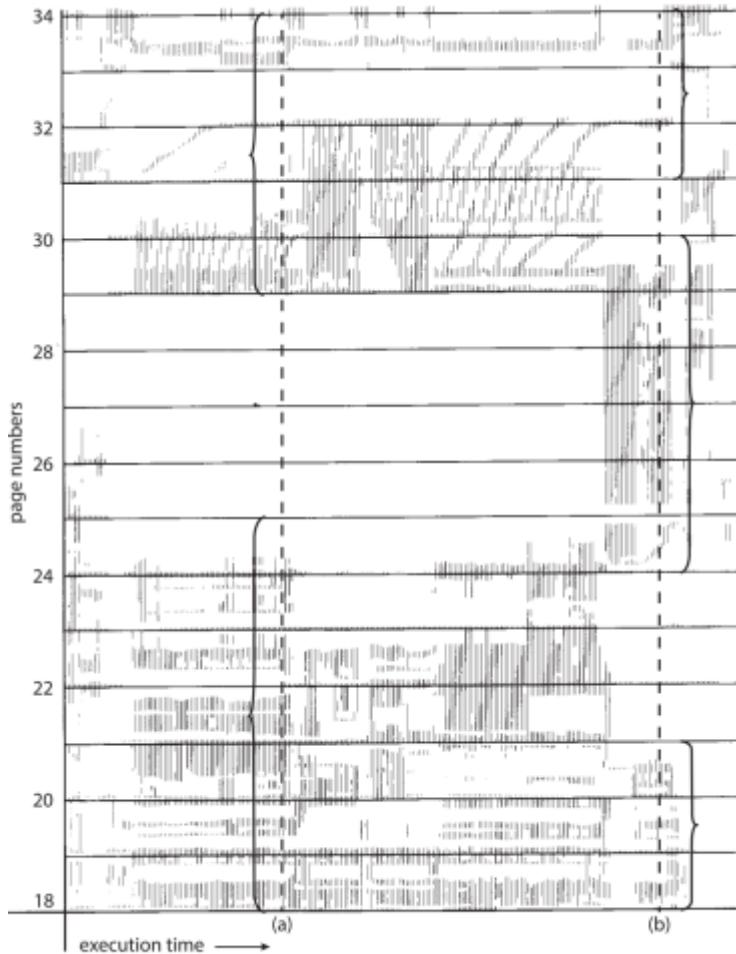
leads to

- low cpu util
- os thinks it needs to increase degree of multiprogramming
- can be exacerbated by more processes being added



## Demand Paging

The solution to thrashing



## Locality In A Memory-Reference Pattern

locality:

- the tendency of processes to reference memory in patterns rather than randomly
- process migrates from one locality to another
- localities may overlap

why does thrashing occur

- total size of locality is greater than total memory size

we can limit the effect by using local or priority page replacement

## Working-Set Model

the working-set model is based on the assumption of locality

$\Delta \equiv$  working-set window  $\equiv$  a fixed number of page references  
Example: 10,000 instructions

$WSS_i$  (working set of Process  $P_i$ ) = total number of pages referenced in the most recent  $\Delta$  (varies in time)

- if  $\Delta$  too small will not encompass entire locality.
- if  $\Delta$  too large will encompass several localities.
- if  $\Delta = \infty \Rightarrow$  will encompass entire program.

$D = \sum WSS_i \equiv$  total demand frames

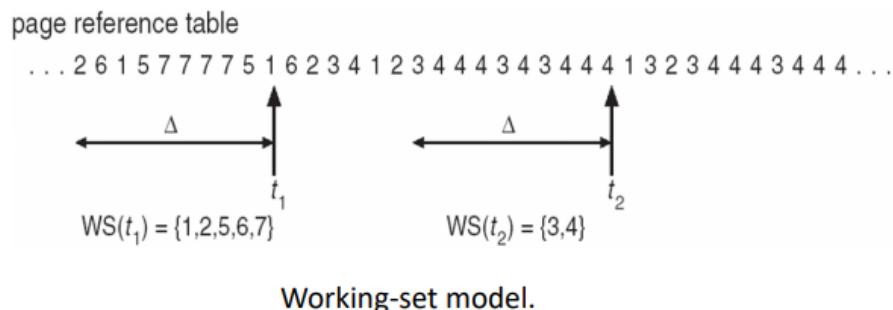
- Approximation of locality.

if  $D > m \Rightarrow$  Thrashing.

- $m$  is the total number of available frames.

Policy if  $D > m$ , then suspend or swap out one of the processes.

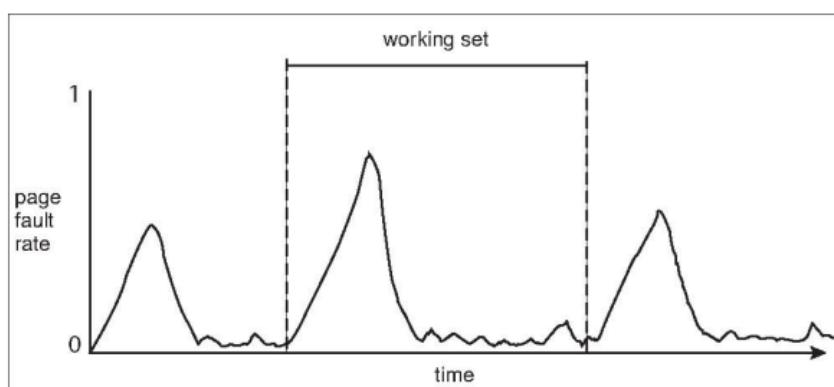
Example,  $\Delta = 10$



Direct relationship between working set of a process and its page-fault rate.

Working set changes over time.

Peaks and valleys over time.



---

Knowledge Check

- Q:

- \_\_\_\_ occurs when a process spends more time paging than executing
- A:
  - a. thrashing
- Q:
  - The \_\_\_\_ is an approximation of a program's locality
- A:
  - b. working set
- Q:
  - if the page fault rate is too high, the process may have too many frames. True or false?
- A:
  - false

## Allocating Kernel Memory

treated differently from user memory, allocating and deallocating memory frequently

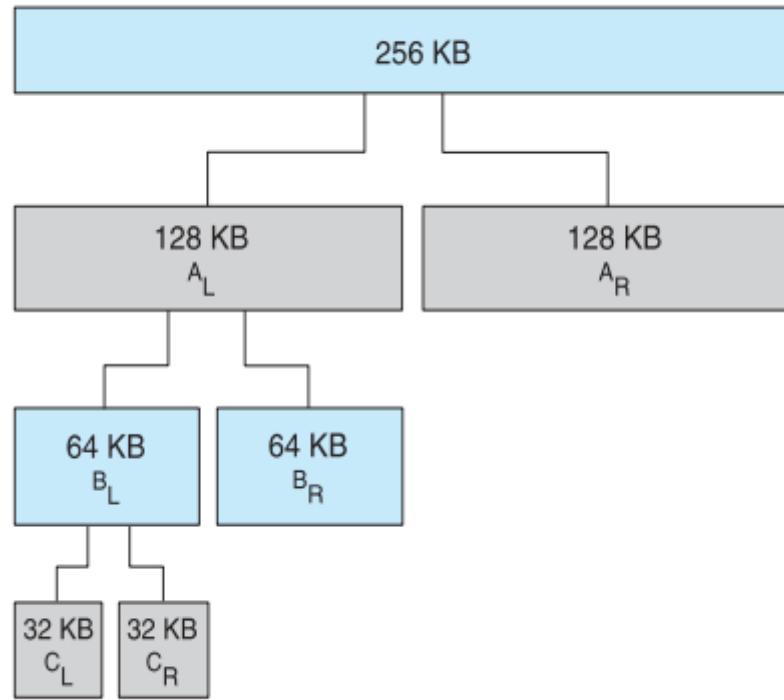
often allocated from a free-memory pool:

- kernel reqs memory for structures of varying sizes
- some kernel memory needs to be contiguous like for device i/o

The “buddy system”

- strategy for managing free memory that is assigned to kernel processes
- allocates memory from fixed-size segments consisting of physically-contiguous pages
- memory allocated using power-of-2 allocator
  - satisfies requests in units sized as power of 2
  - request rounded up to next highest power of 2
  - when smaller allocation needed than is available
    - current chunk split into two buddies of next-lower power of 2
    - continue until the appropriate-sized chunk becomes available
- quickly combining(coalescing) unused chunks into a larger chunk
- can cause internal fragmentation
  - 50% wastage of memory

physically contiguous pages



Buddy system allocation.

example:

- 256kb chunk available
- kernel req 21kb
- split into  $A_L$  &  $A_R$  of 128 kb each
- 1 further divided into  $B_L$  and  $B_R$  of 64kb each
- 1 further divided into  $C_L$  and  $C_R$  of 32 kb each
- finally have a chunk small enough to justify using on the request
  - any smaller would be 16kb

# Storage Management

## Mass-Storage Structure

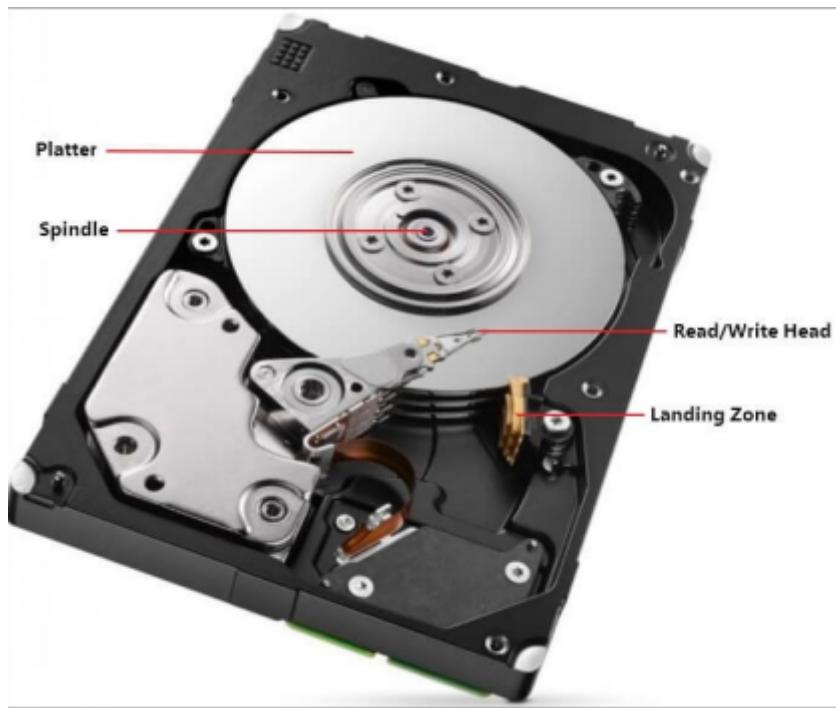
### Overview of Mass Storage Structure

bulk of secondary storage for modern computers provided by

- hard disk drives HDDs
- nonvolatile memory (NVM)
  - semiconductor tech like ram
  - doesn't dump info on power off like ram

- relatively newer

## Hard Disks



Note that there are several read/write heads and magnetically coated platters.

**Landing Zone** is where the read/write head sits when it's not doing anything

HDDs spin platters of magnetically-coated material under moving read-write heads

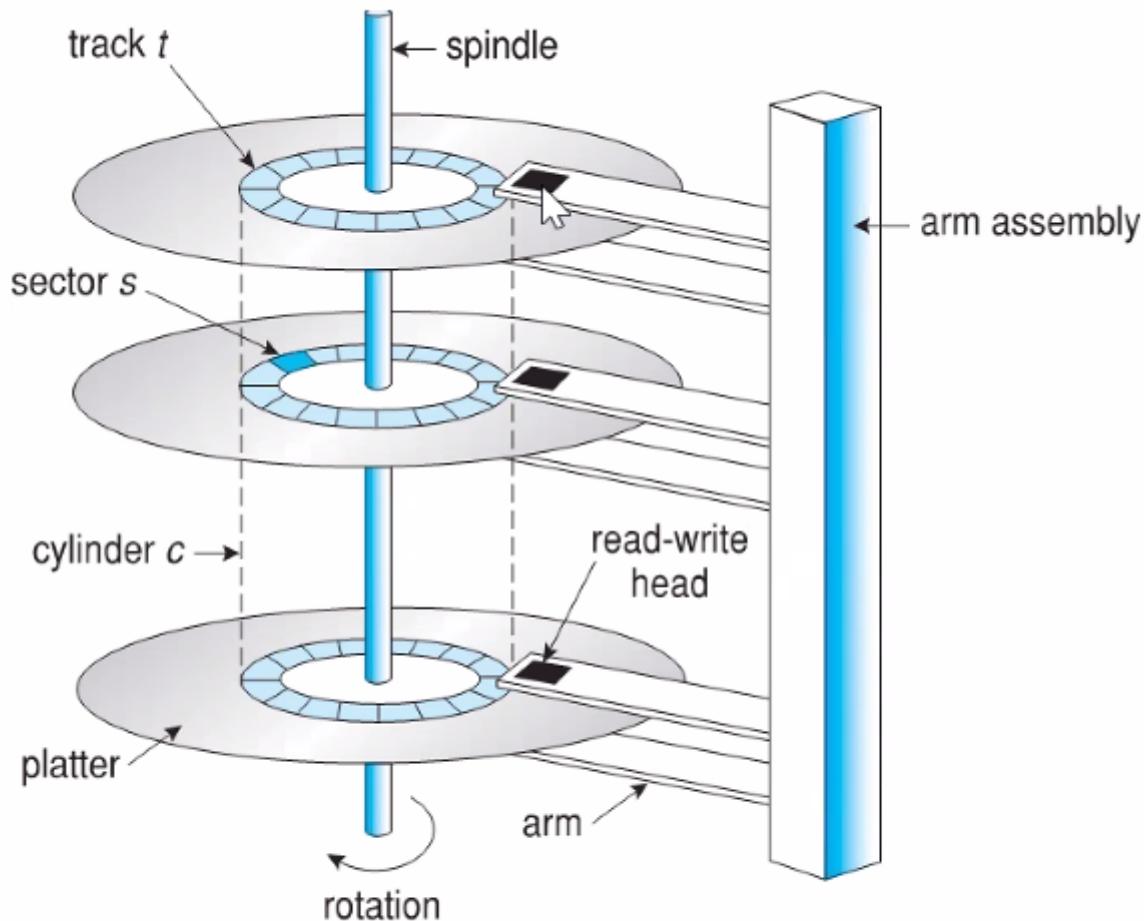
info stored by recording magnetically on platters

read info by detecting magnetic pattern on platters

read-write head “flies” just above each surface of every platter. Head cannot touch the platter or else it causes damage to the disk.

The surface of a platter is logically divided into circular **tracks** which are subdivided into **sectors**.

The set of tracks at a given arm position make up a cylinder.



## HDD moving-head disk mechanism.

The heads themselves cannot move independently of one another as they all share one physical arm assembly but each platter can rotate independently of one another.

Drives rotate at 60-250 times per second (can vary depending on brand and usecase). Rotation speed relates to transfer rates.

Transfer rate is rate at which data flow b/w drive and computer.

This can be improved by other things.

Positioning time (random-access time) is time to

1. move disk arm to desired cylinder (seek time)
2. time for desired sector to rotate under the disk head (rotational latency)

Head crash results from disk head making contact with the disk surface which is really bad. This can be caused by sudden power off. *Prof says and it is true, better to be safe than sorry.*

*But it's also worth noting that a lot of modern drives have mechanisms to prevent head crashes as a result of sudden power loss. Mechanism in place to automatically retract*

the head upon power loss detection. Not infallible mechanisms and not in all drives, especially old ones, but it's worth noting.

Platters range from .85" to 14" (historically). More commonly they will be in 3.5, 2.5, and 1.8 inch sizes

range from 30gb to 3tb per drive we've gotten much larger from then

disks can be removable

performance varies.

## Nonvolatile Memory Devices

electrical rather than mechanical

solid-state-disk (SSD)

- flash-memory based NVM
- used in a disk-drive-like container

USBs also fall under this category

pros:

- can be more reliable than HDDs
- faster
  - no moving parts
    - no seek time
    - no rotational latency
- less power hungry

cons:

- more expensive per mb
- less capacity

## Disk Attachment

A secondary storage device is attached to a computer by the system bus or an I/O bus.

Several busses available, including **advanced technology attachment (ATA)**, **serial ATA (SATA)**, **eSATA**, **serial attached SCSI (SAS)**, **universal serial bus (USB)**, and **fibre channel (FC)**.



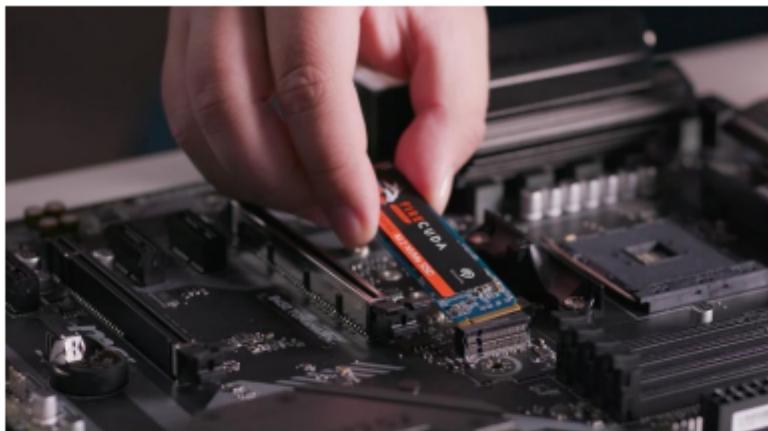
Most common is SATA.

<http://www.cablesonline.com/seratacab.html>

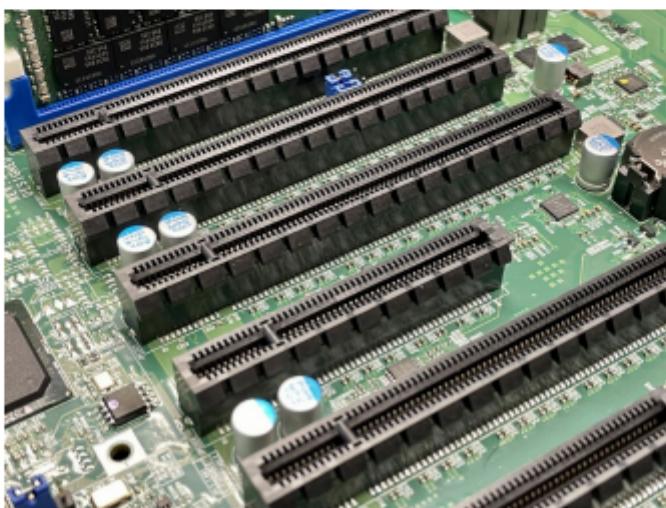
NVM is much faster than HDD

new fast interface for NVM called NVMe (NVMe) connects directly to the PCI (peripheral component interconnect) bus.

PCI provides a high-speed data path b/w cpu and peripher devices such as netowrk adpters, sound cards, graphics cards, and storage controllers



<https://www.seagate.com/ca/en/blog/what-is-an-nvme-ssd/>



<https://www.crystallugged.com/knowledge/what-is-pcie-slots-cards-lanes/>

## Address Mapping

disk drives are addressed as large 1-dimensional array of logical blocks

- 1 block is the smallest unit of transfer
  - low-level formatting creates logical blocks on physical media
  - each logical block maps to a physical sector or semiconductor page
  - logical block address is easier for algos to use than a sector, cylinder, or track
- 

#### Knowledge Check

- read or write performance depends on
  - bit transfer rate, seek time, and rotational latency
- one advantage of HDDs over SSDs is
  - HDDs are cheaper per megabyte than SSDs

Note:

## HDD Scheduling

OS wants to get a fast access time and high disk bandwidth

minimize seek time

- can be minimized by minimizing seek distance

bandwidth

- total bytes transferred / total time b/w first req and completion of last transfer
- bytes transferred / time elapsed for transfer

improve both by managing order in which i/o requests are serviced

**There are many sources of disk I/O request**

- OS.
- System processes.
- Users processes.

I/O request includes input or output mode, disk address, memory address, number of sectors to transfer.

Several algorithms exist to schedule the servicing of disk I/O requests.

- FCFS
- SCAN
- C-SCAN

We illustrate scheduling algorithms with a request queue (0-199) for the following cylinders:

**98, 183, 37, 122, 14, 124, 65, 67**

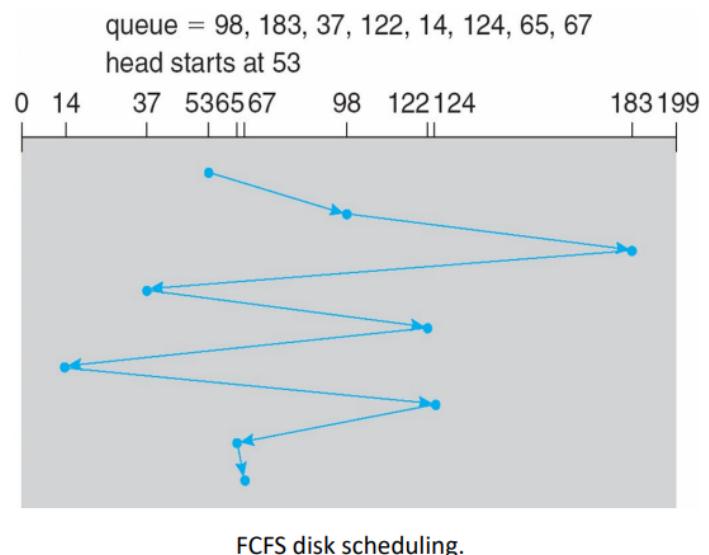
Assume the head pointer is at **53**.

## First Come First Serve

The simplest form of disk scheduling is the first-come, first-served (FCFS) algorithm (or FIFO).

Total head movement of **640** cylinders.

- $45+85+146+85+108+110+59+2$



## SCAN

disk arm starts at one end of the disk, and moves toward the other end, servicing requests until it gets to the other end of the disk.

once it gets to the end of disk it goes back the other way

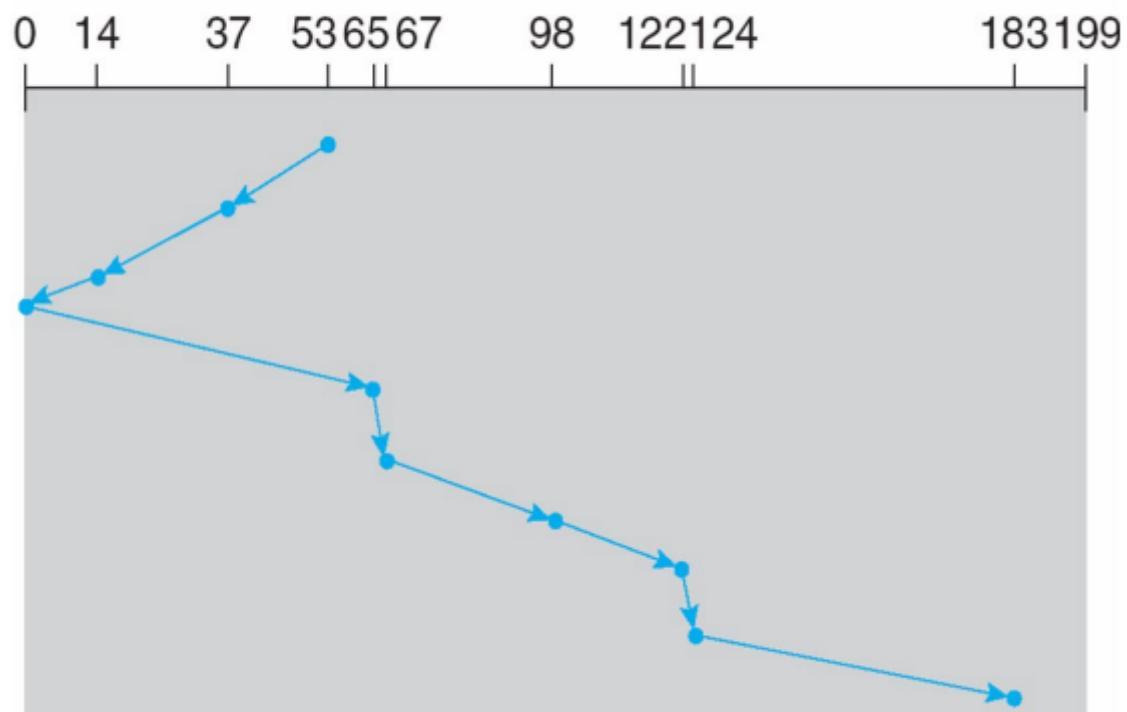
the head continuously scans back and forth across the disk

also known as the elevator algo

prof left this for us to read

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53



SCAN disk scheduling.

## C-SCAN

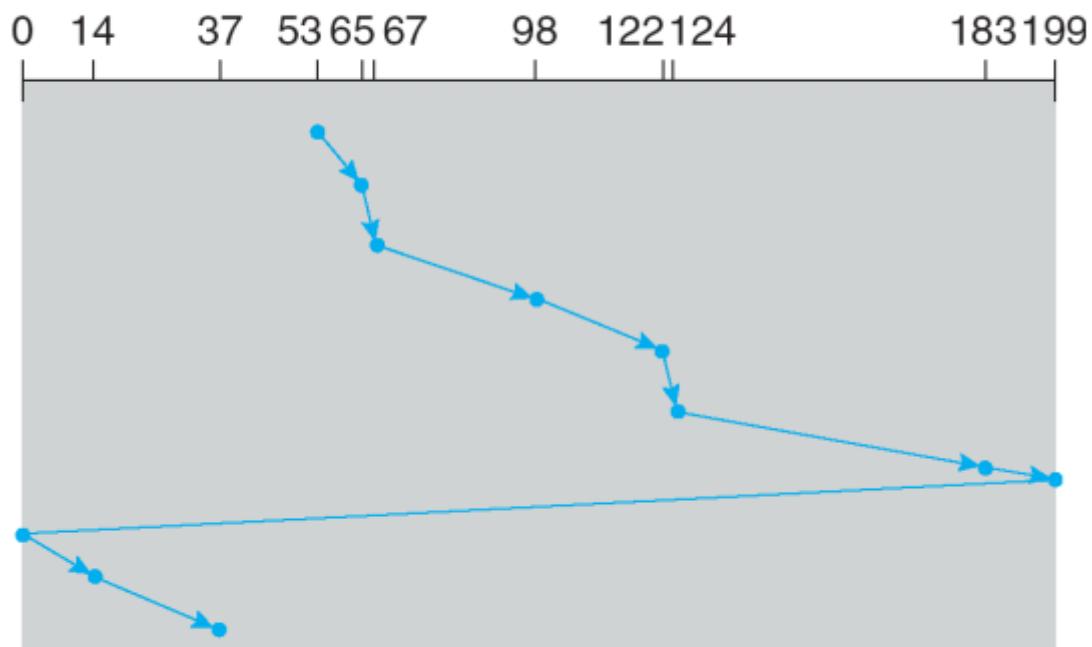
circular scan

higher seek time but more uniform wait time than SCAN

does the same motion of running across the whole disk but it doesn't service on the return trip, just goes back to the start immediately

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53



### C-SCAN disk scheduling.

SCAN and C-SCAN perform better for systems that place a heavy load on the disk

---

## Knowledge Check

- Consider a disk queue holding requests to the following cylinders in the listed order: 116, 22, 3, 11, 75, 185, 100, 87. Using the FCFS scheduling algorithm, what is the order that the requests are serviced, assuming the disk head is at cylinder 88 and moving upward through the cylinders?
  - 100 - 116 - 185 - 87 - 75 - 22 - 11 - 3
  - 87 - 75 - 100 - 116 - 185 - 22 - 11 - 3
  - 100 - 116 - 185 - 3 - 11 - 22 - 75 - 87
  - 116 - 22 - 3 - 11 - 75 - 185 - 100 - 87

d since it's first come first serve

## NVM Scheduling

nvm devices do not contain moving disk heads and commonly use a simple FCFS policy

FCFS is used but adjacent requests are combined for optimization purposes

NVM best perform at random I/O while HDD excels at sequential I/O

**FCFS scheduling is generally used in NVMs, because**

- a) write service time is not uniform in NVMs.
- b) it provides much higher raw throughput compared to HDDs.
- c) NVMs do not contain moving disk heads.
- d) the write performance depends on how “worn” the media is.

c. NVMs do not contain moving disk heads

## **RAID Structure**

Redundant Array of Inexpensive (or independent) Disks

a set of physical disk drives viewed by the operating system as a single logical drive

data is copied onto multiple disks to provide fault tolerance

raid is arranged into different levels *and these different levels are different from each person you may ask lmao*

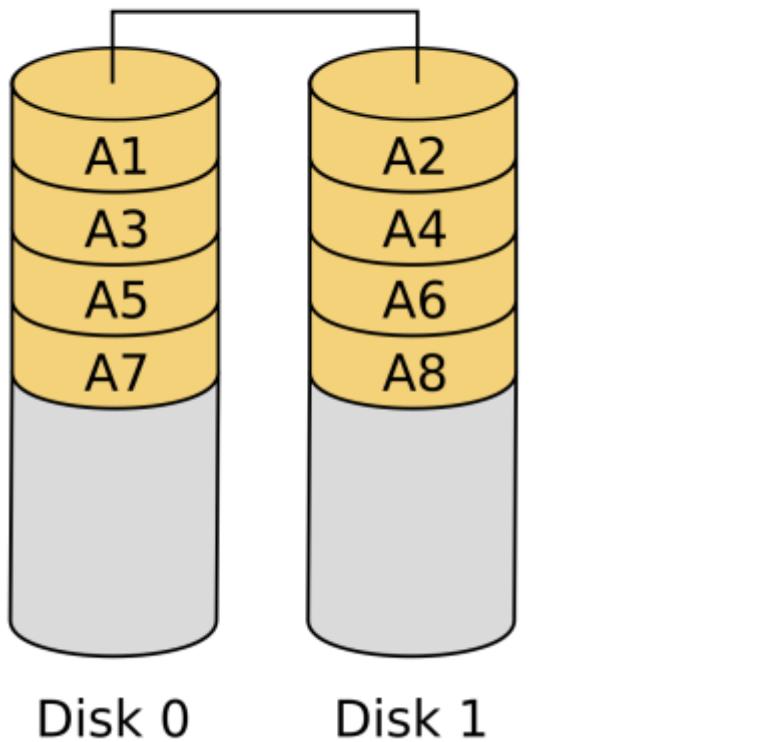
data copying techniques which differ based on raid level

- data mirroring
  - logical disk consists of two physical drives
  - every write carried out on both drives
- data striping
  - splitting the bits of each byte across multiple drives
    - bit-level striping
  - blocks of a file are striped across multiple drives
    - block-level striping

typically a small number of hot-spare disks are left unallocated, automatically replacing a failed disk and having data rebuilt onto them

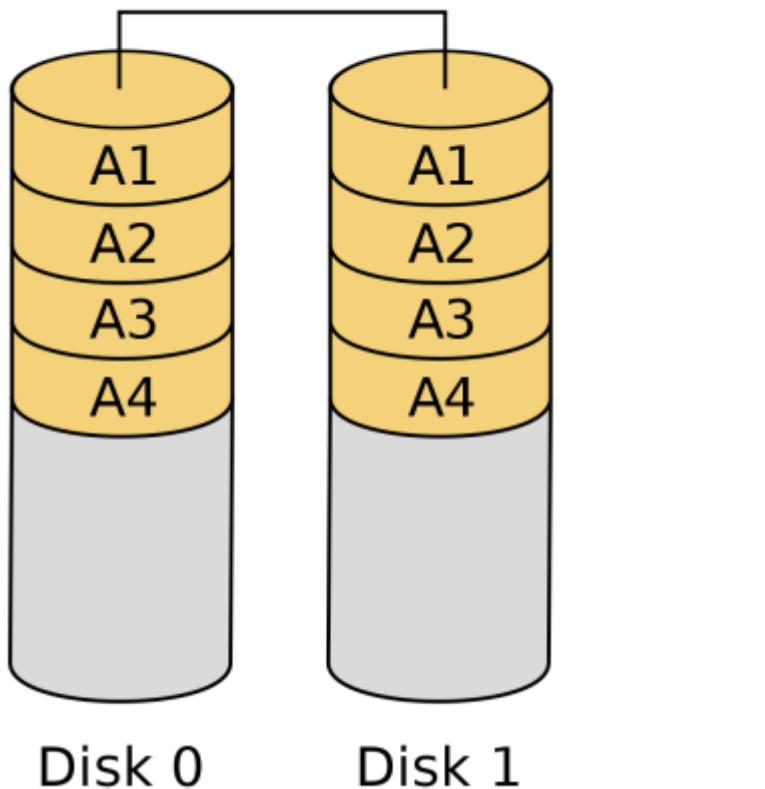
levels and their features (that Ruba thinks is good enough to give us an idea to jump off of)

# RAID 0



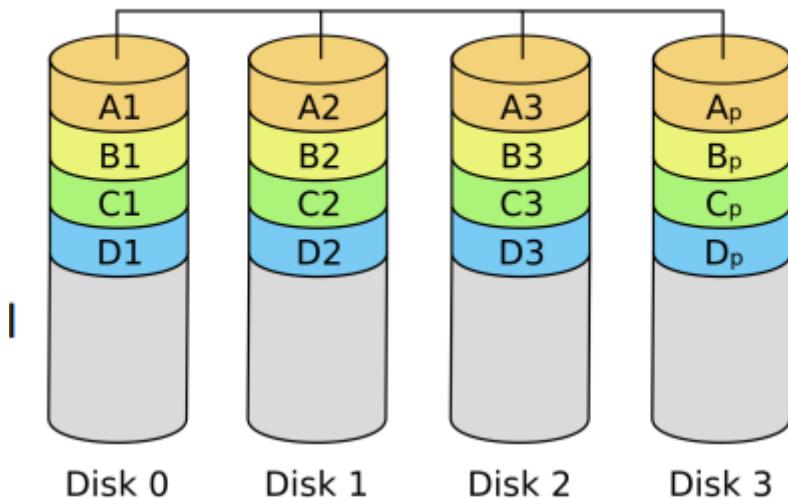
- logical disk is divided into stripes
  - prof is assuming byte level for this
- user and system data are distributed across all of the disks in the array
- doesn't provide fault tolerance
  - failure of one disk causes complete data loss
- improves performance
  - 2 heads instead of 1 makes accessing data faster
- only helps with performance
- might actually be worse for fault tolerance due to the striping since 1 disk failing corrupts the data on the remaining disk

# RAID 1



- drive mirroring present
- when a drive fails the data may still be accessed from the second drive
  - actual fault tolerance
- main disadvantage is cost

## RAID 4



| Input |   | Output  |
|-------|---|---------|
| A     | B | A XOR B |
| 0     | 0 | 0       |
| 0     | 1 | 1       |
| 1     | 0 | 1       |
| 1     | 1 | 0       |

Redundant disk capacity is used to store parity information which guarantees data recoverability in case of a disk failure.

parity is calculated across corresponding blocks on all data disks

- for every stripe of data blocks, a parity block is calculated

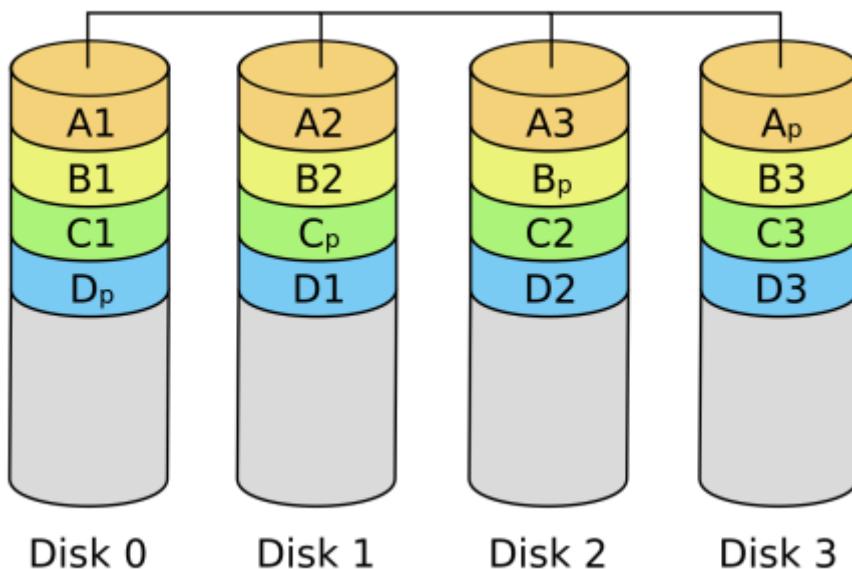
ex:

- an array of three drives where  $A_1 - A_3$  contain data, and  $A_p$  is the parity disk.
- calculate using
- $A_p = A_3 \oplus A_2 \oplus A_1$
- $A_1 = A_p \oplus A_3 \oplus A_2$

If we lose a disk we can recover the data using the parity calculations.

Only 1 disk

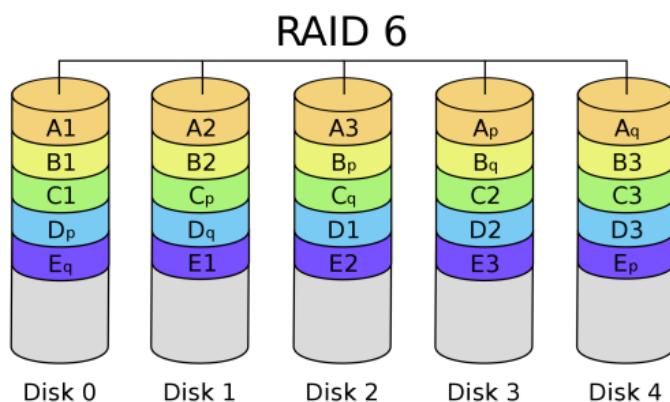
# RAID 5



- distribute parity blocks across all disks
  - interleaved distributed parity
  - typical allocation is a round-robin scheme
- avoids potential overuse of a single parity drive which can occur with RAID 4
- RAID 5 is the most common parity RAID
- can lose any single disk

# RAID Level 6

- Similar to RAID-5 but stores extra redundant information to guard against multiple drive failures.
  - **2 blocks** of redundant data are stored for every 4 blocks of data—compared with 1 parity block in level 5—and the system can tolerate **two drive failures**.
- Striping is done at the block level in RAIDs 4,5, and 6.



## Knowledge Check

- a RAID structure
  - c. stands for redundant arrays of inexpensive disks

- b. is primarily used to ensure higher data reliability
- in most RAID implementations, a hot spare disk is not used for data, but is configured for replacement should any other disk fail
  - true

# I/O Systems

## Overview

I/O management is a major component of operating system design and operation:

- Important aspect of computer operation.
- I/O devices vary greatly.
- Various methods to control them.
- Performance management.
- New types of devices are frequent.

Ports, buses, and device controllers connect to various devices.

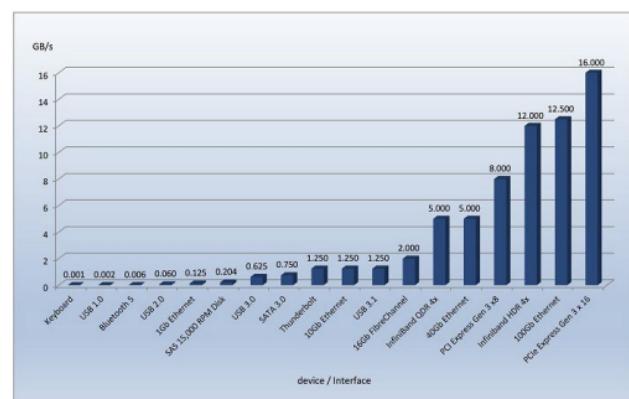
**Device drivers** encapsulate device details:

- Present uniform device-access interface to I/O subsystem, like system calls provided for apps by OS.

# I/O Hardware

## I/O Hardware

- Incredible variety of I/O devices, classified as:
  - Storage (disk, tape).
  - Transmission (network, Bluetooth).
  - Human-interface (monitor, mouse, keyboard, audio).

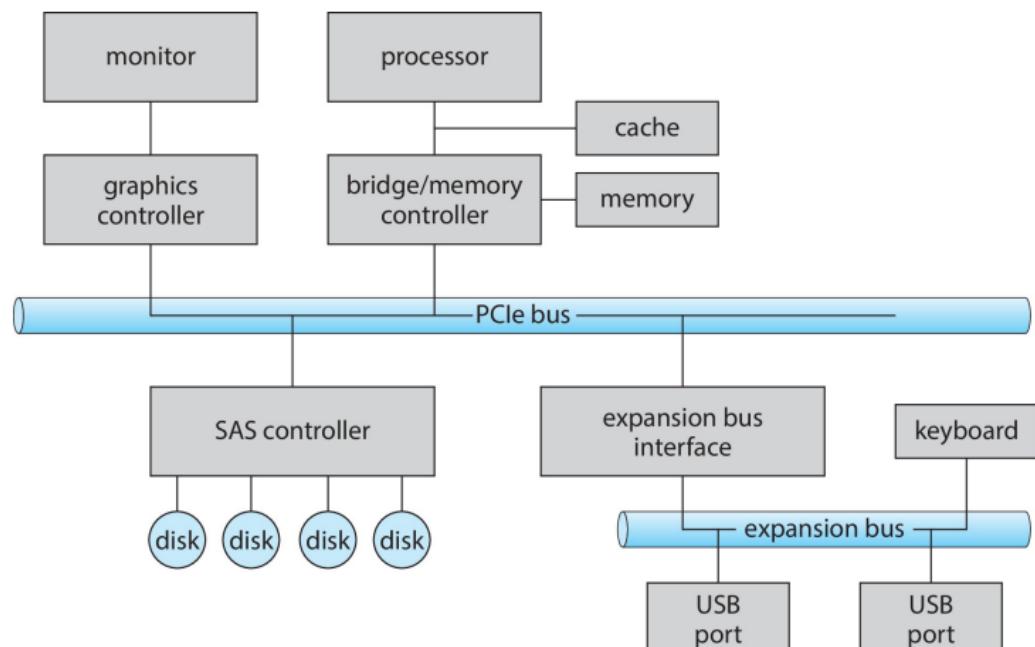


Common PC and data-center I/O device and interface speeds.

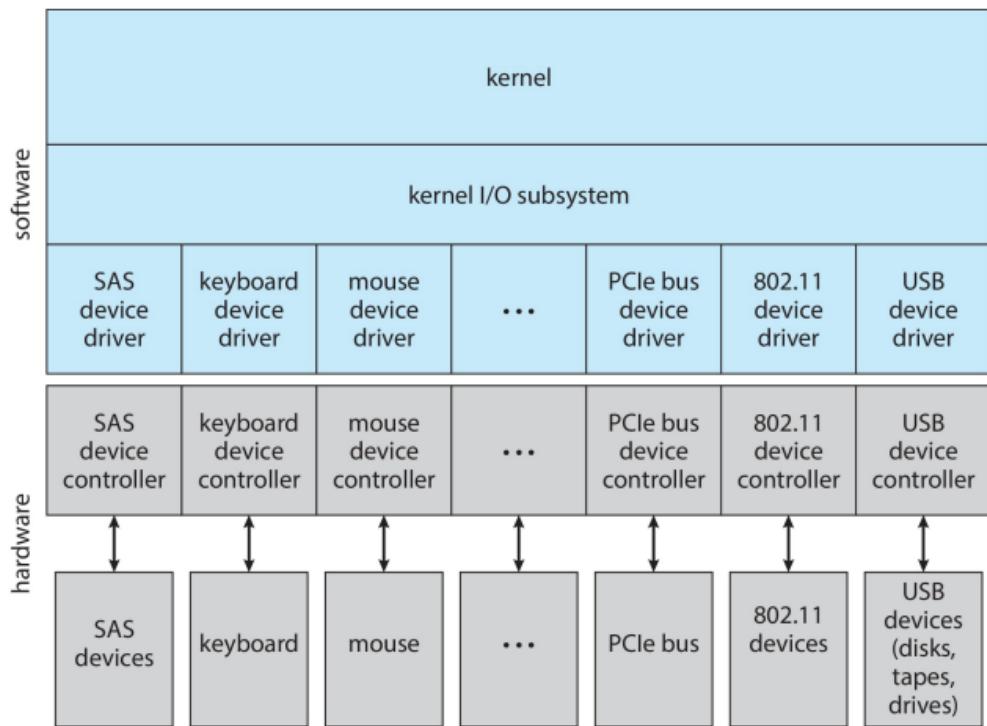
# I/O Hardware Common Concepts

- A device communicates with a computer system by sending signals over a cable or even through the air.
  - **Port:** Connection point for the device.
  - **Bus:** A set of wires and a rigidly defined protocol that specifies a set of messages that can be sent on the wires.
    - **PCI** bus common in PCs and servers, PCI Express (**PCIe**).
    - **Expansion bus** connects relatively slow devices.
    - **Serial-attached SCSI (SAS)** common disk interface.
- **Controller (host adapter)** – electronics that operate port, bus, device.
  - Contains processor, microcode, private memory, bus controller, etc.

## A Typical PC Bus Structure



# A Kernel I/O Structure



next week is security lect and exam review

# Security and Protection

## Security

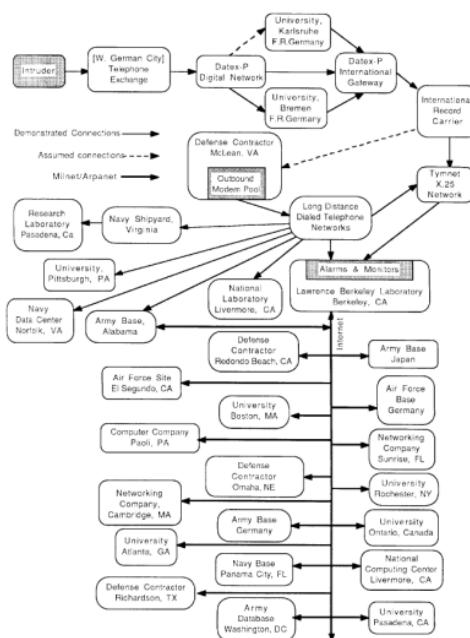
“I will tell you a few stories today. Everybody loves a good story that is perfectly fine but they also have some kind of impact on hardware and software design” - Professor

# The Cuckoo's Egg



10 Months Later

**1986:** A \$0.75 accounting error!



Simplified Connectivity and Partial List of Penetrated Sites

Starting at Lawrence Berkely Laboratory (LBL or Berkely Lab) in California had the 75 cent error. They were performing research for some energy dept. They were connected to many other institutions, army bases, and defense contractors over the combination arpanet and milnet.

How could this happen?

prof wanted to show us a cool little video

Clifford Stoll - Watch from 2:22 to 5:05 - Movie: The KGB, the Computer, and Me available at <https://www.youtube.com/watch?v=PGv5BqNL164>

Clifford Stoll was a grad student at UC Berkely. He pointed out the 75 cent error after taking a summer job there.

Stoll found out that a single account was responsible for this. Surely, just a prankster of u cali. This account was the acct of a prof who was on sabattical after all.



clueless

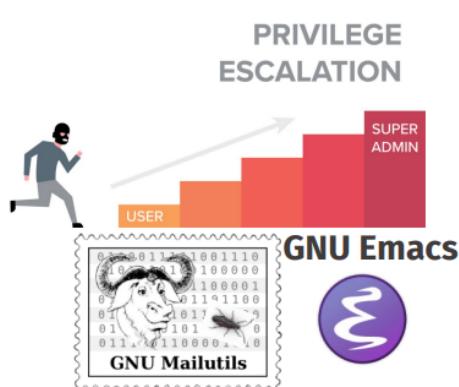
|           |       |
|-----------|-------|
| Starr     | 79987 |
| Vine      | 67088 |
| Treppasso | 59899 |
| Carlson   | 24517 |
| Norman    | 34257 |
| Hunter    |       |
| Benson    | 87459 |
| Johnson   | 56739 |
| Schlar    | 74578 |

|           |       |
|-----------|-------|
| Starr     | 79987 |
| Vine      | 67088 |
| Treppasso | 59899 |
| Carlson   | 24517 |
| Norman    | 34257 |
| Benson    |       |
| Johnson   | 87459 |
| Schlar    | 56739 |
|           | 74578 |

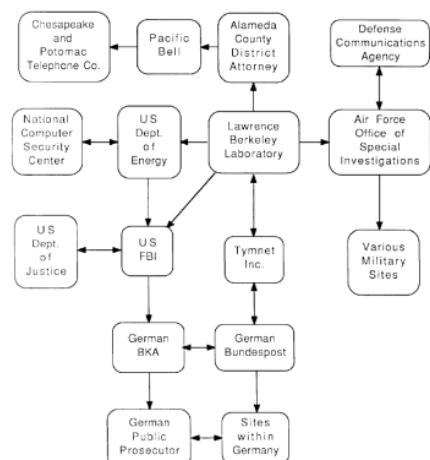
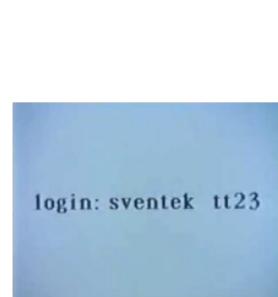
The account from a prof who was on sabatical.

Clifford got a notification that someone was trying to log into docmaster.

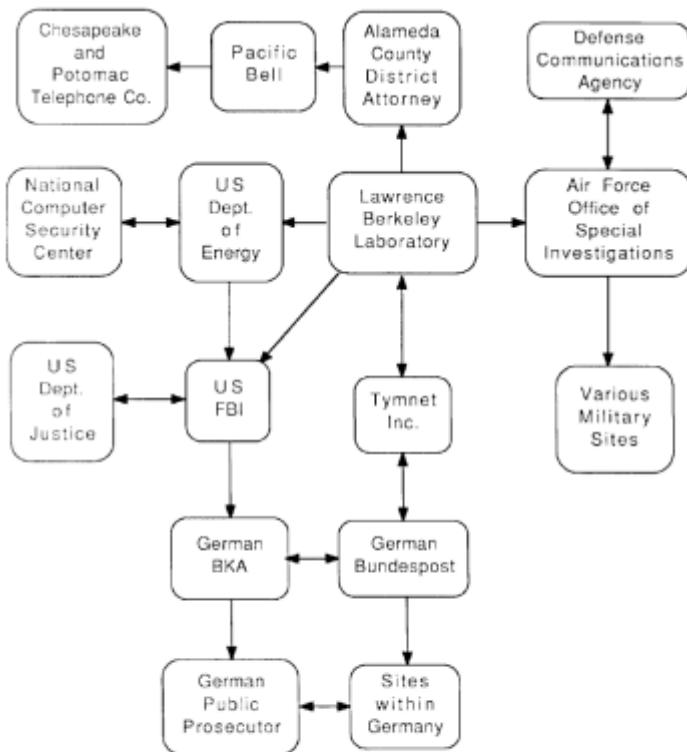
## Privilege Escalation



The attacker used a subtle bug in the Gnu-Emacs program **movemail** to obtain system-manager privileges

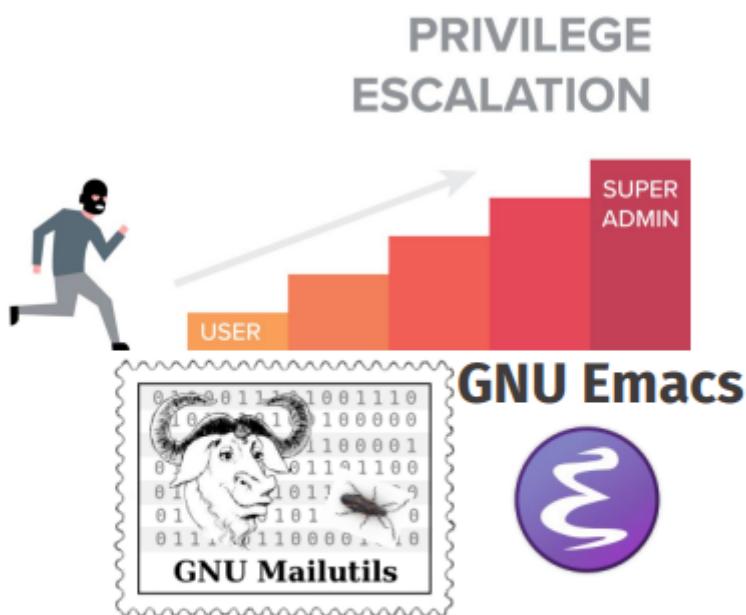


tons of pinging from the person trying to login and creating tons of accounts



## Simplified Communication Paths between Organizations

LBL was a part of a huge network of very important organizations' computers. An attacker started from Tymnet machines.



The attacker used a subtle bug in the Gnu-Emacs program [movemail](#) to obtain system-manager privileges

By gaining the privileges the attacker was able to move between machines.

## movemail

- worked by changing the file ownership and moving it to the designated user's folder
- permanently ran as sudo
  - `movemail` can write into protected system folders
  - given access to any and every folder
  - also had access to `atrun` folder
- the attacker
  - wrote a script
  - used `movemail` to put the script in the `atrun` folder
  - set the script to execute in 5 minutes
  - the script gave his account superuser privileges

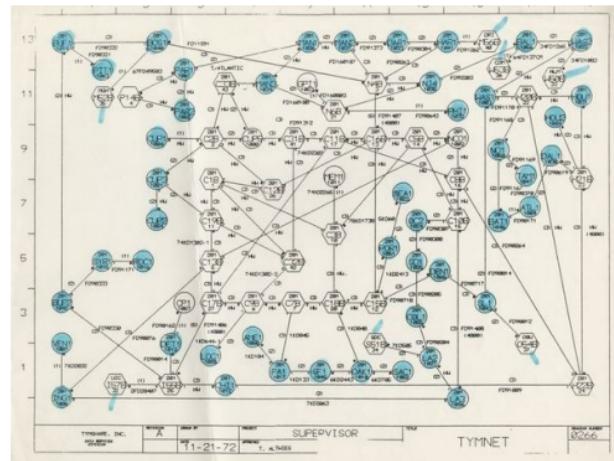
The attacker used the new privileges to make more accounts.

An original thought was to start locking down all those accounts.

But Stoll wanted to get to the root of it. Who is the hacker, what are they doing, and why?

## Tymnet

- Lawrence Berkeley Laboratory (LBL) relies on Tymnet for connectivity, so Stoll asked them for help.
- Tymnet was a "gateway" service that a user called into that routed them to any one of a number of computer systems that also used the service.



Tymnet

Stoll had a gf (SHOCKING) who suggested making some fake docs to lure the attacker to stay in a machine long enough to trace the connection.

This is what would come to be known as a honeypot.

The honeypot of fabricated docs that were supposed described the Strategic Initiative (SDI) aka the Star Wars program.

The honeypot worked and they traced it back to a West German man named Markus Hess.

It took 6 months to get to this point.

3 years later Stoll was flown into Germany to testify.

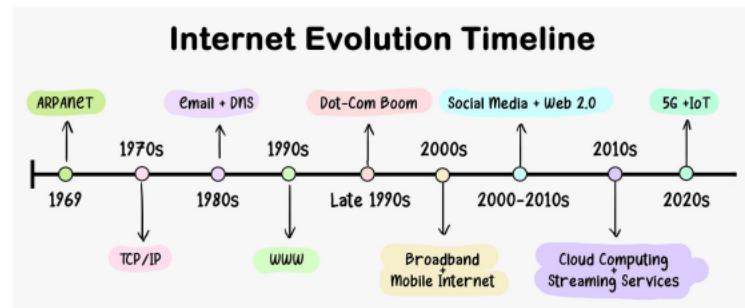
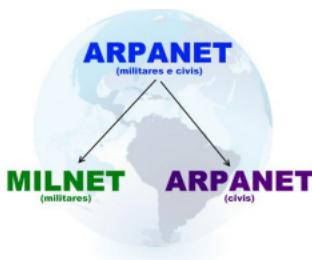
prof likes this summary

Watch The 1985 KGB Hack (6:43 minutes) at  
<https://www.youtube.com/watch?v=dFagMZLMI6o>

There was no huge damage in the end though.

## Sentence

- Hess was able to attack **400** U.S. military computers by using Berkley Lab to "piggyback" to ARPANET and MILNET.
- Hess was found guilty of espionage and was given a 20-month suspended sentence.



## Why Do We Care?

- Stoll spearheaded an investigation that laid bare the specific tactics, techniques, and procedures (TTPs) the adversary used to achieve their goals.
- Many system designers have incorporated lessons that arose from Stoll's article describing the team's efforts, "Stalking the Wily Hacker."
- Paper: Stalking the Wily Hacker :  
<https://dl.acm.org/doi/pdf/10.1145/42411.42412>

**ARTICLES**

### STALKING THE WILY HACKER

An astronomer-turned-sleuth traces a German trespasser on our military networks, who slipped through operating system security holes and browsed through sensitive databases. Was it espionage?

CLIFFORD STOLL

In August 1986 a persistent computer intruder attacked the Lawrence Livermore National Laboratory (LLNL), trying to keep the intruder from us, took the novel approach of allowing him access while we printed out his activities and traced him to his source. This trace back was harder than we expected, requiring nearly a year of work and the cooperation of many organizations. This article tells the story of the break-ins and the trace, and sums up what we learned.

We approached the problem as a short, scientific exercise, dividing the task among several people who tried breaking into our system and document the exploited weaknesses. It became apparent, however, that rather than innocuously playing around, the intruder was using our computer as a hub to reach many others. His main interest was in computers operated by the military and by defense contractors. Targets and keywords suggested that he was attempting espionage by remotely entering sensitive computers and stealing data.

LLNL is a research institute with few military contracts and no classified research (unless our sister laboratory, Lawrence Livermore National Laboratory, which has several classified projects). Our computing environment is typical of a university: widely distributed, heterogeneous, and fairly open. Despite this lack of classified computing, LLNL's management decided to take the intrusion seriously and devoted considerable resources to it, in hopes of gaining understanding and a solution.

The intruder conjured up no new methods for breaking operating systems; rather he repeatedly applied techniques documented elsewhere. Whenever possible, he used known security holes and subtle bugs in different operating systems, including UNIX,\* VMS,\* VM-TSO,\* EMBOSS,\* and SAIL-WAITS. Yet it is a mistake to assume that one operating system is more secure than another. Most of these break-ins were possible because the intruder exploited common blunders.

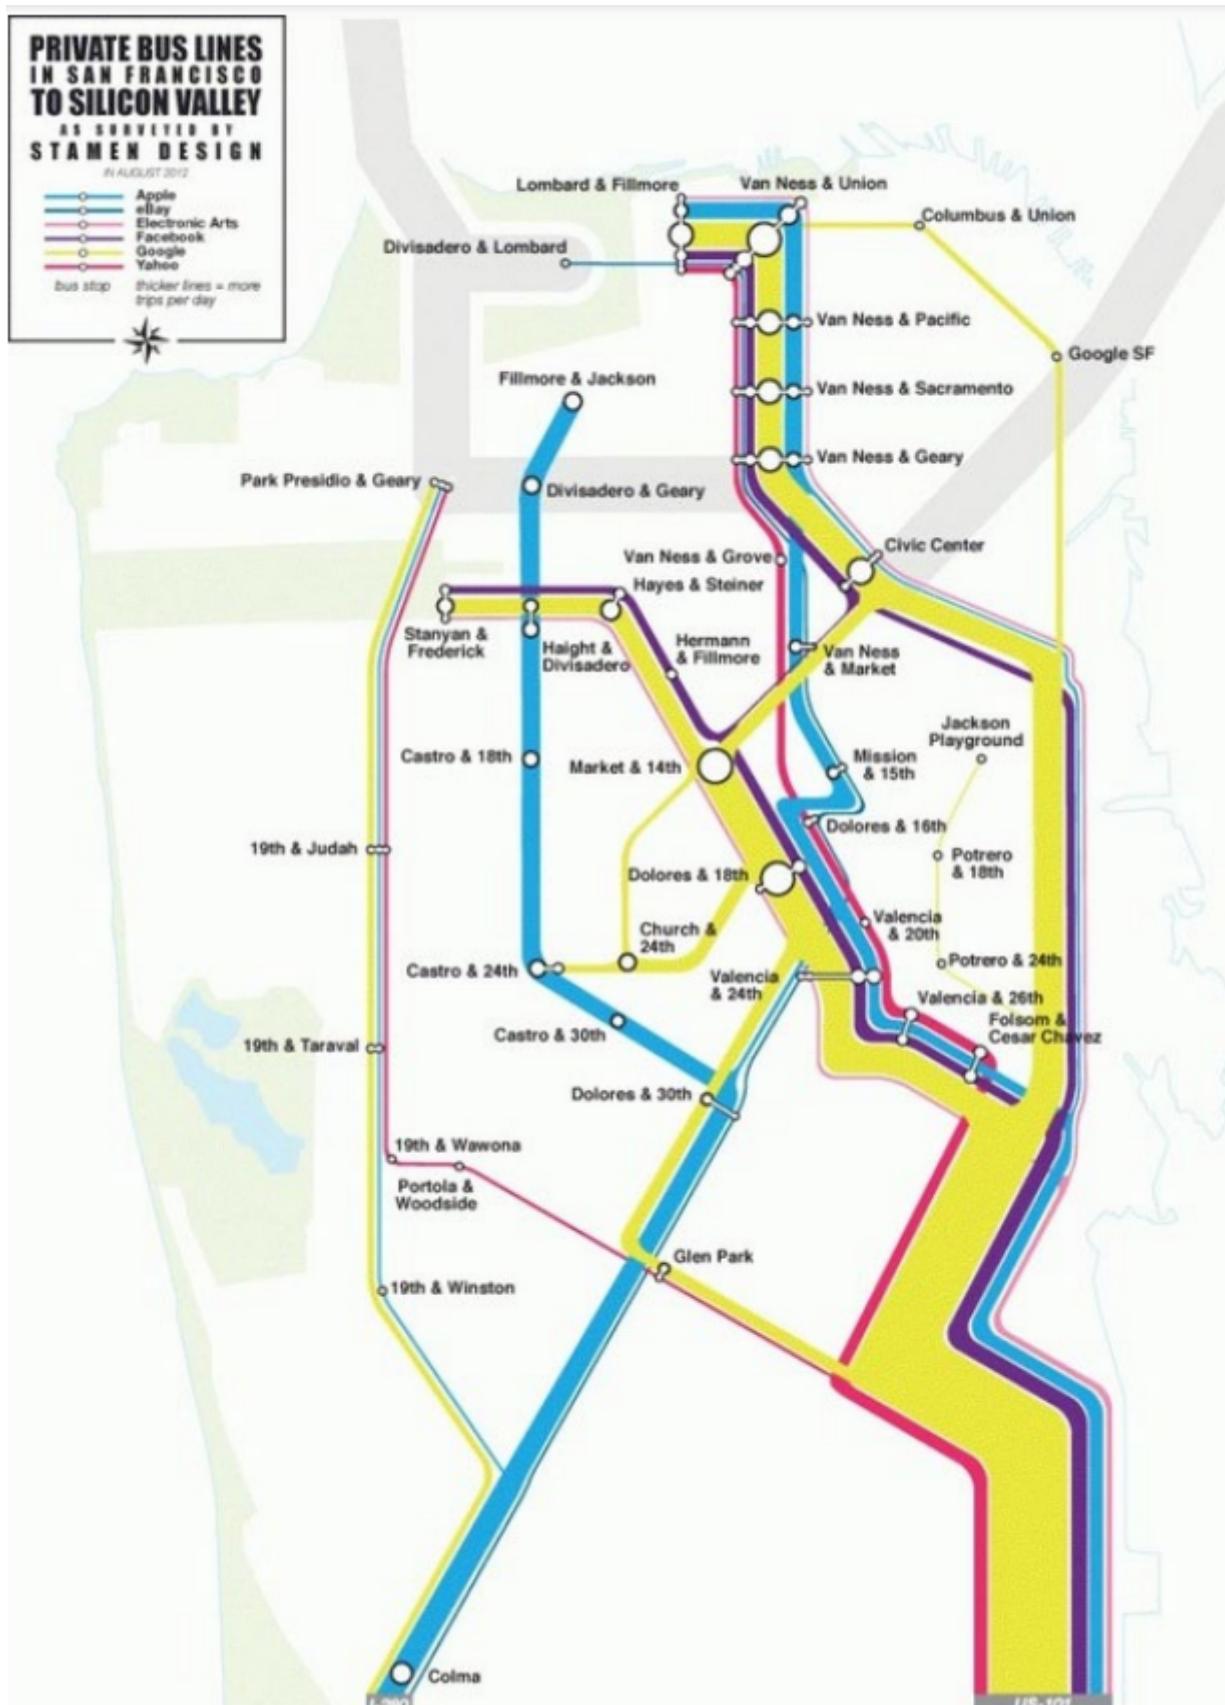
New story: passwords and power drills

This story ends with someone using a power drill to crack open a safe.

In 2012, google had their own internal password manager to store and share secrets for third-party services that don't support better auth mechanisms at the time.

Initially for very few people.

This was used for the auth for shuttlebus wifi for google employees.

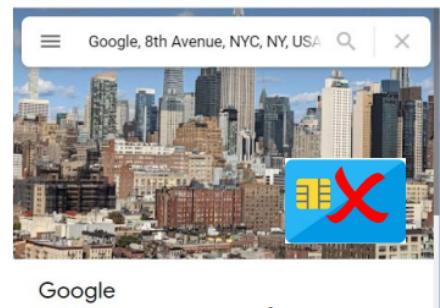


In 2012, Stamen Design tracked the movements of tech shuttles and created a map of them – a map that was never created or released by the companies themselves who keep routes and stops secret.

**September 27, 2012**  
**From:** Google  
transportation team  
**To:** 1000s of employees  
**Subject:** The WiFi  
password has changed .



This crashed the service b/c everyone hopped on to load up the password.



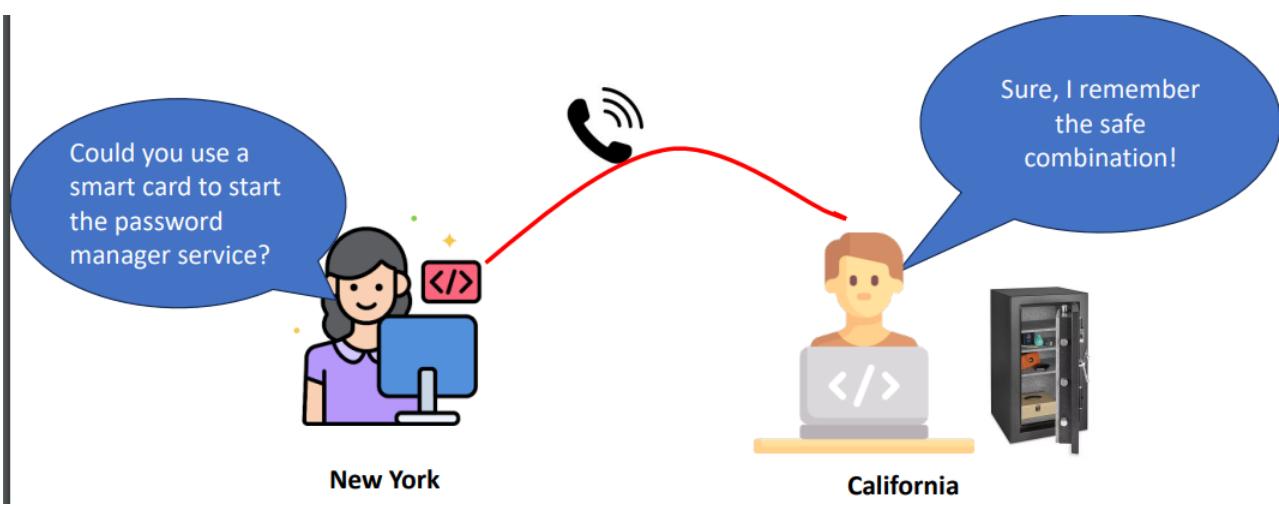
Insert Hardware  
Security Module Smart  
Card.

Could you use a  
smart card to start  
the password  
manager service?



I would love to but....  
The smart card is in  
the safe, and the safe  
combination is in the  
password manager!





▪ One hour later....



It took an hour for everyone to realize that they were putting the card in wrong.

password manager failure was triggered by an availability problem: poor load-balancing.

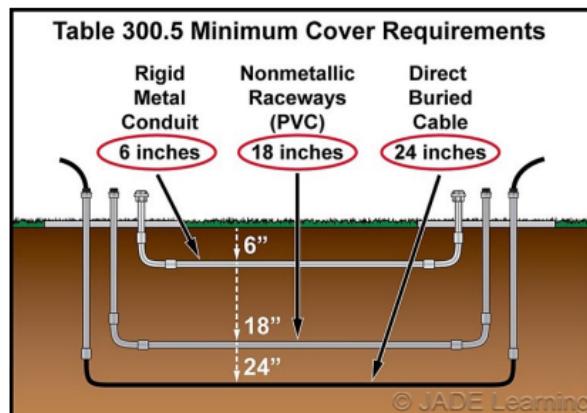
recovery complicated by multiple measures designed to increase system sec

- improve breakglass mechanisms
  - mechanisms to bypass policies to allow engineers to quickly resolve outages

march 2012 there was a power outage at a google datacenter in belgium

turn out a cat destroyed some nearby external psu things

By studying **how complex systems fail** in similar ways, Google has been able to adopt resilient design practices when suspending, burying, and submerging cables around the world.



---

in 2006 an open-source dev noticed that the memory debugger Valgrind was reporting warnings about memory used prior to initialization.

to eliminate the warnings, the dev removed two lines of code.



caused a bug in OpenSSL's pseudo-random number gen. Only got seeded w/ process ID.

lead to things being easily brute force breakable.

understanding the systems' adversaries is critical to building reliable and secure systems

adversaries to reliability are typically benign, routine hardware failures or config changes that make unintended effects

security adversaries are typically humans trying to make the system work in a way that they want

## How Systems Fail

many reasons including

- reliability
  - accidental failures
- usability
  - failures from operating mistakes
- design and goal oversights
  - oversights, errors, and omissions during designing
- security
  - intentional failures created by intelligent parties

all of the above are related

## What is Security

What does **security** mean?

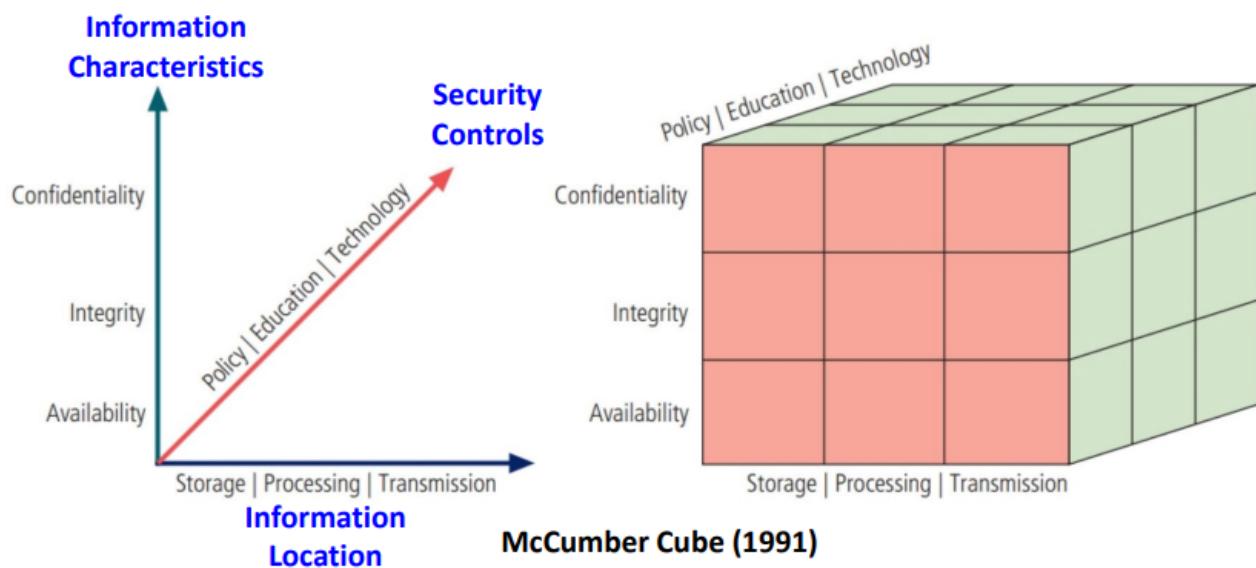
- Often the hardest part of building a secure system is figuring out what security means ("threat modelling").
- Who are the **stakeholders** for which we are considering "security"?
- What are the **assets** to protect?
- What are the **threats** to those assets?
- Who are the **adversaries**, and what are their **resources**?
- What is the **security policy or goals**?

A system is secure if its resources are used and accessed as intended under all circumstances.

CNSS Security Model

- Committee on National Security Systems (CNSS)

- serves as standard for understanding many aspects of infosec



still used today.

ahead of his time

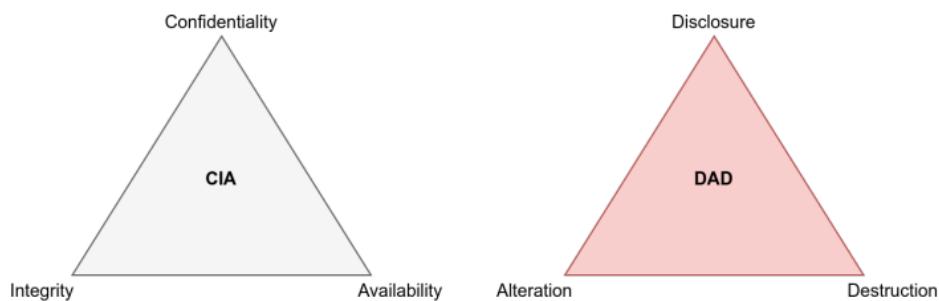
3 dimensions of infosec

- info characteristics
  - confidentiality
  - integrity availability
- security controls
  - policy
  - education
  - technology
- information location
  - storage
  - processing
  - transmission

Each cube cell is a specific 3 way intersection that we should be thinking about

# CIA Triad

- The industry standard for computer security since the development of the mainframe.
- Based on three characteristics that describe the utility of information: Confidentiality, Integrity, and Availability.



Main security violations include breach of confidentiality, integrity, or availability.

- Example: Denial-of-service (DoS) attacks prevent legitimate use of the system.
- Security violations of the system can be categorized as intentional (malicious) or accidental.
- It is easier to protect against accidental misuse than against malicious misuse.

## Information Security Elements

| Confidentiality                                                                                                                                                                            | Integrity                                                                                                                                                          | Availability                                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Information is not made available or disclosed to unauthorized individuals, entities, or processes.<br> | Data is protected from unauthorized modification over its entire lifecycle.<br> | Information and systems must be available when needed by the authorized users.<br> |

- Supporting the CIA triad of information security is the AAA framework:

| Authentication                                                                                                        | Authorization                                                                                                                                                              | Accounting                                                                                                                                               |
|-----------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p>Verifying the user identity.</p>  | <p>Following authentication, a user must gain authorization for doing certain tasks.</p>  | <p>Monitors the resources a user consumes during network access.</p>  |

### AAA processes

- start with identification
  - subject must provide an identity to a system to start the AAA processes
  - identities should be
    - unique
    - securely issued
  - non-descriptive of role
- authentication
  - proves identity of subj that tries to connect
  - subject must supply verifiable creds referred to as factors
    - single-factor
      - just one password
    - multifactor
      - two or more
      - password + duo mobile
      - bank card + pin
  - factor categories
    - knowledge - something you know
      - password, pin, answer to security questions
    - possession - something you have
      - one time auth password, smart cards, ubikey
    - biometrics - something you are
      - fingerprints, eye scan
  - you might notice that these days the message for an incorrect username or password are the same message, this is also a security measure
- authorization
  - determines what a user is allowed to do on a computer system or network
  - the model for this defines how access rights and permissions are granted
    - there are many different models
      - rule based
      - dynamic
      - discretionary

- accounting
  - auditing and monitoring what a user does w/ accessed resources
  - produce audit trail log
  - when the user accessed the resource
  - what the user did with that resource
  - when the user stopped using that resource

## Info Security Strats

don't just hope that it works out for you

off sec (red team) and defensive sec (blue team)

zero trust security: never trust, always verify

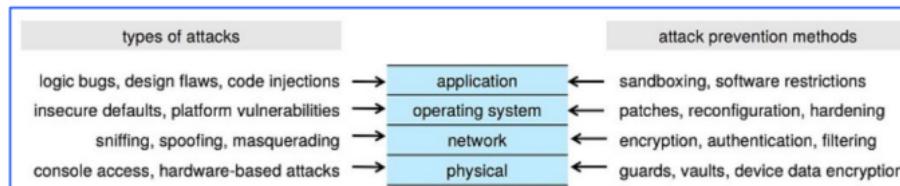
## Offensive Security

- Offensive Security: Break into computer systems, exploit software bugs, and find loopholes in applications to gain unauthorized access to them.
- A **red team** consists of security professionals who act as adversaries to overcome cyber security controls.



# Defensive Security

- Defensive Security: To protect a system, we must take security measures at four levels: physical, network, operating system, and applications.
- A **blue team** consists of security professionals who have an inside out view of the organization. Their task is to protect the organization's critical assets against any kind of threat.



The four-layered model of security.

## Zero Trust Security /1

- A security model that requires **strict identity verification** for every user and device trying to access resources on a private network, regardless of whether they are sitting within or outside of the network perimeter.



Source: seknov

# Zero Trust Security /2

- The Zero Trust model is based on a number of principles, most importantly:
  1. No Implicit Trust.
  2. Verification of Identity and Access.
  3. Principle of Least Privilege.
  4. Continuous Monitoring.
  5. Dynamic Policy Enforcement.



## Protection

### Goal of protection

defend the system from internal or external attacks

concerned with controlling access of processes or users to the resources of the computer system.

Protection problem

- ensure that each object is accessed correctly
- only accessed by those processes that are allowed to do so

## Principles of Protection

protection principles including but not limited to

- principle of least privilege
  - similar to need-to-know basis
  - properly set perms can limit damage in the case of bugs or abuse of perms
  - can be static or dynamic
    - static during life of system or process
    - dynamically changed by process as needed
      - privilege escalation
- compartmentalization
  - derivative concept regarding access to data

- protect each individual system component
- use specific permissions and access restrictions
- ex
  - we can have a virtual testing machine so that if the testing machine gets compromised we can kill it very easily and hopefully the attacker doesn't know to try and jump the hypervisor
  - we don't even have to do a full blown vm for some usecases
- audit trail
  - recording all protection-oriented activities
  - important to understand
    - what happened
    - why
    - what happened that shouldn't have happened

Defense in Depth

- application of multiple, sometimes redundant, defense mechanisms
- if a single control fails then the other controls can still protect the env and assets

# Exam Review

## Exam Housekeeping

Date: Mon 15 apr 2024 @ 0900

location: WSC 106 108 105

Duration: 2 hours

Format: e-class quiz with use of labtest mode

Scope: Everything

bring: id, mfa auth, pen and/or dark pencil, calculator (leave in backpack)

closed book, no

- notes
- book
- e-class resources
- slides

Provided with 5 scratch papers to be left.

Use of the calculator app is permitted.

All questions may be on the exam.

Number of questions, question types, and weight of questions are shown on exam.

Non-sequential, allowed to go back and forth between questions.

protocol:

- Arrive 15 minutes before exam.
- Put device into airplane mode after logging in.
- Leave photo id on desk
- no questions during test
- choose the **best** answer for multiple-choice or true-false questions
- write your assumptions
- don't worry too much about the answer format. just go with what the question is asking for
  - if you're unsure then there will be manual review
- deduct 5 marks if you ask the instructor/proctor to
  - explain a question
    - this will not be answered
  - verify understanding of question
  - verify answer
  - ask about answer format
  - voice concern of missing info in answer
    - do so after the exam

in the case of technical/machine related login failure, restart the machine and try again

All else fails, you will be provided a paper-based copy and retrieve your calculator.

## QnA

Q: are we allowed cheat sheets?

A: No.

Q: Will there be a practice exam?

A: Yes.

Q: Will there be more weight given to post midterm material?

A: There will be more weight but the exam is cumulative so study up on everything.

Q: Will there be any coding questions?

A: No. You won't be creating code to feed to a java or c compiler but we will be looking at snippets that we looked at in class. We may also be asked to write a small piece of code to implement a mutex.

Q: Do we still get to write “I don’t know” for 10%?

A: No

Q: Why not do a paper exam?

A: Prof likes e tests better

Q: Will we be given formulas?

A: Only Amdahl’s law type fomulae. Things like effective access time you need to memorize

Q: What will be the weight of the exam?

A: It’s still 40%. The weight of A2 was placed on to the other completed course materials

Q: Is there any sort of set seating arrangement between the rooms?

A: undecided

Q: Will the difficulty be similar to the midterm?

A; “It depends. Did you think it was Difficult?” “This is a very relative question”

Q: Can we opt-in to paper based?

A: No, it is only an emergency provision. Likely only a single digit number of paper-based tests will be on-site.

## Final Exam Practice

suggested study methods

- go through lecture notes and existing notes
- refer to textbook to attain deeper understanding
  - Operating System Concepts, 10th Edition, by Abraham Silberschatz, Greg Gagne, et al., (Available from York Library in Softcopy)
  - contains practice questions
- practice and understand assignments and tutorials
  - **tutorial, singular**
- use the practice exam given by professor
  - provides answers to the knowledge checks throughout the course
  - new knowledge checks added for chapters that didn’t have many
  - “there are only 117 questions that I encourage you to go through”
  - answer to on the following slide

what have we learned

1. computer system overview

- 2. os structures
- 3. process description and control
- 4. threads
- 5. synchronization and mutex
- 6. deadlock and starvation
- 7. cpu scheduling
- 8. main memory
- 9. virtual memory
- 10. storage management
- 11. Security and Protection

There are also sub-topics