

toc:

- Housekeeping
- Principles of Deadlocks
 - Deadlock
- Deadlock in Multithreaded Applications
- Deadlock Characterization
 - Resource Allocation Graph
- Methods for Handling Deadlocks
 - Prevention
 - Deny Hold and Wait
 - Deny No Preemption
 - Deny Circular Wait
 - Deadlock Avoidance
 - Safe State
 - Avoidance Algos
 - Resource-Allocation Graph Scheme
 - Banker's Algorithm
- Recovery from Deadlock

Housekeeping

Strike won't be affecting us too much

Principles of Deadlocks

Systems have different resources

Threads must:

- request a resource before using it
- release the resource after using it

The number of resources requested cannot exceed the total number of resources available in the system

- a thread cannot request 2 net interface if there is only 1
- can't ask for what we don't have

Request:

- Thread requests the resource
- if you can give it, then give it immediately
- If it can't be given right now then the requesting thread must wait until it can acquire the resource

Use:

- the thread can operate on the resource
- ex: use the mutex lock to access a process' critical section

Release:

- resource is released and made usable again

Reusable Resources:

- can be used safely by one process at a time
- not depleted by that use
- not consumable
- ex: processors, I/O channels, memory (main and secondary), I/O devices, data structures (files, databases, and semaphores)

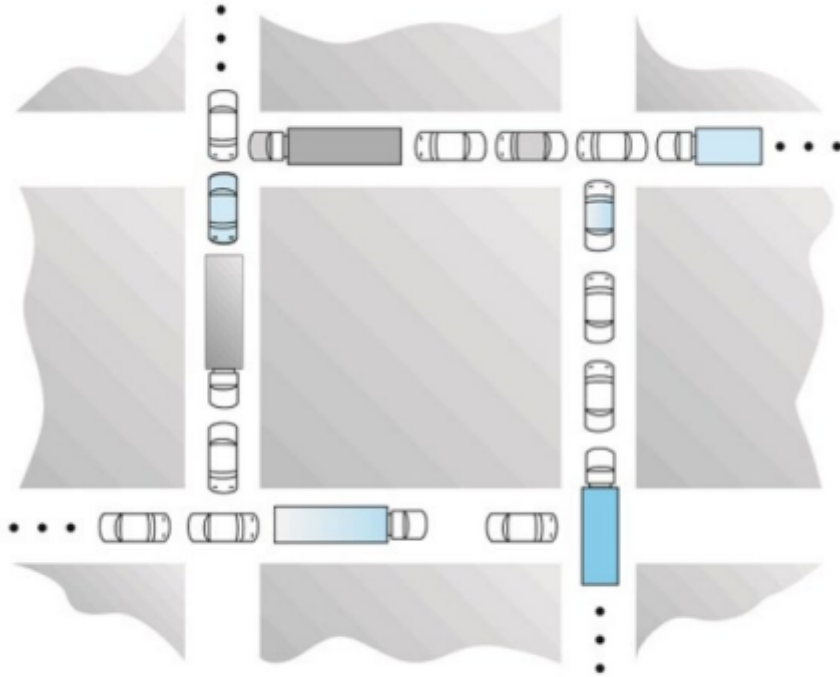
Consumable Resources:

- created/produced and destroyed/consumed
- ceases to exist after being acquired by the consuming process
- ex: interrupts, signals, messages, info in I/O buffers

Deadlock

A deadlock is the **permanent** block.

Two or more processes are waiting for the other process to release a shared resource.




Traffic Deadlock

A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set.

The block is permanent because none of the events ever get triggered.

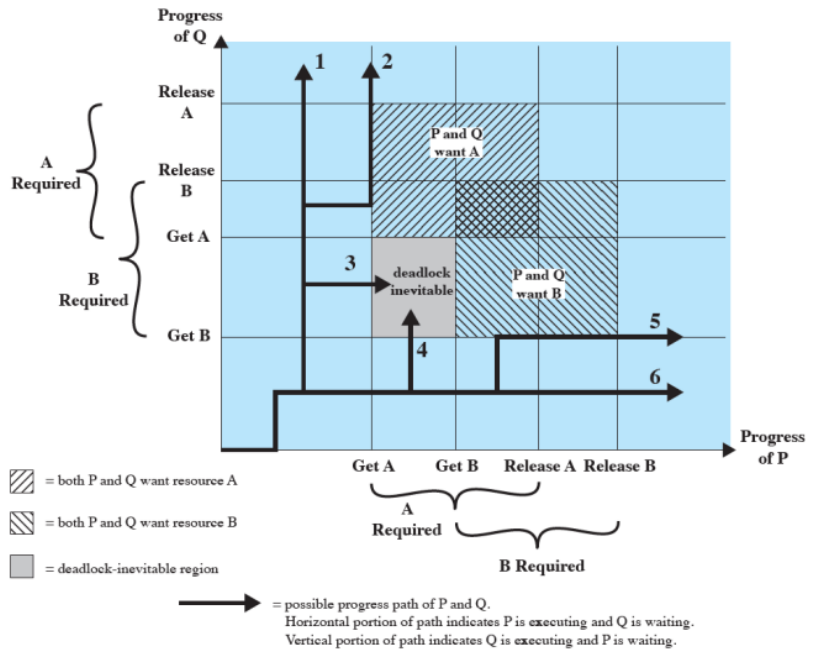
Illustration of Deadlock

- All deadlocks involve conflicting needs for resources by two or more processes.



Joint Progress Diagram – Deadlock

Process P	Process Q
...	...
Get A	Get B
...	...
Get B	Get A
...	...
Release A	Release B
...	...
Release B	Release A
...	...



Trace 1 sees process *Q* execute all of its desired operations before process *P* can perform any single operation

Trace 2 sees process *Q* get A then make process *P* wait until it is done.

Trace 3 sees *Q* get B and *P* get A. 4 does the same but in reverse order

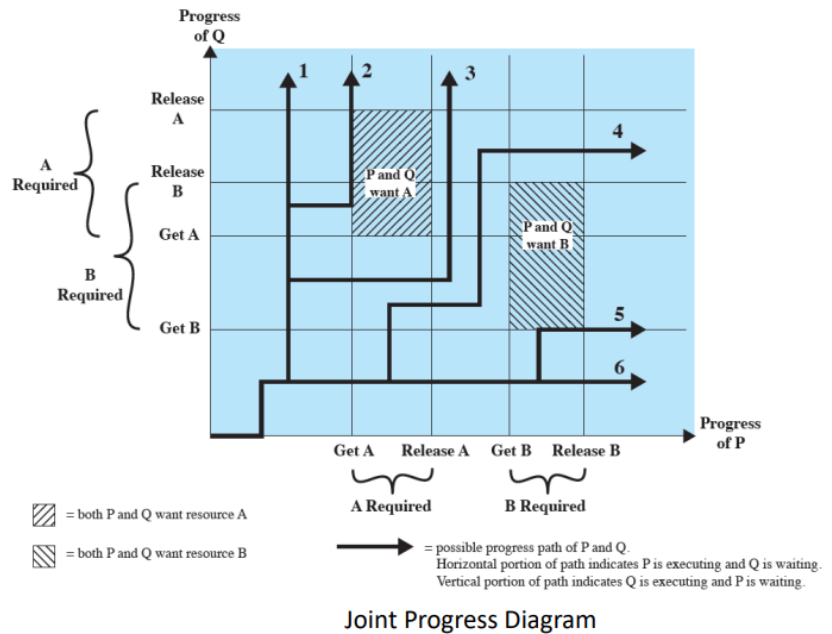
Traces 3 & 4 eventually lead to a deadlock as each process will want the other's acquired resource that they can't release until they get the resource they want.

Trace 5 sees process *P* get B then make process *Q* wait until it's done.

Trace 6 sees *P* execute everything before *Q* gets to make some meaningful progress.

Joint Progress Diagram – No Deadlock

Process P	Process Q
...	...
Get A	Get B
...	...
Release A	Get A
...	...
Get B	Release B
...	...
Release B	Release A
...	...



“I would like to leave that for the banking algorithm” - Prof, no one knows what she meant by that

This rearrangement is not always a possible solution since we don't know what the order of the instructions/operations will be beforehand.

Traces 1, 2, 5, and 6 are unaffected but now traces 3 and 4 don't lead to deadlock.

instead with Trace 3

- Q gets B
- P gets A then releases A
- Q gets A, releases B, then releases A
- P gets B then releases B

Similarly with trace 4

- 2 processes that compete for exclusive access to a disk file **D** and a tape drive **T**.
- **Questions:** Are these reusable or consumable resources? Will there be a deadlock? What sequence of execution will result in a deadlock?

Step	Process P Action	Step	Process Q Action
P ₀	Request (D)	Q ₀	Request (T)
P ₁	Lock (D)	Q ₁	Lock (T)
P ₂	Request (T)	Q ₂	Request (D)
P ₃	Lock (T)	Q ₃	Lock (D)
P ₄	Perform function	Q ₄	Perform function
P ₅	Unlock (D)	Q ₅	Unlock (T)
P ₆	Unlock (T)	Q ₆	Unlock (D)

These are reusable resources.

There will be a deadlock.

$p_0q_0p_1q_1$ introduces the possibility of a deadlock

$p_0q_0p_1q_1p_2q_2$ is where the deadlock happens.

- Memory Space is available for allocation of **200Kbytes**, and the following sequence of requests occur.
- **Question:** Are these reusable or consumable resources? When will there be a deadlock?

P1	P2
...	...
Request 80 Kbytes;	Request 70 Kbytes;
...	...
Request 60 Kbytes;	Request 80 Kbytes;

These are reusable resources.

There will be a deadlock.

P1 and P2 can make their requests then the next requests will cause a deadlock.

P1 or P2 can make both requests but then upon the next request from the other process we will run into deadlock. That is unless there is some kind of release condition that we don't see beyond the processes' second requests.

The cause of the deadlock being that we don't have enough resources to give.

- Each process is attempting to receive a message from the other process and **then** send a message to the other process:
- **Question:** Are these reusable or consumable resources? Will there be a deadlock?

P1	P2
...	...
Receive (P2);	Receive (P1);
...	...
Send (P2, M1);	Send (P1, M2);

Consumable resource.

There is a deadlock as they are both waiting on each other.

If receive is not blocking - i.e. the processes just make themselves open to receiving and don't wait - then there is no deadlock.

It depends on the nature of the receive operation.

Prof will tell us if it's blocking or not blocking.

Deadlock in Multithreaded Applications

2 mutex locks are created and initialized:

```
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;

pthread_mutex_init(&first_mutex, NULL);
pthread_mutex_init(&second_mutex, NULL);
```

2 threads are created and both threads have access to both mutex locks.

```

/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}

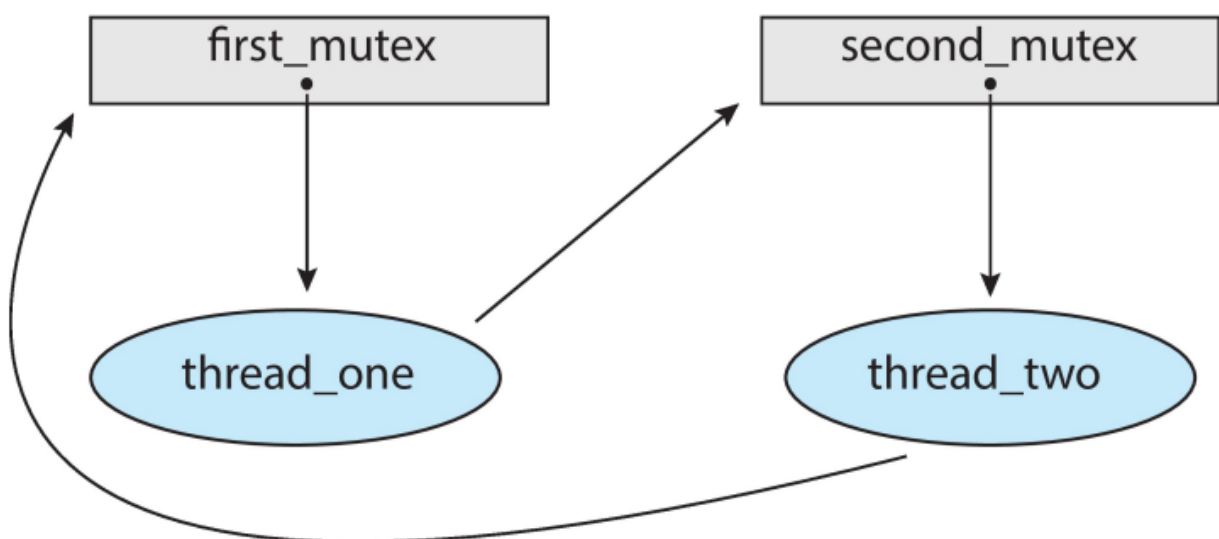
/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}

```

Deadlock is possible if thread 1 acquires `first_mutex` and thread 2 acquires `second_mutex`

Each thread then waits for the other's mutex.



The order depends on how the threads are scheduled by the CPU scheduler. Hard to test for deadlocks as they may only occur under certain scheduling circumstances.

this is the livelock of waiting continually trying to acquire

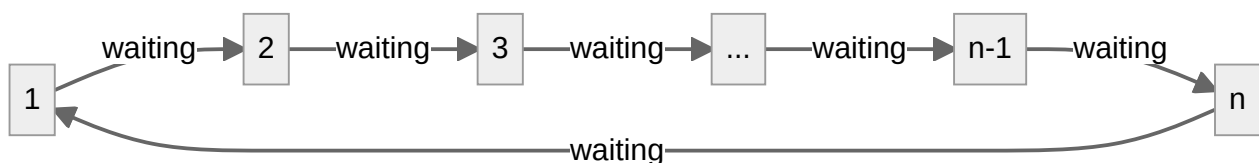
knowledge check:

- a deadlocked state occurs whenever
 - every process in a set is waiting for an event that can only be caused by another process in the set
- Deadlock occurs when every thread in a set is blocked waiting for an event that can be caused only by another thread in the set, while livelock occurs when a thread continuously attempts an action that fails.
 - true
- in the dining philosophers problem, there is a possibility of deadlock but not livelock
 - false
 - everyone picks a fork up then there is deadlock
 - if they keep picking it up and putting it down then there is livelock

Deadlock Characterization


deadlock can arise if 4 conditions hold at the same time:

- mutual exclusion
 - only one process at a time can use a resource
- hold and wait
 - a process holding at least one resource is waiting to acquire additional resources held by other processes
- no preemption (for resources not the process)
 - resource only released by the process holding it after it's done its task
 - resource cannot be preempted
 - processor can't take resources from process and give it to other processes arbitrarily
- circular wait
 - process 1 is waiting on process 2
 - process 2 is waiting on process 3
 - ...
 - process n-1 is waiting on process n
 - process n is waiting on process 1



- **All four conditions must hold** for a deadlock to occur.

Possibility of Deadlock	Existence of Deadlock
1. Mutual exclusion	1. Mutual exclusion
2. No preemption	2. No preemption
3. Hold and wait	3. Hold and wait
	4. Circular wait



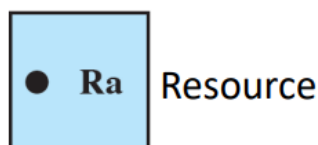
Resource Allocation Graph

- Deadlocks can be described more precisely in terms of a directed graph called a system **resource-allocation graph**.
 - A set of vertices V and a set of edges E .
 - V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the **processes** in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all **resource** types in the system

A digraph with vertices V partitioned into P processes and R resources

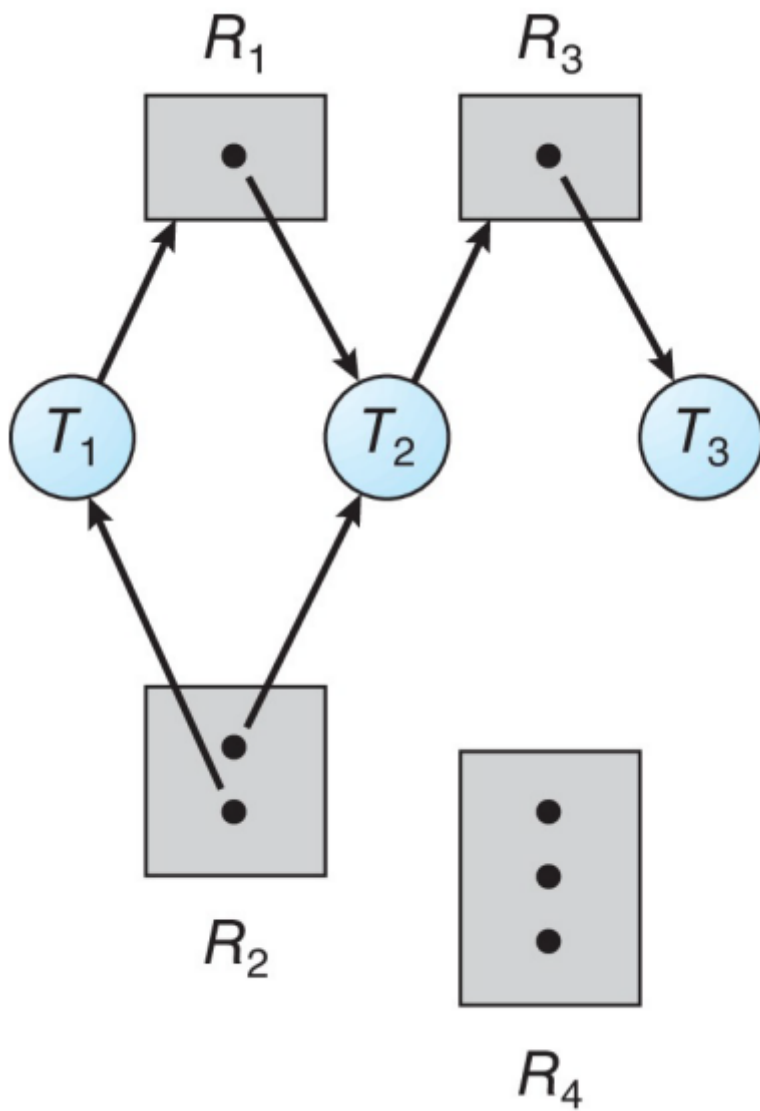
request edge – directed edge $P_i \rightarrow R_j$

assignment edge – directed edge $R_j \rightarrow P_i$



● 1 Instance of the resource





There are 3 threads and 4 resources.

.	T_1	T_2	T_3
R_1	req	assigned	
R_2	assigned	assigned	
R_3		req	assigned
R_4			

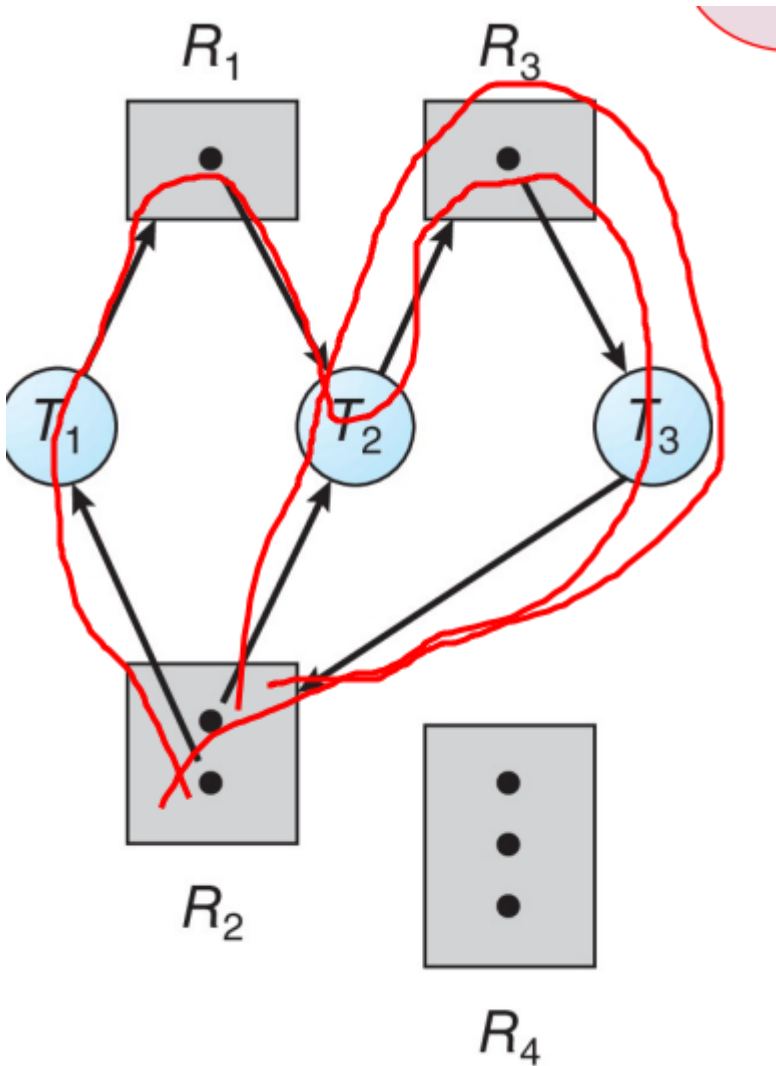
There are 1 instance of R_1 and R_2

2 of R_2 and 3 of R_4

There are no cycles in the graph.

Therefore there are no deadlocks

If there is no cycle then there is no deadlock.



There are 2 cycles.

T_1 is assigned R_2 and wants R_1

T_2 is assigned R_1+R_2 and wants R_3

T_3 is assigned R_3 and wants R_2 which it can't get since there's no more instances of R_2

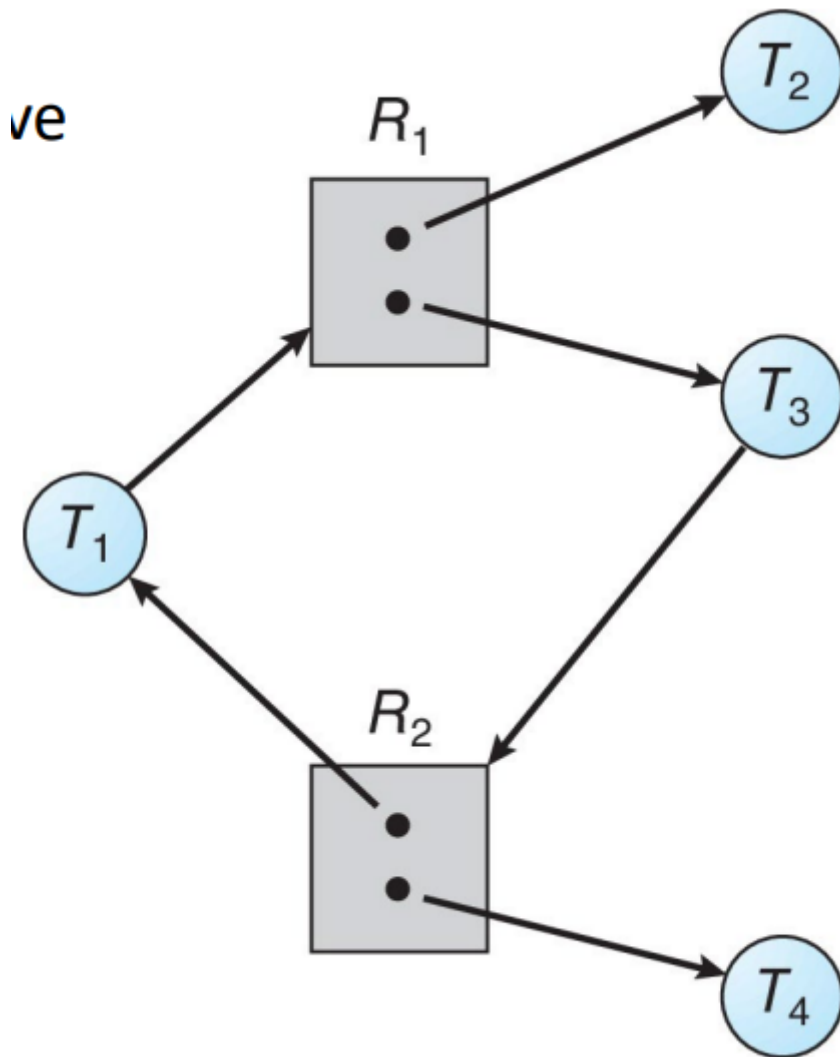
It might seem like there's no deadlock if we look at the smaller cycle.

T_2 is assigned R_2 and wants R_3

T_3 is assigned R_3 and wants R_2 which it can get since there's another instance of R_2

T_3 finishes with R_3 and gives it up for T_2 .

This is a static situation so we have a deadlock



There is a cycle with the middle 4 nodes.

There is no deadlock.

The outer threads of T_2 and T_4 can finish and release their instances of R_1 and R_2 respectively.

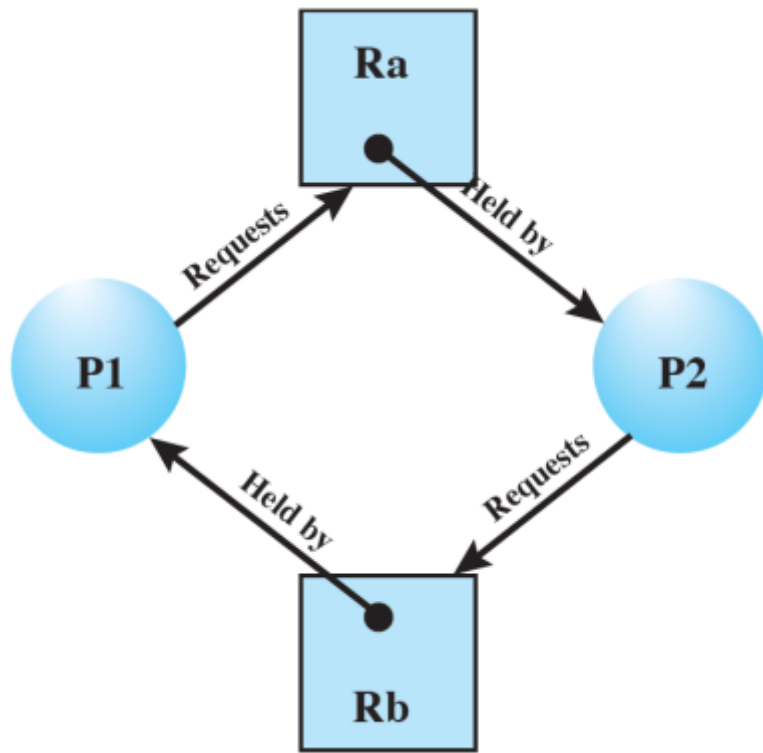
Then T_1 and T_3 can get hold of the resources and work on them.

If there is no cycle, then there is no deadlock

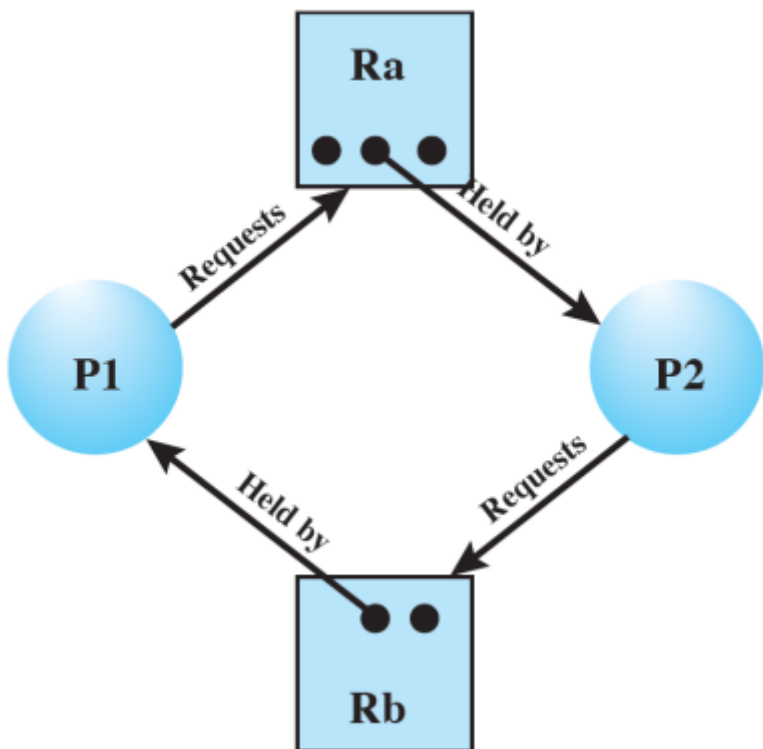
If there is a cycle then there may or may not be a deadlock

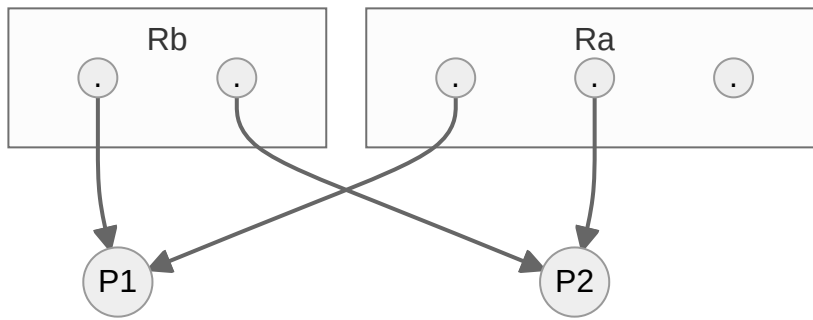
- if only one instance per resource type, then deadlock
- if there are several instances per resource type, possibility of deadlock

deadlock:




no deadlock:





Resource Allocation graph for the car deadlock example showing circular wait



We can take the numbers on the cars to be the numbers of the processes.

Car 1 is sitting lane **Ra** and wants to cross lane Rb

Car 2 is sitting lane Rb and wants to cross lane Rc

Car 3 is sitting lane Rc and wants to cross lane Rd

Car 4 is sitting lane Rd and wants to cross lane **Ra**

Knowledge Check

- one necessary condition for deadlock is ____, which states taht a process must be holding one resource and waiting to acquire additional resources
 - a: hold and wait
- a cycle in a resource-allocation graph is ____
 - d. a necessary and sufficient condition for a deadlock in the case that each resource has exactly one instance
- if a resource-allocation graph has a cycle, the system must be in a deadlocked state.
 - false

Methods for Handling Deadlocks

In general, there are 3 ways to deal with deadlocks:

1. prevention and avoidance
2. detection and recovery
3. ignoring it

most OSes including windows and linux just pretend that deadlocks don't happen ???

Prevention

In order to prevent deadlocks we have to invalidate one of the 4 necessary conditions for deadlock

as a refresher

deadlock can arise if 4 conditions hold at the same time:

- mutual exclusion
 - only one process at a time can use a resource
- hold and wait
 - a process holding at least one resource is waiting to acquire additional resources held by other processes
- no preemption (for resources not the process)
 - resource only released by the process holding it after it's done its task
 - resource cannot be preempted
 - processor can't take resources from process and give it to other processes arbitrarily
- circular wait
 - process 1 is waiting on process 2
 - process 2 is waiting on process 3
 - ...
 - process n-1 is waiting on process n
 - process n is waiting on process 1

Though there are shareable resources, we cannot invalidate mutual exclusion as a whole since some resources just cannot be shared and used at the same time.

Deny Hold and Wait

We can invalidate hold and wait by ensuring that whenever a process requests a resource it won't be holding any other resources. Put another way, If the process is holding a resource then they won't be making any requests.

how:

- require process to request and allocate all resources before execution
 - **dynamic nature of resource requesting makes this impractical**

- only allow resource requests when process has no resource allocated
 - must release resources before requesting for more

cons:

- low resource utilization
 - resources allocated but not used
- starvation possible
 - a thread may have to wait indefinitely if it needs several popular resources
 - those resources will always end up allocated to processes that will only need those resources
 - I imagine that even with bounded wait time considerations, priority will cause this to happen as well

Deny No Preemption

reminder: non-preemptive is when a process enters into the processor and only leaves when it's finished or voluntarily leaves when waiting on something

if

- a process is holding some resources
- the process requests another resources
- that resource cannot be immediately allocated to it

then **all resources currently being held are preempted**, released by the processor.

Preempted resources are added to the list of resources that the process is waiting for.

The process starts again once it regains all of the resources it's requesting (the old preempted ones and the new ones).


This is often *selectively* applied to resources whose state can be easily saved and restored later, such as cpu registers and database transactions.

cons:

- cannot be applied to resources like mutex locks and semaphores which are the type of resources where deadlock occurs most commonly
 - it's hard to save the condition of a mutex when a process wants to lock for a critical section
 - not easy to store and retrieve → not easy to preempt
 - we can't do this for the shit we actually care about bruhhhhhh

Deny Circular Wait

This is the most common approach



how:

- assign each resource a unique number
- resources must be acquired in the order of their unique number

Knowledge check:

- to handle deadlocks, operating systems most often ____
 - a. pretend that deadlocks never occur
- both deadlock prevention and deadlock avoidance techniques ensure that the system will never enter a deadlocked state
 - true
- most operating systems choose to ignore deadlocks, because
 - d. all of the above
 - handling is expensive, they occur infrequently, livelock recovery methods can be use on deadlocks

Deadlock Avoidance

Dynamically made decision to see if an allocation request will potentially lead to a deadlock if granted.

System considers

- currently available resources
- currently allocated resources
- and the future requests and releases of each thread

requires knowledge of future process requests.

basics:

- state
 - reflects current allocation of resources to processes
- safe state
 - state where there is at least one execution path where all the processes will finish
 - there will be no deadlock
- unsafe state
 - state resulting in possibility of deadlock
- avoidance

- ensuring system will never enter an unsafe state

Safe State


When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.

System is in **safe state** if:

- There exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$

That is:

- If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on..
-
- If a system is in **safe state** \Rightarrow **no deadlocks**
 - If a system is in **unsafe state** \Rightarrow **possibility of deadlock**
 - **Avoidance** \Rightarrow ensure that a system will never enter an unsafe state.



- A system has 12 resources and 3 threads T_0 , T_1 , and T_2 shown below.
- At t_0 , T_0 is holding 5 resources, T_1 holding two, and T_2 holding two.
- **Question:** Does the sequence $\langle T_1, T_0, T_2 \rangle$ satisfies the safety condition?
 - T_1 : $2 + 2 \rightarrow$ done, release back 4 \rightarrow 5 available.
 - T_0 : $5 + 5 \rightarrow$ done, release back 10 \rightarrow 10 available.
 - T_2 : $2 + 9 \rightarrow$ done, release back 9 \rightarrow 12 available.

	<u>Maximum Needs</u>	<u>Current Needs</u>
T_0	10	5
T_1	4	2
T_2	9	2

In the above example we only have 3 resources available.

For T_1 we give it 2 then release back 4 once it's done.

Then the above steps just go as shown.

So we are in a safe state

- A system has 12 resources and 3 threads T_0 , T_1 , and T_2 shown below.
- At t_1 , thread T_2 requests and is allocated one more resource.
- **Question:** Is the system now in a safe state?
 - T_2 : 3, may request 6 \rightarrow at t_1 : 2 available.
 - T_0 : 5, may request 5 \rightarrow at t_1 : 2 available.
 - T_1 : $2 + 2 \rightarrow$ done, release back 4 \rightarrow at t_2 : 4 available.

	<u>Maximum Needs</u>	<u>Current Needs</u>
T_0	10	5
T_1	4	2
T_2	9	3

This is the same example as seen above but now T_2 needs 1 more resource, requests the resource and receives it, resulting in only 2 resources available at the beginning.

We can start the sequence with T_1 and allocate 2 then get back 2 once done.

We end up with only 4 available resources which isn't enough to satisfy the maximum needs of T_0 and T_2

Avoidance Algos

There are 2 algos:

1. using a resource-allocation when there's a single instance of a resource type
2. use the Banker's Algo when there's multiple instances of a resource type

Resource-Allocation Graph Scheme

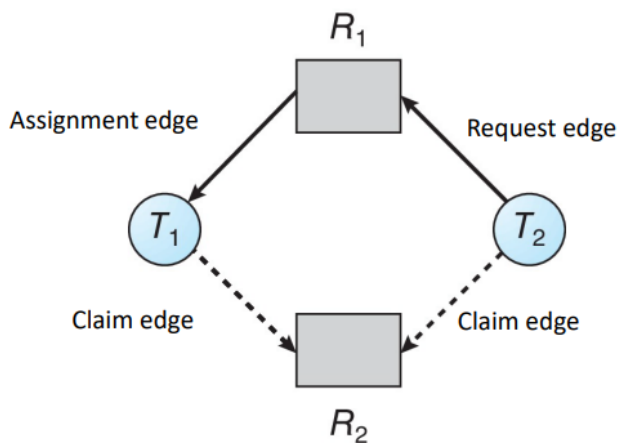
Claim edge $P_i \rightarrow R_j$ indicated that process P_i **may request** resource R_j represented by a dashed line

- Claim edge converts to **request edge** when a process requests a resource
- Request edge converted to an **assignment edge** when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge

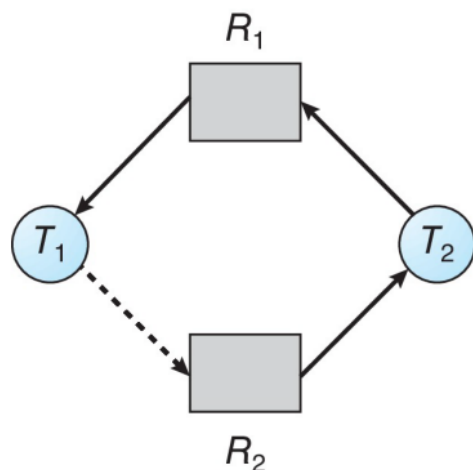
Resources must be claimed *a priori* in the system.

a priori = ahead of time

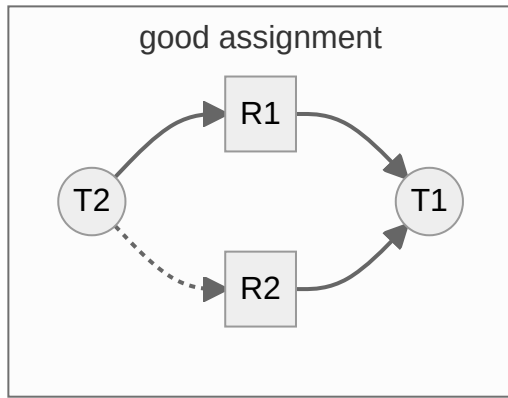
claims turn into requests which turn into assignments before turning back into claim edges



Although R2 is currently free,
we cannot allocate it to T2.



An unsafe state in a resource-allocation graph.



Suppose that process P_i requests a resource R_j

The request can be granted **only** if converting the request edge to an assignment edge **does not result in the formation of a cycle** in the resource allocation graph.

make sure there is no cycles forming if you convert a request edge to an assignment edge

Banker's Algorithm

prof misspells this is as Baker at some point. It's Banker.

conditions:

- multiple instances of resources
- each process must have a prior claim maximum use
- when a process requests a resource, it may have to wait
- when a process gets all its resources it must return them in a finite amount of time

Let n = number of processes, and m = number of resources types.

- **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- **Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- **Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- **Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$

Safety Algorithm

1. Let **Work** and **Finish** be vectors of length m and n , respectively. Initialize:
Work = **Available**
Finish [i] = **false** for $i = 0, 1, \dots, n-1$
2. Find an index i such that both:
(a) **Finish** [i] = **false**
(b) **Need** _{i} ≤ **Work**
If no such i exists, go to step 4
3. **Work** = **Work** + **Allocation** _{i}
Finish [i] = **true**
go to step 2
4. If **Finish** [i] == **true** for all i , then the system is in a **safe** state otherwise we are in **unsafe** state.

Request _{i} = request vector for process P_i If **Request** _{i} [j] = k then process P_i wants k instances of resource type R_j

1. If **Request** _{i} ≤ **Need** _{i} go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
 2. If **Request** _{i} ≤ **Available**, go to step 3. Otherwise, P_i must wait, since resources are not available.
 3. Pretend to allocate requested resources to P_i by modifying the state as follows:
Available = **Available** – **Request** _{i} ;
Allocation _{i} = **Allocation** _{i} + **Request** _{i} ;
Need _{i} = **Need** _{i} – **Request** _{i} ;
 - If **safe** ⇒ the resources are allocated to P_i (*use Safety alg. in the prev slide*).
 - If **unsafe** ⇒ P_i must wait, and the old resource-allocation state is restored.
-

Example of Banker's Algorithm /1

- 5 processes P_0 through P_4 ;
- 3 resource types:
 A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

	<u>Need</u>
	$A \ B \ C$
P_0	7 4 3
P_1	1 2 2
P_2	6 0 0
P_3	0 1 1
P_4	4 3 1

Q: Does the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfy safety criteria?

A: yes

we start with $\langle 3, 3, 2 \rangle$

P_1 can be allocated and give back $\langle 5, 3, 2 \rangle$

P_3 can be allocated and give back $\langle 7, 4, 3 \rangle$

P_4 can be allocated and give back $\langle 7, 4, 5 \rangle$

P_2 can be allocated and give back $\langle 10, 4, 7 \rangle$, making all instances of A and C available

P_0 can be allocated and give back $\langle 10, 5, 7 \rangle$, making all instances of all resources available.

Example of Banker's Algorithm /3

- Does the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfy safety criteria?

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	3 3 2
P_1	2 0 0	1 2 2	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ **satisfies** safety requirement.

you do the same process with this.

A: the sequence satisfies safety requirement

Example: P_1 Request (1,0,2)

- Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow$ true

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Is the new state safe?
- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ **satisfies** safety requirement.
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?

prof will pick this up next time

Recovery from Deadlock

