

this covers the **virtual memory** slide deck

toc:

- Background
- Demand Paging
- Copy on Write
- Page Replacement
 - Page Replacement Algorithms
 - Page Replacement Algo Evaluation
 - Different Types
 - FIFO
- Allocation of Frames
- Thrashing
- Allocating Kernel Memory

Background

code needs to be in physical memory to execute but entire program is rarely used

- error code
 - used to handle unusual error conditions
 - almost never executed
- large data structures
 - often allocated more mem than needed
- unusual routines
 - lesser used options and features of a program

we only have to load what we need for execution


ability to execute partially-loaded program

- no longer limited by physical mem
- each program takes less memory while running
 - more programs can run at the same time
 - cpu util and throughput increase
 - response time and turnaround time the same
- less i/o needed to load or swap programs into mem
 - each user program runs faster

virtual mem

- separation of user logical mem from physical mem
- only part of the program needs to be in mem for execution
- logical address space can be larger than physical address space


- allows address spaces to be shared by several processes
- allows for more efficient process creation
- more programs running concurrently
- less i/o needed to load or swap processes



Virtual Memory That is Larger Than Physical Memory

Virtual Address Space

- **Virtual address space** – logical view of how process is stored in memory:
 - Usually start at address 0, contiguous addresses until end of space.
 - Meanwhile, physical memory organized in page frames.
 - MMU must map logical pages to physical page frames in memory.
- Logical memory is also known as virtual address space.



Virtual address space of a process in memory.

The large blank space (or hole) between the heap and the stack is part of the virtual address space.

- Will require actual physical pages only if the heap or stack grows.

Virtual address spaces that include holes are known as **sparse address spaces**.


- Beneficial because the holes can be filled as the stack or heap segments grow or if we wish to dynamically link libraries during program execution.

Shared Library Using Virtual Memory

- Virtual memory allows files and memory to be shared by two or more processes through **page sharing**.

- Examples:

- Library mapped as read-only.
- Shared memory for 2 processes to communicate.



Shared library using virtual memory.

Logical Memory vs. Virtual Memory vs. Real Memory

- Logical memory is the memory space perceived by a process.
- Real memory is the physical RAM installed in the system.
- Virtual memory is a memory management technique used by the operating system to extend available memory beyond physical RAM.

- logical memory
 - memory space perceived by a process
- real/physical memory

- physical ram sticks in you're machine
 - virtual memory
 - memory management technique
 - used by operating system
 - extends beyond available physical ram
-

Knowledge check:

- which of the following is a benefit of allowing a program that is only partially in memory to execute?
 - d. all of the above
 - programs can be written to use more memory than is available in physical memory
 - cpu util and throughput is incdreaded
 - less io needed to load or swap each user program into mem
- in general virtual memory decreases the degree of multiprogramming in a system
 - false
 - why would these even effect each other???


Demand Paging

virtual memory can be implemented via

- demand paging
- demand segmation

demand paging:

- bring a page into memeory ONLY when it is needed
 - less required
 - I/O
 - memory
 - faster response
 - more users
- similar to a paging system with swapping

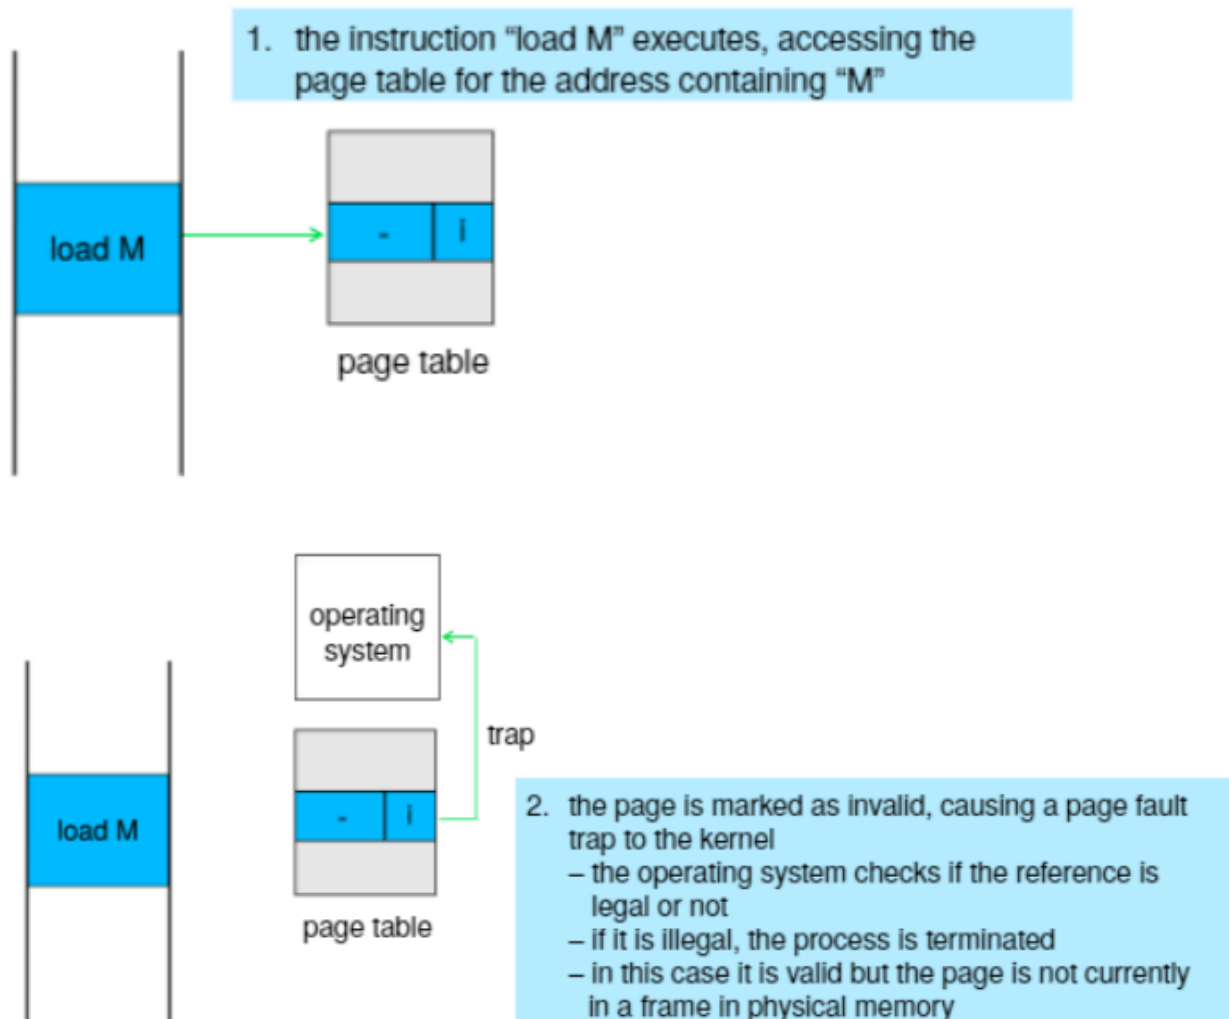


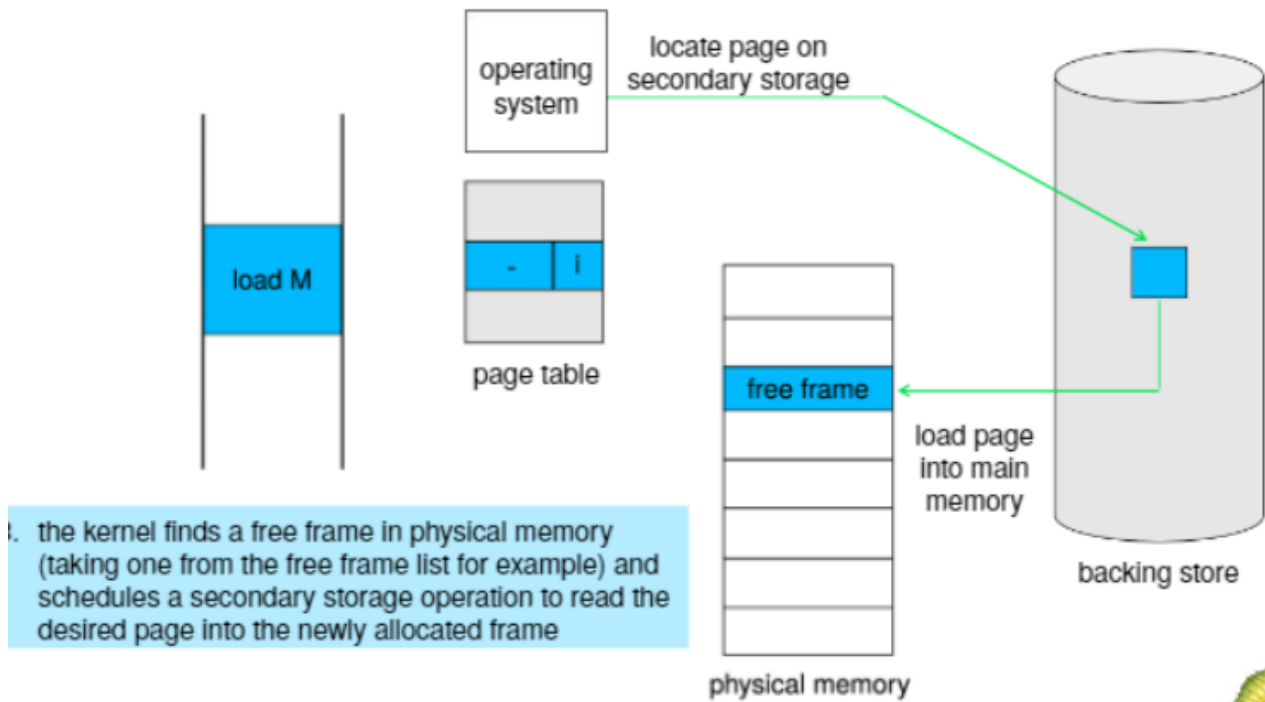
Swapping with paging.

valid-invalid bit

- with each page entry there is a valid-invalid bit
 - v → in-memory - memory resident
 - i → not-in-memory
- initially valid-invalid bit is set to **i** on all entries
 - all invalid
- during MMU address translation
 - if bit is invalid for a page entry then there's a page fault

Steps in Handling a Page Fault





valid

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical memory

valid-invalid bit	
frame	
0	4 v
1	i
2	6 v
3	i
4	i
5	9 v
6	i
7	i

page table

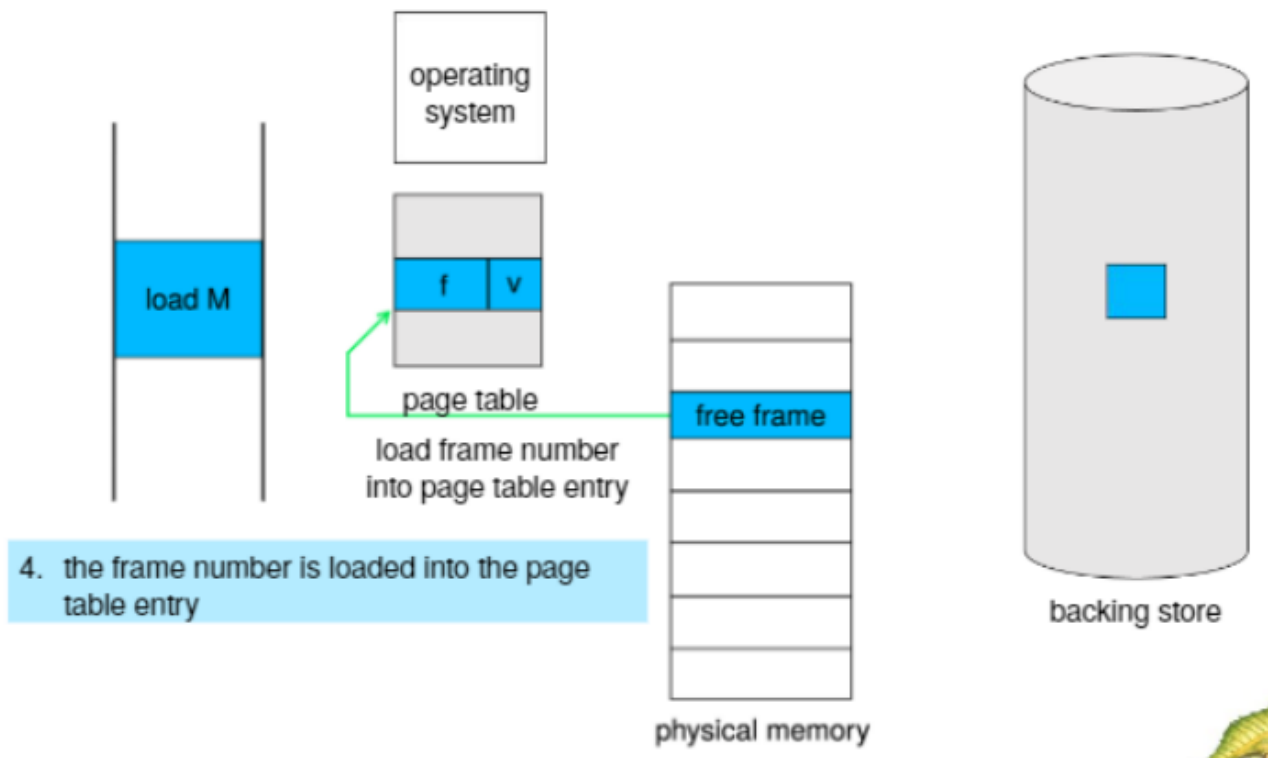
0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

physical memory

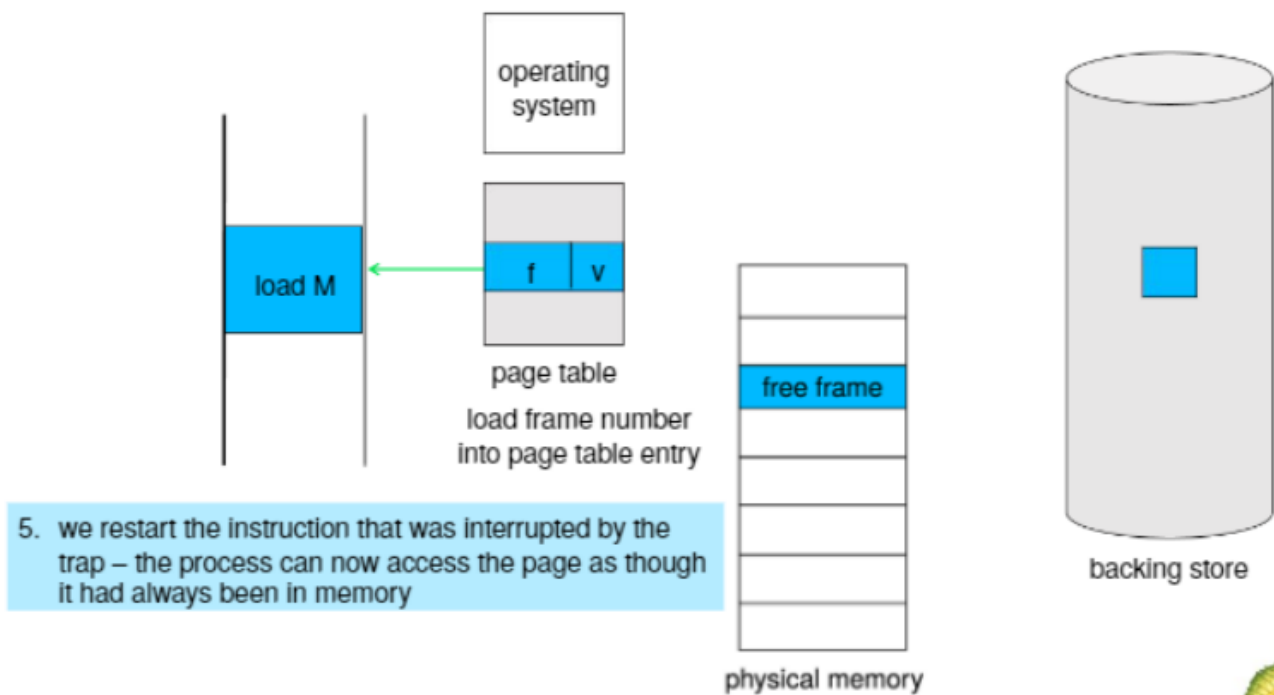


d-
e

Page Table When Some Pages Are Not in Main Memory



Steps in Handling a Page Fault




there are some more slides that are the same thing

free-frame list

- a pool of free frames for satisfying such requests
- maintained by most operating systems

- when a page fault occurs, the operating system must bring the desired page from secondary storage into main memory
- OS typically allocate free frames using zero-fill-on-demand
 - we flush out the frame
 - the content of the frames zeroed-out before being allocated
- when a system starts up, all available memory is placed on the free-frame list
 - this lets us start allocating frames as they are demanded



performance of demand paging

- 3 major activities of servicing a page-fault
 - service the page-fault interrupt
 - read in the page
 - this is where most of the time is spent
 - restart the process
- page fault rate $0 \leq p \leq 1$
 - $p = 0 \rightarrow$ no page faults
 - $p = 1 \rightarrow$ every reference is a fault
- Effective Access Time (EAT) =
 - $(1 - p) \times \text{memory access} + p(\text{page fault overhead} + \text{swap page out} + \text{swap page in})$
 - we're basically averaging it out

Example of Demand Paging

memory access time = 200 ns

avg page-fault service time = 8 ms

$$\begin{aligned}
 \text{EAT} &= (1 - p) \times 200 + p(8\text{ms}) \\
 &= ((1 - p) \times 200) + (p \times 8,000,000) \\
 &= 200 + (p \times 7,999,800)
 \end{aligned}$$

if 1 in 1000 accesses causes a page fault ($p = 0.001$) then EAT = 8.2 microseconds. a slowdown by a factor of 40

if we want a performance degradation (slowdown) less than 10 percent

- 10% of 200 (effective access time we want) = 20 ns
- $220 > 200 + 7,999,800 \times p$
- $20 > 7,999,800 \times p$
- $p < 0.0000025$
- $p <$ one page fault in every 400,000 memory accesses

This needs hardware support, it can't all be done by the OS.

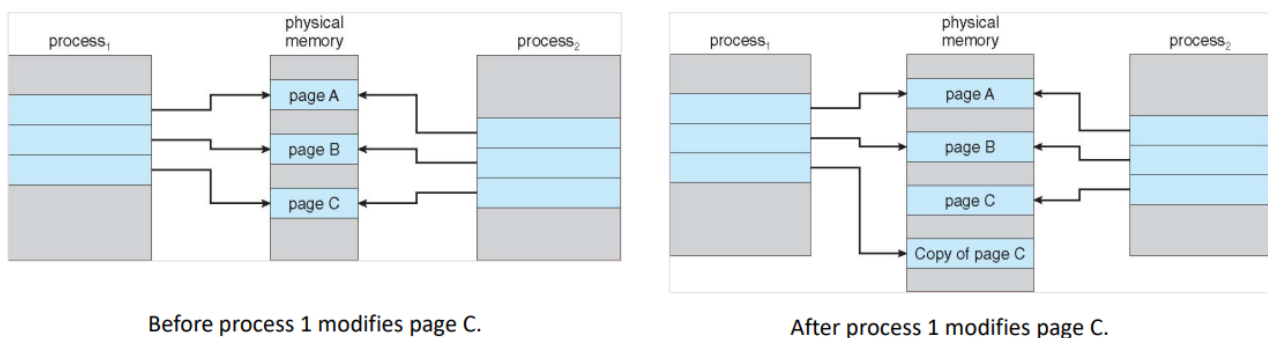
Knowledge Check

- on a system with demand-paging, a process will experience a high page fault rate when the process begins execution
 - true
 - similar principle to caching where we need to slowly fill things up in the memory that the process asks for
 - since the page table is invalid to start with we end up faulting a ton as we're setting things up for the process
 - as it starts going on it'll start hitting the pages we already loaded for it
- a page fault occurs when
 - c. a process tries to access a page that is not loaded in memory
- if memory access time is 250 ns and average page fault service time 10 ms the probability of page faults must be less ____ to keep the performance degradation less than 20%
 - options
 - a) 0.0000025
 - b) 0.000005
 - c) 0.0000075
 - d) 0.00001
 - takehome exercise

Copy on Write

copy-on-write (COW):

- allows both parent and child processes to initially share the same pages in memory
- the page is copied once either process modifies a shared page
 - read: write
- allows more efficient process creation as only modified pages are copied



Knowledge Check

- ____ allows the parent and child processes to initially share the same pages, but when either process modifies a page, a copy of the shared page is created.

- a. copy-on-write

Page Replacement

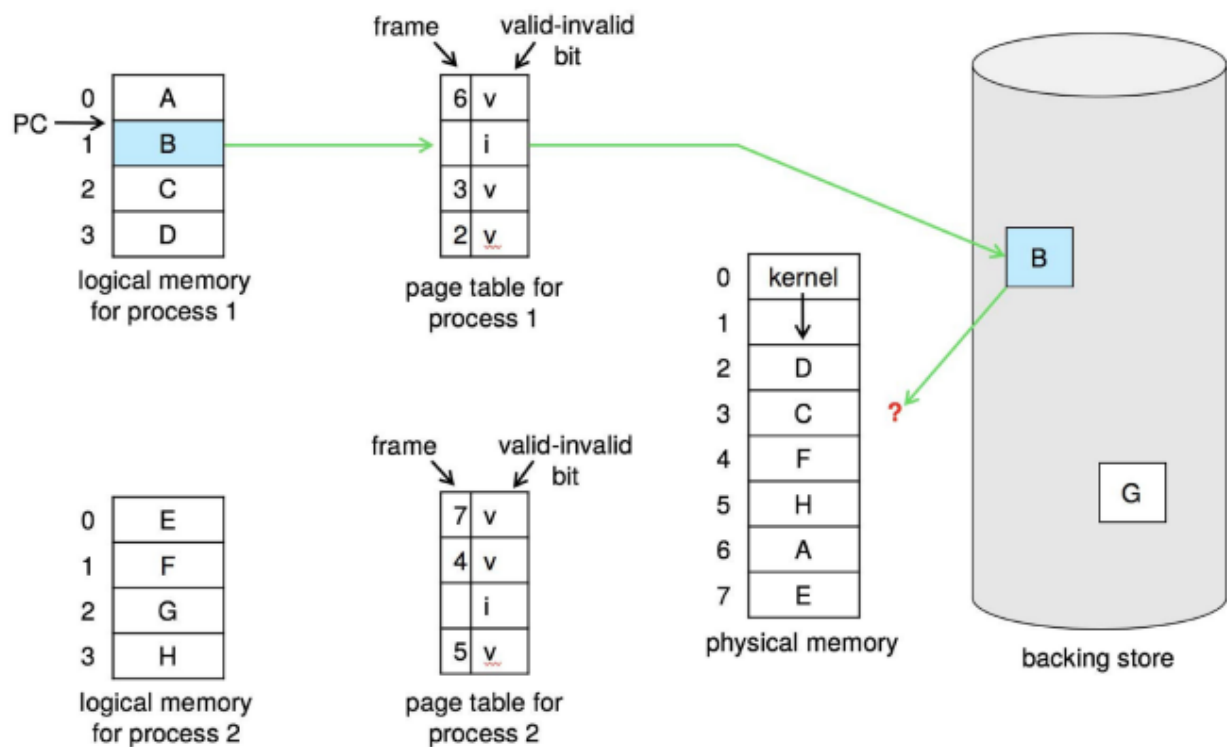
sometimes there is no free frame

how:

- memory is used up by process pages but also in demand from the kernel, I/O buffers, etc
- the page is a hot commodity that everyone needs access to

scenario

- process executing
- page fault occurs
- OS determines where the desired page is residing on secondary storage
- but there are no free frames on the free-frame list
 - all mem is in use



need for page replacement

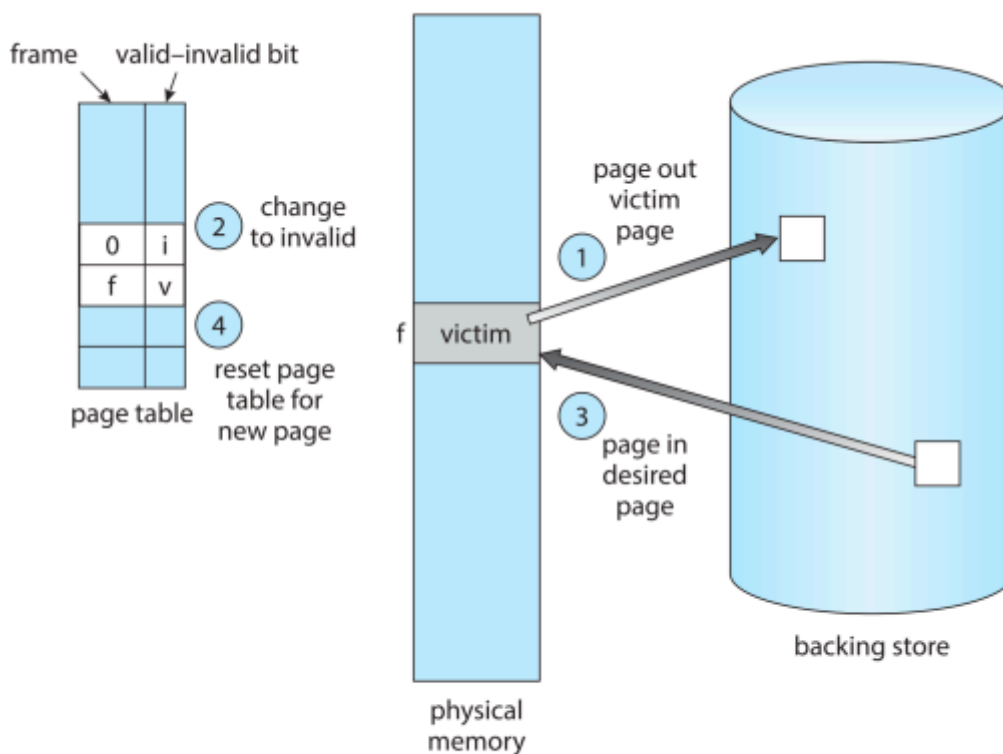
- algorithms
 - do we terminate it
 - do we swap out the whole process till there are more free frames
 - do we replace the page

- we can find some page in memory but not really in use, page it out
- most operating systems now combine swapping pages with page replacement

we want an algo that has the minimum number of page faults

basic page replacement:

- prevent over-allocation of mem
 - modify page-fault service routine to include page replacement
- use modify (dirty) bit to reduce the overhead of page transfers
 - only modified pages are written to disk
 - unmodified pages have no need to be written back to backing store since it wouldn't make a difference
- page replacement completes separation b/w logical mem and phys mem
 - large virt mem can be provided on a smaller phys mem



Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame:
 - a) If there is a free frame, use it
 - b) If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - c) Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables.
4. Continue the process from where the page occurred.

frame allocation algo determines how many frames given to each process

- some small processes are smaller
- are frames reserved
- etc

page-replacement algo

- which pages to replace
- want lowest page-fault rate on both first access and re-access

Page Replacement Algorithms

Page Replacement Algo Evaluation

evaluate algo by running it on a particular string of memory reference (reference string) and computing the number of page faults on that string


- string is just page numbers, not full addresses
- if we have a reference to a page p , then any references to page p that immediately follow will never cause a page fault
- results depend on the number of frames available

- Example: if we trace a particular process, we may record the address sequence:

0100, 0432, 0101, 0612, 0102, 0103, 0104, 0101, 0611, 0102, 0103, 0104, 0101, 0610, 0102, 0103, 0104, 0101, 0609, 0102, 0105

- This sequence is reduced to the following reference string:

1, 4, 1, 6, 1, 6, 1, 6, 1, 6, 1



Graph of Page Faults Versus The Number of Frames (in general)

Different Types

algos:

- FIFO
- Optimal
- Least Recently Used (LRU)
- LRU-Approximation
 - Second-Chance
 - Enhanced Second-Chance
- Counting-Based

In all our examples, the reference string of referenced page numbers is: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

FIFO

first-in-first-out algo

Reference string: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

3 frames (3 pages can be in memory at a time per process)

How to track ages of pages? Just use a FIFO queue.


7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2	2	4	4	4	0			0	0			7	7	7
	0	0	0		3	3	3	2	2	2			1	1			1	0	0
		1	1		1	0	0	0	3	3			3	2			2	2	1

page frames

FIFO: 15 page-faults

FIFO Illustrating Belady's Anomaly

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
 - Adding more frames can cause more page faults! → **Belady's Anomaly**



Page-fault curve for FIFO replacement on a reference string.

Allocation of Frames

Thrashing

Allocating Kernel Memory